

FUSING APACHE SPARK AND LUCENE FOR NEAR-REALTIME PREDICTIVE MODEL BUILDING

Debasish Das
Principal Engineer
Verizon

Pramod Lakshmi Narasimha
Principal Engineer
Verizon

Contributors

Platform: Pankaj Rastogi, Venkat Chunduru, Ponrama Jegan, Masoud Tavazoei

Algorithm: Santanu Das, Debasish Das (Dave)

Frontend: Altaff Shaik, Jon Leonhardt



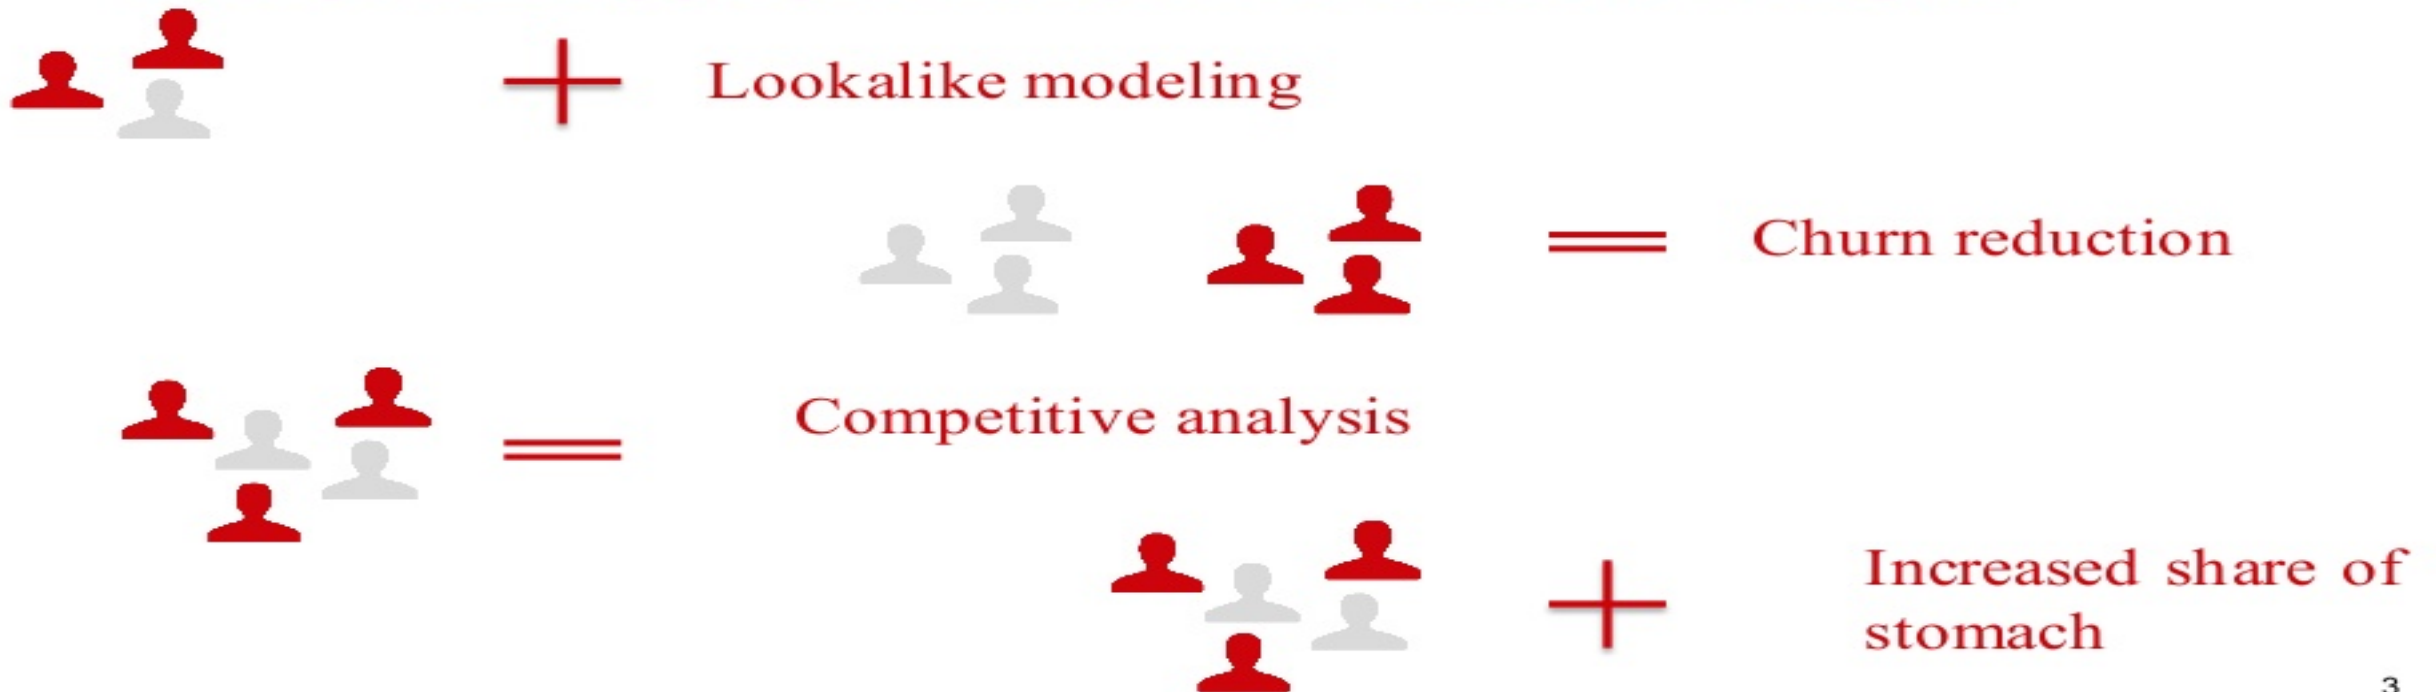
Data Overview

- **Location data**
 - Each srcIp defined as unique row key
 - Provides approximate location of each key
 - Timeseries containing latitude, longitude, error bound, duration, timezone for each key
- **Clickstream data**
 - Contains clickstream data of each row key
 - Contains startTime, duration, httpHost, httpUri, upload/download bytes, httpMethod
 - Compatible with IPFIX/Netflow formats



Marketing Analytics

- Anonymous aggregate analysis for customer insights



Data Model

- Dense dimension, dense measure

Schema: srcip, date, hour, tld, zip, tldvisits, zipvisits
Data: 10.1.13.120, d1, H2, [macys.com](#), 94555, 2, 4

- Sparse dimension, dense measure

Schema: srcip, date, tld, zip, clickstreamvisits, zipvisits
Data: 10.1.13.120, d1, {[macys.com](#), [kohls.com](#)}, {94555, 94301}, 10, 15

- Sparse dimension, sparse measure

Schema: srcip, date, tld, zip, tldvisits, zipvisits
Data: 10.1.13.120, d1, {[macys.com](#), [kohls.com](#)}, {94555, 94301}, {[macys.com](#):4, [kohls.com](#):6}, {94555:8, 94301:7}
Schema: srcip, week, tld, zip, tldvisits, zipvisits
Data: 10.1.13.120, week 1, {[macys.com](#), [kohls.com](#)}, {94555, 94301}, {[macys.com](#):4, [kohls.com](#):6}, {94555:8, 94301:7}

- Sparse dimension, sparse measure, **last N days**

Schema: srcip, tld, zip, tldvisits, zipvisits
Data: 10.1.13.120, {[macys.com](#), [kohls.com](#)}, {94555, 94301}, {[macys.com](#):4, [kohls.com](#):6}, {94555:8, 94301:7}

- Competing technologies: PowerDrill, Druid, LinkedIn Pinot, EssBase



Document Dataset Representation

- Example

Schema: srcip, tld, zip, tldvisits, zipvisits

Data: 10.1.13.120, {macys.com, kohls.com}, {94555, 94301}, {macys.com:4, kohls.com:6}, {94555:8, 94301:7}

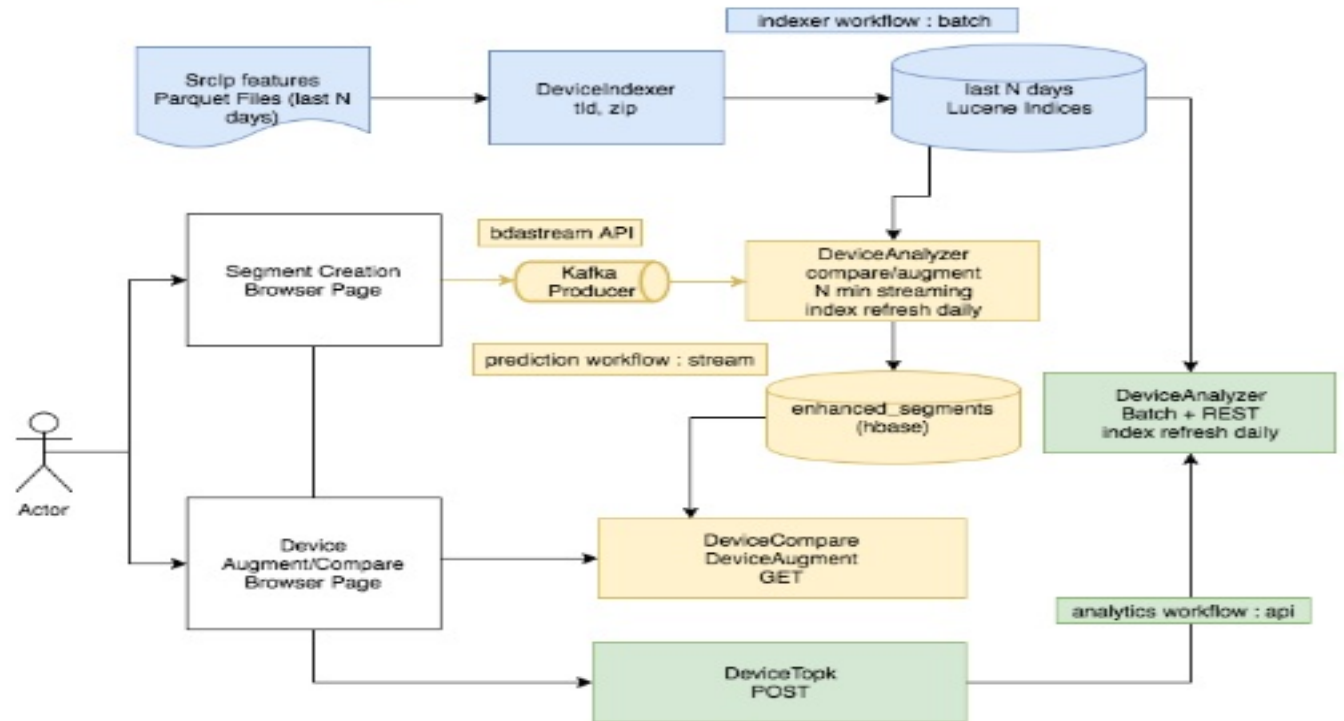
- DataFrame row to Lucene Document mapping

<i>Store/schema</i>	<i>Row</i>	<i>Document</i>
<i>srcip</i>	<i>primary key</i>	<i>docId</i>
<i>tld</i> <i>zip</i>	<i>String</i> <i>Array[String]</i>	<i>SingleValue/MultiValue</i> <i>Indexed Fields</i>
<i>tldvisits</i> <i>zipvisits</i>	<i>Double</i> <i>Map[String, Double]</i>	<i>SparseVector</i> <i>StoredField</i>

- Distributed collection of srcip as RDD[Document]
 - ~100M srcip, 1M+ terms (sparse dimensions)

DeviceAnalyzer

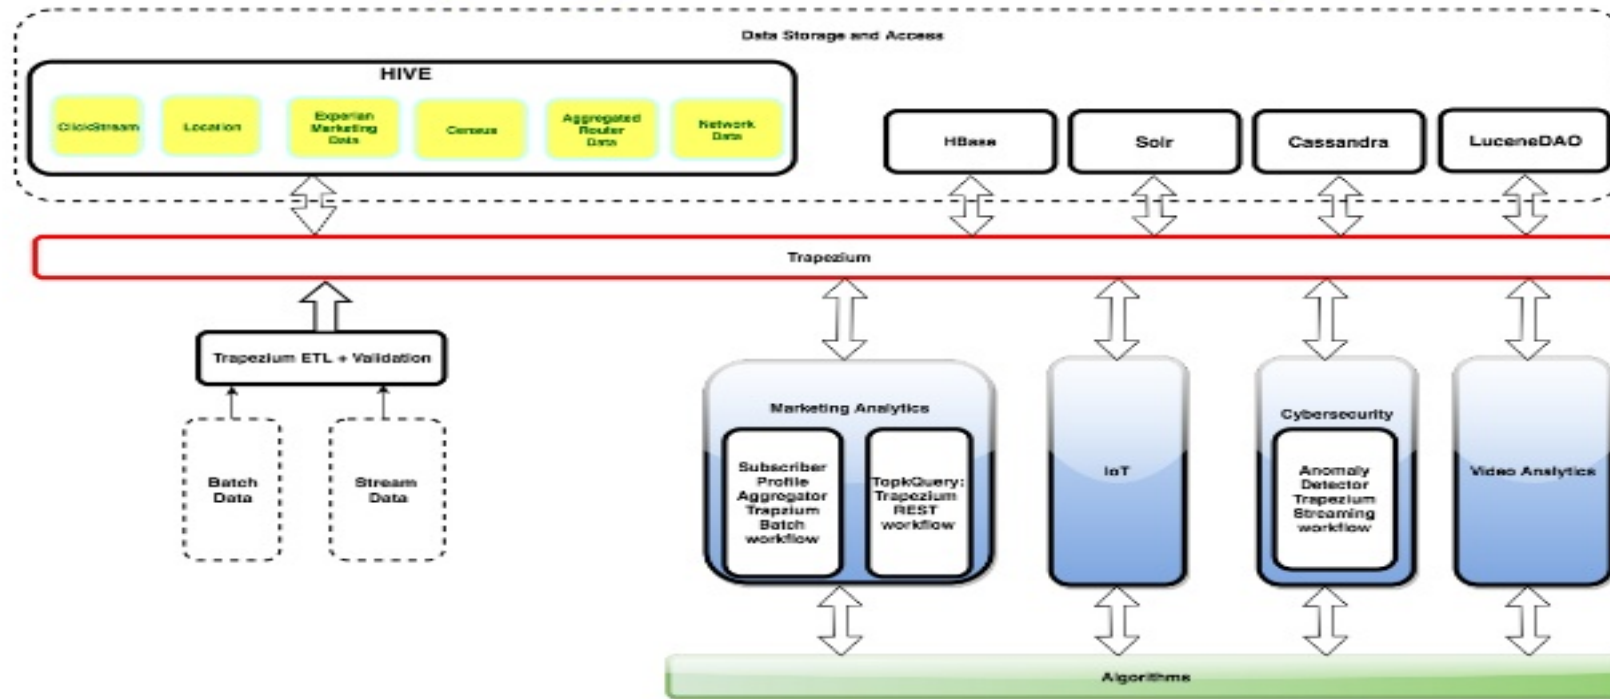
- DeviceAnalyzer goals
 - Search and retrieve devices that matched query
 - Generate statistical and predictive models on retrieved devices




What is Trapezium ?

DAIS Open Source framework to build batch, streaming and API services

<https://github.com/Verizon/trapezium>



Trapezium Architecture

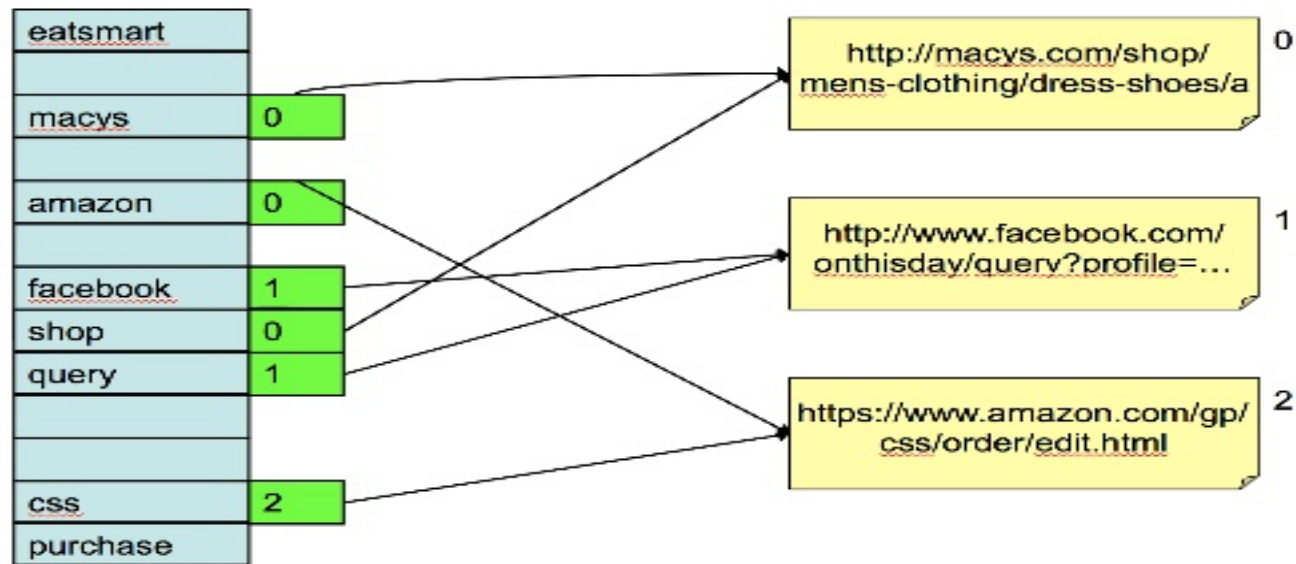


© Verizon 2016 All Rights Reserved

Information contained herein is provided AS IS and subject to change without notice. All trademarks used herein are property of their respective owners.

Lucene Overview

- Scalable, full-text search library
- Focus: Indexing + searching documents



Trapezium LuceneDAO

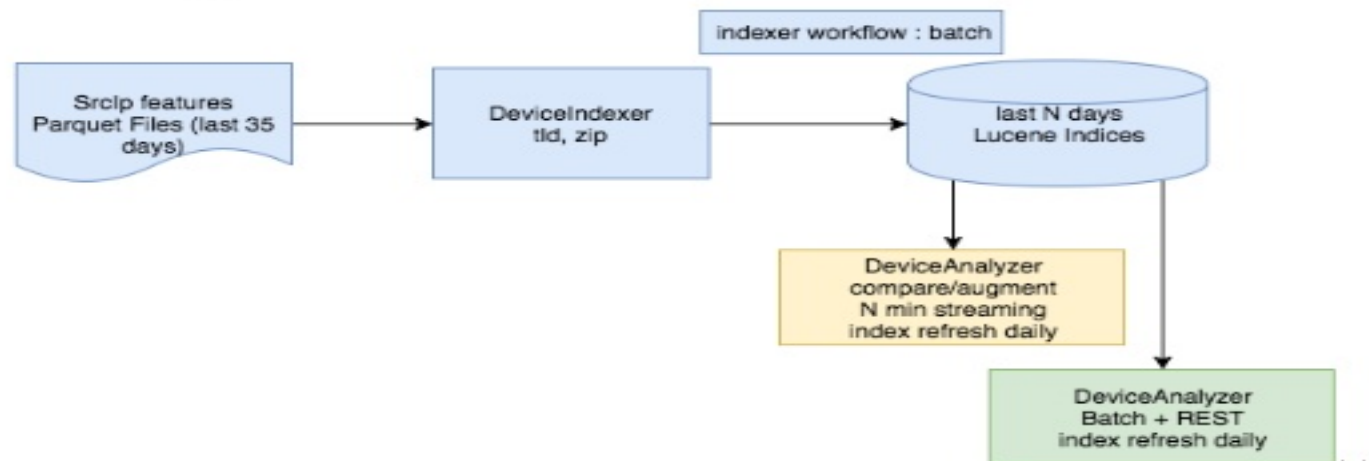
- SparkSQL and MLlib optimized for full scan, **column indexing not supported**
- Why Spark + Lucene integration
 - Lucene is battle tested Apache Licensed Open Source Project
 - Adds column search capabilities to Spark
 - Adds spark operators (treeAggregate, treeReduce, map) to Lucene
- LuceneDAO features
 - Build distributed lucene shards from Dataframe
 - Save shards to HDFS for QueryProcessor (CloudSolr)
 - Access saved shards through LuceneDAO for ML pipelines



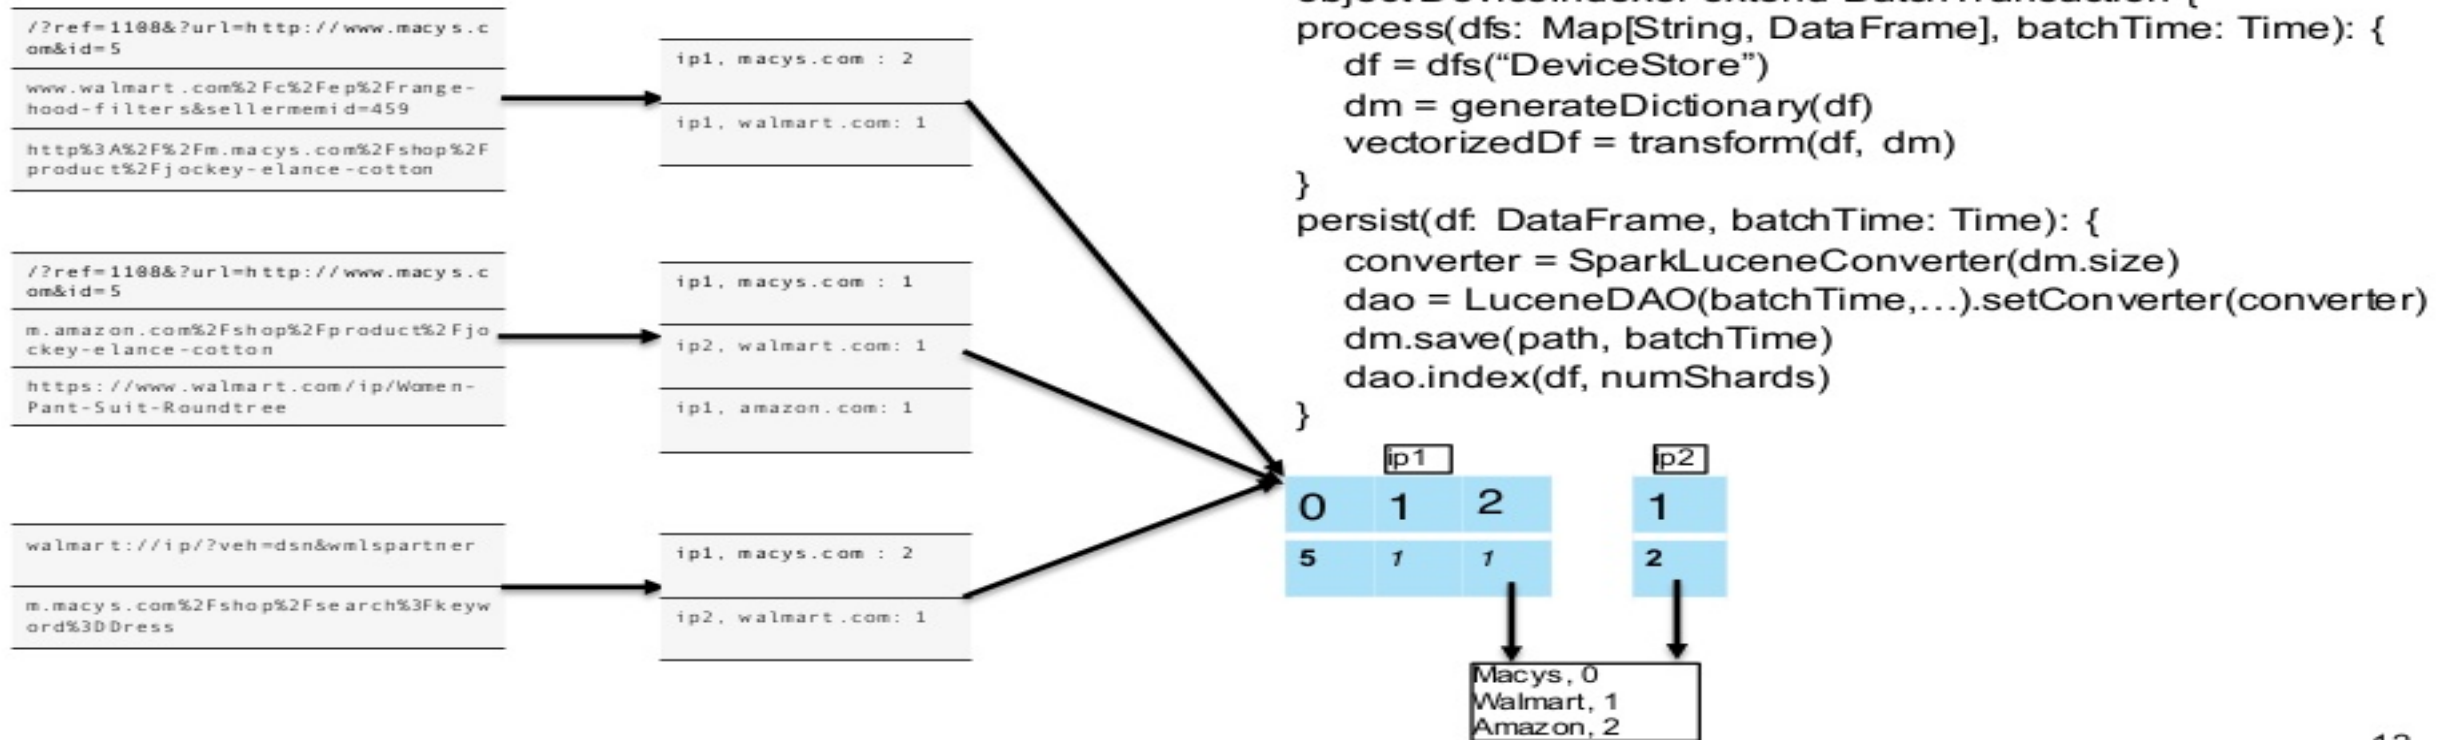
Trapezium Batch

```
runMode = "BATCH"
dataSource = "HDFS"
dependentWorkflows = {
  workflows=[aggregate]
  frequencyToCheck=100
}
hdfsFileBatch = {
  batchSize = 86400
  timerStartDelay = 1
  batchInfo = [{
    name = "DeviceStore"
    dataDirectory = {saiph-devqa=/aggregates}
    fileFormat = "parquet"
  }]
}
```

```
transactions = [{
  transactionName = "DeviceIndexer"
  inputData = [{name = "DeviceStore"}]
  persistDataName = "indexed"
}]
```

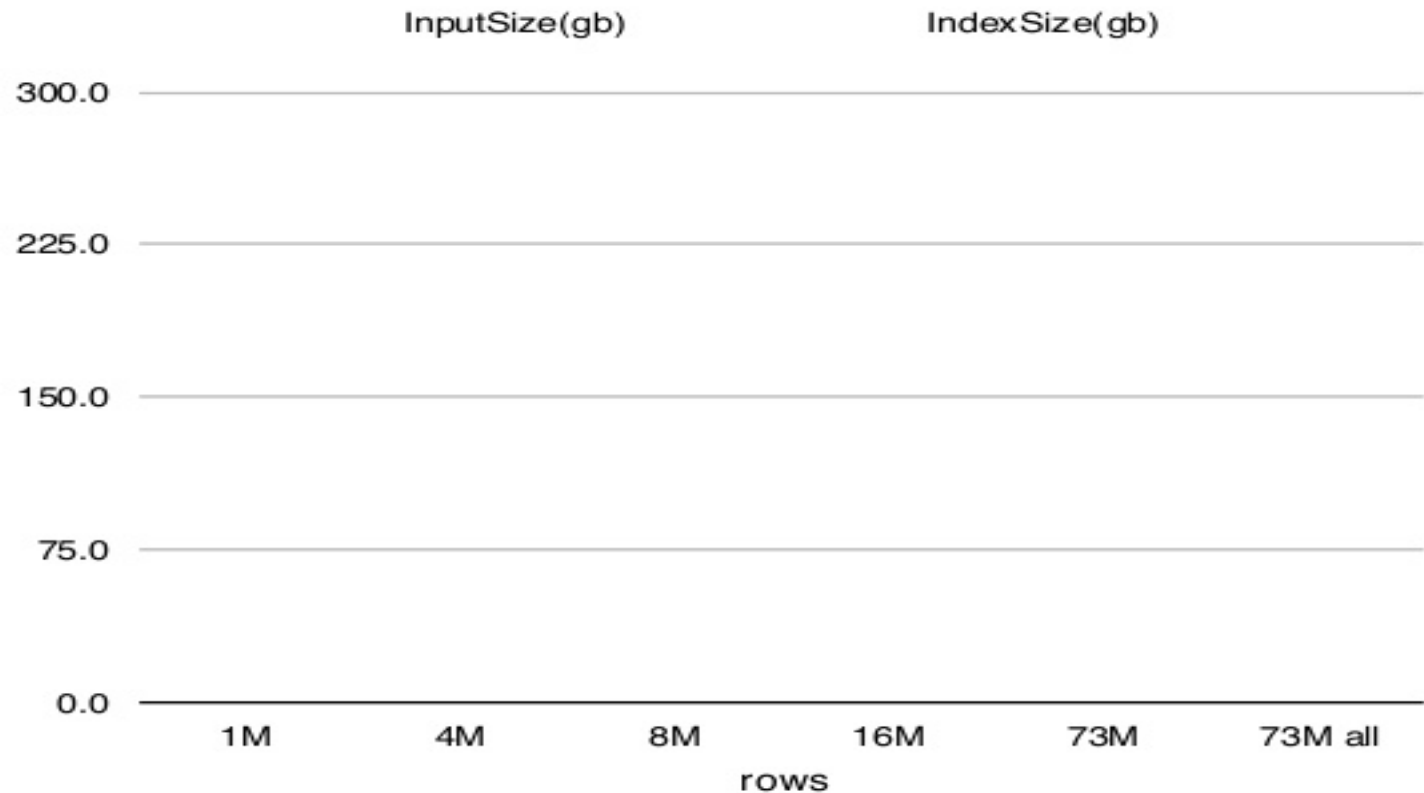


DeviceAnalyzer: Indexing



LuceneDAO Index Size

rows	InputSize(gb)	IndexSize(gb)
1M	4.0	5.1
4M	14.4	19.0
8M	27.9	35.7
16M	58.8	63.2
73M	276.5	228.0
73M all	276.5	267.1

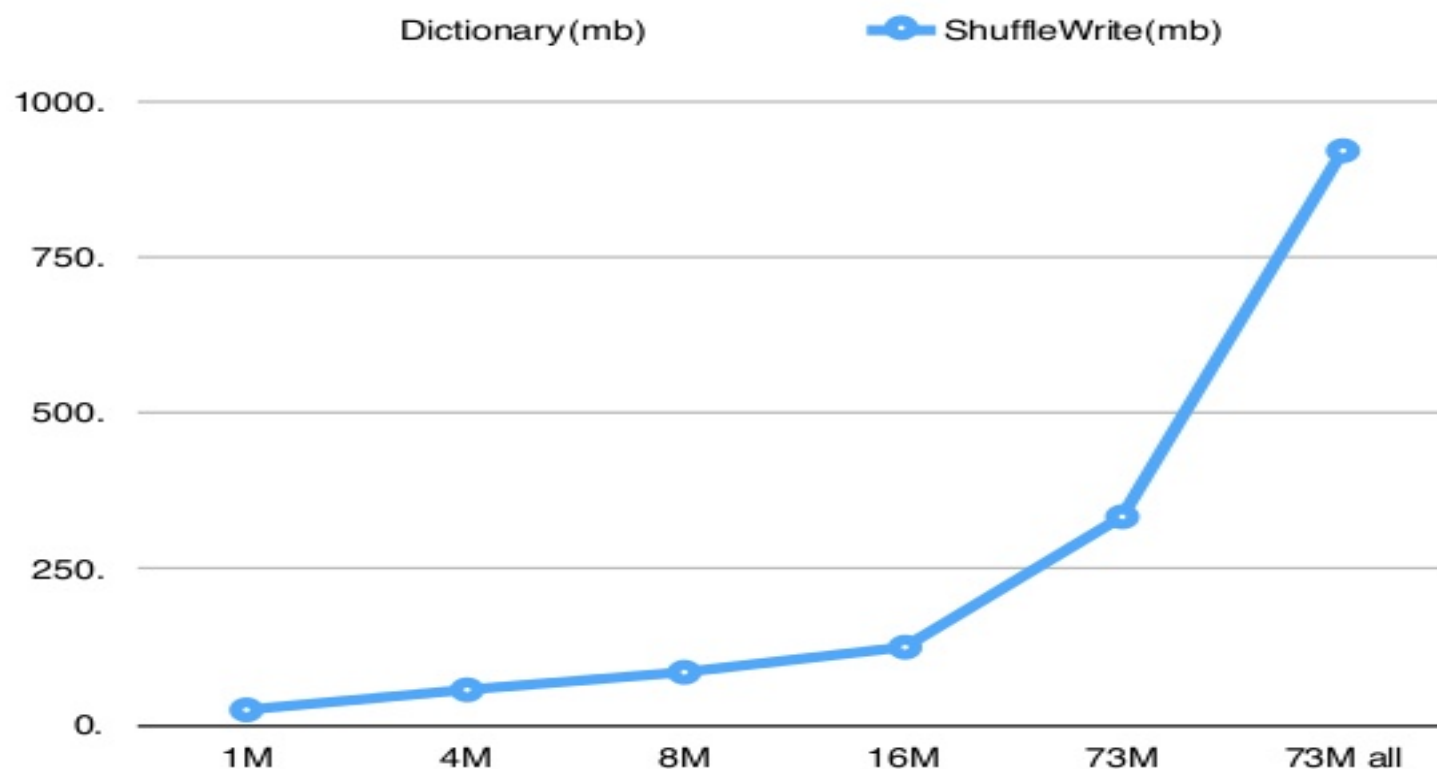


© Verizon 2016 All Rights Reserved

Information contained herein is provided AS IS and subject to change without notice. All trademarks used herein are property of their respective owners.

LuceneDAO Shuffle Size

rows	ShuffleWrite(mb)	Dictionary(mb)
1M	25	22.0
4M	56	30.0
8M	85	31.6
16M	126	32.2
73M	334	32.4
73M all	921	146.5



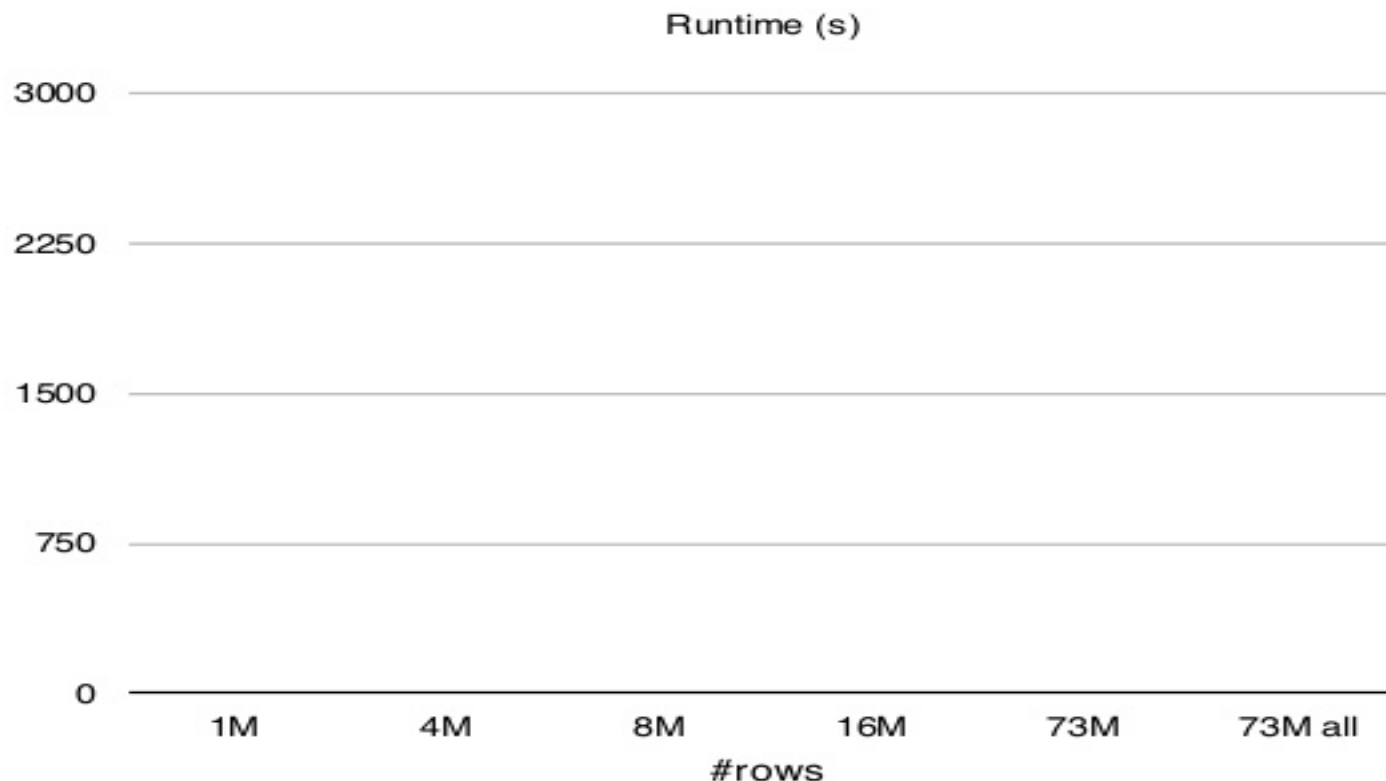
© Verizon 2016 All Rights Reserved

Information contained herein is provided AS IS and subject to change without notice. All trademarks used herein are property of their respective owners.

LuceneDAO Index Runtime

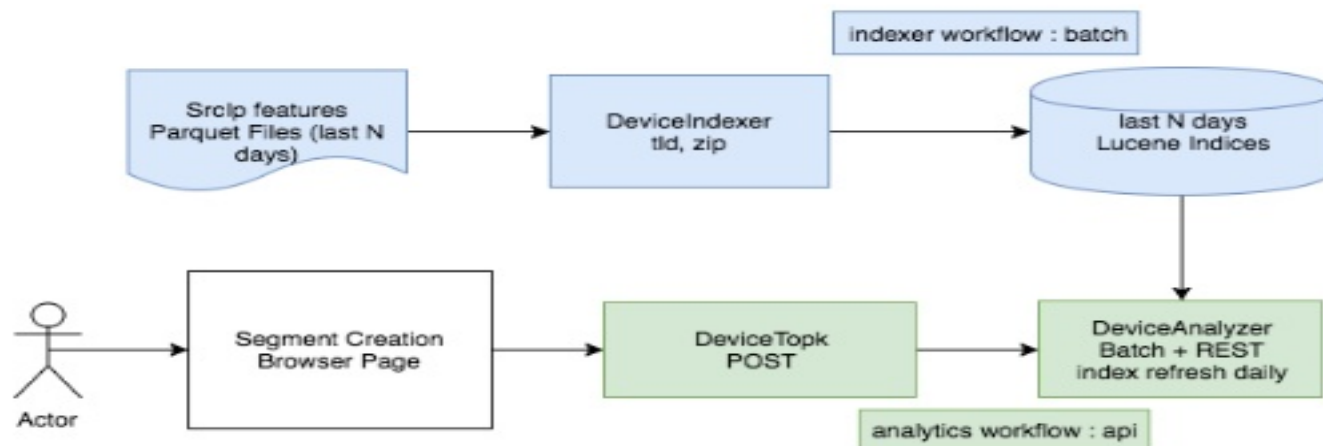
20 executors 16 cores
Executor RAM 16 GB
Driver RAM 8g

rows	Runtime (s)
1M	135
4M	228
8M	434
16M	571
73M	1726
73M all	2456



Trapezium Api

```
runMode = "BATCH"
dataSource = "HDFS"
httpServer = {
  provider = "akka"
  hostname = "localhost"
  port = 19999
  contextPath = "/"
  endPoints = [{
    path = "analyzer-api"
    className = "TopKEndPoint"
  }]
}
```



DeviceAnalyzer: Topk

- Given a query `select * from devices where tld='macys.com' OR 'nordstorm.com' AND (city='SanFrancisco' OR 'Brussels') AND (device='Android') ...`
 - ML: Find topk dimensions highly correlated with selected device
 - BI: group by tld order by `sum(visits)` as `tldVisits` limit topk

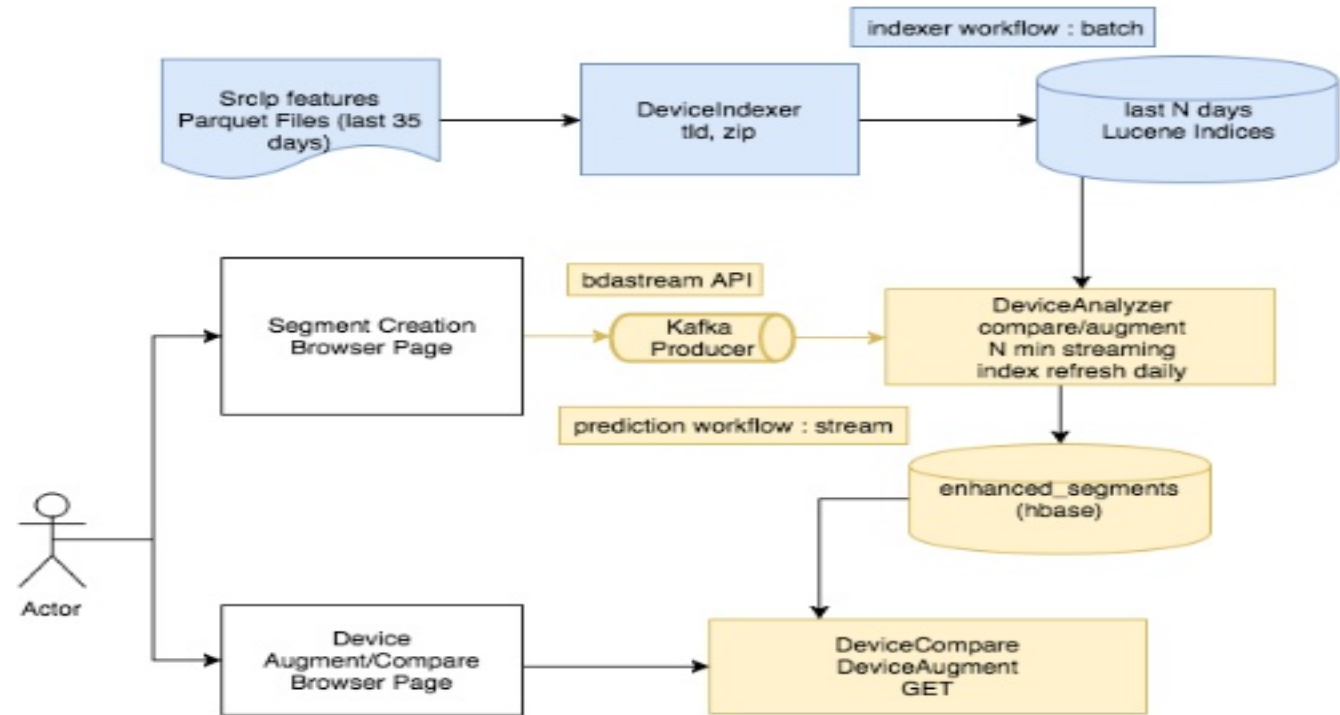
```
class TopkController(sc: SparkContext) extends SparkServiceEndPoint(sc) {  
  override def route : topkRoute  
  converter = SparkLuceneConverter(dm.size)  
  batchTime = Trapezium.getSyncTime("indexer")  
  dao = LuceneDAO(batchTime...)  
    .setConverter(converter).load(sc, indexPath)  
  dict = loadDictionary(sc, indexPath, batchTime)  
  def topkRoute : {  
    post { request => {  
      devices = dao.search(request)  
      response = getCorrelates(devices, dict, topk)  
    }  
  }  
}
```

df[deviceId, vector]

sum, support
mean, median, stddev

Trapezium Stream

```
runMode = "STREAM"
dataSource = "KAFKA"
kafkaTopicInfo = {
  consumerGroup = "KafkaStreamGroup"
  maxRatePerPartition = 970
  batchSize = "5"
  streamsInfo = [{
    name = "queries"
    topicName = "deviceanalyzer"
  }]
}
transactions = [{
  transactionName = "DeviceAnalyzer"
  inputStreams = [{name: "queries"}]
  persistStreamName = "deviceanalyzer"
  isPersist = "true"
}]
```



DeviceAnalyzer: Compare

- Given two queries

```
select * from Devices where  
tld='macys.com' OR 'nordstorm.com' AND  
(city='SanFrancisco') AND (device='Android')  
select * from Devices where  
tld='macys.com' OR 'nordstorm.com' AND  
(city='Brussels') AND (device='Android')
```

- Find the dimensions that discriminate the devices associated with two groups

```
def processStream(streams: Map[String,  
DStream[Row]], workflowTime: Time): {  
  streams("queries").collect().map{ requests =>  
    group1 = dao.search(requests(0))  
    group2 = dao.search(requests(1))  
    response = runLDA(aud1, aud2, dict)  
  }  
}
```

- Sparse weighted least squares using Breeze QuadraticMinimizer
- L1 Regularized logistic regression

```
def persistStream(responses: RDD[Row],  
batchTime: Time) {  
  HBaseDAO.write(responses)  
}
```



DeviceAnalyzer: Augment

- Given a query

```
select * from Devices where  
tld='macys.com' OR 'nordstorm.com'  
AND (city='SanFrancisco' OR 'Brussels')  
AND (device='Android')...
```

- Find devices similar to seed as lookalikes
- Find dimensions that represent lookalikes

```
object DeviceAnalyzer extends StreamingTransaction {  
  converter = SparkLuceneConverter(dm.size)  
  batchTime = Trapezium.getSyncTime("indexer")  
  dao = LuceneDAO(batchTime...)  
    .setConverter(converter).load(sc, indexPath)  
  dict = loadDictionary(sc, indexPath, batchTime)  
  all = dao.search("**:*")  
  def processStream(streams: Map[String, DStream[Row]]) :  
  {  
    streams("queries").collect().map{ request =>  
      audience = dao.search(request)  
      response = getLookalikeDimensions(all, audience, dict)  
    }  
  }
```

- Sparse weighted least squares using Breeze QuadraticMinimizer
- L2 regularized linear regression



FastSummarizer

- Statistical and predictive operators
 - sum: sum over numeric measures
 - support: sum over distinct docID
 - sumSquared: L2 norm
 - gram: Uses BLAS sspr
 - solve: Uses BreezeQuadraticMinimizer to support L1
- Implemented using Array[Float] for shuffle opt
- Scala/Java for Level1 operations
- OpenBLAS for Level3 operations



Sync API Benchmark

73M rows 1M+ search terms

1 measure on 250K sparse dimensions


20 executors 8 cores

32 GB driver RAM 16 GB executor RAM

akka-http cores: 24 default

topk

qps	runtime(s)
1	1.389
5	1.663
10	3.214
20	5.992
40	12.174



Async API Benchmark

73M rows, 1M+ search terms

1 measure on 250K sparse dimensions

20 executors 8 cores

32 GB driver RAM 16 GB executor RAM


forkjoinpool = 40

Kafka Fetch + compare/augment + HBase Persist

compare

augment

predictions		
qps	compare(s)	augment(s)
1	9	16
5	13	36
10	23	70
20	42	142



topk tld + apps

Behaviors

Search for Sites, Apps and Keywords

Audience Definition -

Filter Audience +

Behaviors

Demographics

Devices

Locations

Behaviors

Rank	Site	Audience %	Rank	App	Audience %
31	realtor.com	5.58%	11	com.apple.mobilesafari	2.58%
32	usatoday.com	5.33%	12	Find My Friends	1.39%
33	stumbleupon.com	5.16%	13	Instagram	1.28%
34	disney.com	5.11%	14	com.zillow.zillowmap	1.21%
35	zillow.com	5.04%	15	Starbucks	0.84%
36	dailymail.co.uk	5.01%	16	Find My iPhone	0.72%
37	starwars.com	4.85%	17	com.acmeasom.android.radarym	0.72%
38	latimes.com	4.83%	18	com.facebook.katana	0.7%
39	truste.com	4.75%	19	com.apple.geod	0.66%
40	nypost.com	4.7%	20	com.hearst.kcraiphone	0.63%

Previous

1

2

3

4

5

Next

Previous

1

2

3

4


5

Next


© Verizon 2016 All Rights Reserved

Information contained herein is provided AS IS and subject to change without notice. All trademarks used herein are property of their respective owners.

Augment: Auto Enthusiastic

 **Look Alike Analysis**

Palomar has predicted the most meaningful and accurate new features to add so you can successfully increase your audience reach with customers that resemble your seed audience.

 Create Snapshot

AutoEnthusiastsAudE1

AutoEnthusiastsAud

Top Prediction Features

The number of 'Reach' subscribers is the remainder, or difference, between the Seed Audience counts and the Look-alike Audience counts.

Min

Max

Apply

All Features

Rank	Feature	Weight
1.	Site: autotrader.com	
2.	Site: showroomlogic.com	
3.	Site: carsforsale.com	
4.	Site: honda.com	
5.	Site: bmwusa.com	
6.	Site: automobilemag.com	
7.	Site: cars.com	
8.	Site: dealerfire.com	
9.	Site: carfax.com	
10.	Site: nissanusa.com	

Look-Alike Audience

Create This Audience Save This Audience Show Seed Audience Stats

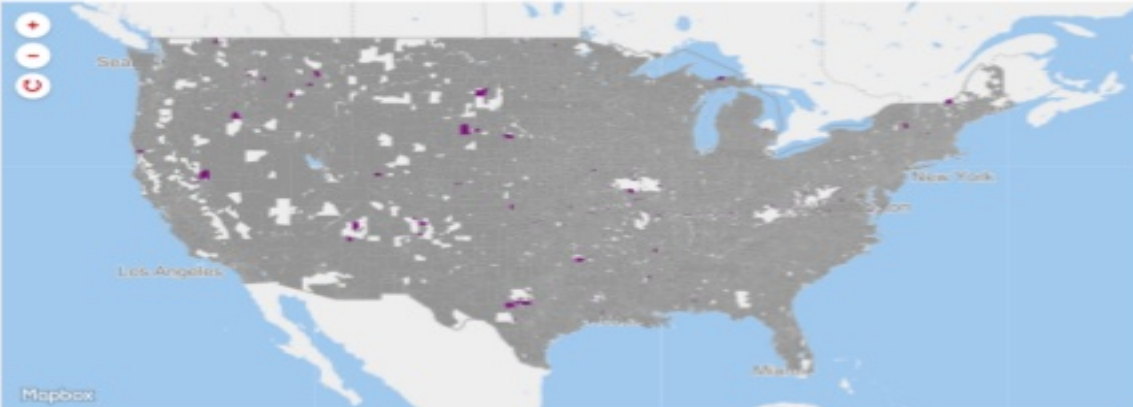
Look-Alike Audience: 6.03M
Seed Audience: 4.22M

Seed Audience: AutoEnthusia ...
Seed Audience Definition: Site: autoweek.com, caranddriver.com, edmunds.com, topspeed ...

Top Prediction Features

Behaviors 1. autotrader.com 2. showroomlogic.com 3. carsforsale.com 4. honda.com 5. bmwusa.com 6. aut...


Locations



© Verizon 2016 All Rights Reserved

Information contained herein is provided AS IS and subject to change without notice. All trademarks used herein are property of their respective owners.

Augment Model Performance



Compare: Leisure vs Business Travellers

LeisureVSBusinessTravelersAudE1

LeisureTravelersAud Vs BusinessTravelersAud

LeisureTravelersAud

Audience Definition:

Site: orbitz.com, expedia.com, kayak.com, cheapcaribbean.com, bookit.com, fun...

BusinessTravellerAud

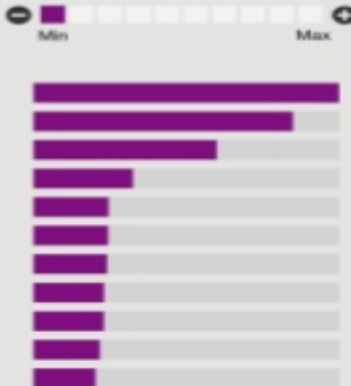
Audience Definition:

Site: expedia.com, travelocity.com, kayak.com, united.com, aa.com, virginamer...

Behavioral Differentials

Sites & Apps

App : Priceline - Book Che...
App : com.hotwire.hotwire
Site : estara.com
Site : applevacations.com
Site : capecodonline.com
Site : aircanada.com
Site : microsofttranslator.c...
Site : cheapflights.co.uk
Site : winetasting.com
Site : vegas.com
Site : displays2go.com



Sites & Apps

Site : marvel.com
Site : efinancialnews.com
Site : containerstore.com
Site : geico.com
App : Marriott International
Site : nbc12.com
Site : simon.com
Site : chick-fil-a.com
Site : justanswer.com
Site : military.com
Site : qutitrics.com



Location Differentials



© Verizon 2016 All Rights Reserved

Information contained herein is provided AS IS and subject to change without notice. All trademarks used herein are property of their respective owners.

THANK YOU. Q&A

Join us and make machines intelligent
Data & Artificial Intelligence Systems
499 Hamilton Ave, Palo Alto
California

verizon✓

