

# Deep Dive Into Catalyst: Apache Spark's Optimizer

Herman van Hövell

Spark Summit Europe, Brussels  
October 26<sup>th</sup> 2016



# Who is Databricks

## Why Us

- Created Apache Spark to enable big data use cases with a single engine.
- Contributes 75% of Spark's code




## Our Product

- Bring Spark to the enterprise: The just-in-time data platform.
- Fully managed platform powered by Apache Spark.
- A unified solution for data science and engineering teams.



# Overview



# Why structure?

- By definition, structure will *limit* what can be expressed.
- In practice, we can accommodate the vast majority of computations.

**Limiting the space of what can be expressed enables optimizations.**

# Why structure?

## **RDD**

```
pdata.map { case (dpt, age) => dpt -> (age, 1) }  
      .reduceByKey { case ((a1, c1), (a2, c2)) => (a1 + a2, c1 + c2) }  
      .map { case (dpt, (age, c)) => dpt -> age / c }
```


## **Dataframe**

```
data.groupBy("dept").avg("age")
```

## **SQL**

```
select dept, avg(age) from data group by 1
```

# Why structure?



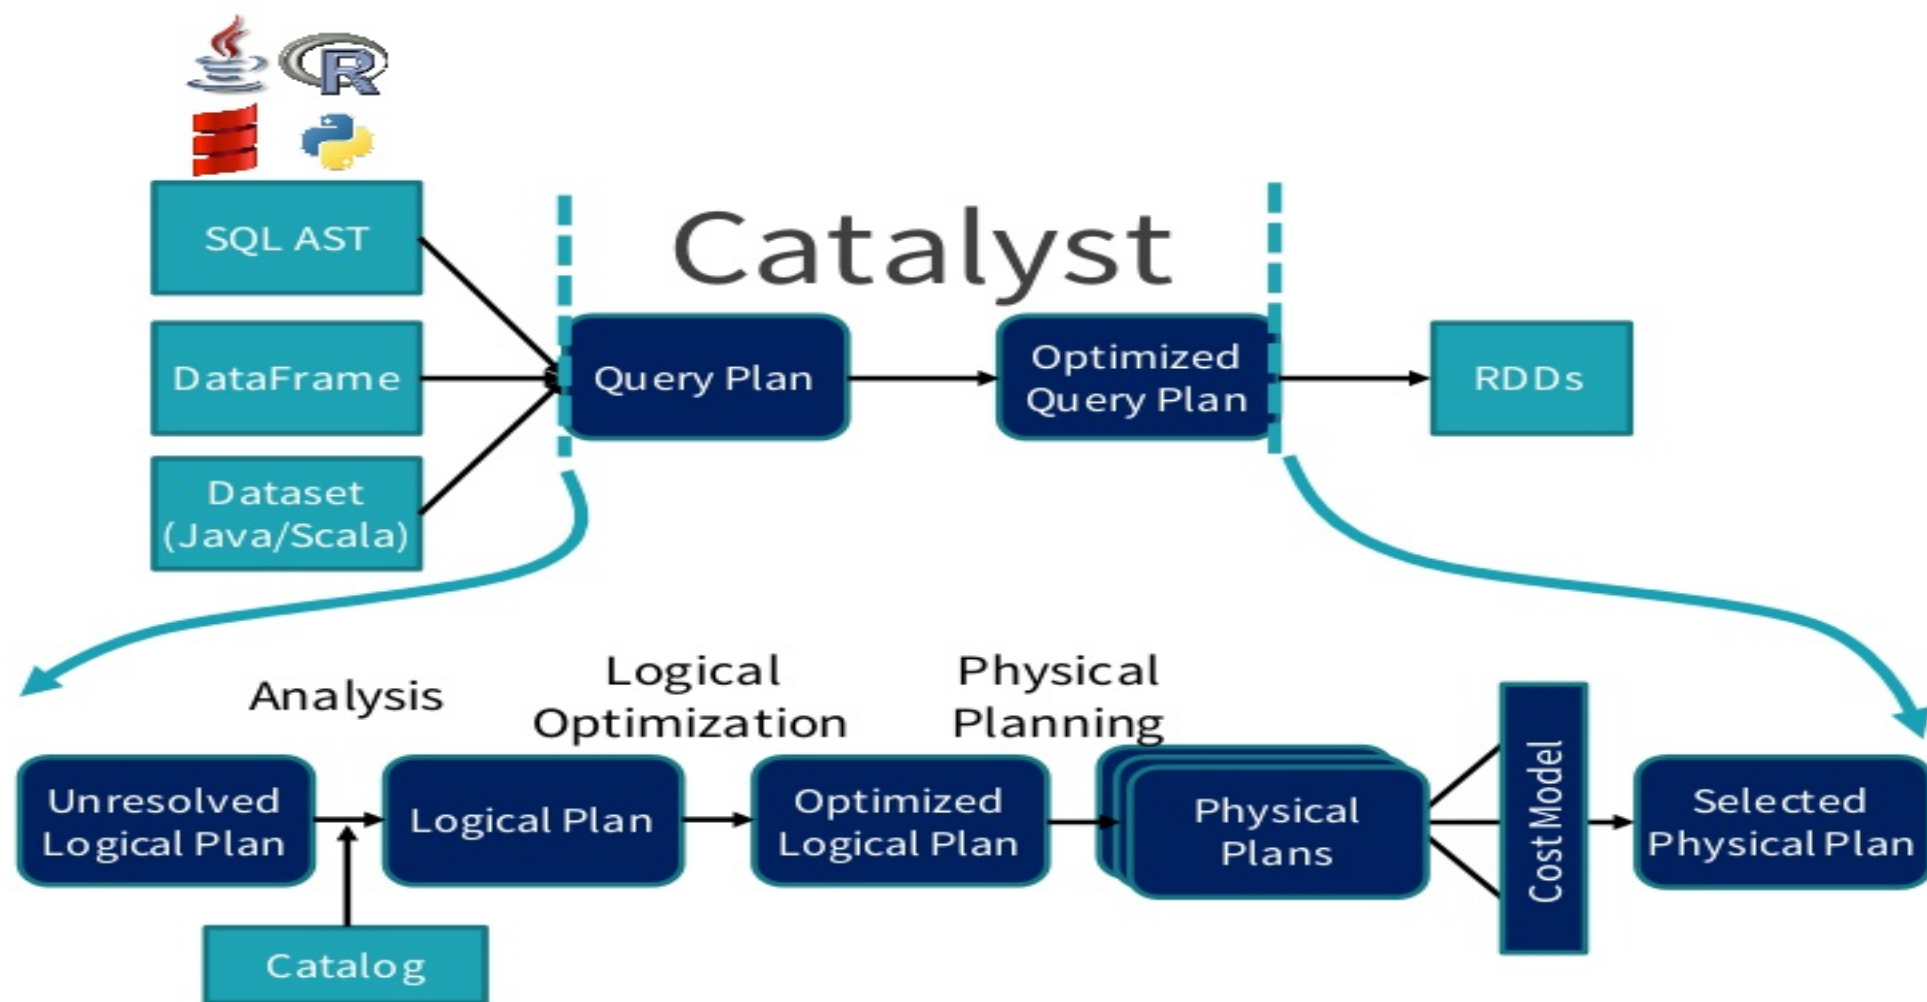


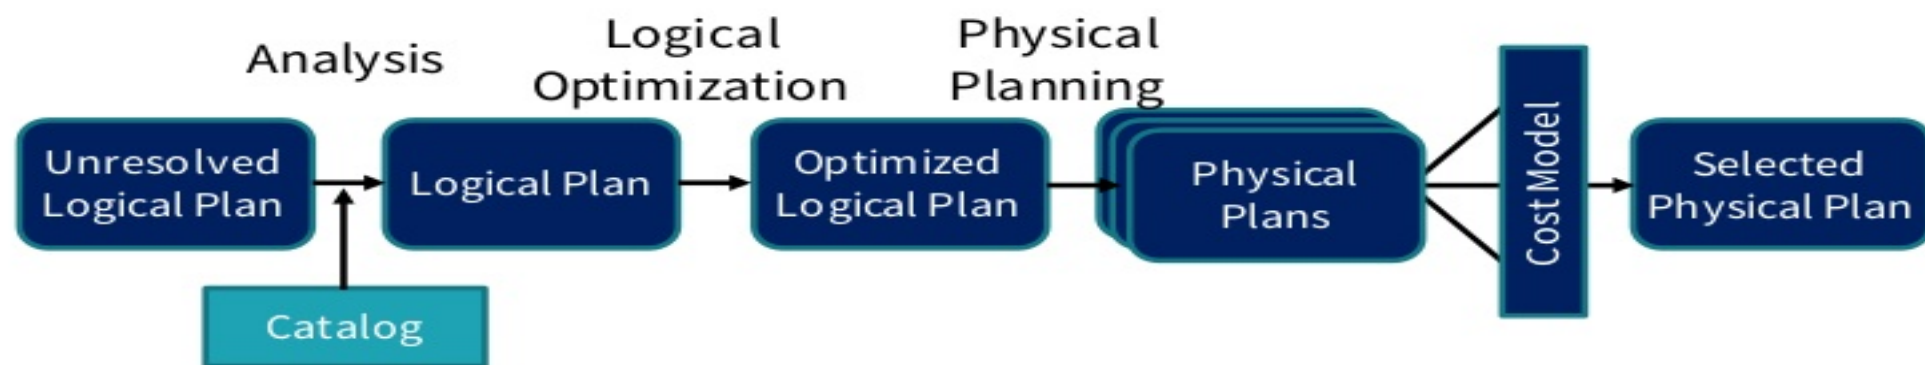
# How?

- Write programs using high level programming interfaces
  - Programs are used to describe what data operations are needed without specifying how to execute those operations
  - High level programming interfaces: SQL, DataFrames, and Dataset
- Get an optimizer that **automatically** finds out the most efficient plan to execute data operations specified in the user's program

# Catalyst: Apache Spark's Optimizer

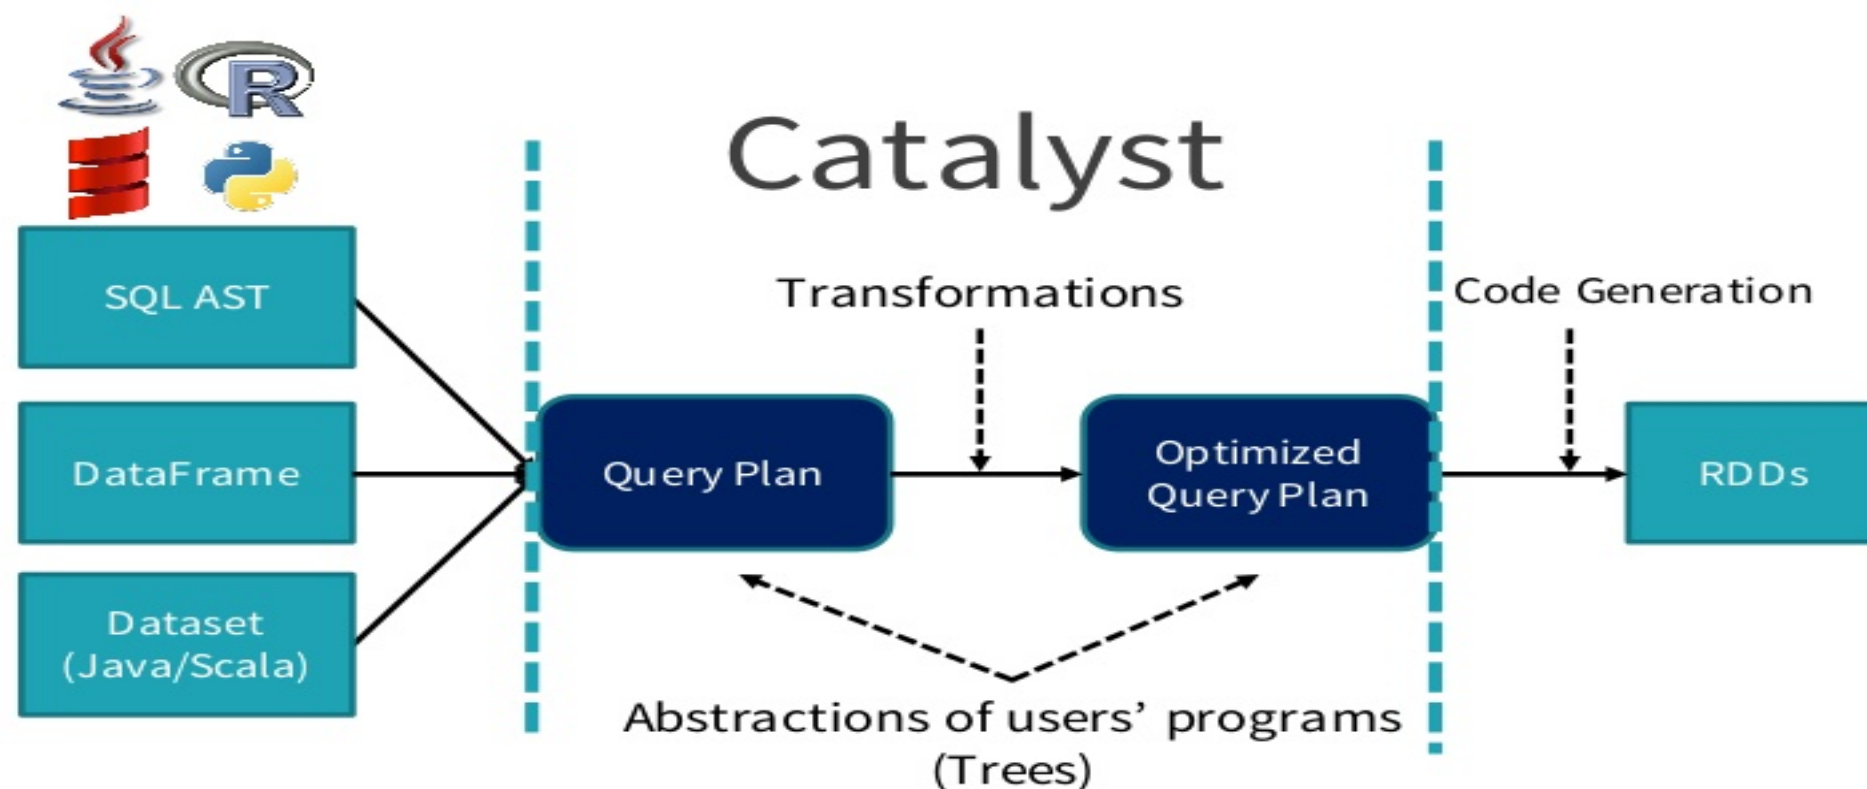




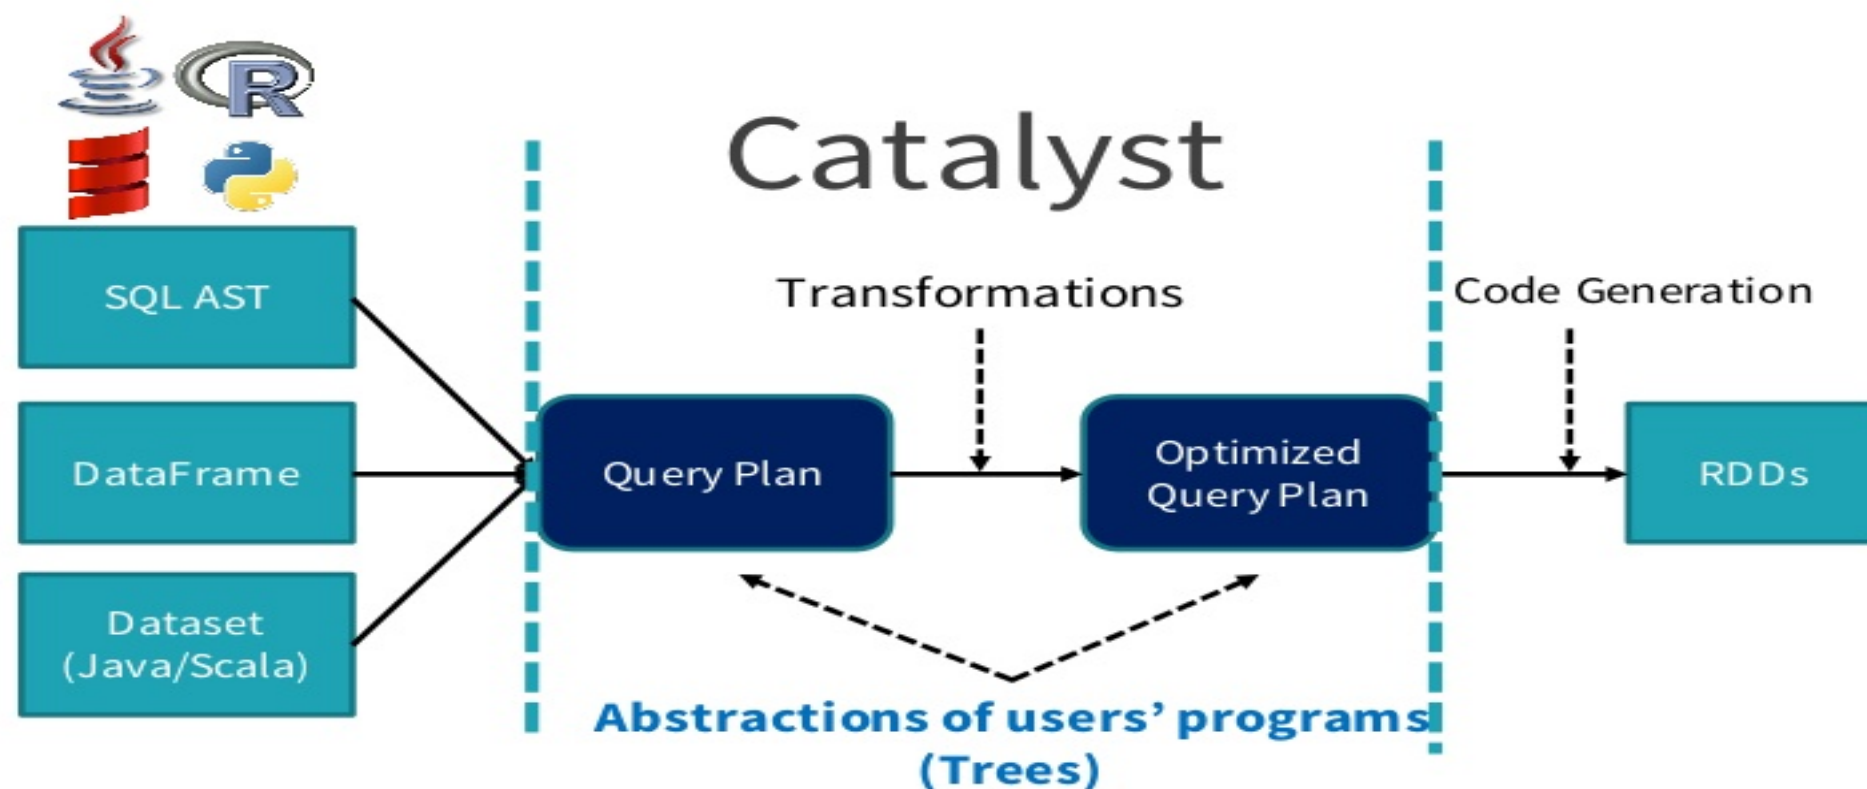


- **Analysis (Rule Executor):** Transforms an Unresolved Logical Plan to a Resolved Logical Plan
  - Unresolved => Resolved: Use Catalog to find where datasets and columns are coming from and types of columns
- **Logical Optimization (Rule Executor):** Transforms a Resolved Logical Plan to an Optimized Logical Plan
- **Physical Planning (Strategies + Rule Executor):** Transforms a Optimized Logical Plan to a Physical Plan

# How Catalyst Works: An Overview



# How Catalyst Works: An Overview



# Trees: Abstractions of Users' Programs

```
SELECT sum(v)
FROM (
  SELECT
    t1.id,
    1 + 2 + t1.value AS v
  FROM t1 JOIN t2
  WHERE
    t1.id = t2.id AND
    t2.id > 50 * 1000) tmp
```

# Trees: Abstractions of Users' Programs

## Expression

```
SELECT sum(v)
FROM (
  SELECT
    t1.id,
    1 + 2 + t1.value AS v
  FROM t1 JOIN t2
  WHERE
    t1.id = t2.id AND
    t2.id > 50 * 1000) tmp
```


- An expression represents a new value, computed based on input values
  - e.g. `1 + 2 + t1.value`
- Attribute: A column of a dataset (e.g. `t1.id`) or a column generated by a specific data operation (e.g. `v`)



# Trees: Abstractions of Users' Programs

## Query Plan


```
SELECT sum(v)
FROM (
  SELECT
    t1.id,
    1 + 2 + t1.value AS v
  FROM t1 JOIN t2
  WHERE
    t1.id = t2.id AND
    t2.id > 50 * 1000) tmp
```






# Logical Plan

- A Logical Plan describes computation on datasets **without** defining how to conduct the computation
- **output**: a list of attributes generated by this Logical Plan, e.g. [**id**, **v**]
- **constraints**: a set of invariants about the rows generated by this plan, e.g. **t2.id > 50 \* 1000**

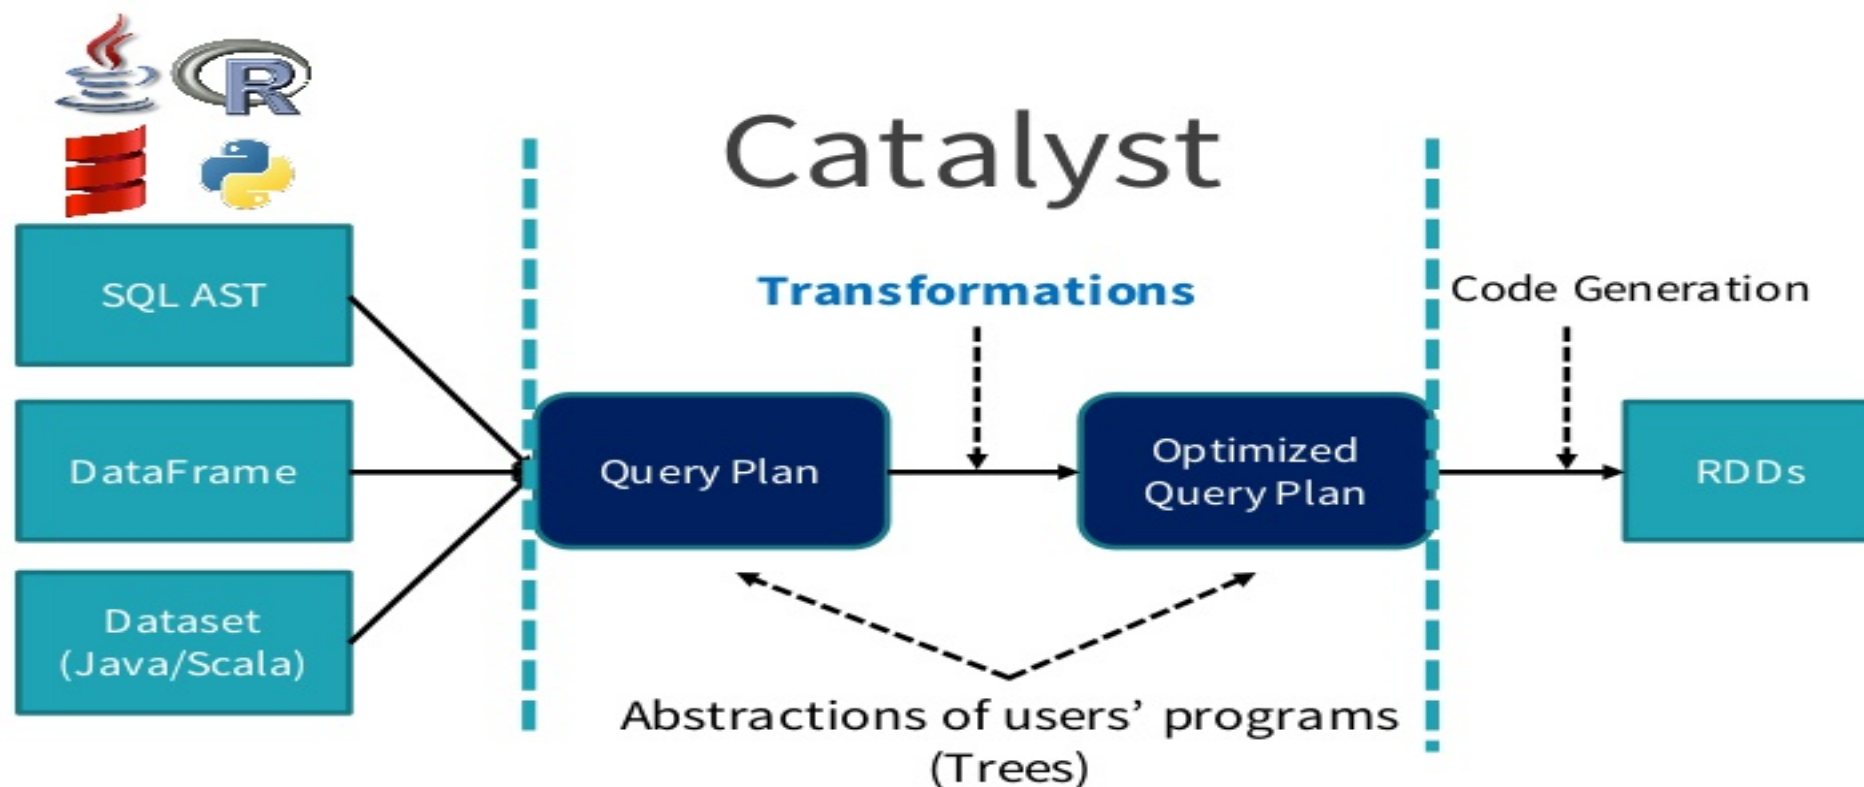


# Physical Plan

- A Physical Plan describes computation on datasets with specific definitions on how to conduct the computation
- A Physical Plan is executable



# How Catalyst Works: An Overview



# Transformations


- Transformations without changing the tree type (Transform and Rule Executor)
  - Expression  $\Rightarrow$  Expression
  - Logical Plan  $\Rightarrow$  Logical Plan
  - Physical Plan  $\Rightarrow$  Physical Plan
- Transforming a tree to another kind of tree
  - Logical Plan  $\Rightarrow$  Physical Plan

# Transform

- A function associated with every tree used to implement a single rule

$1 + 2 + t1.value$


Evaluate  $1 + 2$   
for every row



Evaluate  $1 + 2$  once




$3 + t1.value$



# Transform

- A transformation is defined as a Partial Function
- Partial Function: A function that is defined for a subset of its possible arguments

```
val expression: Expression = ...  
expression.transform {  
  case Add(Literal(x, IntegerType), Literal(y, IntegerType)) =>  
    Literal(x + y)  
}
```

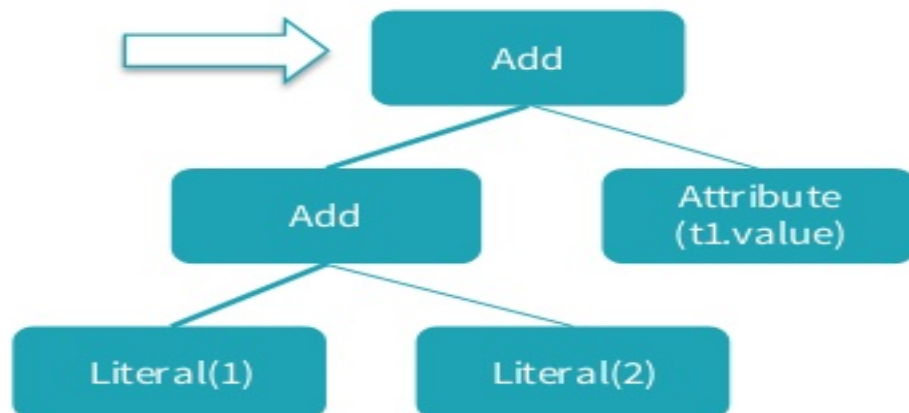


Case statement determine if the partial function is defined for a given input

# Transform

```
val expression: Expression = ...  
expression.transform {  
  case Add(Literal(x, IntegerType), Literal(y, IntegerType)) =>  
    Literal(x + y)  
}
```

1 + 2 + t1.value






# Transform

```
val expression: Expression = ...  
expression.transform {  
  case Add(Literal(x, IntegerType), Literal(y, IntegerType)) =>  
    Literal(x + y)  
}
```

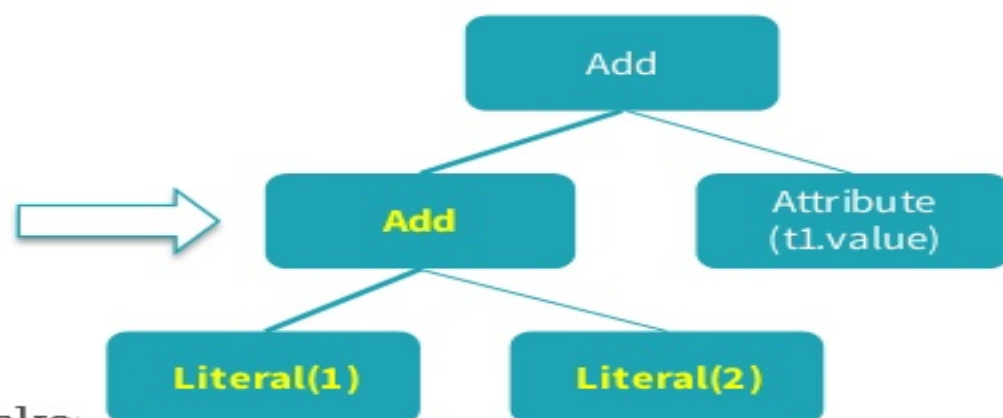
1 + 2 + t1.value



# Transform

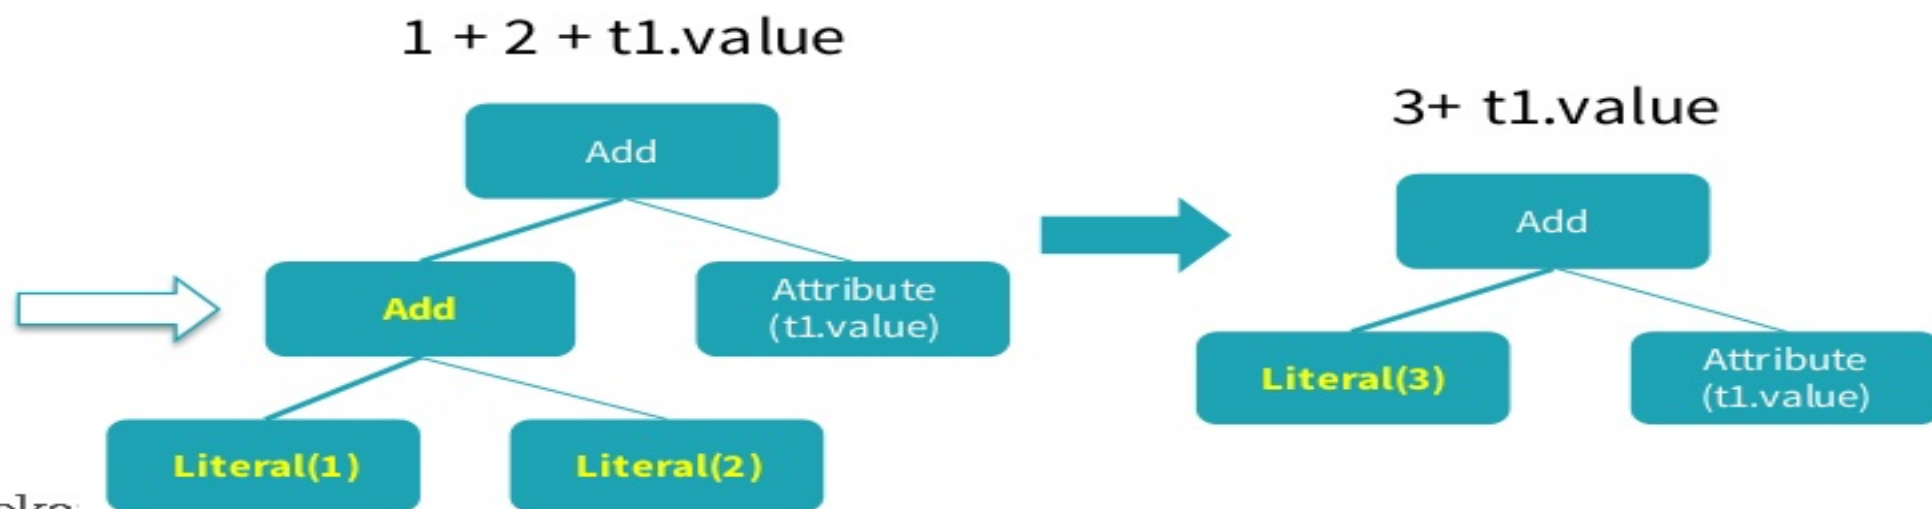
```
val expression: Expression = ...  
expression.transform {  
  case Add(Literal(x, IntegerType), Literal(y, IntegerType)) =>  
    Literal(x + y)  
}
```

1 + 2 + t1.value



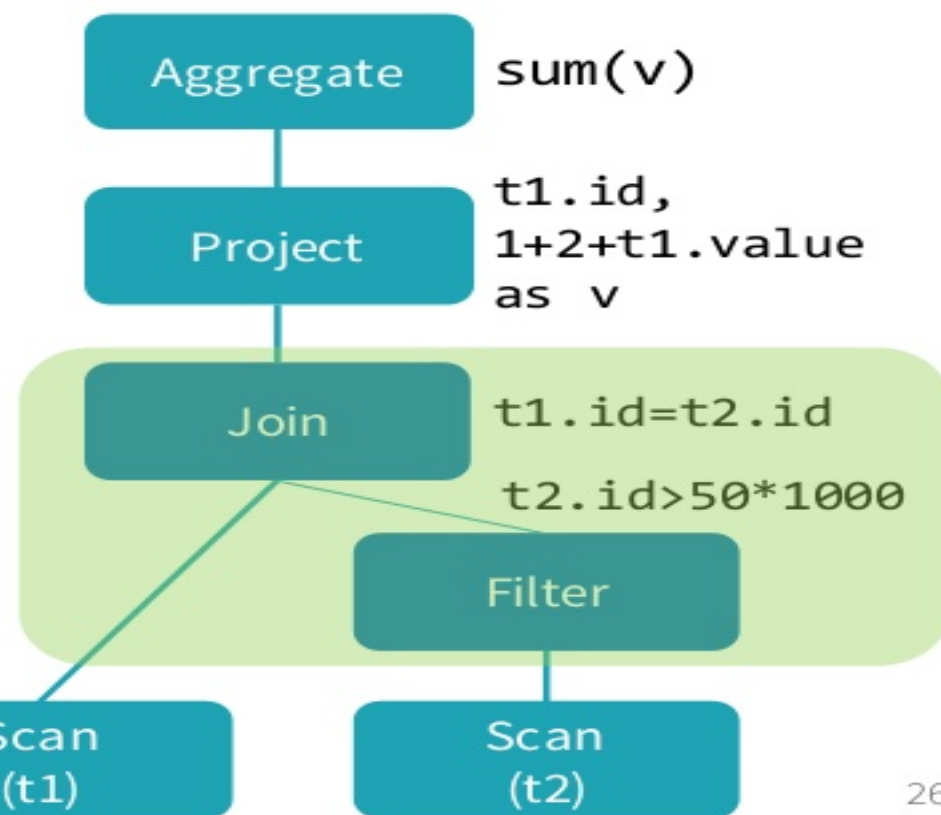


# Transform

```
val expression: Expression = ...  
expression.transform {  
  case Add(Literal(x, IntegerType), Literal(y, IntegerType)) =>  
    Literal(x + y)  
}
```



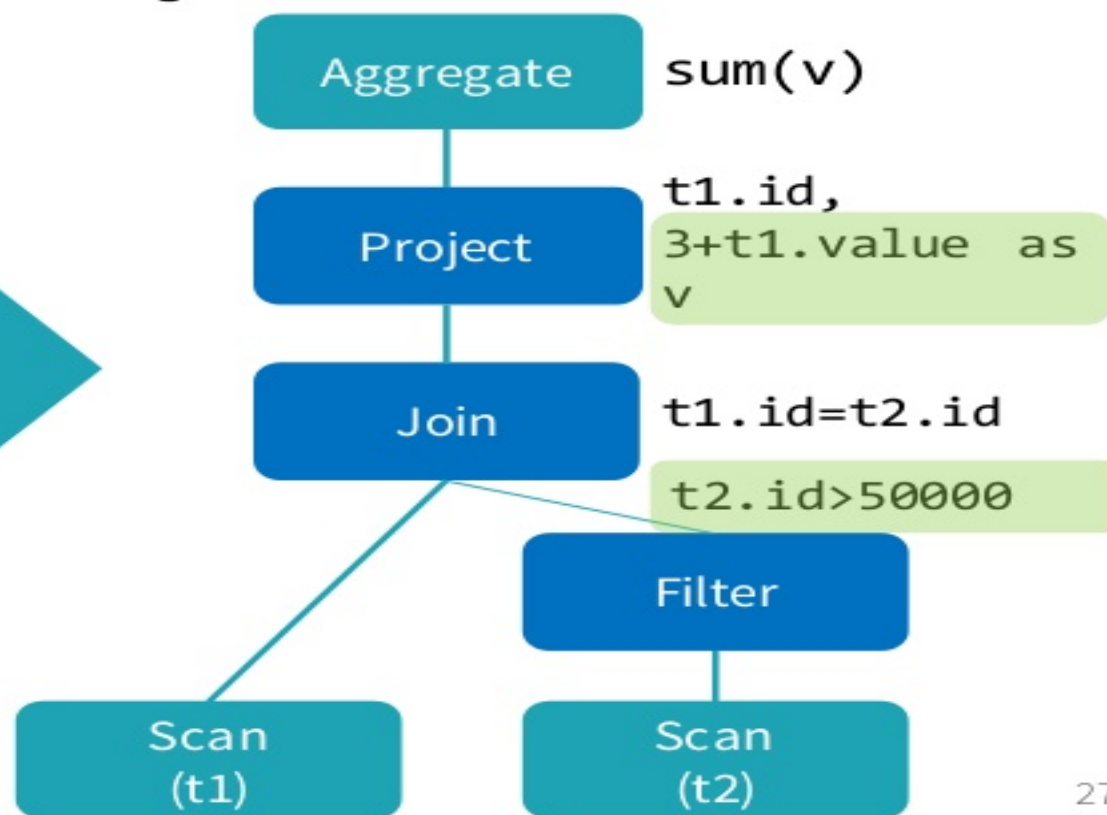
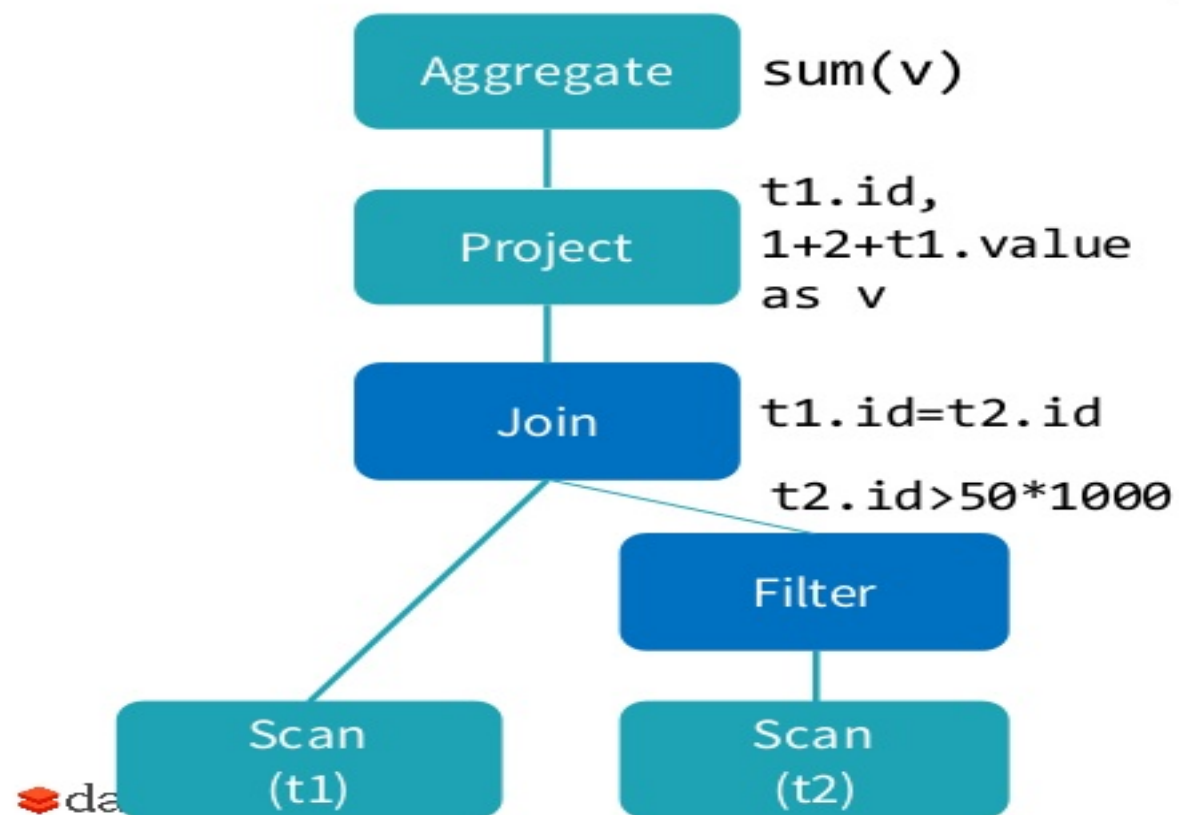
# Combining Multiple Rules

Predicate Pushdown



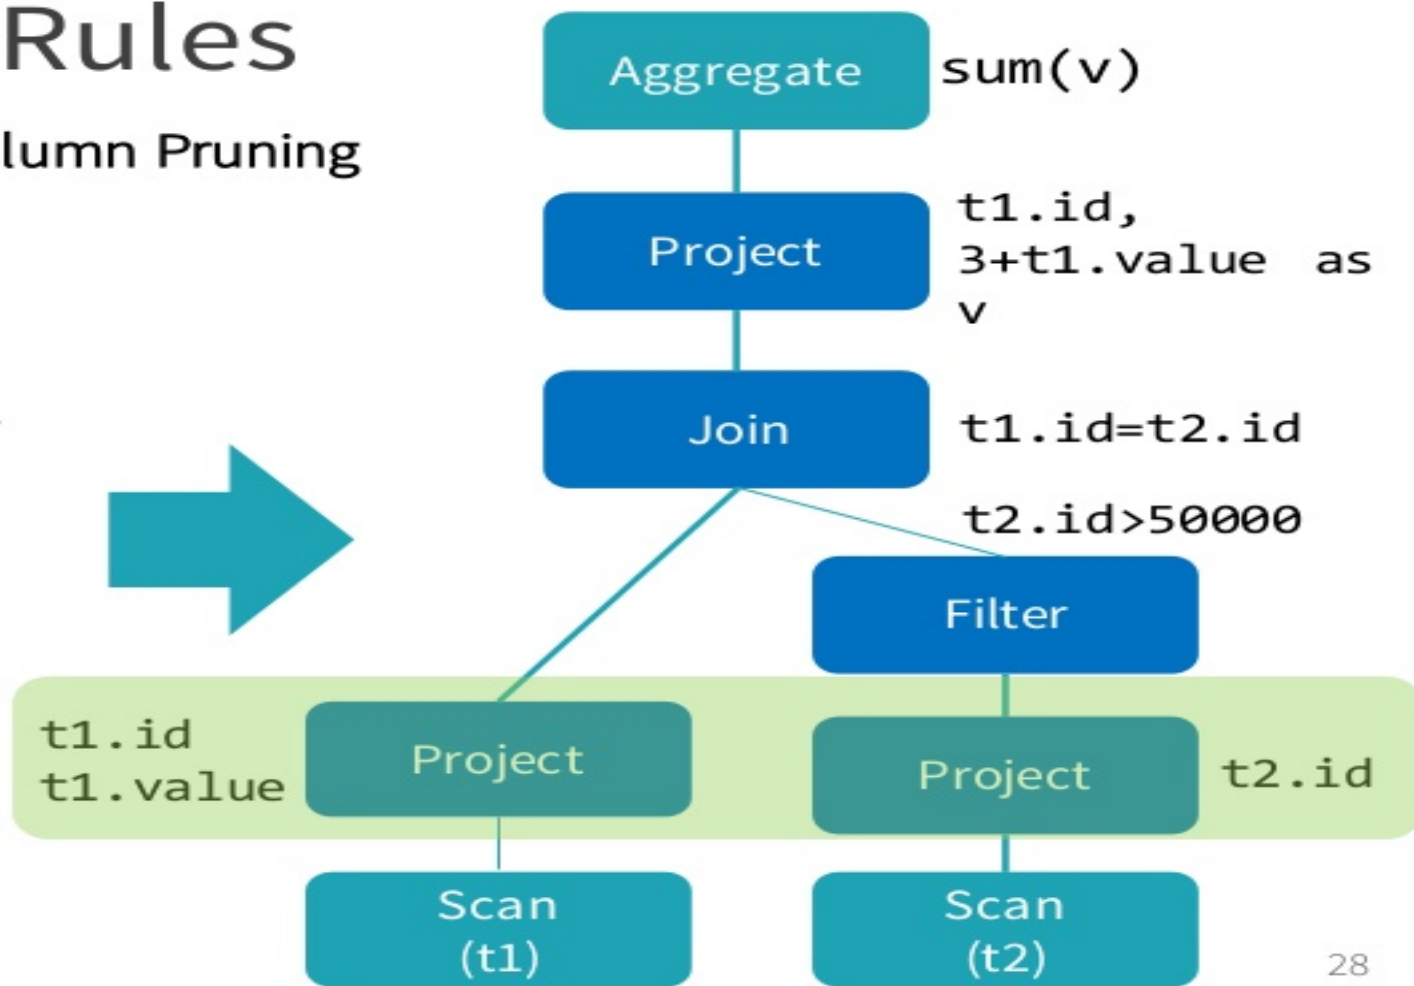
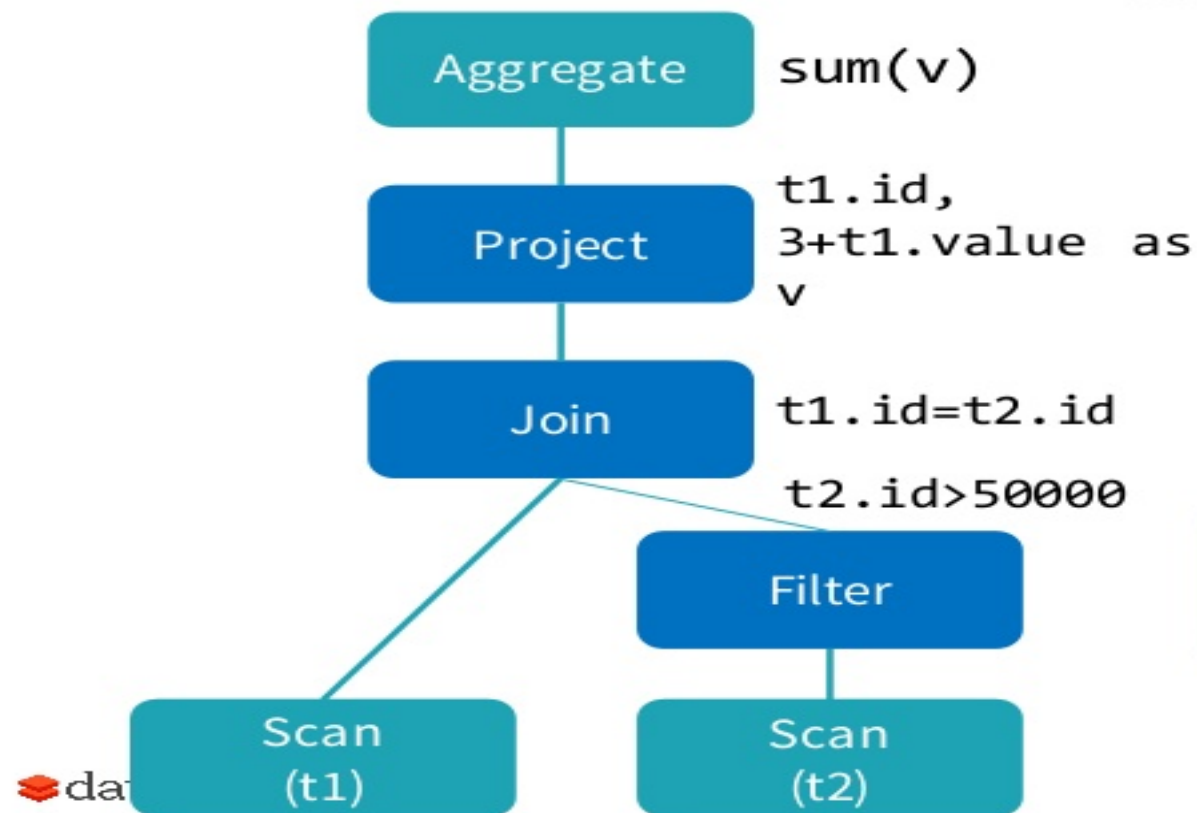
# Combining Multiple Rules

Constant Folding





# Combining Multiple Rules

Column Pruning

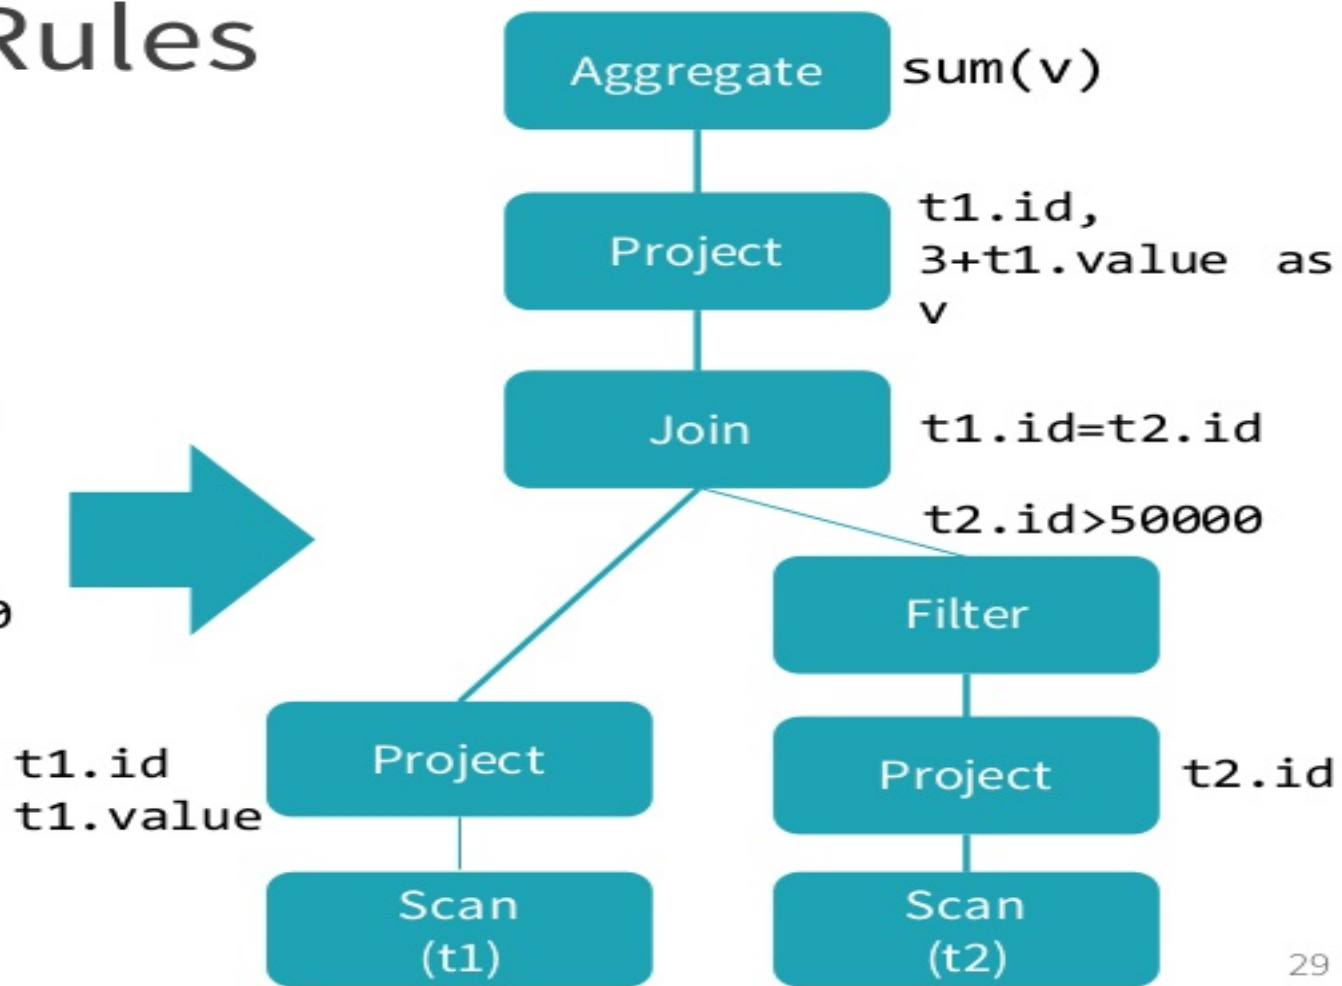


# Combining Multiple Rules

Before transformations



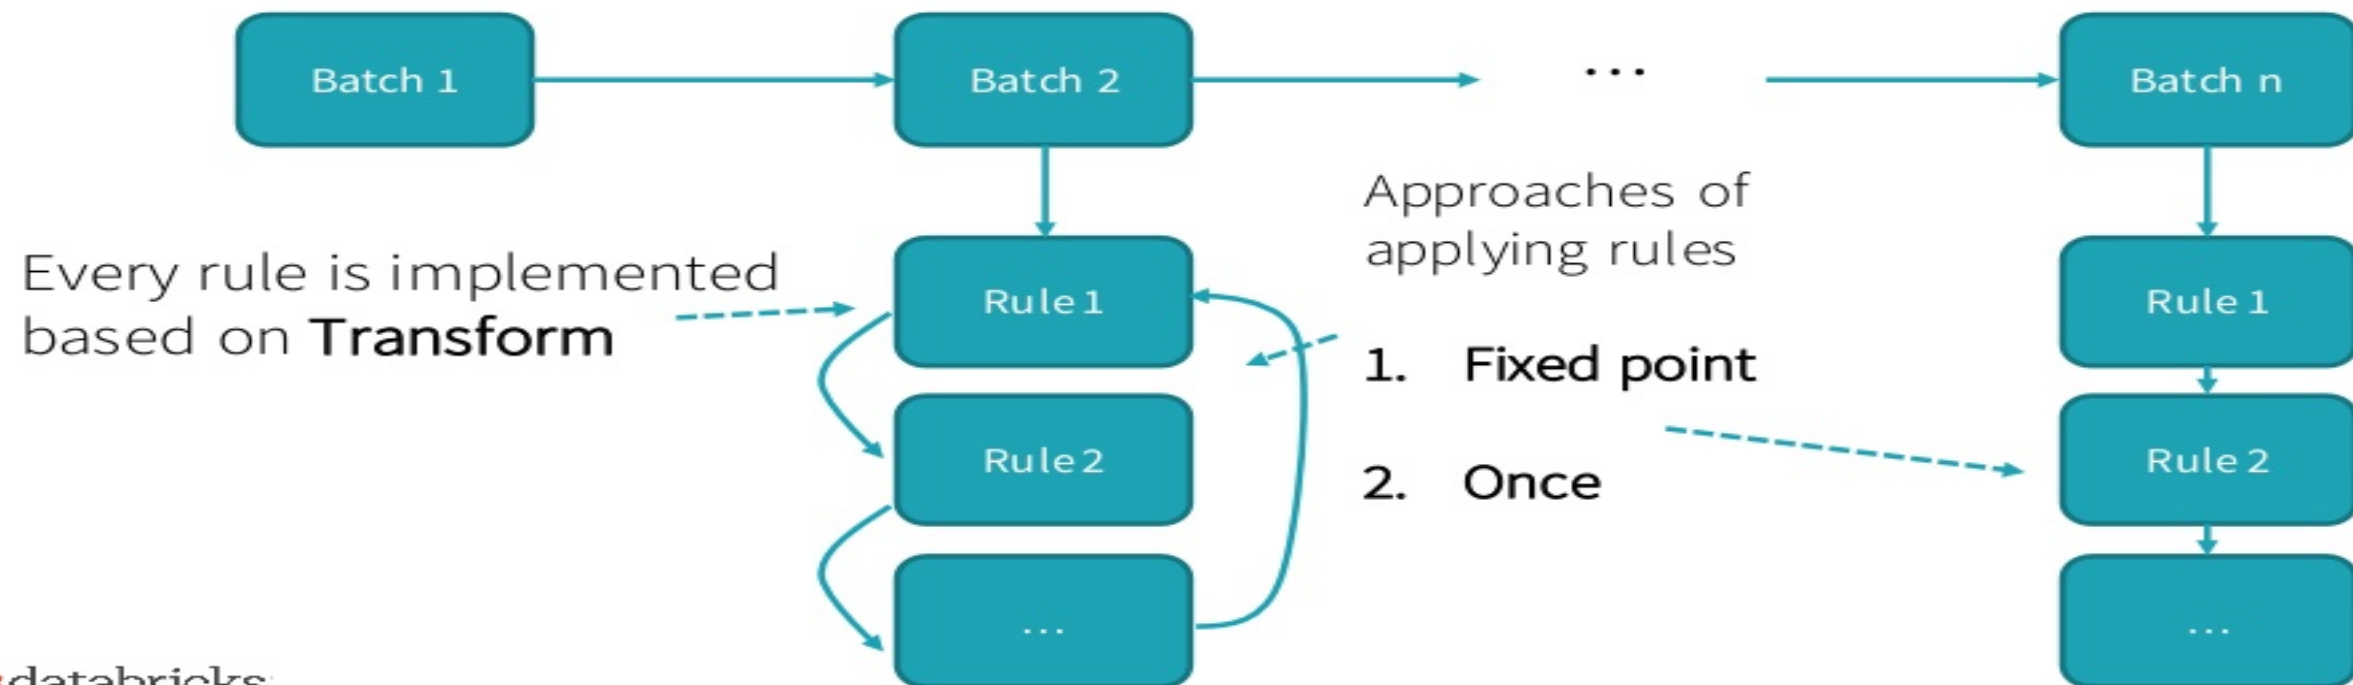
After transformations





# Combining Multiple Rules: Rule Executor

A Rule Executor transforms a Tree to another same type Tree by applying many rules defined in batches



# Transformations

- Transformations without changing the tree type (Transform and Rule Executor)
  - Expression  $\Rightarrow$  Expression
  - Logical Plan  $\Rightarrow$  Logical Plan
  - Physical Plan  $\Rightarrow$  Physical Plan
- Transforming a tree to another kind of tree
  - Logical Plan  $\Rightarrow$  Physical Plan

# From Logical Plan to Physical Plan

- A Logical Plan is transformed to a Physical Plan by applying a set of **Strategies**
- Every Strategy uses pattern matching to convert a Tree to another kind of Tree

```
object BasicOperators extends Strategy {  
  def apply(plan: LogicalPlan): Seq[SparkPlan] = plan match {  
    ...  
    case logical.Project(projectList, child) =>  
      execution.ProjectExec(projectList, planLater(child)) :: Nil  
    case logical.Filter(condition, child) =>  
      execution.FilterExec(condition, planLater(child)) :: Nil  
    ...  
  }  
}
```

Triggers other Strategies


# Spark's Planner

- 1st Phase: Transforms the Logical Plan to the Physical Plan using Strategies
- 2nd Phase: Use a Rule Executor to make the Physical Plan ready for execution
  - Prepare Scalar sub-queries
  - Ensure requirements on input rows
  - Apply physical optimizations

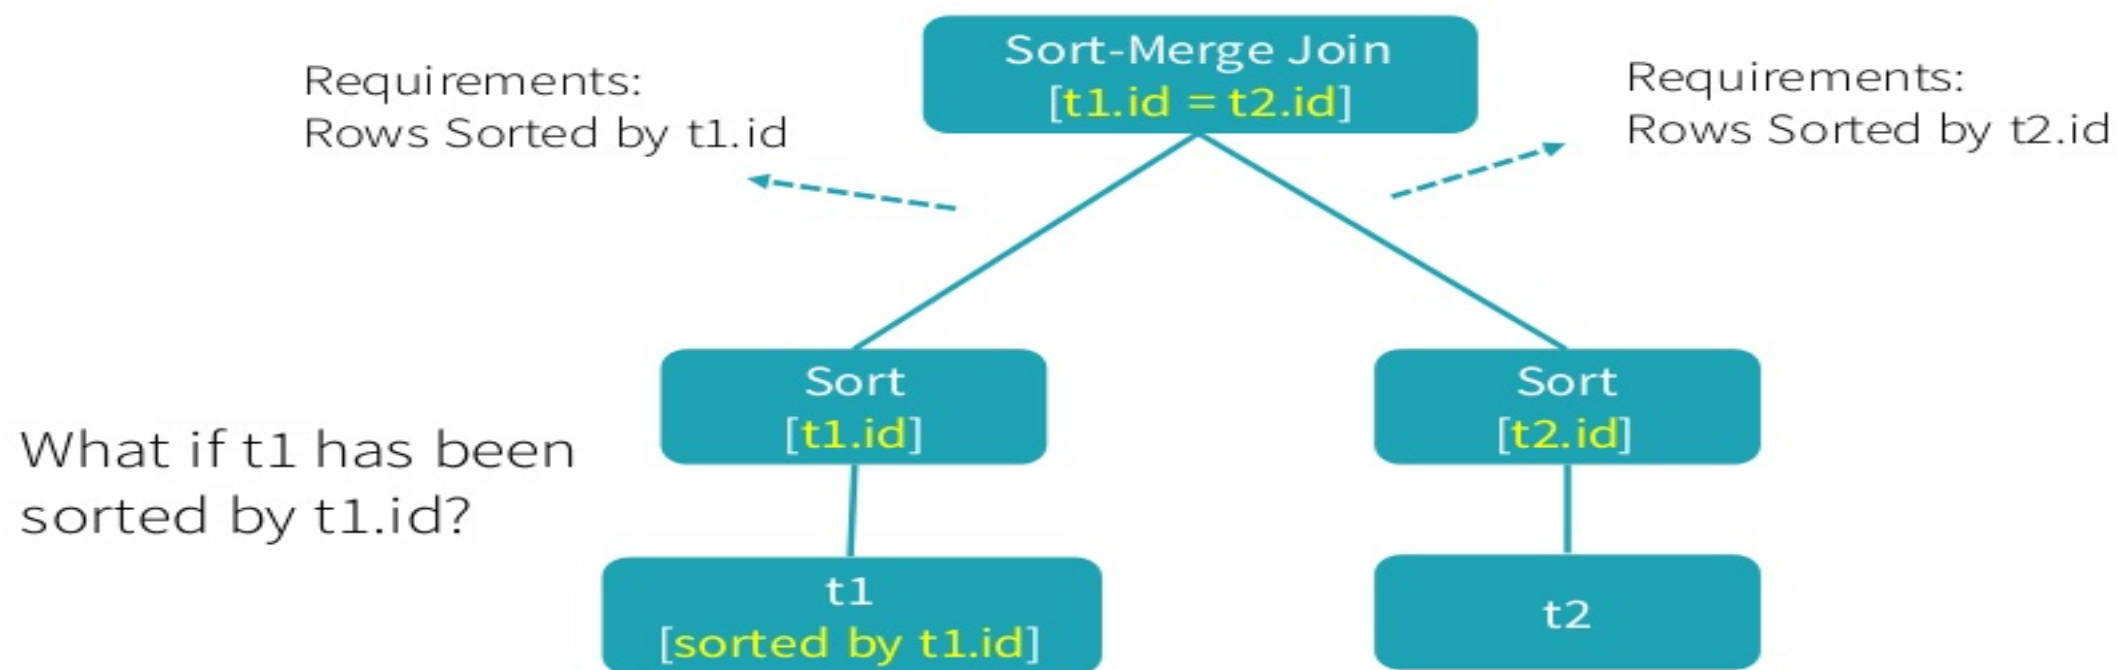
# Ensure Requirements on Input Rows



# Ensure Requirements on Input Rows




# Ensure Requirements on Input Rows





# Ensure Requirements on Input Rows



# Roll your own Planning Rule

## Roll your own Planner Rule

```
import org.apache.spark.sql.functions._

val tableA = spark.range(1000000000).as('a')
val tableB = spark.range(1000000000).as('b')

val result = tableA
  .join(tableB, $"a.id" === $"b.id")
  .groupBy()
  .count()
result.count()
```

*This takes ~22 Seconds on Databricks Community edition*

Roll your own Planner Rule

Can we do better?


# Roll your own Planner Rule - Analysis

```
== Physical Plan ==
*HashAggregate(keys=[], functions=[count(1)], output=[count#43L])
+- Exchange SinglePartition
  +- *HashAggregate(keys=[], functions=[partial_count(1)], output=[count#48L])
    +- *Project
      +- *SortMergeJoin [id#21L], [id#25L], Inner
        :- *Sort [id#21L ASC], false, 0
        :   +- Exchange hashpartitioning(id#21L, 200)
        :     +- *Range (0, 100000000, step=1, splits=Some(8))
        +- *Sort [id#25L ASC], false, 0
        :   +- Exchange hashpartitioning(id#25L, 200)
        :     +- *Range (0, 100000000, step=1, splits=Some(8))
```

# Roll your own Planner Rule

Exploit the structure of the problem

We are joining two intervals; the result will be the intersection of these intervals



## Roll your own Planner Rule - Matching

```
case object IntervalJoin extends Strategy with Serializable {  
  def apply(plan: LogicalPlan): Seq[SparkPlan] = plan match {  
    case Join(  
      Range(start1, end1, 1, part1, Seq(o1)),  
      Range(start2, end2, 1, part2, Seq(o2)),  
      Inner,  
      Some(EqualTo(e1, e2)))  
      if ((o1 semanticEquals e1) && (o2 semanticEquals e2)) ||  
          ((o1 semanticEquals e2) && (o2 semanticEquals e1)) =>  
        // Rule...  
    case _ => Nil  
  }  
}
```



## Roll your own Planner Rule - Body

```
if astart1 <= end2) && (end1 >= end2)) {  
  val start = math.max(start1, start2)  
  val end = math.min(end1, end2)  
  val part = math.max(part1.getOrElse(200), part2.getOrElse(200))  
  val result = RangeExec(Range(start, end, 1, part, o1 :: Nil))  
  val twoColumns = ProjectExec(  
    Alias(o1, o1.name)(exprId = o1.exprId) :: Nil,  
    result)  
  twoColumns :: Nil  
} else {  
  Nil  
}
```

# Roll your own Planner Rule

Hook it up with Spark

```
spark.experimental.extraStrategies = IntervalJoin :: Nil
```

Use it

```
result.count()
```

*This now takes 0.46 seconds to complete*

## Roll your own Planner Rule

```
== Physical Plan ==
*HashAggregate(keys=[], functions=[count(1)],
output=[count#43L])
+- Exchange SinglePartition
   +- *HashAggregate(keys=[], functions=[partial_count(1)],
output=[count#48L])
      +- *Project
         +- *Project [id#21L AS id#21L]
            +- *Range (0, 100000000, step=1, splits=Some(8))
```

# Community Contributed Transformations

SPARK-3462 push down filters and projections into Unions #2345

**Closed** koeninger wants to merge 3 commits into `apache:master` from `mediacrossinginc:SPARK-3462`

Conversation 15 Commits 3 Files changed 2 +110 -0

Showing 2 changed files with 110 additions and 0 deletions. Unified Split

110 line patch took this user's query from  
“never finishing” to 200s.

Overall 200+ people have contributed to the analyzer/optimizer/planner in the last 2 years.

# Where to Start

- Source Code:
  - Trees: [TreeNode](#), [Expression](#), [Logical Plan](#), and [Physical Plan](#)
  - Transformations: [Analyzer](#), [Optimizer](#), and [Planner](#)
- Check out previous pull requests
- Start to write code using Catalyst
- Open a pull request!

# Try Apache Spark with Databricks

- Try latest version of Apache Spark

<http://databricks.com/try>



## FULL-PLATFORM TRIAL

Put Apache® Spark™ to work

- Unlimited clusters
- Notebooks, dashboards, production jobs, RESTful APIs
- Interactive guide to Spark and Databricks
- Deployed to your AWS VPC
- BI tools integration

14-Day Free Trial

START YOUR TRIAL TODAY



## COMMUNITY EDITION

Learn Apache® Spark™ for free

- Min 5GB cluster for learning Spark
- Interactive notebooks and dashboards
- Online learning resources
- Public environment to share your work

Free

SIGN UP



# Questions?

I will be available in the Databricks booth (D1) afterwards

@Westerflyer

