

EXTENDING SPARK WITH JAVA AGENTS

Jaroslav Bachorik
Adrian Popescu



unravel



SPARK SUMMIT
EUROPE 2016

Meet The Speakers

Jaroslav

- Backend engineer
- Hands-on experience with JVM
 - OpenJDK Reviewer
- Strong focus on performance
 - VisualVM Contributor
 - BTrace Maintainer and Lead Developer

Adrian

- Backend engineer
- 5+ years of experience in performance monitoring & modeling
- Focusing on tuning and optimization of Big Data apps

Unravel

Performance Intelligence for Big Data Stack



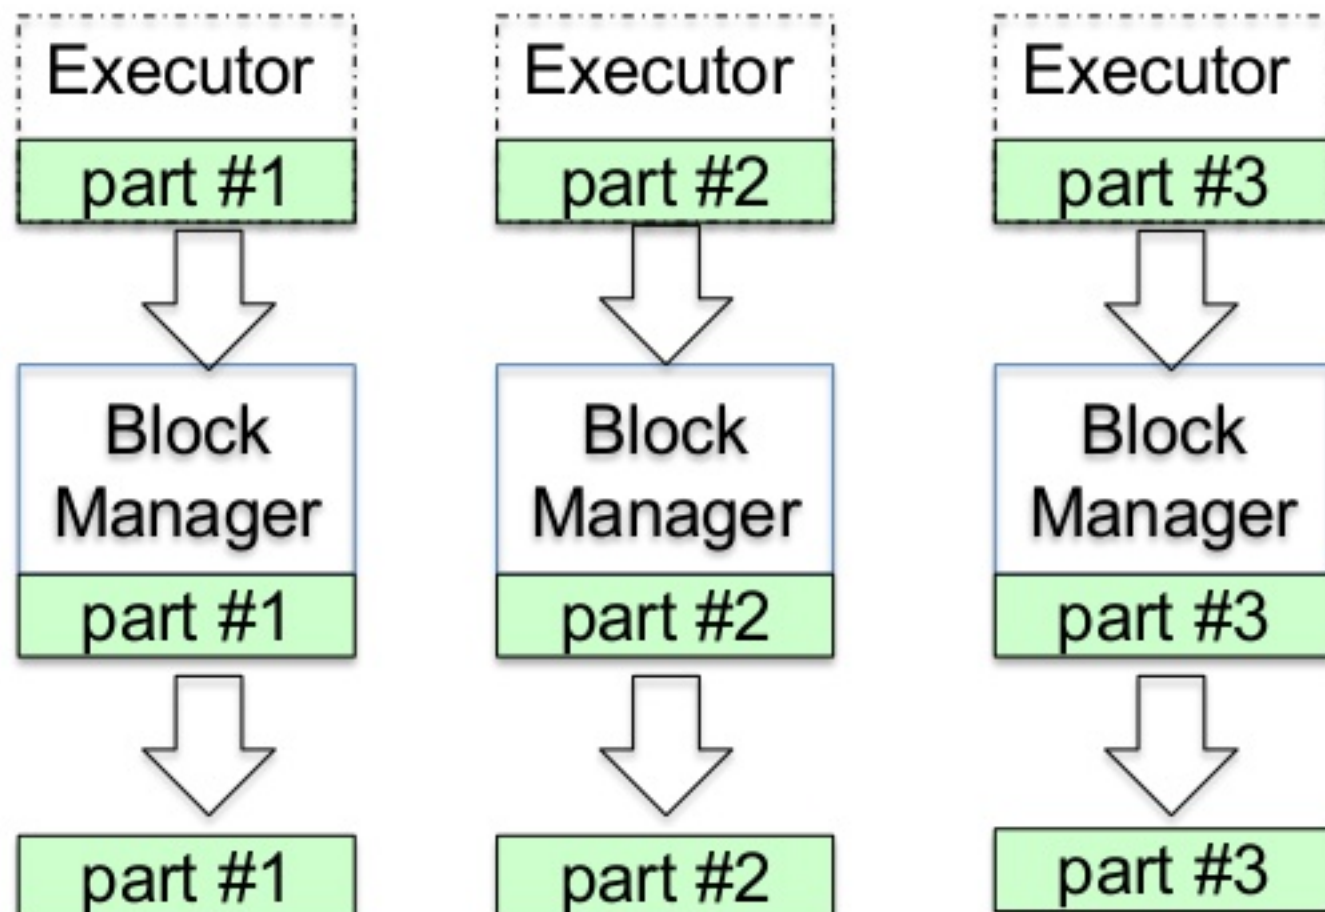
Caching in Spark

```
...  
rdd.cache()  
rdd.count()  
...  
  
rdd.reduceByKey(  
  (x,y) =>  
    x+y).count()
```

1. Evaluate

2. Caching

3. Fetch or
re-evaluate



Avoiding re-evaluation for cached RDD blocks

Spark Caching Problem

- **RDD Caching**

- Benefits: speedup, saved resources
- Too much caching may waste resources



- **Spot the opportunity**

- Which RDDs to cache (for most benefit)?
- How to prioritize when memory is scarce?
- Where to cache?
- Benefit depends on memory state



- <http://unraveldata.com/to-cache-or-not-to-cache/>

Challenging to decide **when**, **what**, and **where** to cache

Algorithm

1. **Find** stages that share RDDs

2. **Cached RDDs:** Measure benefit
Block *hitRate* & *time saved* for every Stage

3. **RDDs not cached:** Potential benefit
Approximate time saved & RDD storage size

Insufficient visibility
into Spark structures

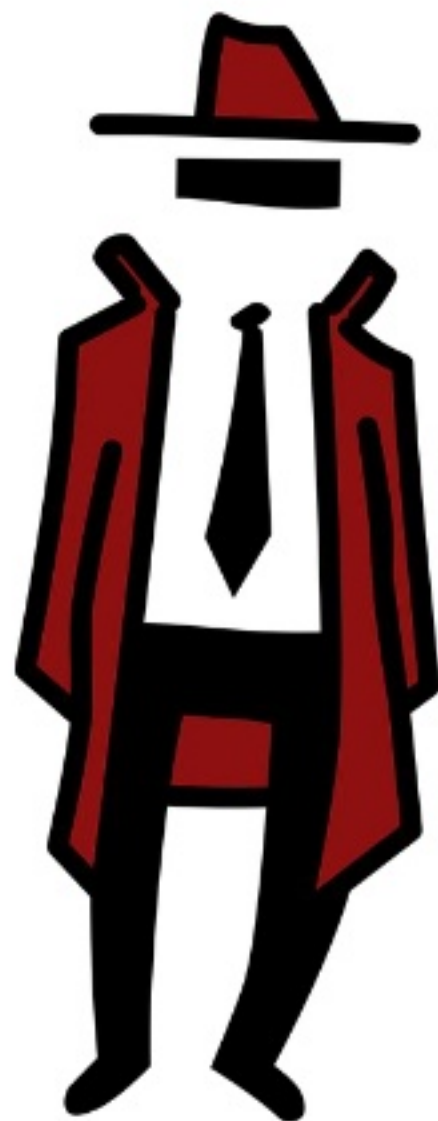
e.g., only block info in
BlockStatusListener

4. **Suggest** which RDDs to cache

Need to collect additional metrics



Meet Java Agent



Meet Java Agent

Launcher

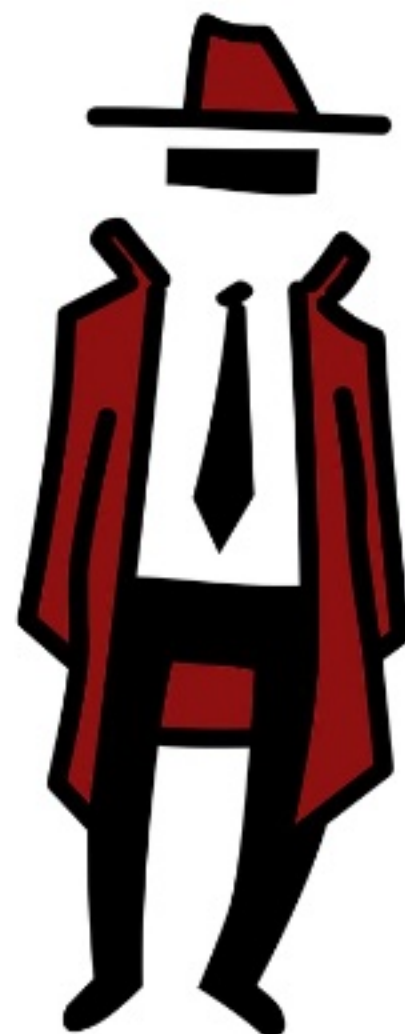
```
java -javaagent:myfancyagent.jar=<args>
```

- Runs JVM with an agent
- Agent jar
 - agent class
 - required manifest entries
- Arguments
 - any string understood by agent class
- Multiple agents can be specified

```
Premain-Class: com.myorg.agents.FancyAgent
```

```
Can-Redefine-Classes: true
```

```
Can-Transform-Classes: true
```



Meet Java Agent

Launcher

```
java -javaagent:myfancyagent.jar=<args>
```

FancyAgent.java

```
public static void premain(String args, Instrumentation inst) {
```

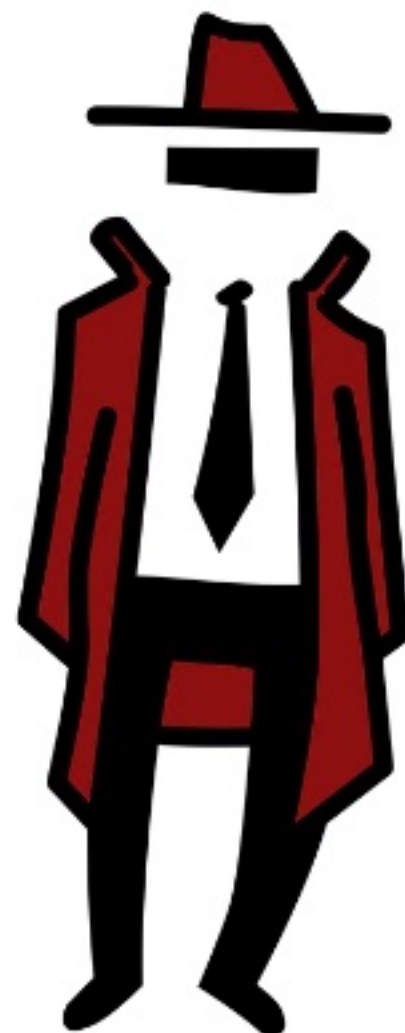
- Executed at JVM startup
- **Instrumentation** provides access to JVM internals

META-INF/manifest.mf

```
Premain-Class: com.myorg.agents.FancyAgent
```

```
Can-Redefine-Classes: true
```

```
Can-Transform-Class: true
```



Meet Java Agent

Launcher

```
java -javaagent:myfancyagent.jar=<args>
```

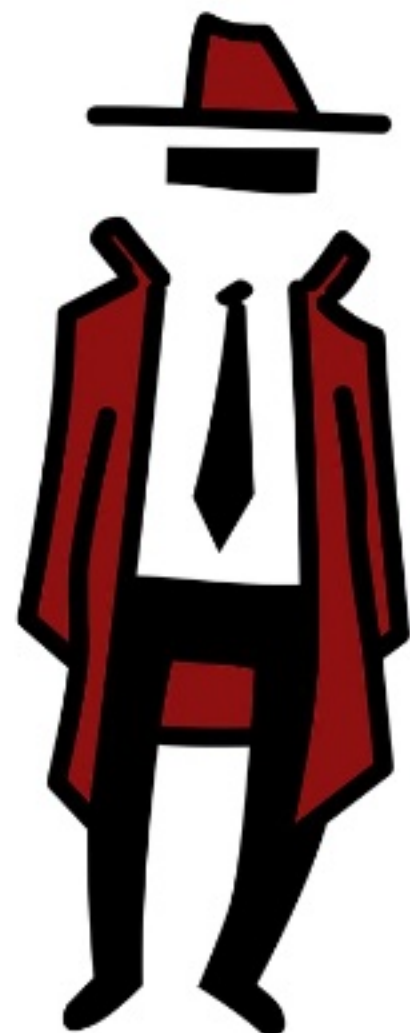
FancyAgent.java

```
public static void premain(String args, Instrumentation inst) {  
    // do the magic stuff here  
}
```

META-INF/manifest.mf

```
Premain-Class: com.myorg.agents.FancyAgent
```

- Agent class
- Must be present in the agent jar



Meet Java Agent

Launcher

```
java -javaagent:myfancyagent.jar=<args>
```

FancyAgent.java

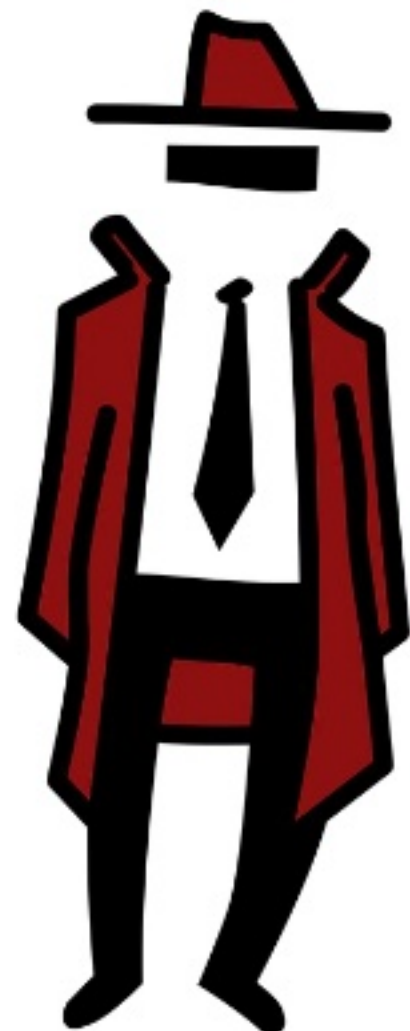
```
public static void premain(String args, Instrumentation inst) {  
    // do the magic stuff here  
}
```

META-INF/manifest.mf

Transform classes upon load

Can-Redefine-Classes: true

Can-Retransform-Classes: true



Meet Java Agent

Launcher

```
java -javaagent:myfancyagent.jar=<args>
```

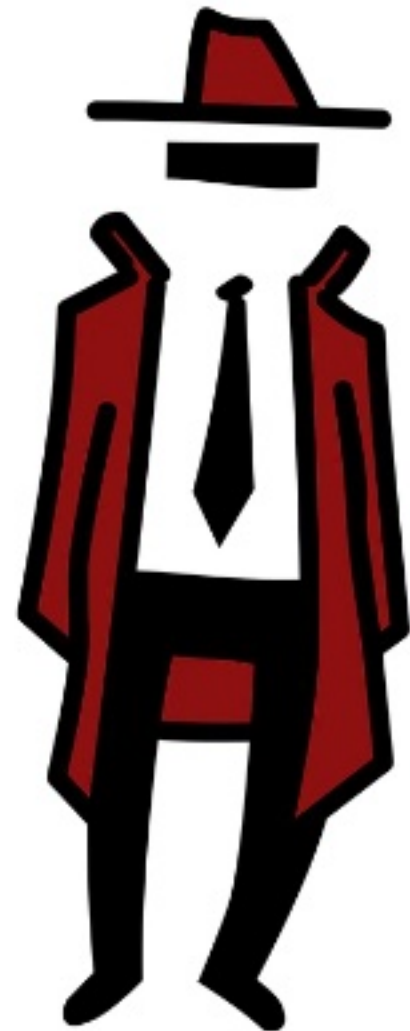
FancyAgent.java

```
public static void premain(String args, Instrumentation inst) {  
    // do the magic stuff here
```

- Agent capabilities and settings
 - retransform classes
 - redefine classes
 - native method prefix

<https://docs.oracle.com/javase/7/docs/api/java/lang/instrument/package-summary.html>

Can-Retransform-Classes: true



Class Transformers

```
class SpecialLibraryClass {  
    public void doStuff() {  
        System.out.println(  
            "Doing stuff #" + id(System.currentTimeMillis())  
        );  
    }  
    private int id(int input) {  
        return input % 42;  
    }  
}
```

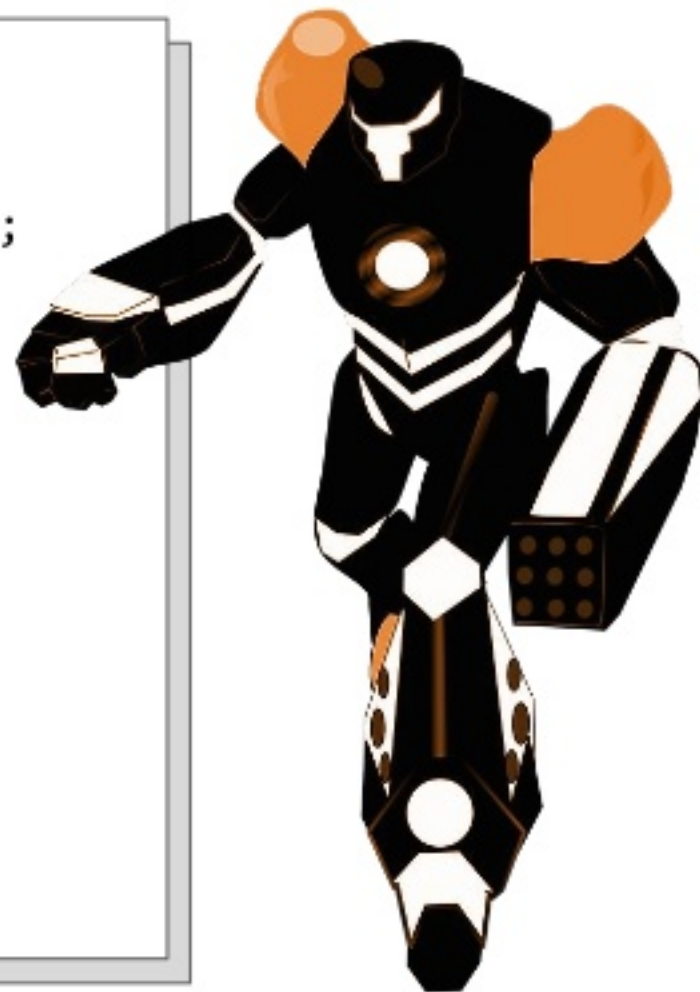
Private method in 3rd party package
private method.

Patch and recompile? Reflection?



Class Transformers

```
class SpecialLibraryClass {  
    public void doStuff() {  
        System.out.println("Doing stuff #" + id(System.currentTimeMillis()));  
    }  
    private int id(int input) {  
        AgentCallbacks.onEnter("SLC", "id", "input = " + input);  
        int tmp = input % 42;  
        AgentCallbacks.onReturn("SLC", "id", tmp);  
        return tmp;  
    }  
}
```



Class Transformers

```
java.lang.instrument.ClassTransformer
```

```
byte[] transform(ClassLoader l, String name, Class<?> cls,  
                ProtectionDomain pd, byte[] classfileBuffer)
```

- Inspect and modify class data
 - without recompilation
 - even on-the-fly
- Powerful
 - introduce callbacks
 - patch problematic code
- Not a simple task
 - class data structure
 - constant pool
 - stack frame maps



<http://wikipedia.org>

BTrace

<https://github.com/btraceio/btrace>

- Open source dynamic instrumentation framework for JVM
- Java annotations based DSL for class transformations
- Optimized for high performance
- Used eg. in VisualVM
(<http://visualvm.github.io>)



btrace.io

Spark Config Tuning

- Intercept **SparkContext** instantiation at Spark Driver
- Modify the associated **SparkConf** object
 - before **SparkContext** is fully instantiated
- Any Spark property can be changed
 - eg. “**spark.executor.cores**” based on system load
- Modified **SparkConf** will be applied to the current job
 - and distributed to to all executors

Config Tuning Transformation

```
@OnMethod(clazz = "org.apache.spark.SparkContext",  
          method = "<init>",  
          location = @Location(Kind.RETURN)  
)  
  
public static void sparkContextConstructor(@Self SparkContext scObj) {  
    SparkConf conf = scObj.conf();  
    if (conf == null)  
        return;  
    int recommendedCores = AdviceSystem.getRecommendedCores(scObj);  
    conf.set("spark.executor.cores", recommendedCores);  
}
```

Agents on Spark?



Agents on Spark!

- Spark is a JVM application
- Any Java agent can be used with Spark



*<http://apache.org>
<http://pixabay.com>*

Register JVM Agent for Spark

```
java -javaagent:myfancyagent.jar=<args>
```

Spark configuration parameters

```
spark.executor.extraJavaOptions=  
    "-javaagent:/opt/myagent/myagent.jar=opt.load:true"  
spark.driver.extraJavaOptions=  
    "-javaagent:/opt/myagent/myagent.jar=track.conf:true"
```

Store the options in **spark-defaults.conf** to apply system-wide.



<http://arthurandsage.blogspot.cz/>

Let's Put It All Together!



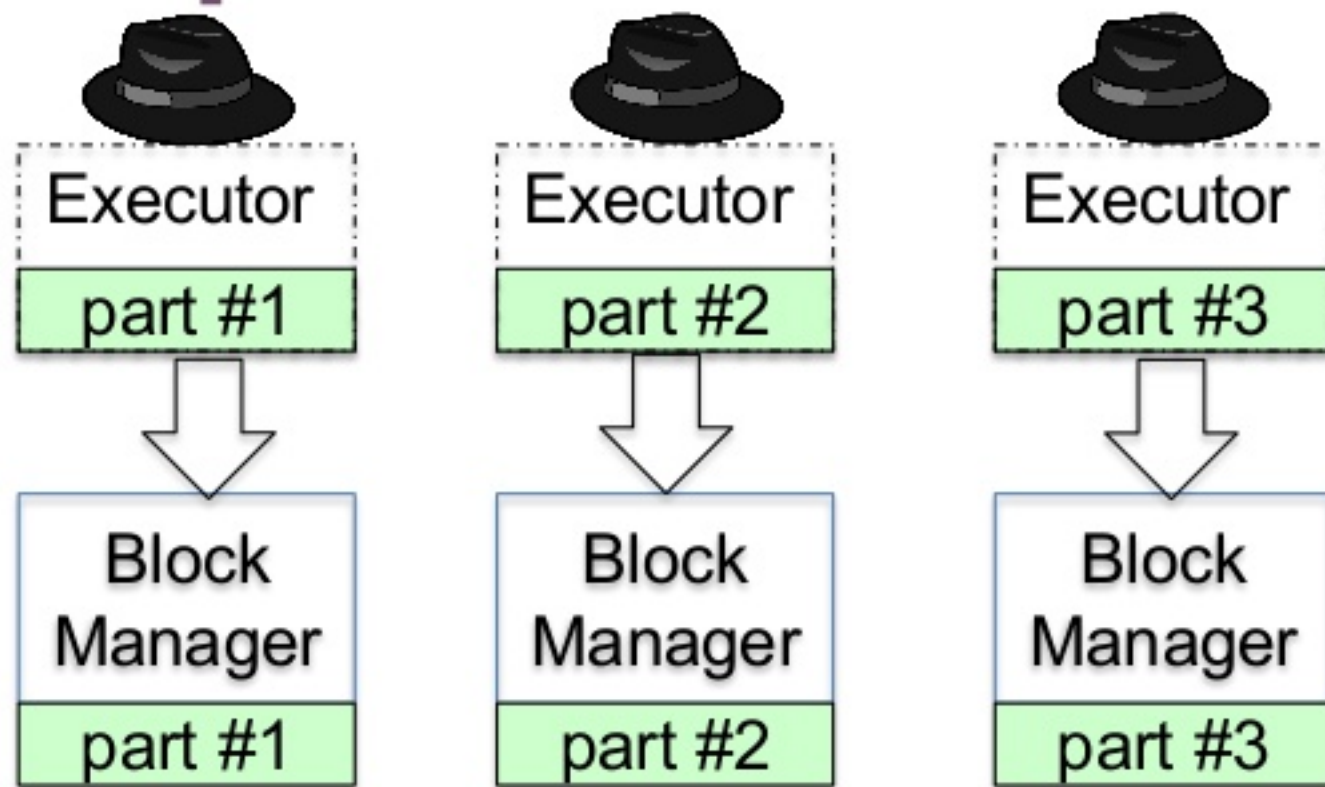
Identify Required Data

...
`rdd.cache()`
`rdd.count()`
...

`rdd.reduceByKey((x,y) => x+y).count()`
...

1. Evaluation

2. Caching



Agents
allow:

- Access to Spark internal structures
- Collecting: RDD block usage, cache hit, potential benefit
- Run algorithm and suggest which RDDs to cache

Transform Classes

```
def get(blockId: BlockId): Option[BlockResult] = {  
  val local = getLocal(blockId)  
  if (local.isDefined) {  
    logInfo(s"Found block $blockId locally")  
    return local  
  }  
  val remote = getRemote(blockId)  
  if (remote.isDefined) {  
    logInfo(s"Found block $blockId remotely")  
    return remote  
  }  
  None  
}
```



```
def get(blockId: BlockId): Option[BlockResult] = {  
  AgentRuntime.useBlock(blockId, currentStage)  
  val local = getLocal(blockId)  
  if (local.isDefined) {  
    logInfo(s"Found block $blockId locally")  
    return local  
  }  
  val remote = getRemote(blockId)  
  ...  
  ...  
}
```


Unravel Spark Cache Events

SPARK application_1469434457872_8057

Identify the costs of not-caching

OPPORTUNITY FOR RDD CACHING

22.0 minutes spent recomputing RDDs in the application

Give very precise advice

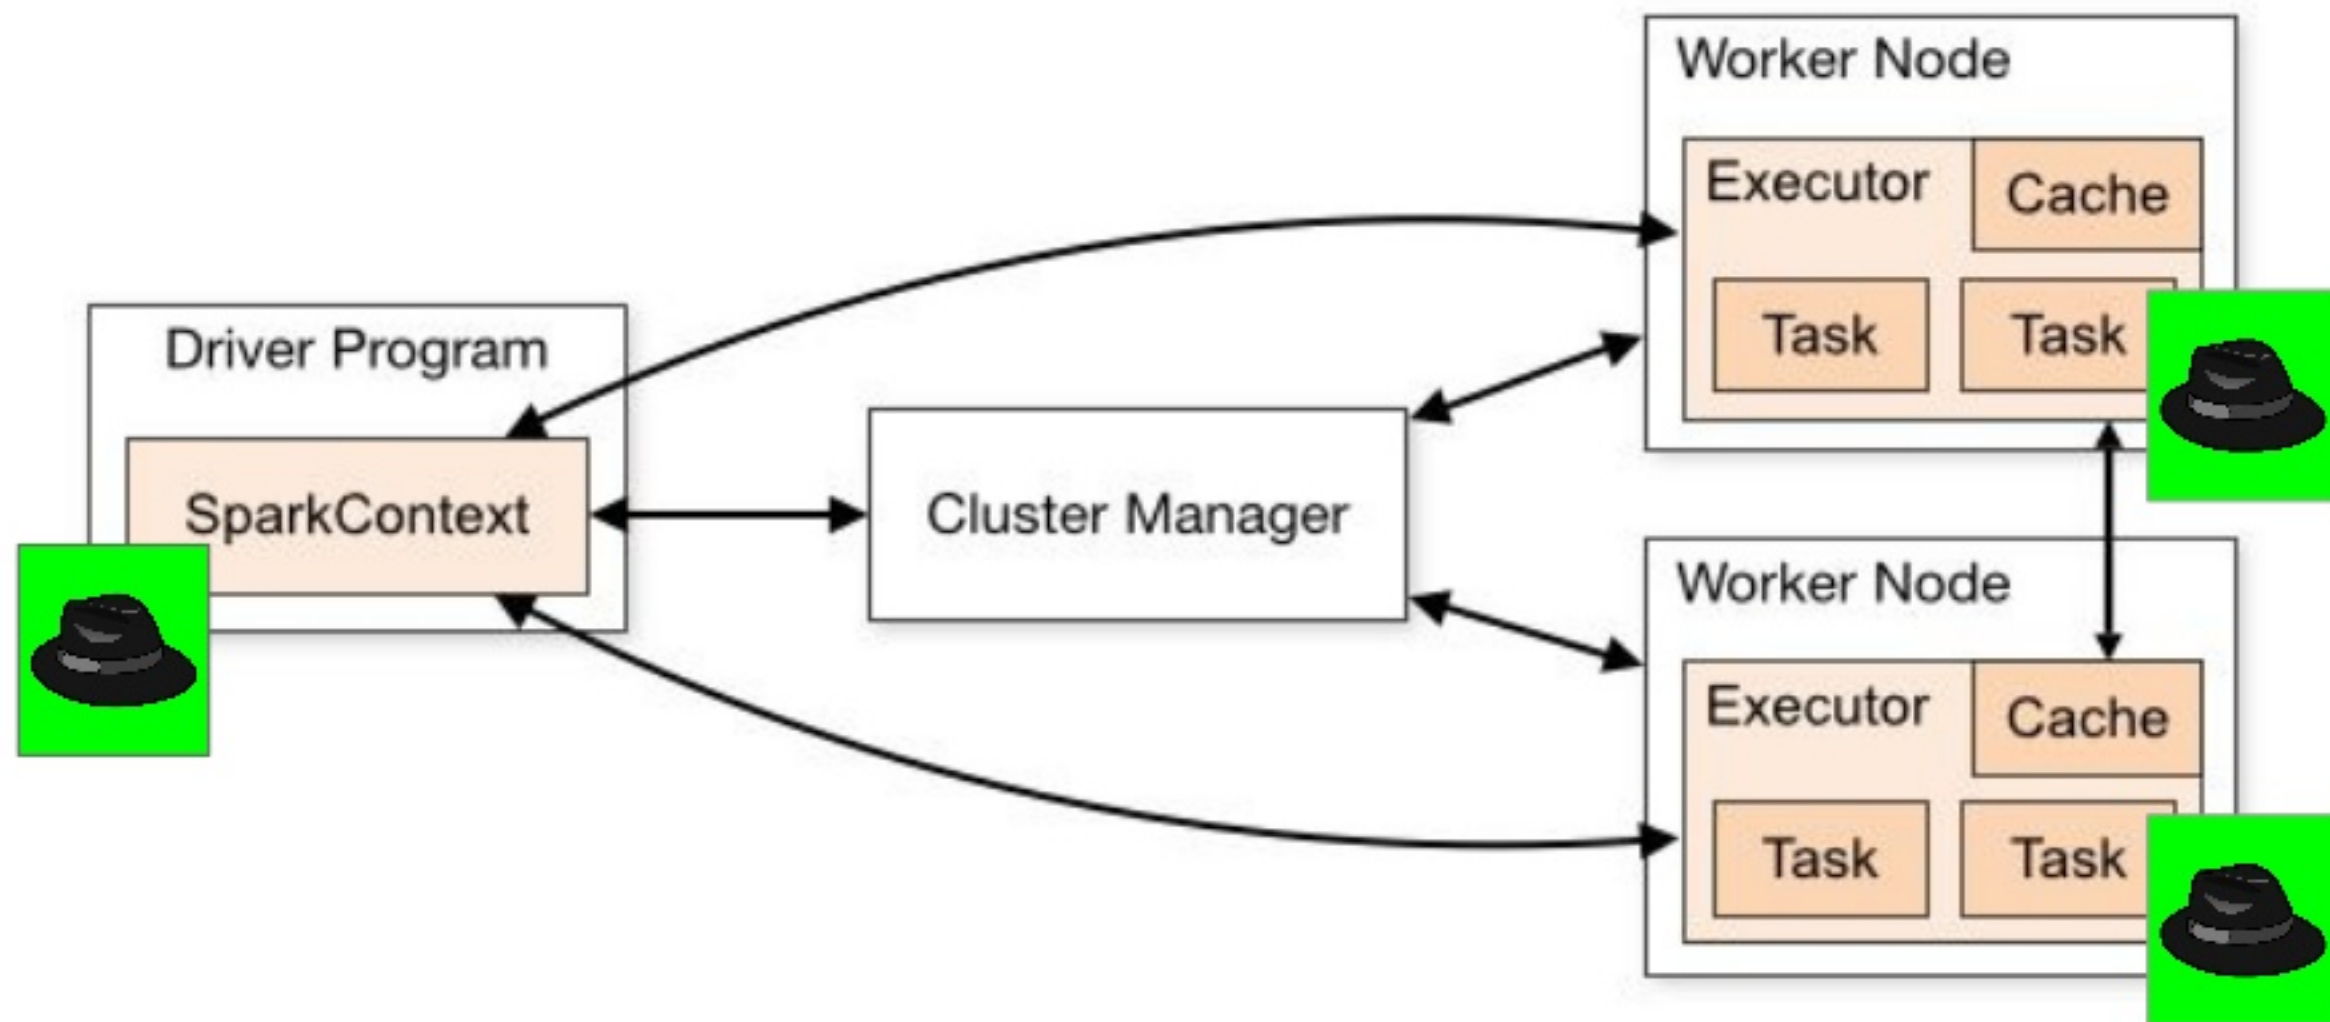
Adding a `cache()` statement before `count at UniformGroupByTest.scala:77` can save up to 22.0 minutes

Caching with `StorageLevel.MEMORY_AND_DISK_SER` is recommended

Put It On Cluster



Agents In Cluster



Distributing Agents for Spark

Make sure all nodes have the latest and same copy of agent binaries

- Manually copy agents to all nodes
- Script that copy process
- Put that script in CRON!



Yarn Localization Service

- Built-in resource distribution system
- Cluster resources
 - linked into executor working directories
 - kept up-to-date
- Supported resource types
 - files
 - archives
- <https://spark.apache.org/docs/1.5.2/running-on-yarn.html>



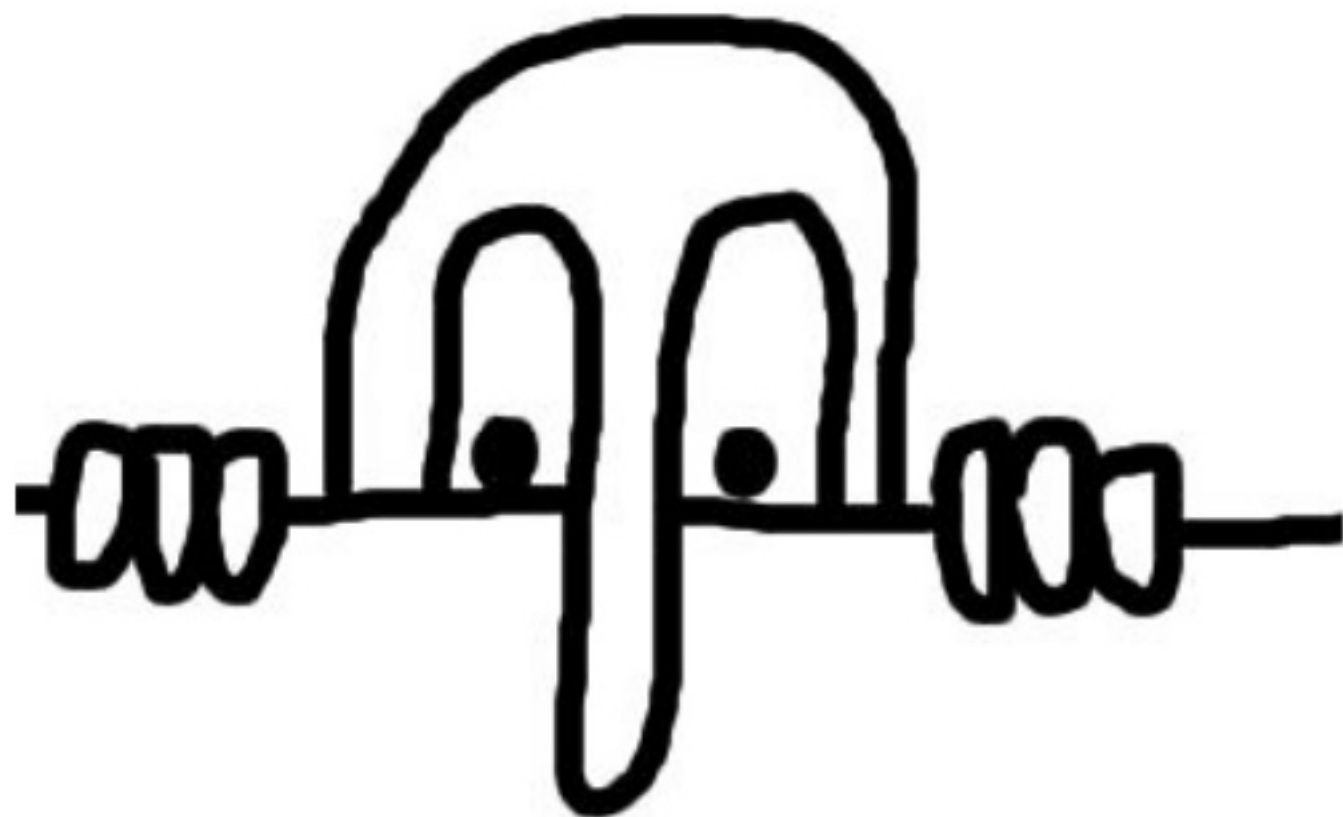
Takeaways

- Sometimes even rich set of Spark metrics, logs and events is not sufficient
- Java agents are powerful tool
 - extend application beyond the original intent
 - hot-patch problems or enhancements
 - may be removed when not needed
- Try Unravel to see the power of Java agents unleashed

Resources

- To simplify application and operations management with Spark, try out Unravel
 - <http://unraveldata.com/free-trial/>
- Start with Java agents and try out BTrace
 - <https://github.com/btraceio/btrace>
 - contributors most sincerely welcome!
- Spark Configuration
 - <http://spark.apache.org/docs/latest/configuration.html>
 - <https://spark.apache.org/docs/1.5.2/running-on-yarn.html>
- Java Agents
 - <https://docs.oracle.com/javase/7/docs/api/java/lang/instrument/package-summary.html>
 - <http://zeroturnaround.com/rebellabs/how-to-inspect-classes-in-your-jvm/>

Q&A



Unravel Free Trial: <http://unraveldata.com/free-trial/>

THANK YOU.

Jaroslav Bachorik, jaroslav@unraveldata.com

Adrian Popescu, adrian@unraveldata.com

