

# DYNAMIC RESOURCE ALLOCATION, DO MORE WITH YOUR CLUSTER

Luc Bourlier  
Lightbend




- Dynamic Resource Allocation
- ^^ in Spark
- External Shuffle Service
- Configuration
- Demo
- Spark Streaming

# Dynamic Resource Allocation


# Dynamic Resource Allocation




Cluster



# Dynamic Resource Allocation




Cluster



# Dynamic Resource Allocation




Cluster



# Dynamic Resource Allocation



Cluster






# Dynamic Resource Allocation



Cluster






# Dynamic Resource Allocation



Hmm, actually, I don't need all this power anymore.

Cluster



# Dynamic Resource Allocation



Cluster




# Dynamic Resource Allocation

Why?

- Shared cluster
- Optimization of resource usage


When?

- variable load job

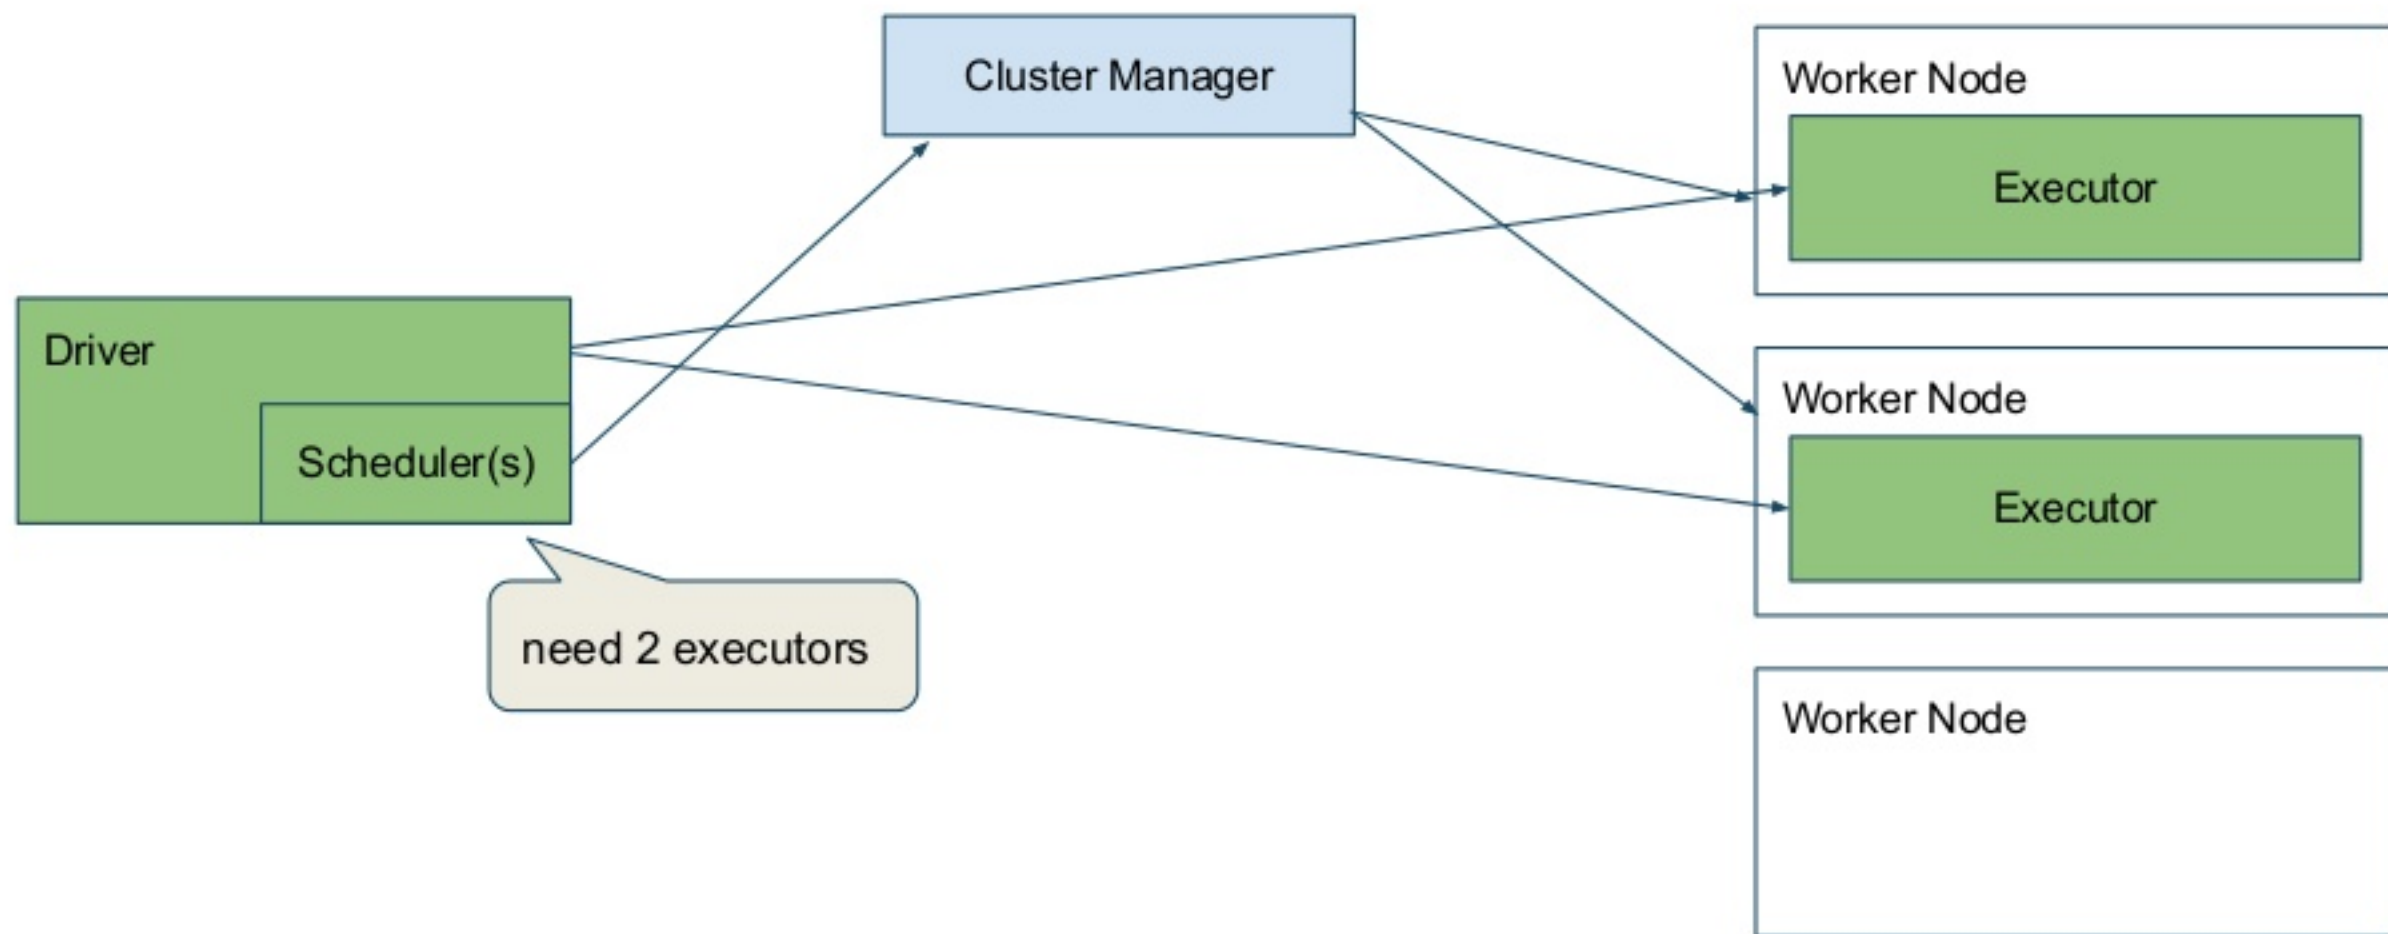


# Spark Cluster Architecture

# Spark Cluster Architecture




# Spark Dynamic Allocation

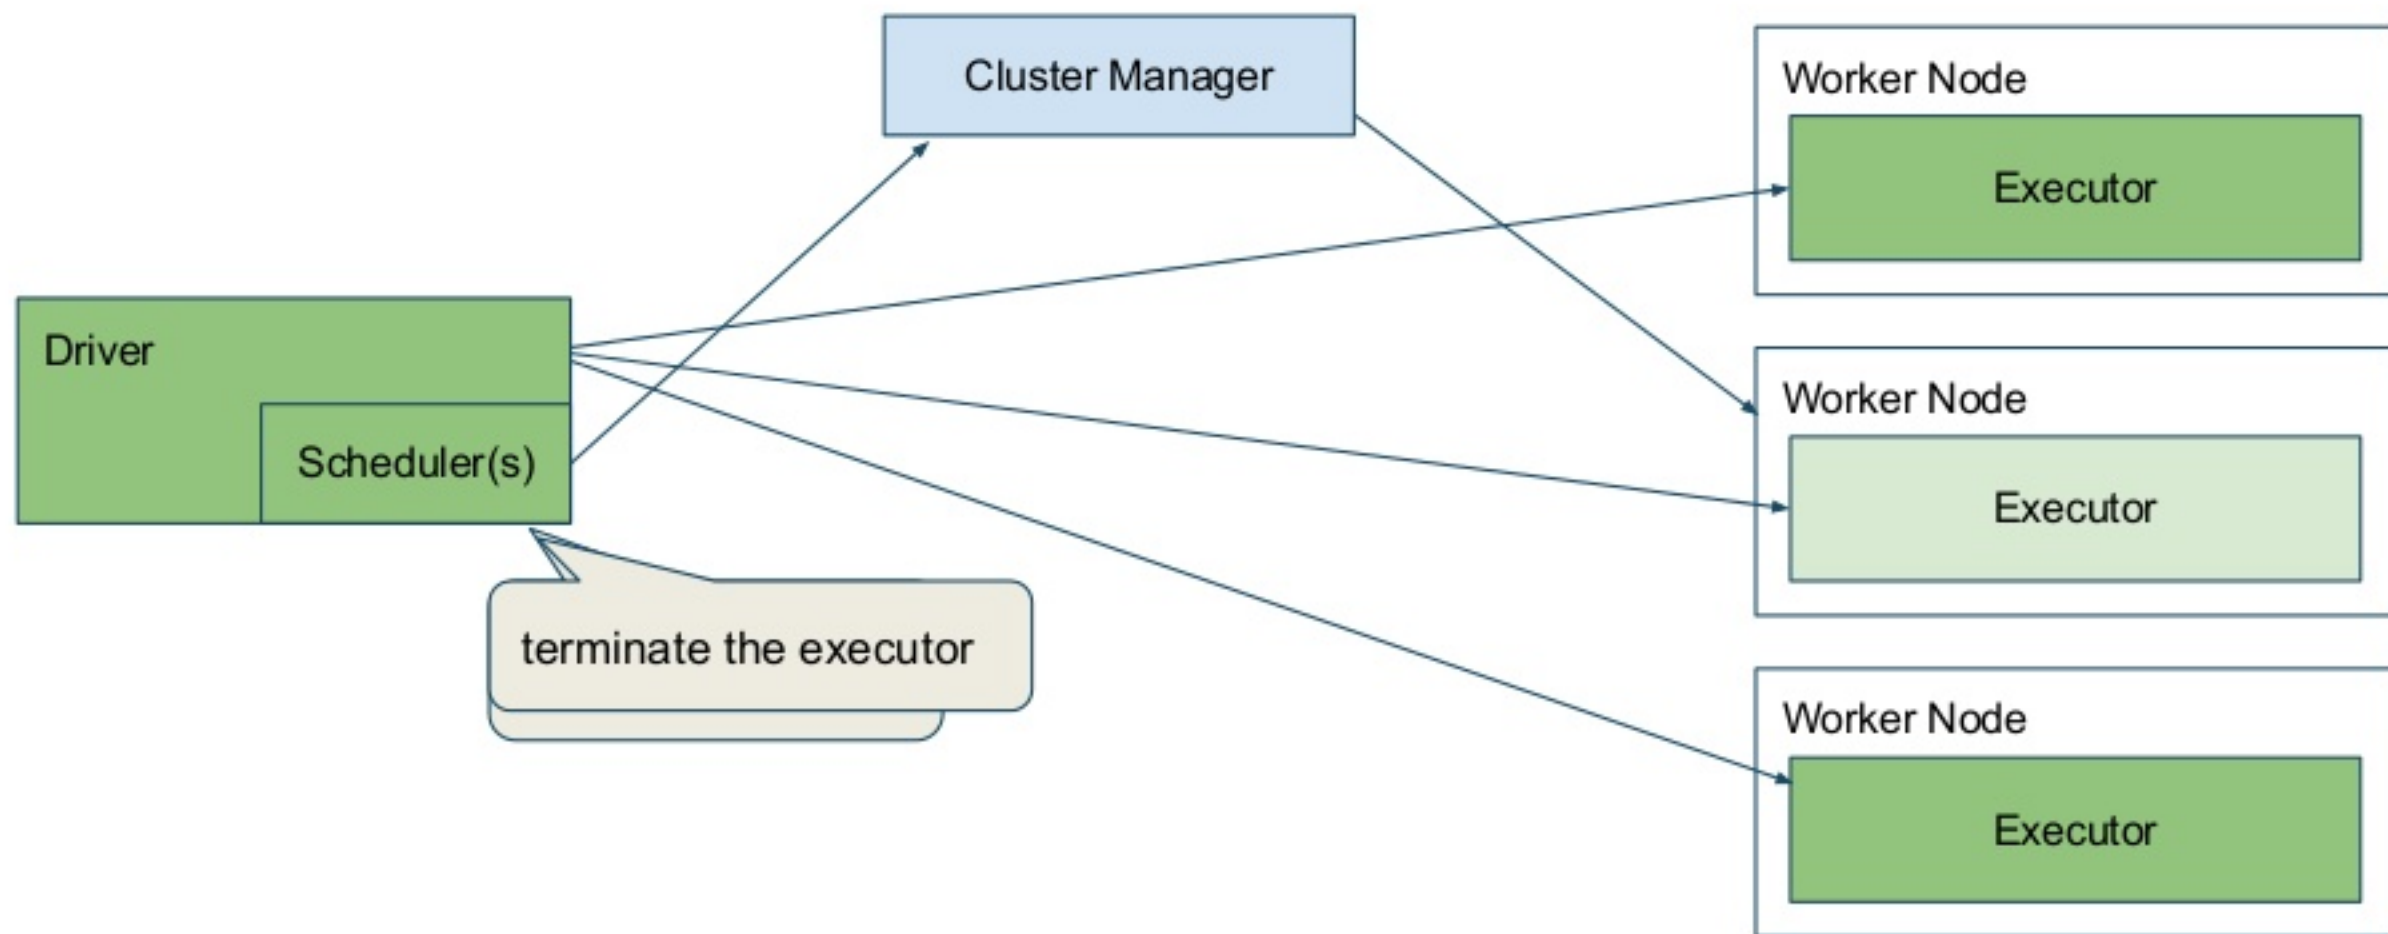




# Spark Dynamic Allocation



# Spark Dynamic Allocation

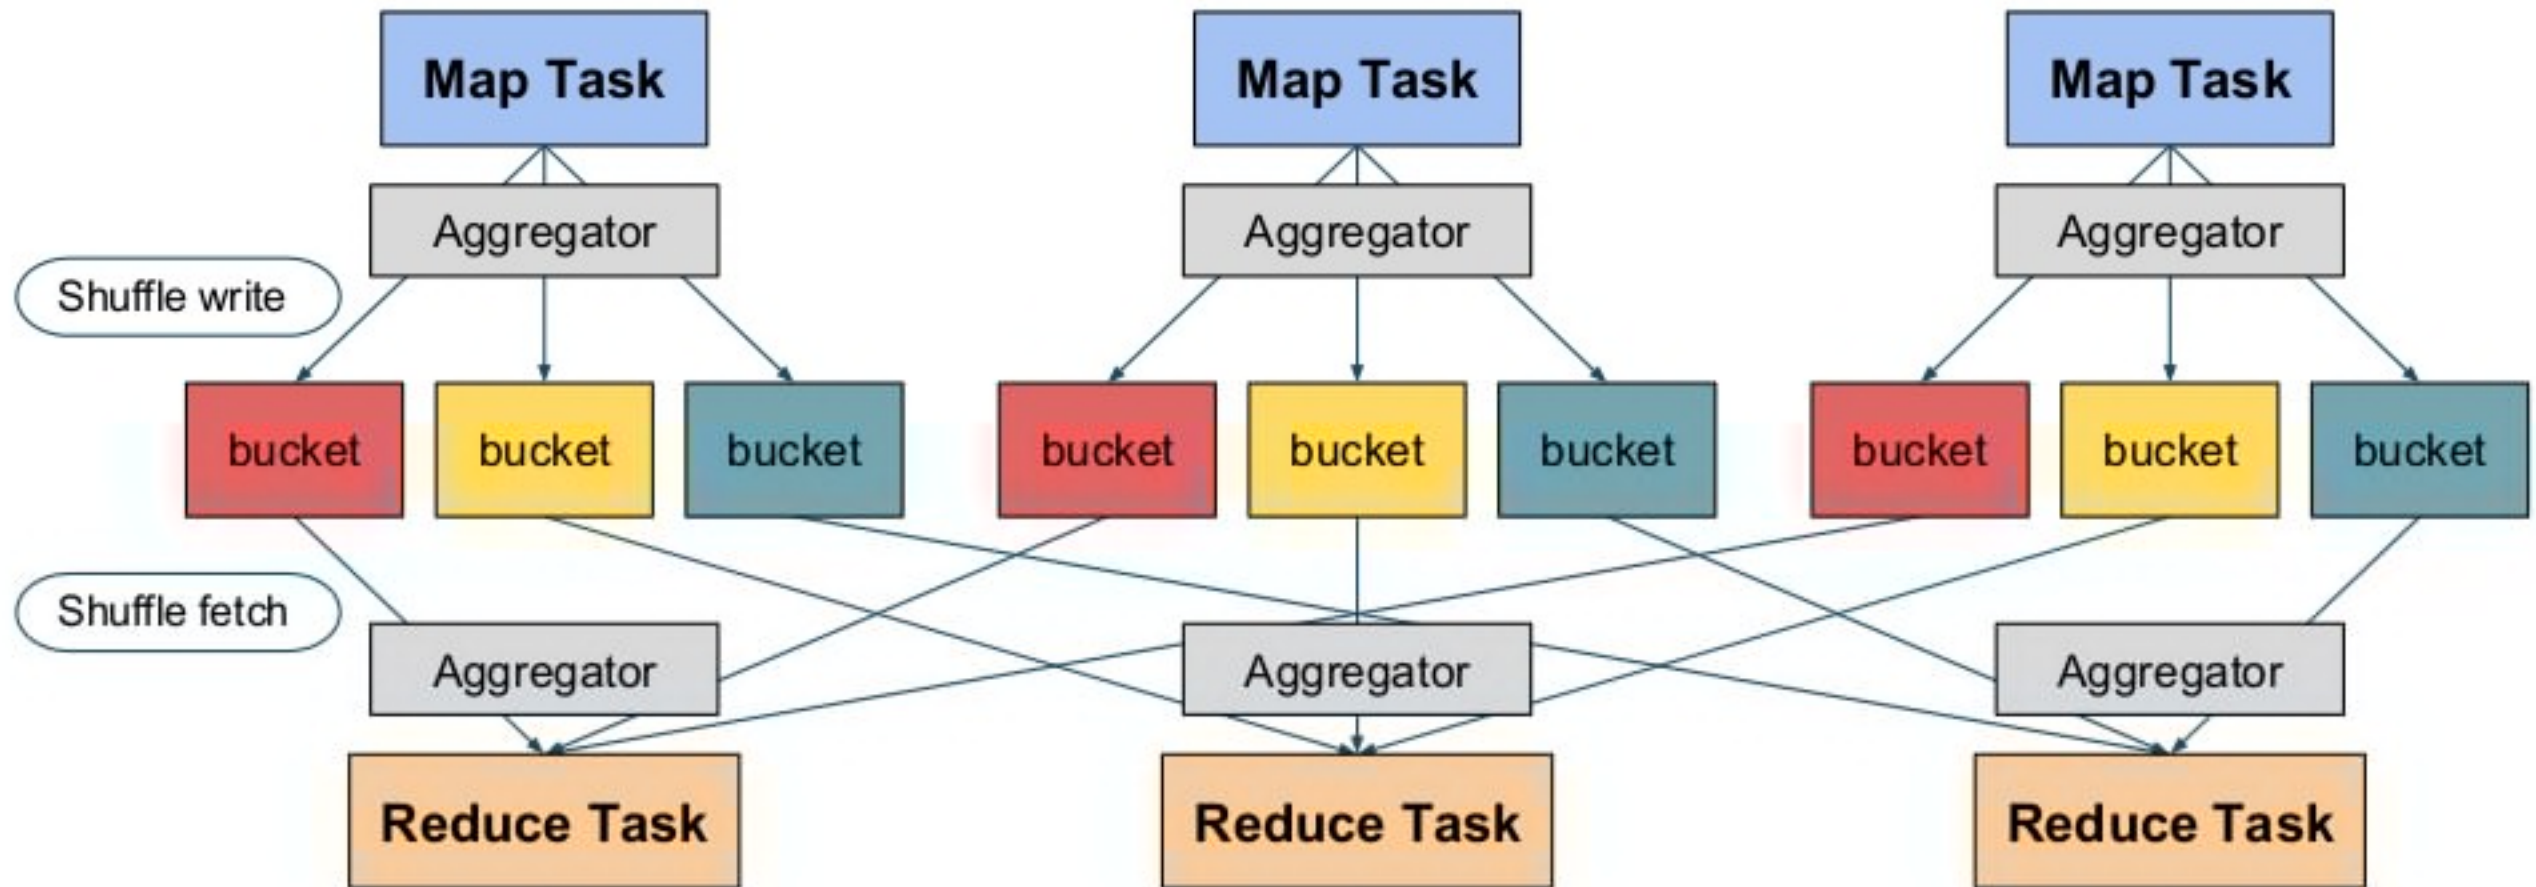


# External Shuffle Service

# External Shuffle Service

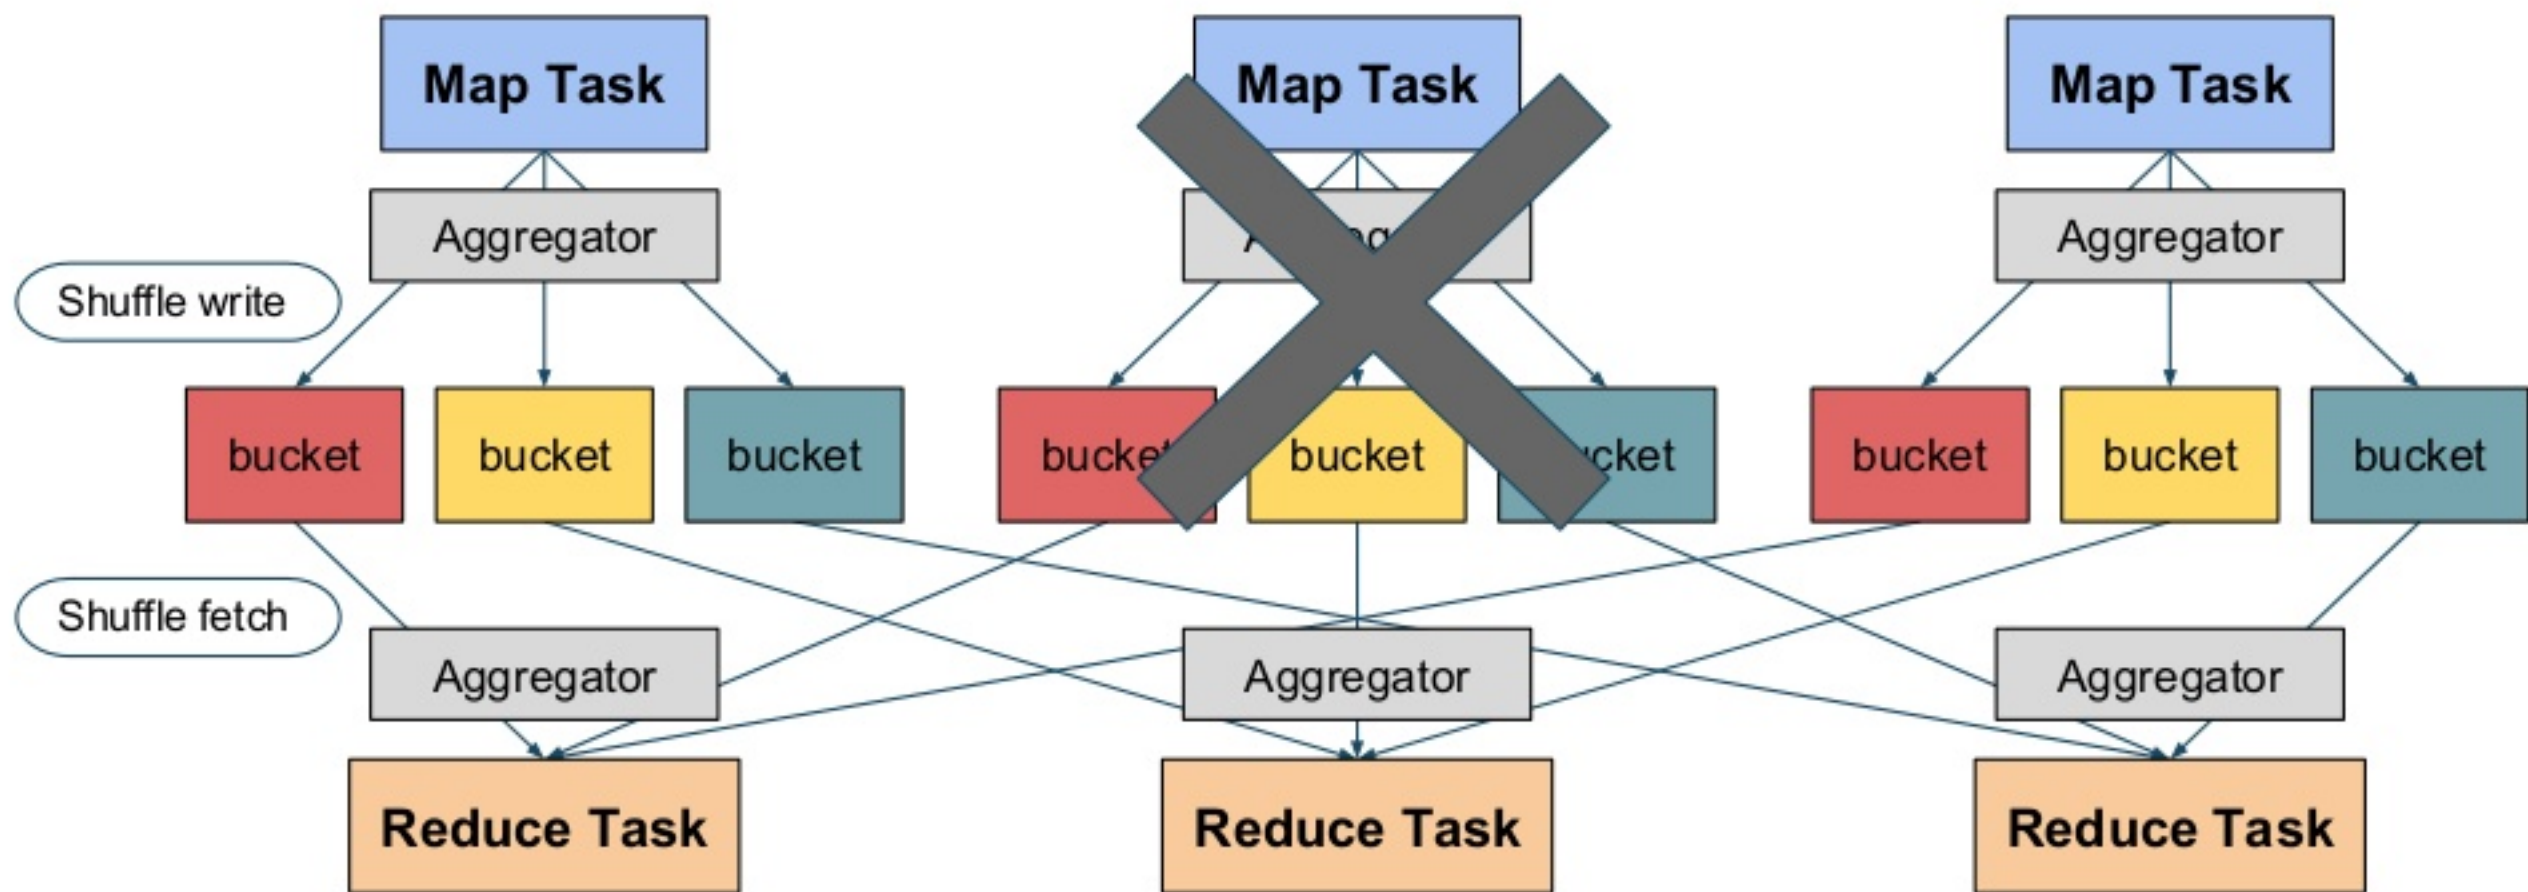
- Did we lose any data?

# External Shuffle Service



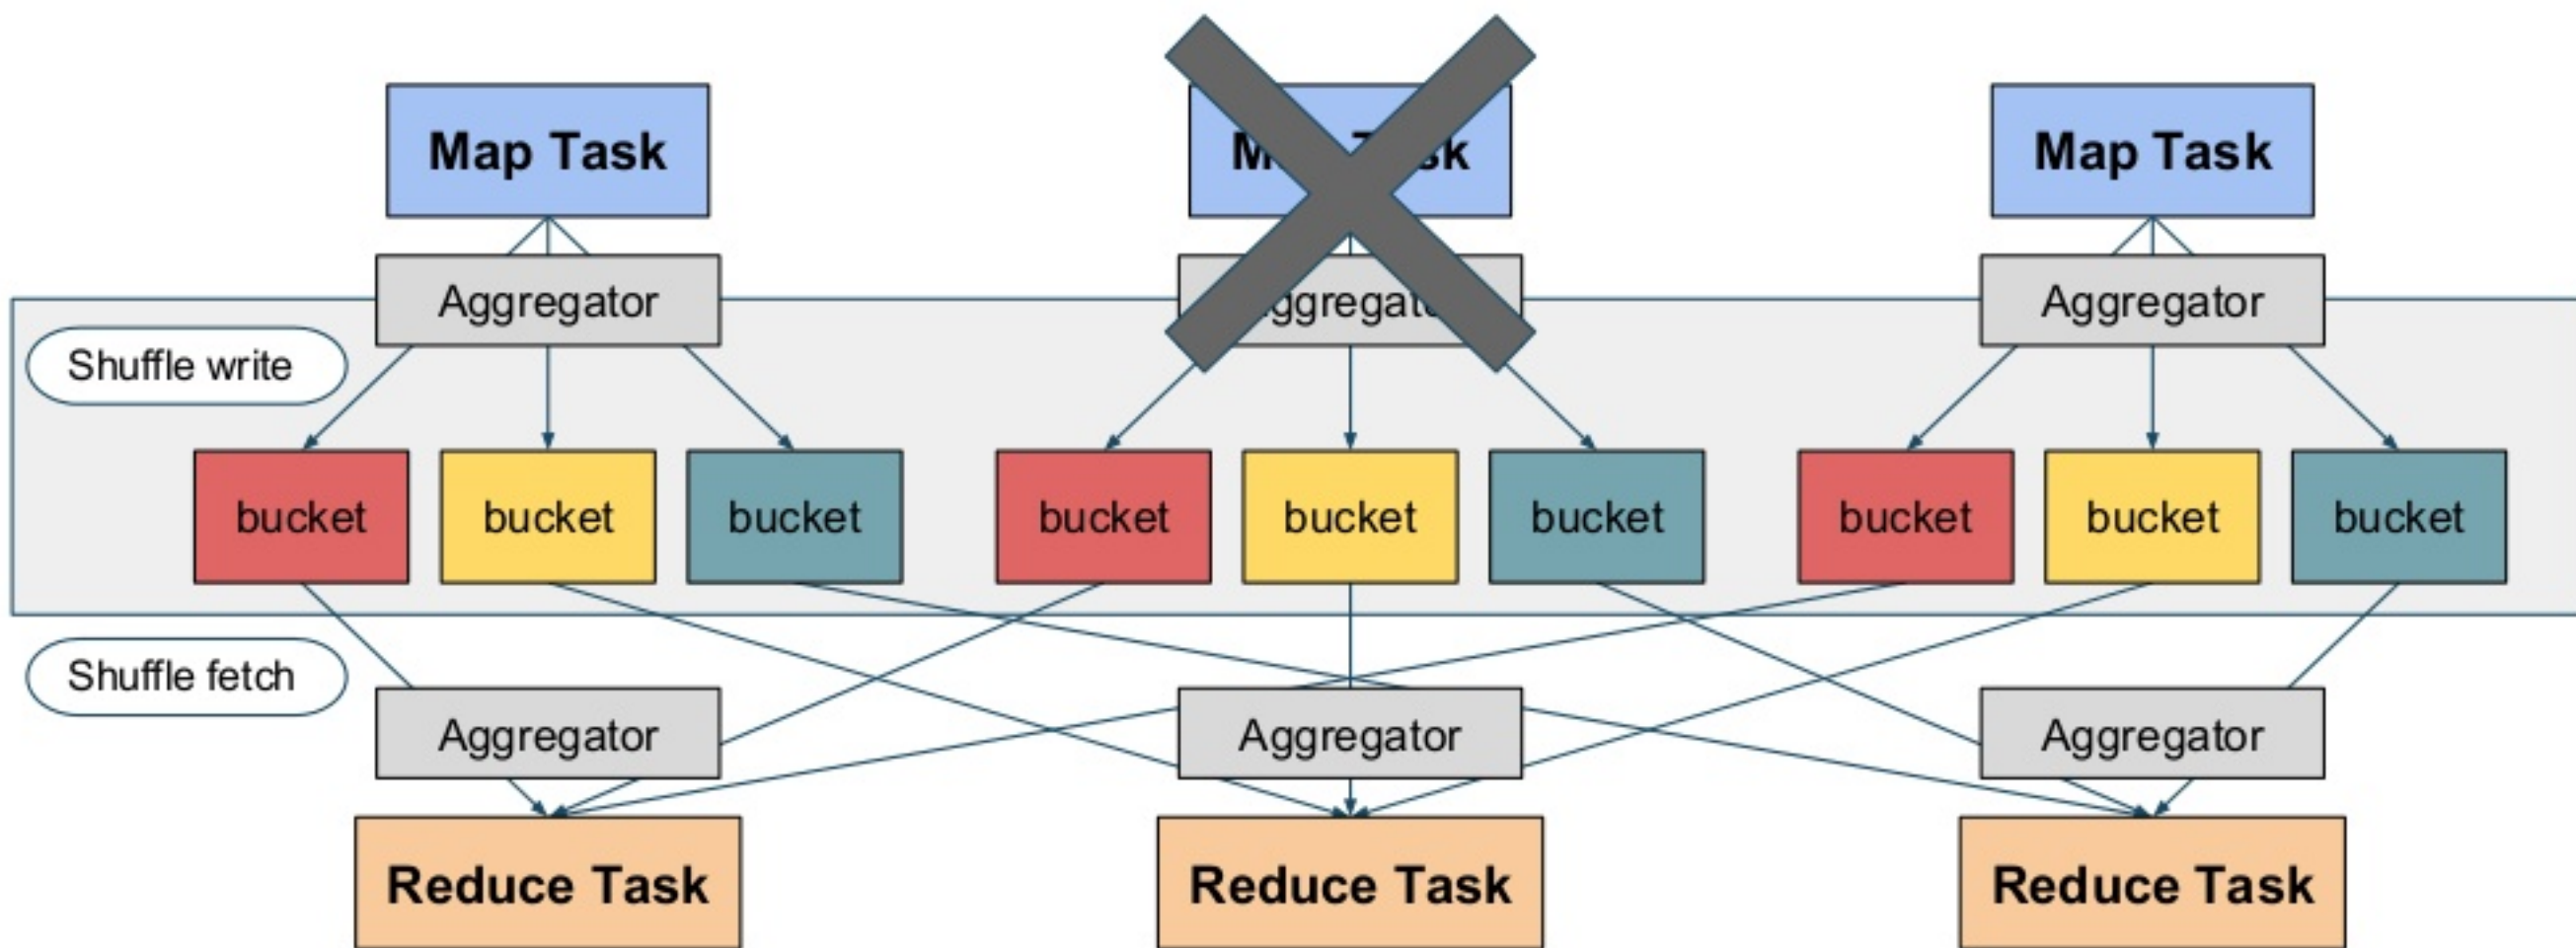


# External Shuffle Service





# External Shuffle Service



# External Shuffle Service

- Extracted from executor
- Manage the local aggregated data for the shuffle operations
- Maintain the data until the application is done.

# Configuration

# Configuration

- Dynamic Allocation
  - `spark.dynamicAllocation.enabled`
  - `spark.dynamicAllocation.initialExecutors`
  - `spark.dynamicAllocation.maxExecutors`
  - `spark.dynamicAllocation.minExecutors`

# Configuration

- Dynamic Allocation

- `spark.dynamicAllocation.schedulerBacklogTimeout`
- `spark.dynamicAllocation.executorIdleTimeout`
- `spark.dynamicAllocation.sustainedSchedulerBacklogTimeout`

# Configuration

- External Shuffle Service
  - `spark.shuffle.service.enabled`
  - `spark.shuffle.service.port`



# Demo

## Dynamic Allocation in Action

# Configuration values?

# Configuration Values?

It depends ....

No, seriously

# Configuration Values

- `spark.dynamicAllocation.initialExecutors`
- `spark.dynamicAllocation.maxExecutors`
- `spark.dynamicAllocation.minExecutors`

Depends on workload and how many resources are potentially available to you.

# Configuration Values

- `spark.dynamicAllocation.schedulerBacklogTimeout`

Too short, might trigger for short burst of tasks.

Too long, might be less effective.

- `spark.dynamicAllocation.sustainedSchedulerBacklogTimeout`

Executor start duration.

Default set to `schedulerBacklogTimeout`.

# Configuration Values

- `spark.dynamicAllocation.executorIdleTimeout`

Relative to the duration of the longer task.

No big drawback on being too long, except cost.



# Spark Streaming

# Spark Streaming



<http://spark.apache.org/docs/latest/streaming-programming-guide.htm>

# Spark Streaming

- In most case, `schedulerBacklogTimeout` longer than batch interval.
- `executorIdleTimeout` a portion of batch interval.
- Should allow to manage processing delay.
- Not compatible with the dynamic rate estimator.

# More Dynamic?

# More Dynamic?

<https://github.com/twosigma/Cook>

‘Fair’ job scheduler for Spark on top of Mesos

- Not a recommendation, just a suggestion.
- Some assembly required.

# THANK YOU.

[github.com/skyluc/tree/master/talks/sparksummit-eu-2016](https://github.com/skyluc/tree/master/talks/sparksummit-eu-2016)





# External Shuffle Service

