

SPARK AND COUCHBASE

AUGMENTING THE OPERATIONAL DATABASE WITH
SPARK

Michael Nitschinger
Couchbase



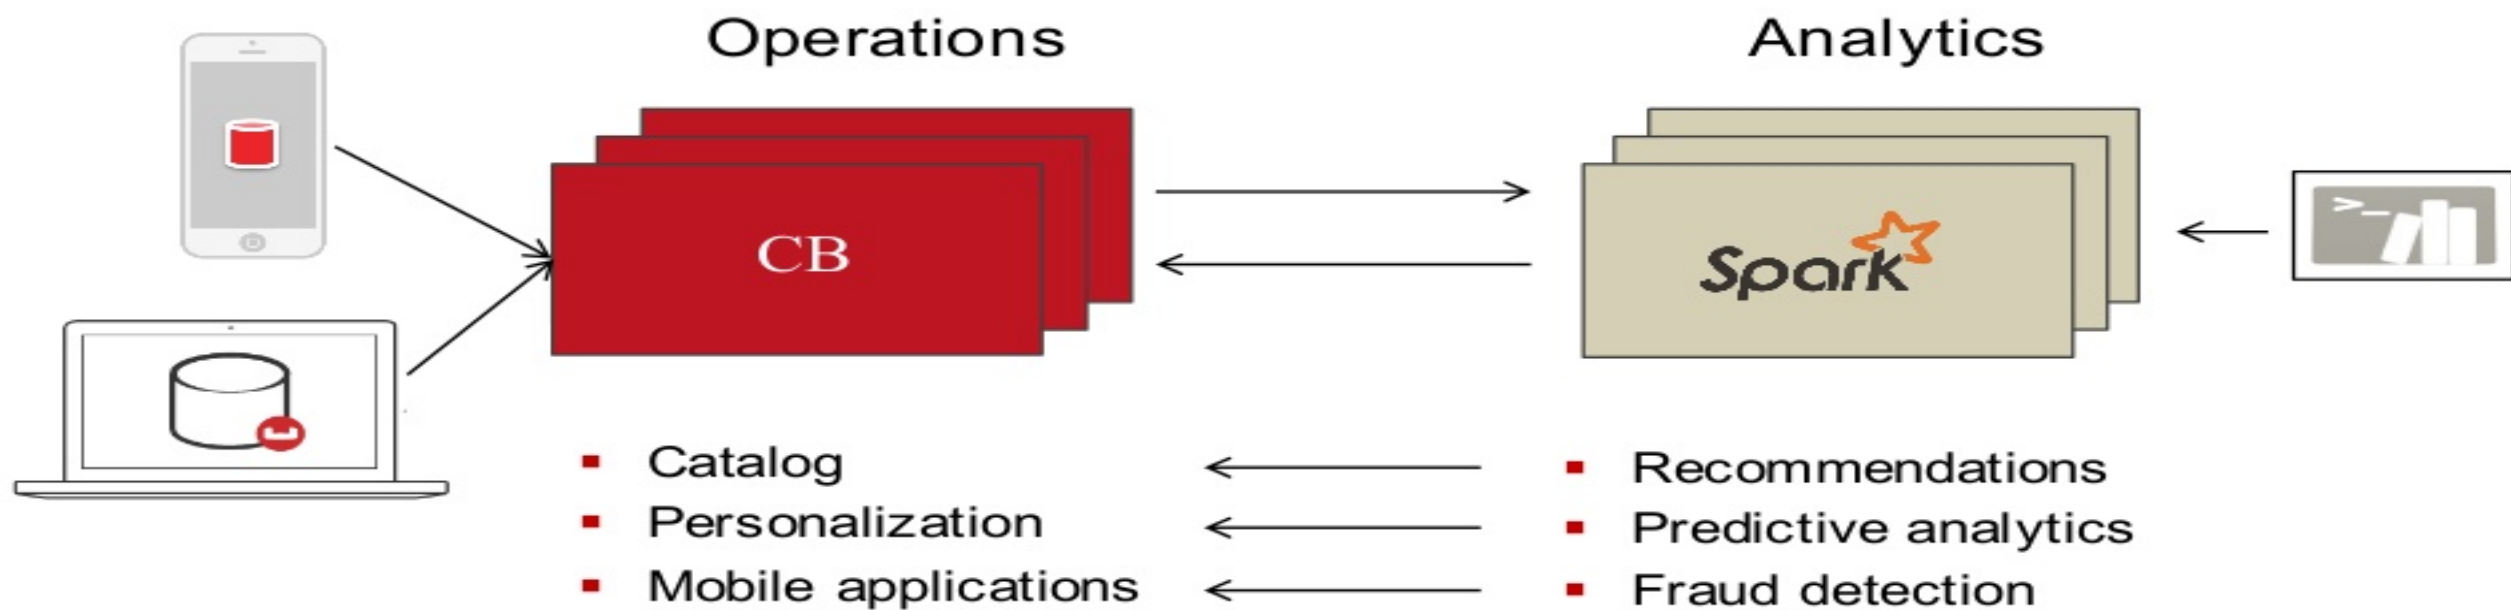


Couchbase

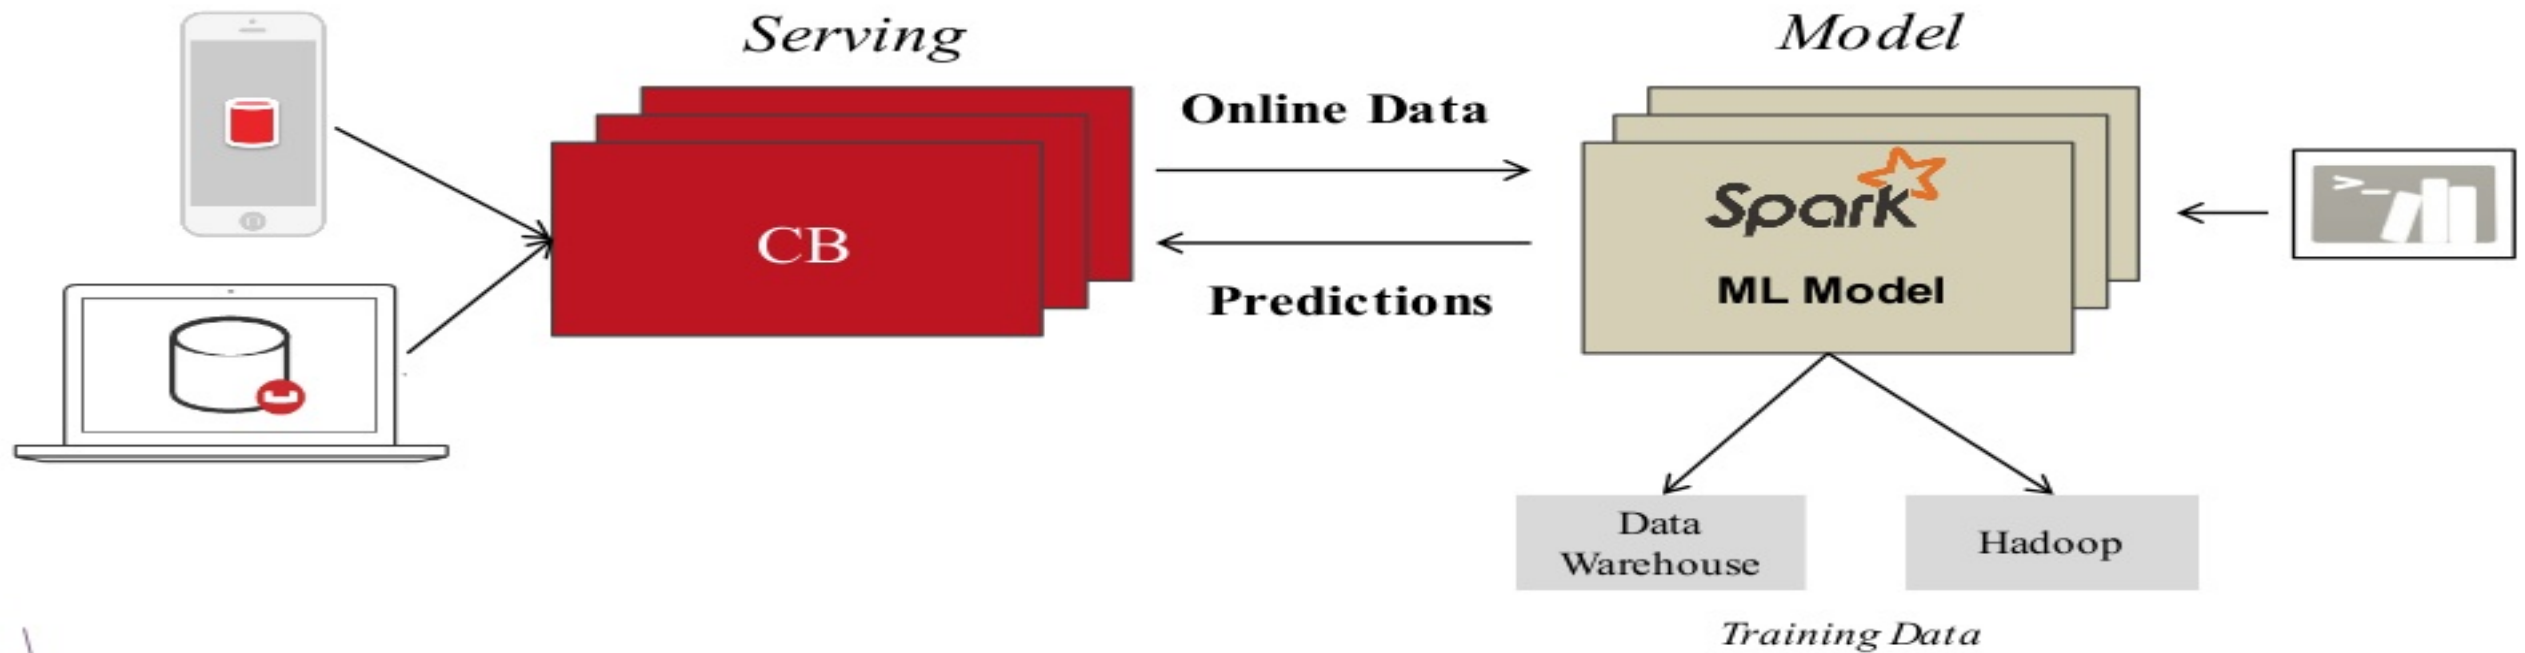
Overview & Use-Cases

WHY SPARK AND COUCHBASE

Use Cases




Use Case: Operationalize Analytics / ML






Use Case: Data Integration



Standalone Deployment



Side-By-Side Deployment



From Spark to Couchbase and Back Again

ACCESS PATTERNS

Key-Value

Fetch/Store by
Document ID

Key-Value

Fetch/Store by
Document ID

N1QL Query

Fetch by
Criteria "SQL"

Key-Value

Fetch/Store by
Document ID

N1QL Query

Fetch by
Criteria "SQL"

Map-Reduce Views

Materialized Indexes
(Aggregation)

Key-Value

Fetch/Store by
Document ID

N1QL Query

Fetch by
Criteria "SQL"

Map-Reduce Views

Materialized Indexes
(Aggregation)

Streaming

Mutation Streams
For Processing

Key-Value

Fetch/Store by
Document ID

N1QL Query

Fetch by
Criteria "SQL"

Map-Reduce Views

Materialized Indexes
(Aggregation)

Streaming

Mutation Streams
For Processing

Full Text

Search on Freeform
Text

Key-Value

RDD

N1QL Query

RDD &
Spark SQL


Map-Reduce
View

RDD

Stream

DStream &
Structured
Streaming

Couchbase Data Partitioning



Data Locality

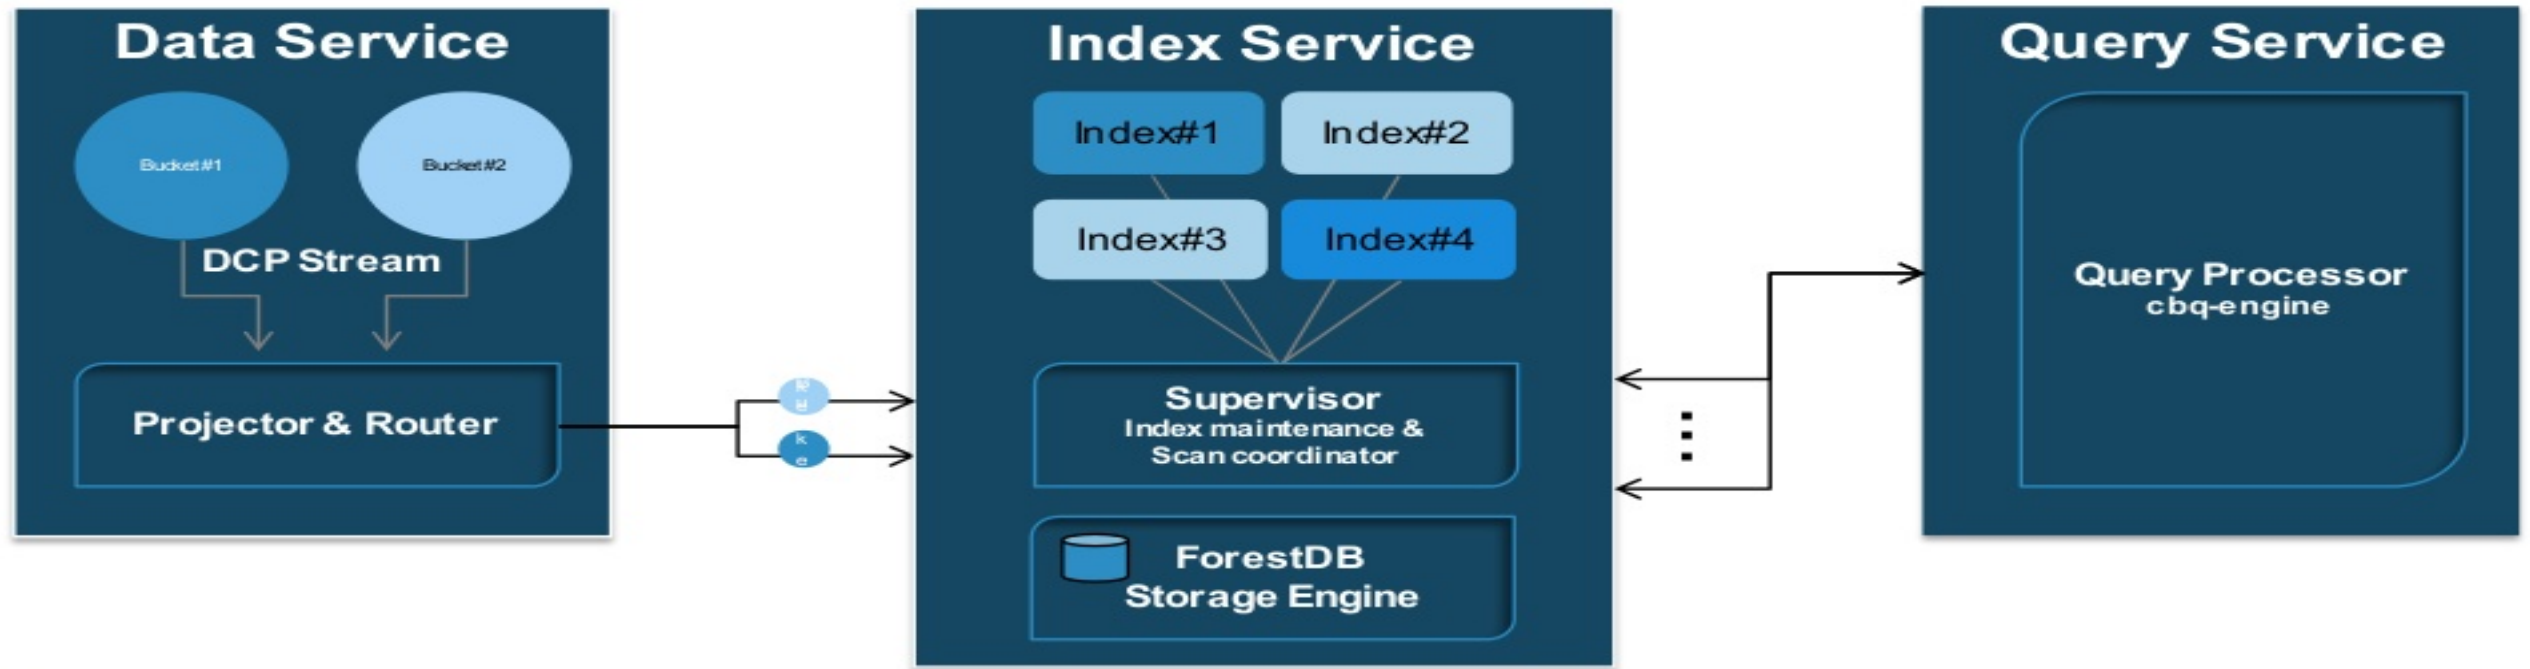
- RDD Location Hints based on the Cluster Map
- Not available for N1QL or Views
 - Round robin - can't give location hints
 - Back end is scatter gather with 1 node responding

N1QL Query

- N1QL is a **SQL** service with **JSON** extensions
- Uses Couchbase's Global Secondary Indexes
- Can run on any nodes within the cluster
- Nodes with differing services can be added and removed as needed on the fly



Couchbase Query Architecture



Spark SQL Sources

TableScan

Scan all of the data and return it

PrunedScan

Scan an index that matches only relevant data to the query at hand.

PrunedFilteredScan

Scan an index that matches only relevant data to the query at hand.

Predicate Conversion

```
filter match {  
  case EqualTo(attr, value) => s" ${attrToFilter(attr)} = " + valueToFilter(value)  
  case GreaterThan(attr, value) => s" ${attrToFilter(attr)} > " + valueToFilter(value)  
  case GreaterThanOrEqualTo(attr, value) => s" ${attrToFilter(attr)} >= " + valueToFilter(value)  
  case LessThan(attr, value) => s" ${attrToFilter(attr)} < " + valueToFilter(value)  
  case LessThanOrEqualTo(attr, value) => s" ${attrToFilter(attr)} <= " + valueToFilter(value)  
  case IsNull(attr) => s" ${attrToFilter(attr)} IS NULL"  
  case IsNotNull(attr) => s" ${attrToFilter(attr)} IS NOT NULL"  
  case StringContains(attr, value) => s" CONTAINS(${attrToFilter(attr)}, '$value')"  
  case StringStartsWith(attr, value) =>  
    s" ${attrToFilter(attr)} LIKE ' " + escapeForLike(value) + "%' "
```


Schema Inference

```
val spark = SparkSession
  .builder()
  .config("spark.couchbase.nodes", "127.0.0.1")
  .config("spark.couchbase.bucket.travel-sample", "")
  .getOrCreate()

val airlines = spark.read.couchbase(EqualTo("type", "airline"))

airlines
  .select("name", "callsign")
  .sort(airlines("callsign").desc)
  .show(10)
```

Schema Inference


N1QLRelation:28 - Inferring schema from bucket travel-sample with query `'SELECT META(`travel-sample`).id as `META_ID`, `travel-sample`.* FROM `travel-sample` WHERE `type` = 'airline' LIMIT 1000'`

N1QLRelation:28 - Executing generated query: `'SELECT `name`,`callsign` FROM `travel-sample` WHERE `type` = 'airline''`


Schema Inference

```
root
|-- META_ID: string (nullable = true)
|-- callsign: string (nullable = true)
|-- country: string (nullable = true)
|-- iata: string (nullable = true)
|-- icao: string (nullable = true)
|-- id: long (nullable = true)
|-- name: string (nullable = true)
|-- type: string (nullable = true)
```

DCP and Spark Streaming



Structured Streaming Source



(Un)Structured Streaming?

```
val schema = StructType(  
    StructField("META_ID", StringType) ::  
    StructField("type", StringType) ::  
    StructField("name", StringType) :: Nil  
)
```

Structured Streaming Source

```
val records = spark.readStream  
    .format("com.couchbase.spark.sql")  
    .schema(schema)  
    .load()
```

Structured Streaming Sink

```
val query = wordCounts.writeStream
  .outputMode("complete")
  .option("checkpointLocation", "mycheckpointlocation")
  .option("idField", "value")
  .format("com.couchbase.spark.sql")
  .start()

query.awaitTermination()
```

Couchbase Spark Connector 1.2.1

- Spark 1.6.x support, including Datasets
- DCP Flow Control
- Enhanced Java APIs

Couchbase Spark Connector 2.0.0

- Spark 2.0.x Support
- Enhanced DCP Client
- Experimental Structured Streaming

Resources

- **Spark Packages** <https://spark-packages.org/package/couchbase/couchbase-spark-connector>
- **Docs** <http://docs.couchbase.com>
- **Source** <https://github.com/couchbase/couchbase-spark-connector>
- **Bugs** <https://issues.couchbase.com/browse/SPARKC>

THANK YOU.

Michael Nitschinger

@daschl

michael.nitschinger@couchbase.com

