

Boosting Spark Performance on Many-Core Machines

Qifan Pu

Sameer Agarwal (Databricks)

Reynold Xin (Databricks)

Ion Stoica



Me

Ph.D. Student in AMPLab, advised by Prof. Ion Stoica

- Spark-related research projects
 - e.g., how to run Spark in a geo-distributed fashion
- Big data storage (e.g., Alluxio)
 - how to do memory management for multiple users

Intern at Databricks in the past summer

- Spark SQL team: aggregates, shuffle

This project

Spark performance on many-core machines

- Ongoing research, feedbacks are welcome
- Focus of this talk:
 - understand shuffle performance
 - Investigate and implement In-memory shuffle

Moving beyond research

- Hope is to get into Spark (but no guarantee yet)

Why do we care about many-core machines?



Spark started as a **cluster** computing framework

Spark on cluster


- Spark's one big success has been high scalability
 - Largest cluster known is 8000
 - Winner of 2014 Daytona GraySort Benchmark
- Typical cluster sizes in practice:

“Our observation is that companies typically experiment with cluster size of under 100 nodes and expand to **200 or more nodes** in the *production* stages.”

-- Susheel Kaushik (Senior Director at Pivotal)

Increasingly powerful single node machines

- More cores packed on single chip
 - Intel Xeon Phi: 64-72 cores
 - Berkeley FireBox Project: ~100 cores
- Larger instances in the Cloud
 - Various 32-core instances on EC2, Azure & GCE
 - EC2 X-Instance with 128 cores, 2TB (May 2016)





Cost of many-core nodes

	<i>Memory (GB)</i>	<i>vCPUs</i>	<i>Hourly Cost (\$)</i>	<i>Cost/100GB</i>	<i>Cost/8vCPU</i>
x1.32xlarge	1952	128	13.338	0.68	0.83
g2.2xlarge	15	8	0.65	4.33	0.65
i2.2xlarge	61	8	1.705	2.80	1.70
m3.2xlarge	30	8	0.532	1.78	0.53
c3.2xlarge	15	8	0.42	2.80	0.42

1 x1.32xlarge instance is a small cluster
(with more memory, fast inter-core network)

Spark's design was based on many nodes

- Data communication (a.k.a. shuffle)  Focus of this talk
 - Store intermediate data on disk
 - Serialization/deserialization needed across nodes
 - Now: much memory to spare, intra-node shuffle
- Resource management  Ongoing work
 - Designed to handle moderate amount on each node
 - Now: 1 executor for 100 cores + 2TB memory?

Can we improve shuffle on single, multi-core machines?


- 1, Memory is fast
- 2, We can use memory for shuffle
- 3, Therefore,
Shuffle will be fast

Will this “common sense” work?

Put in practice...

- Spark: write to a file stream, and save all bytes to **disk**
- Solution: ..., to **memory (bytes on heap)**

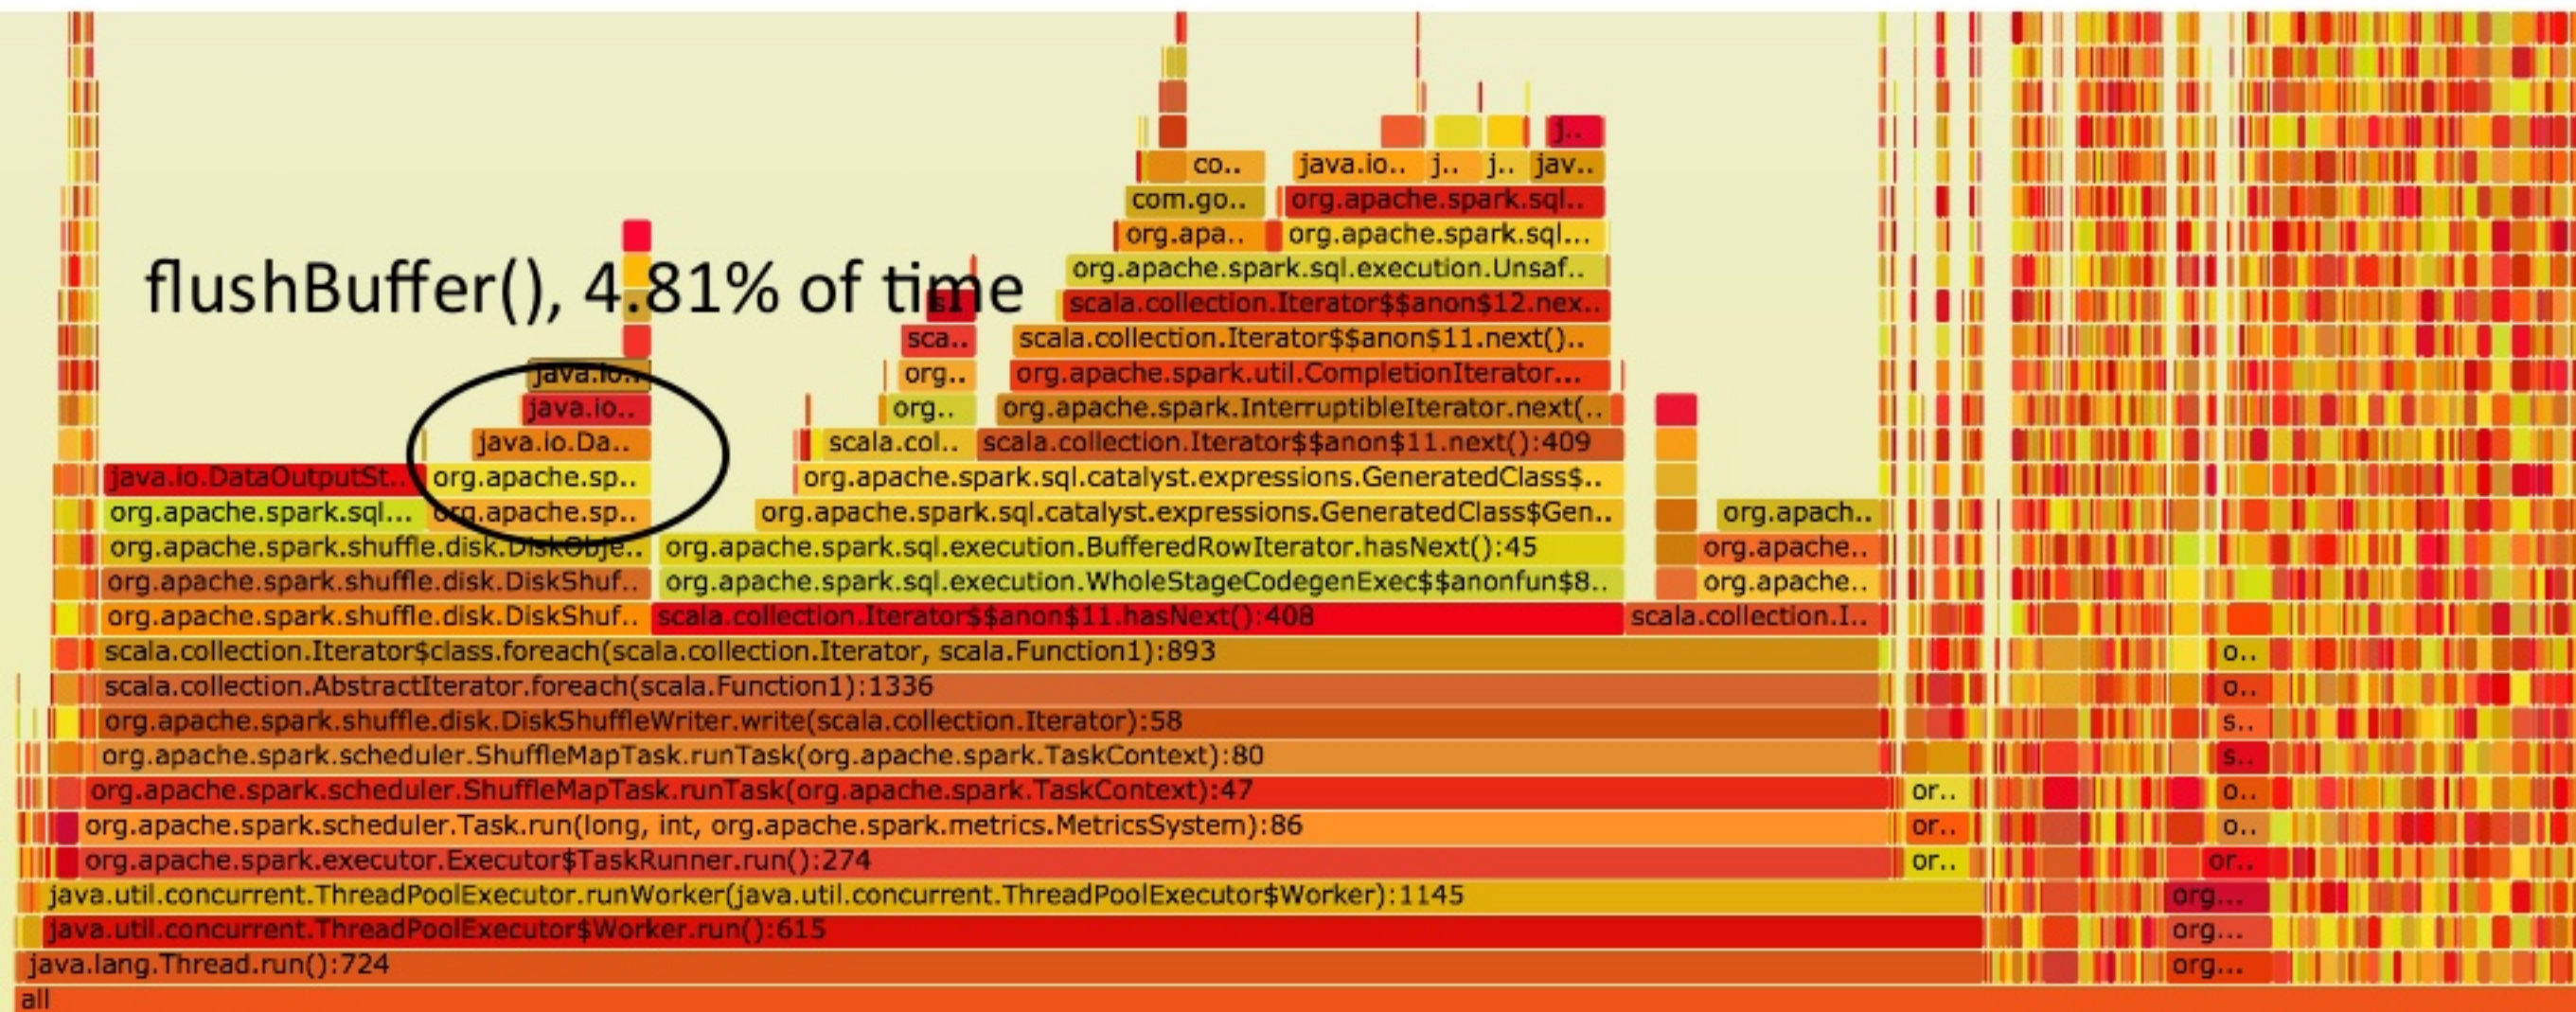
```
spark.range(16M).repartition(1).selectExpr("sum(id)").collect()
```



Why zero improvement?

Why zero improvement?

flushBuffer(), 4.81% of time



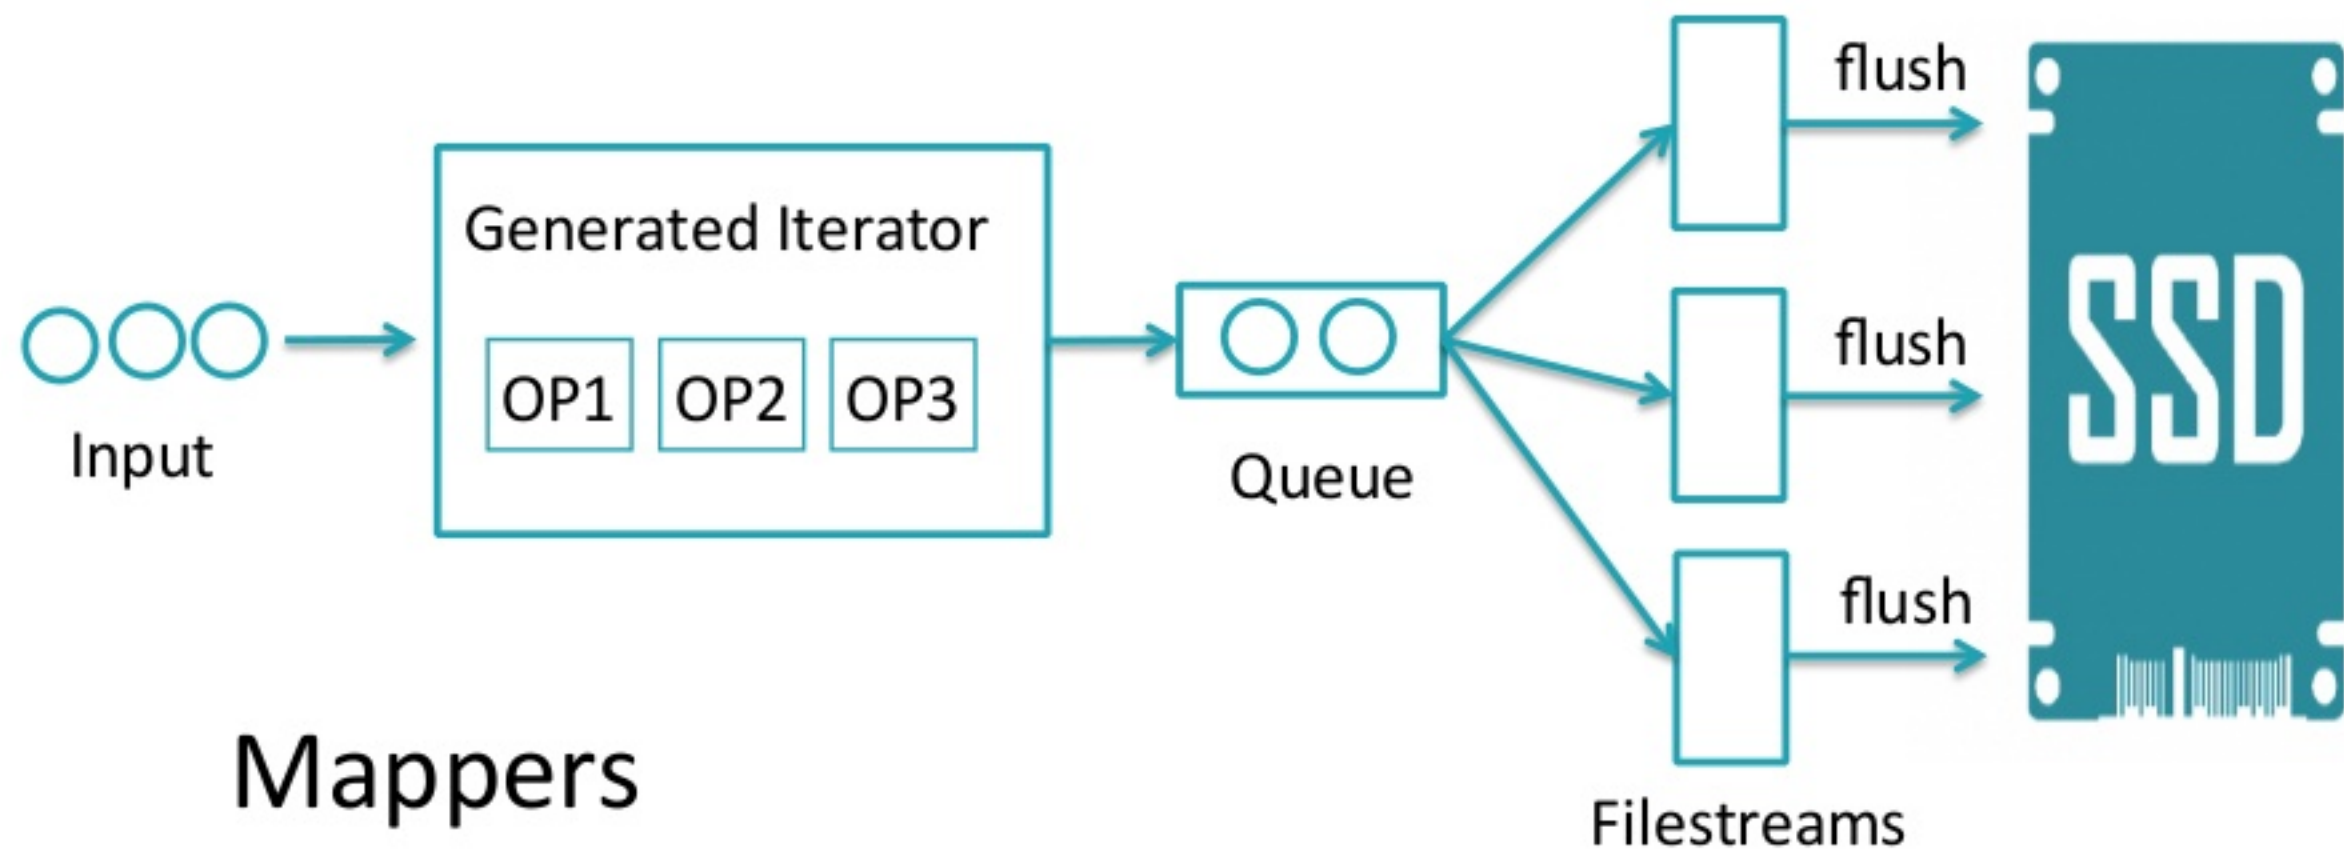
Function: java.io.BufferedOutputStream.flushBuffer():82 (44 samples, 4.81%)

Entire duration of a job (100%)

Why zero improvement?

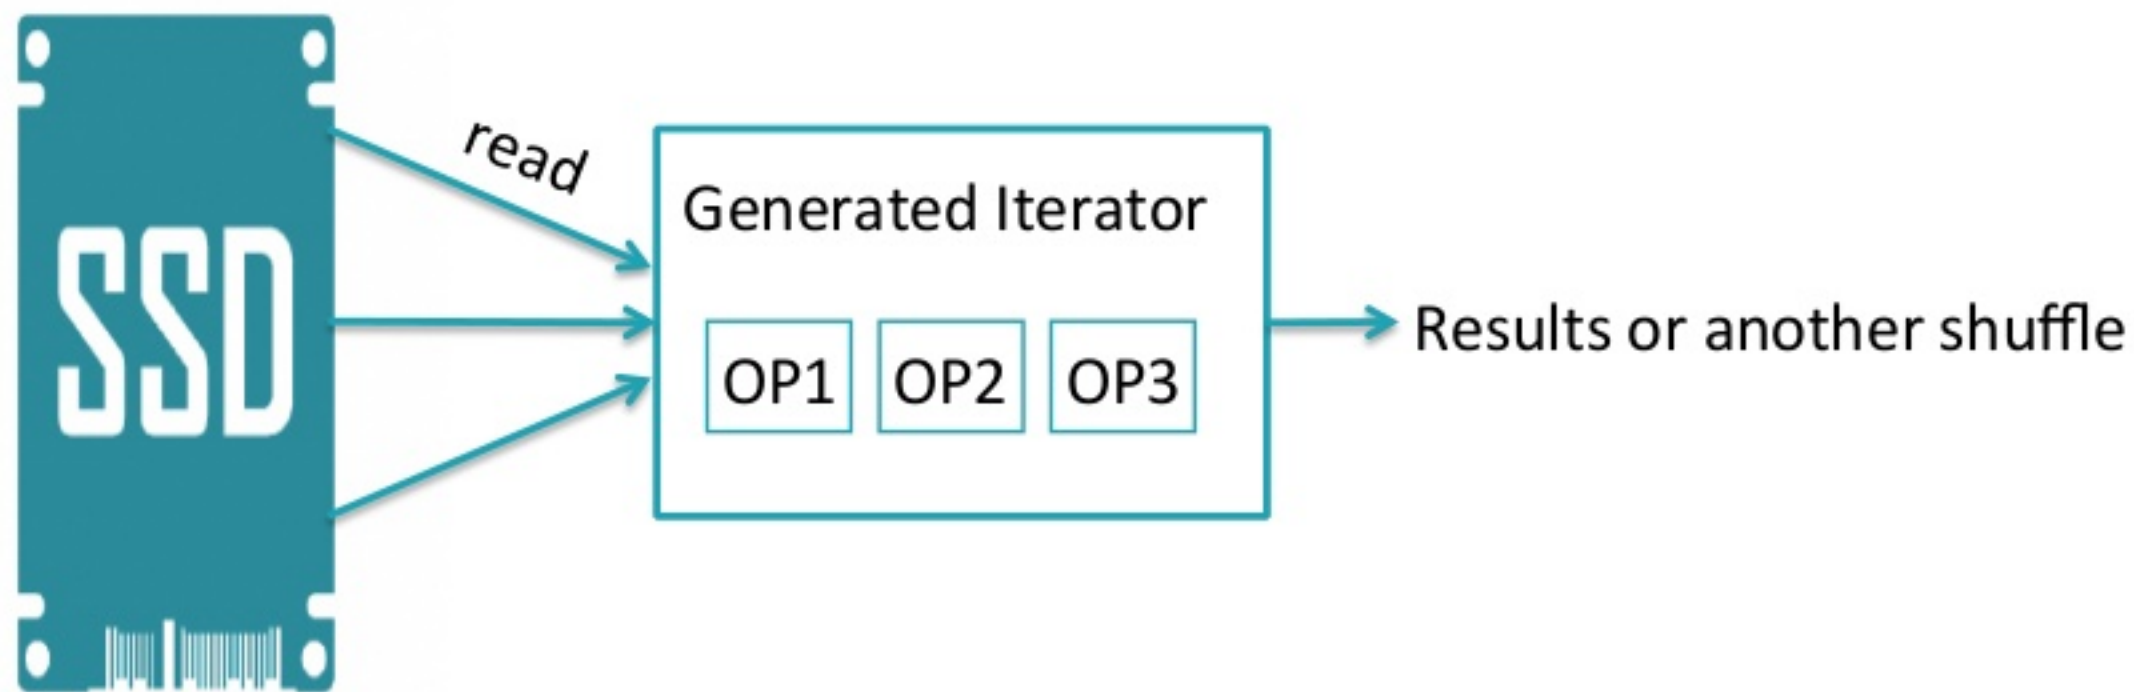
- 1, I/O throughput is not the bottleneck (in this job)
- 2, Buffer cache:
memory is being exploited by disk-based shuffle

Understanding Shuffle Performance



Understanding Shuffle Performance

Reducers




More complications

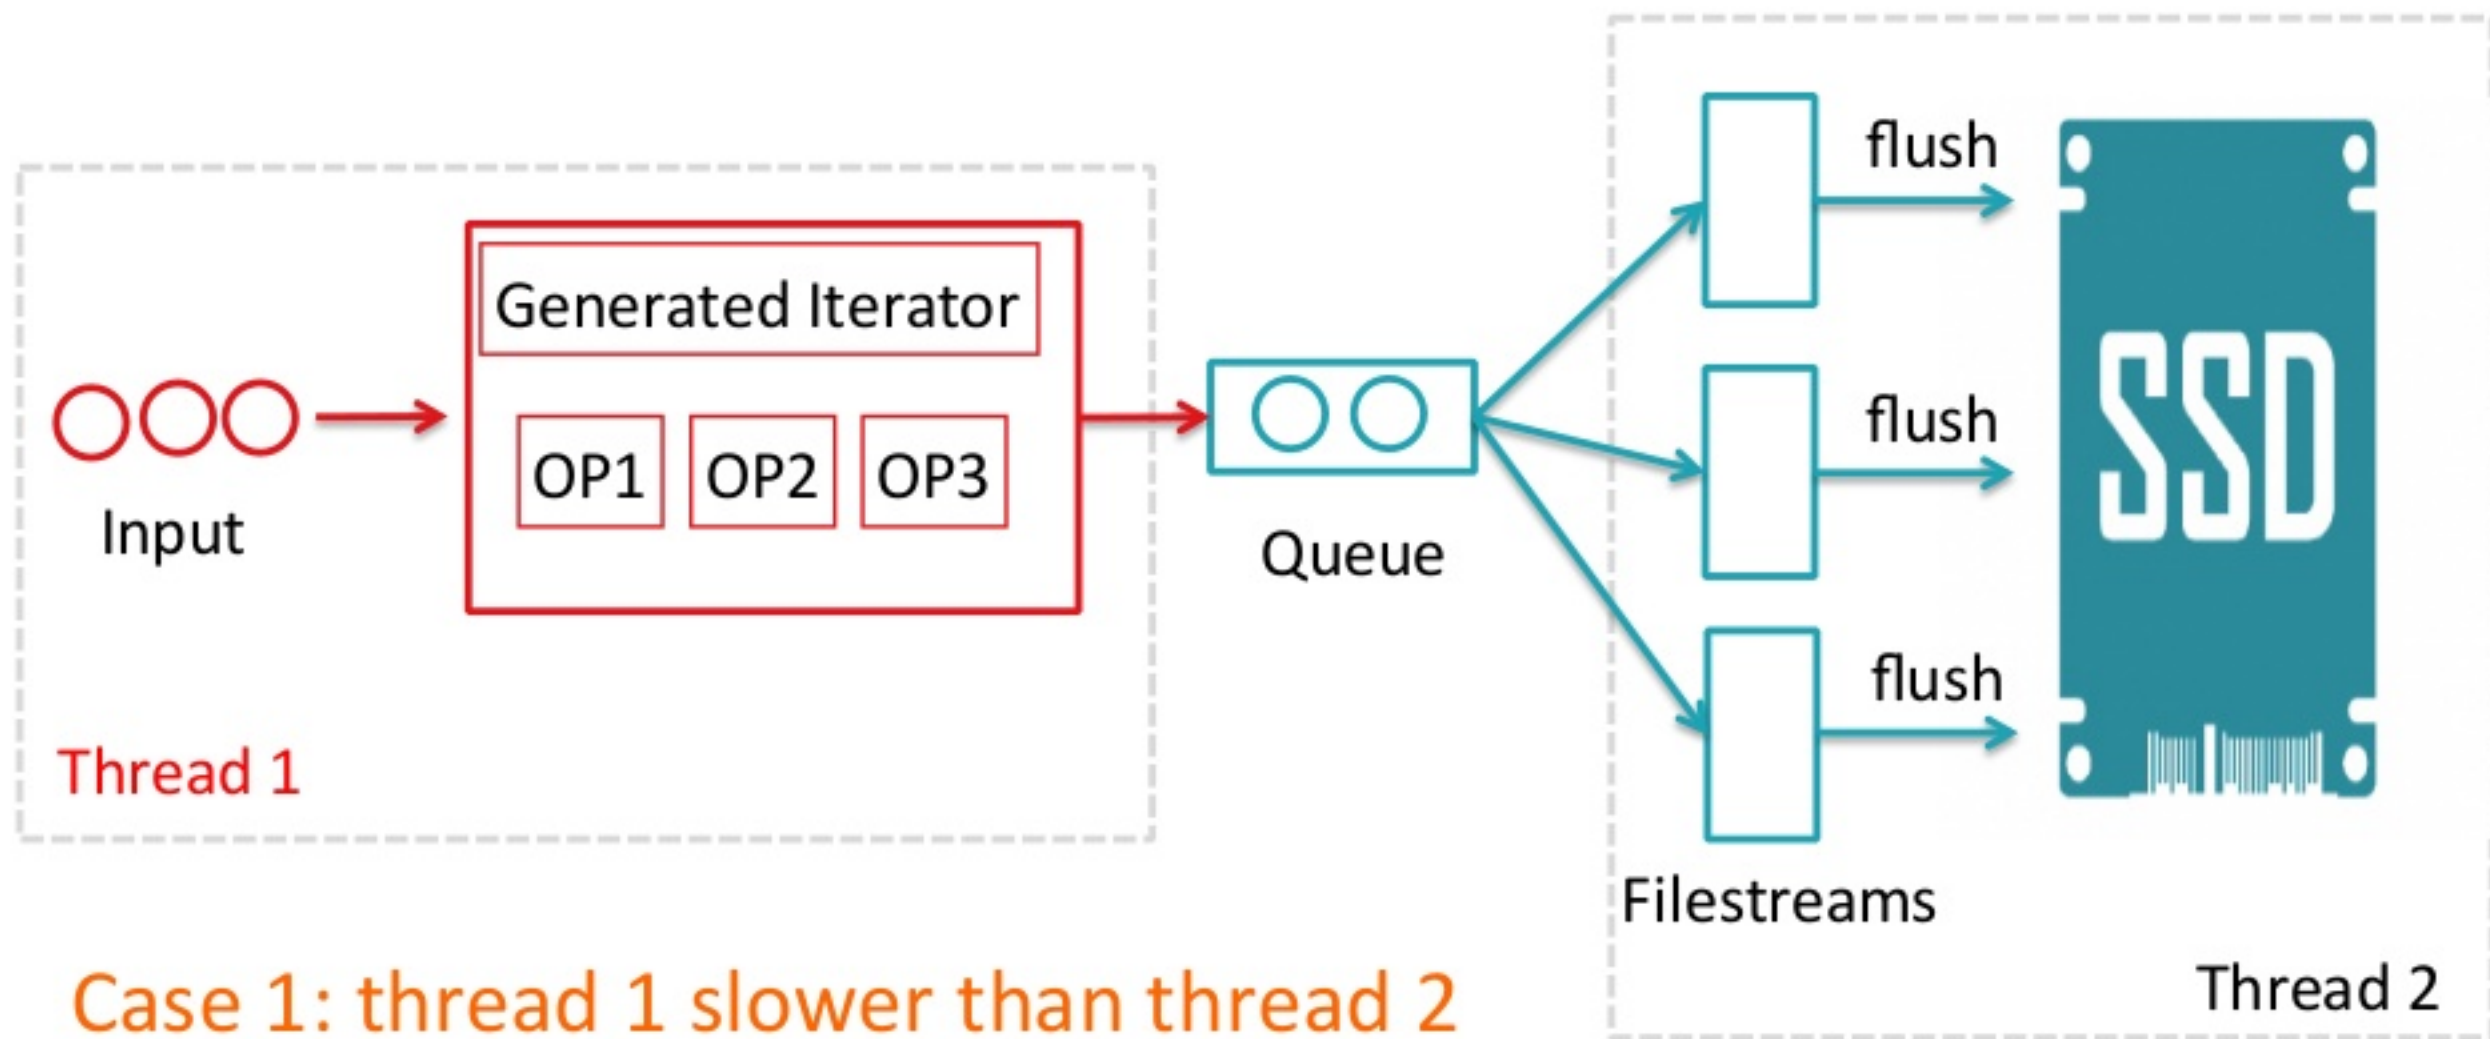
- Sort vs. hash-based shuffle
- Spill when memory runs out
- Clear data after shuffle

...


Case 1



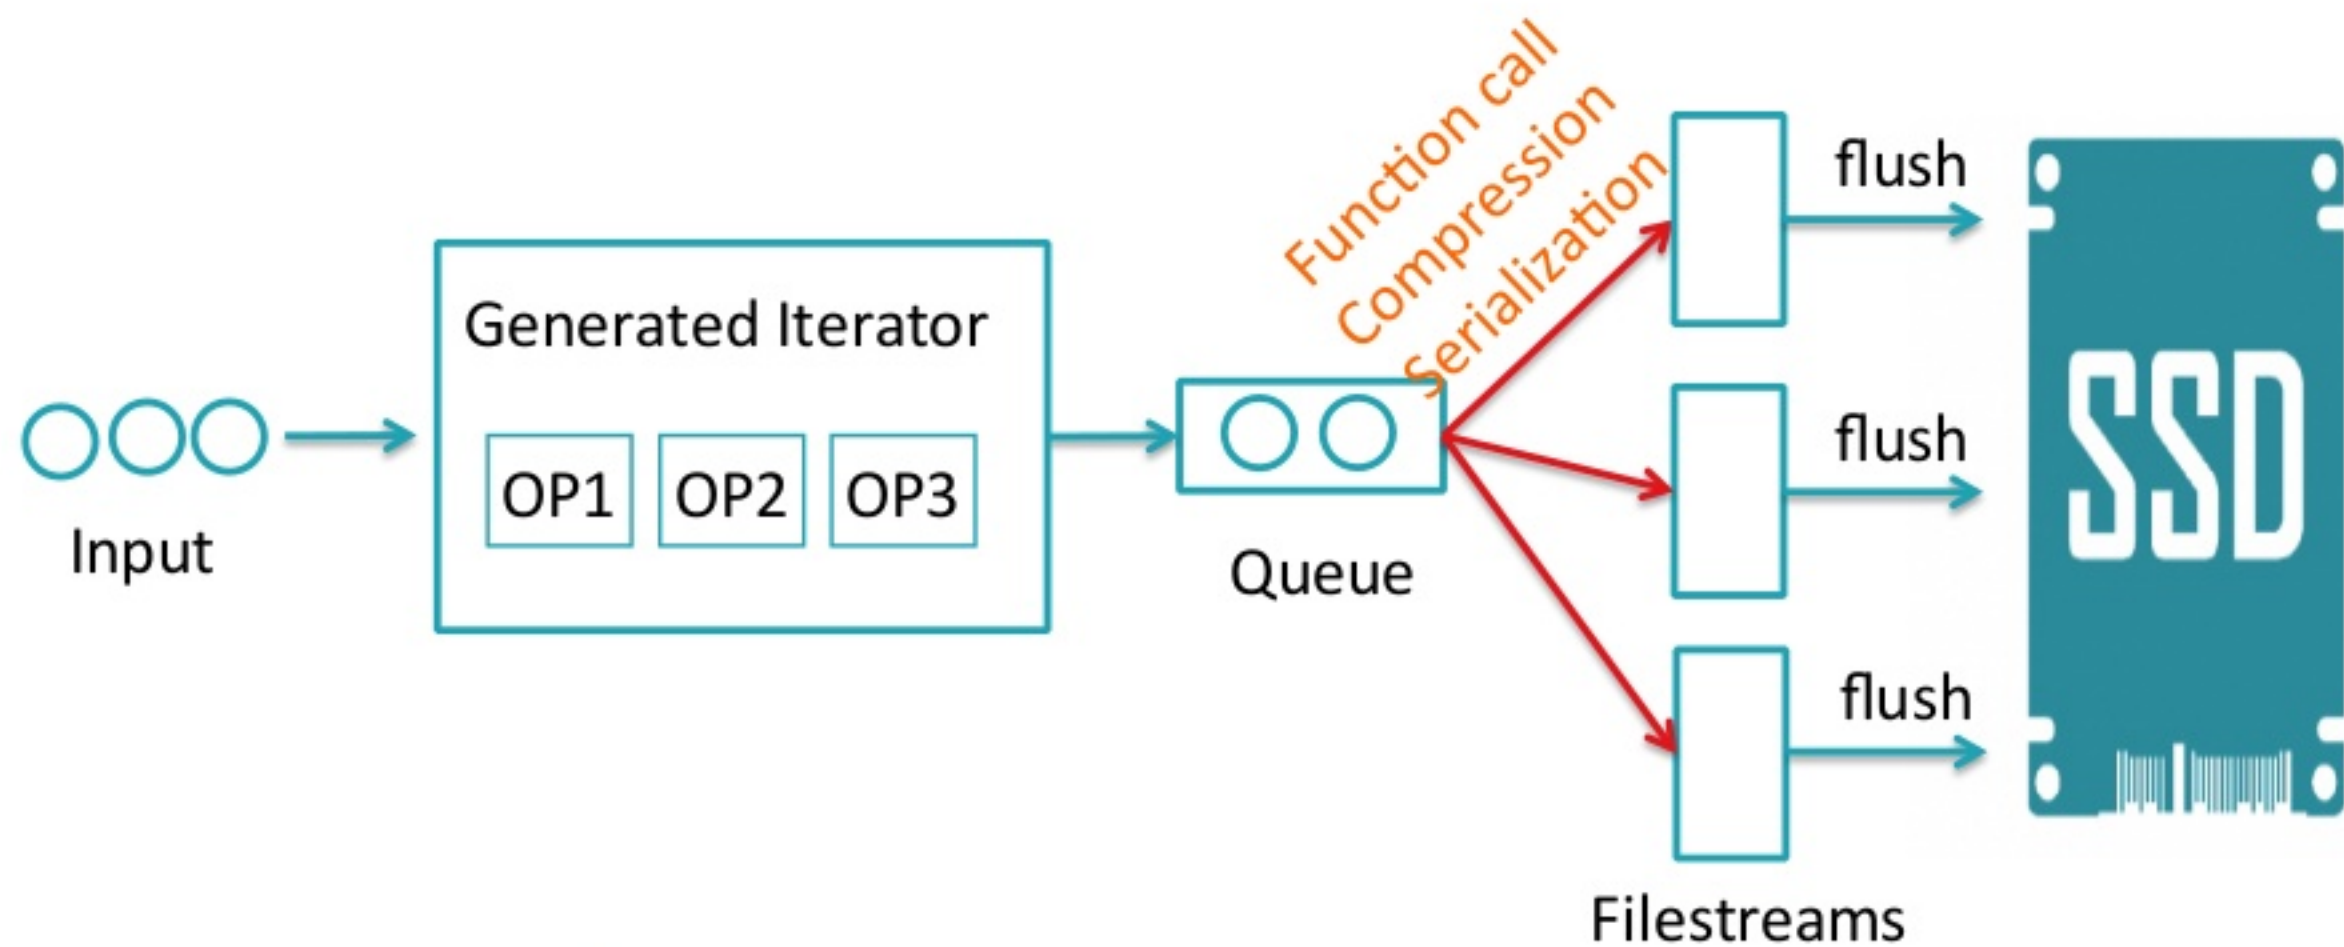
Case 1



Case 2



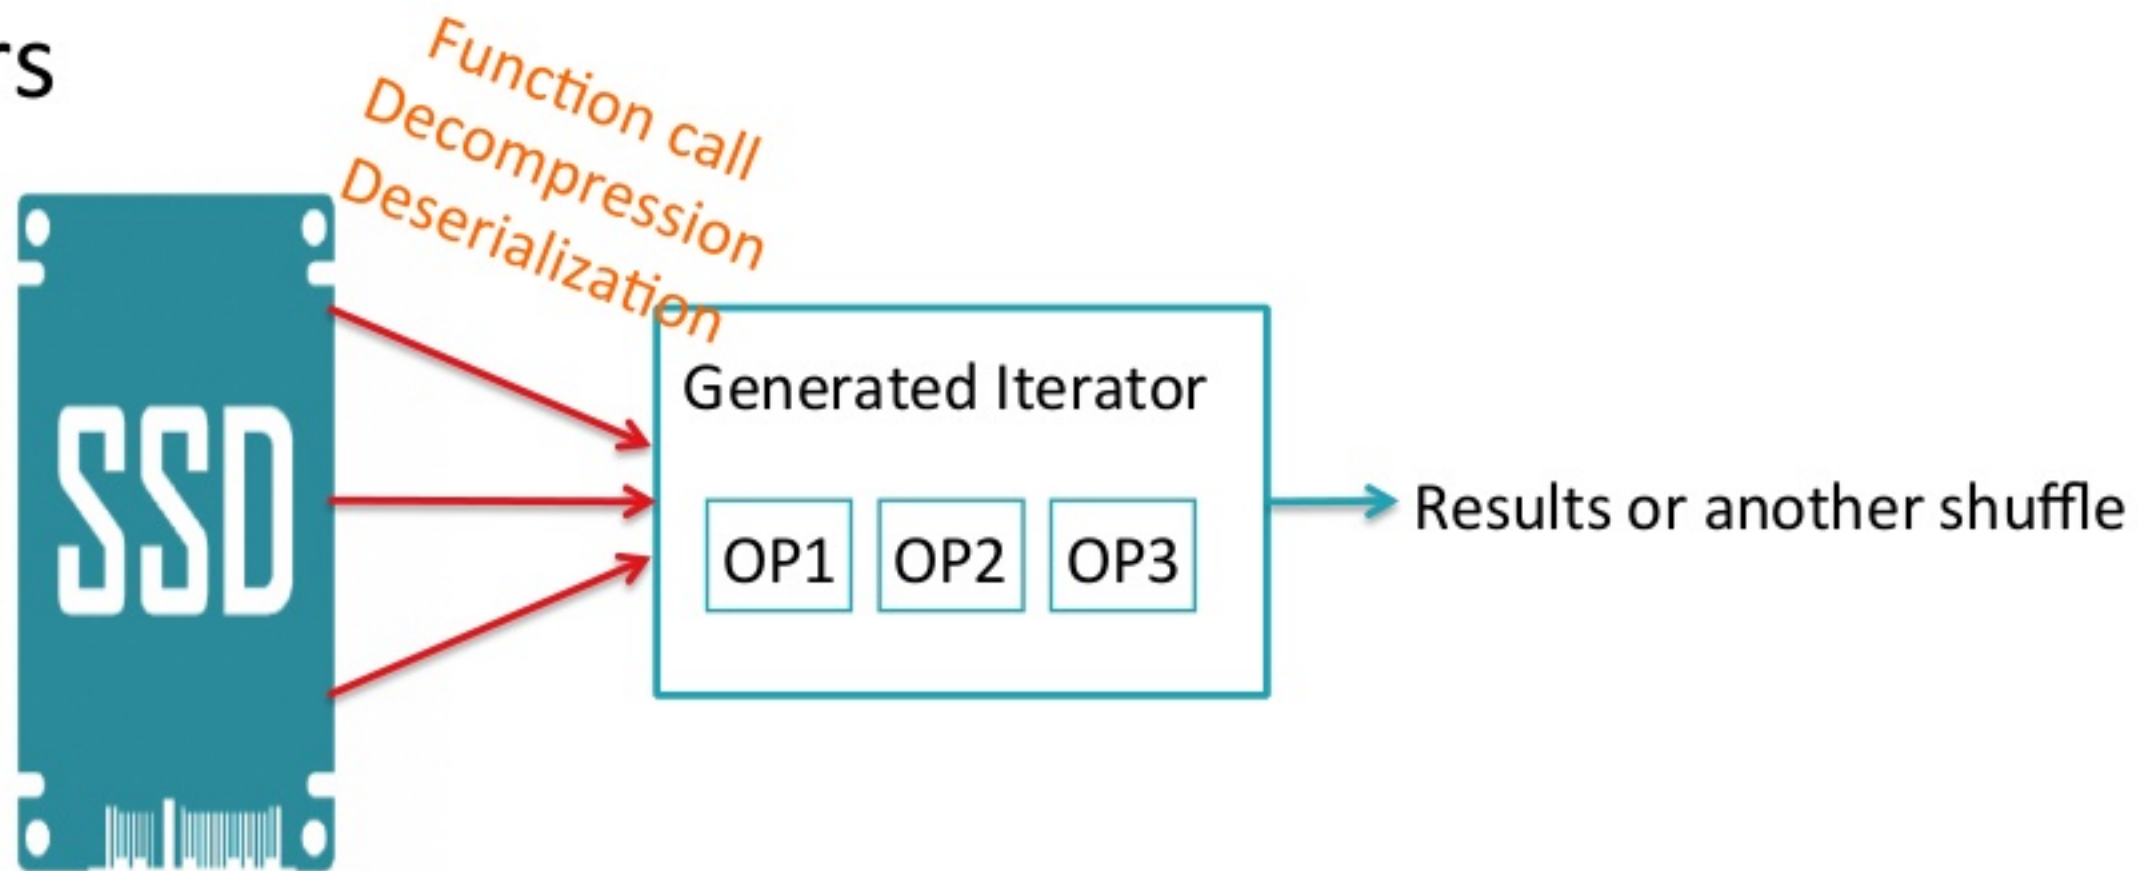
Case 2



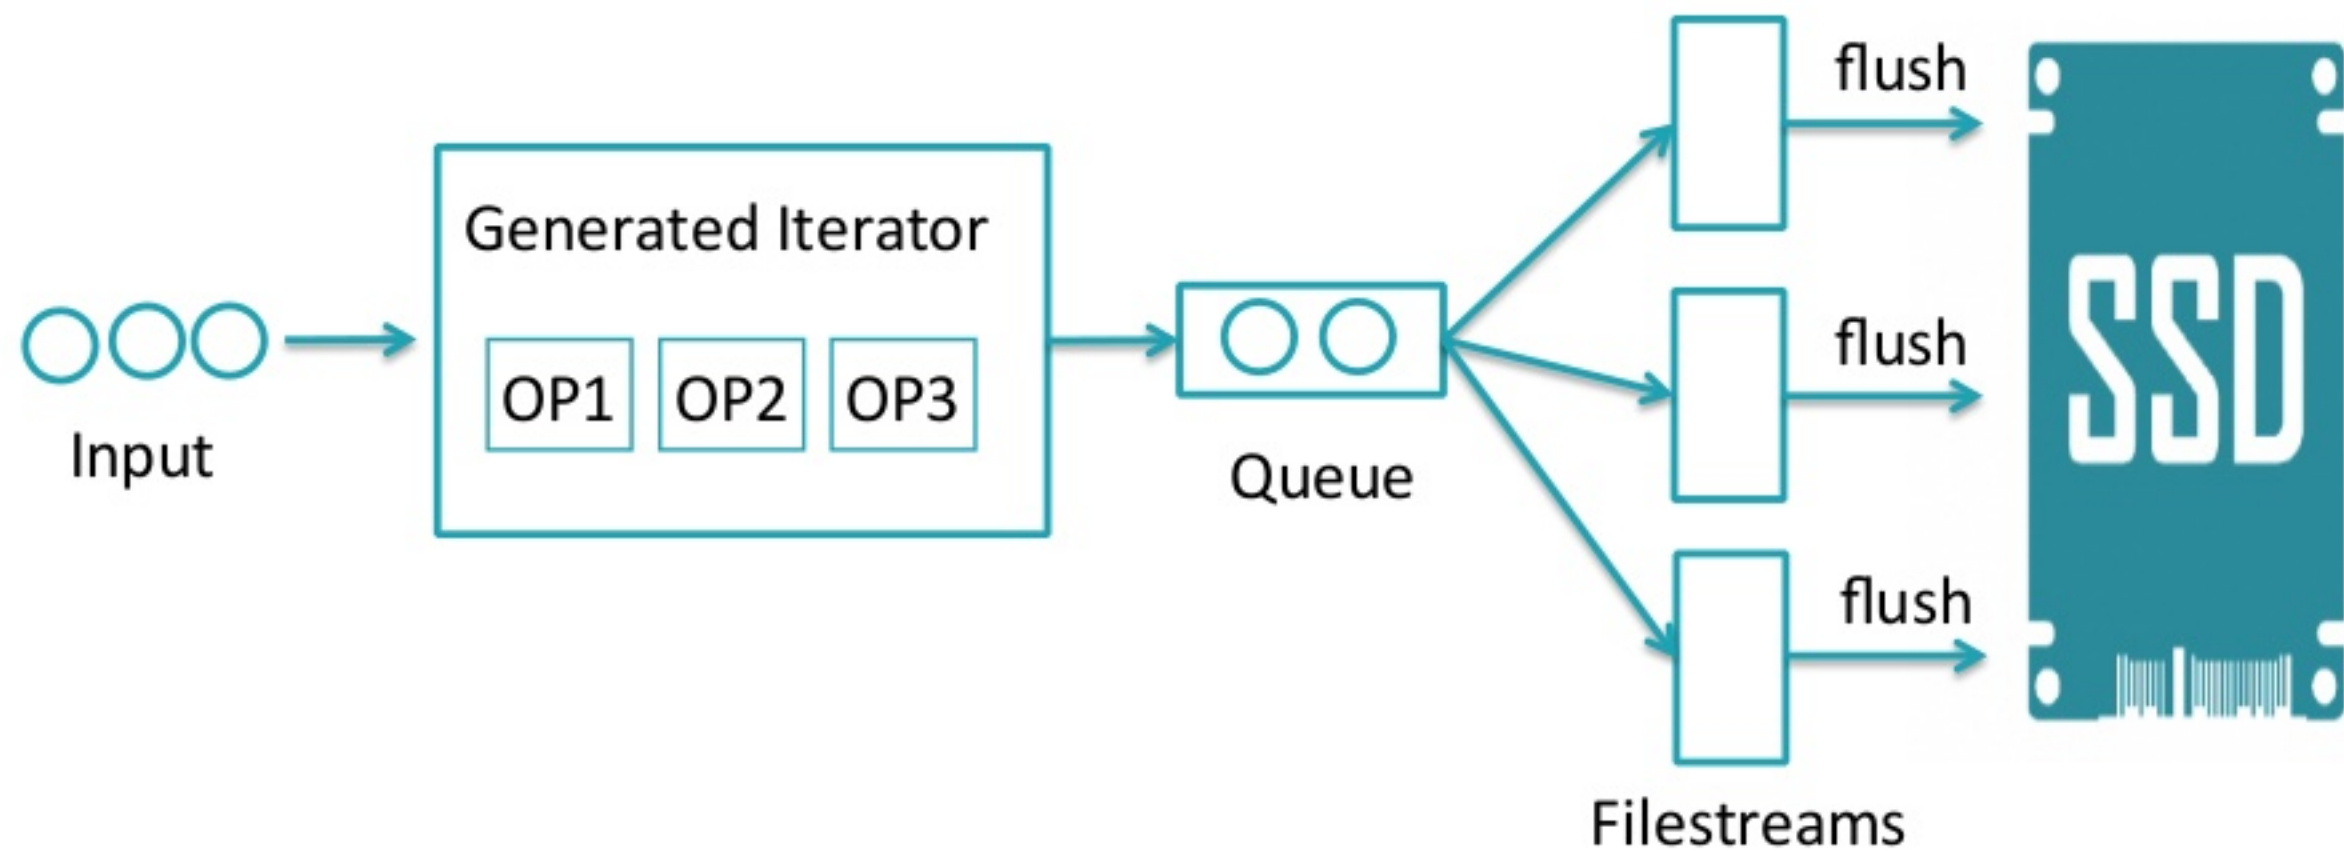
Case 2: writing/reading file streams is slow

Case study 2

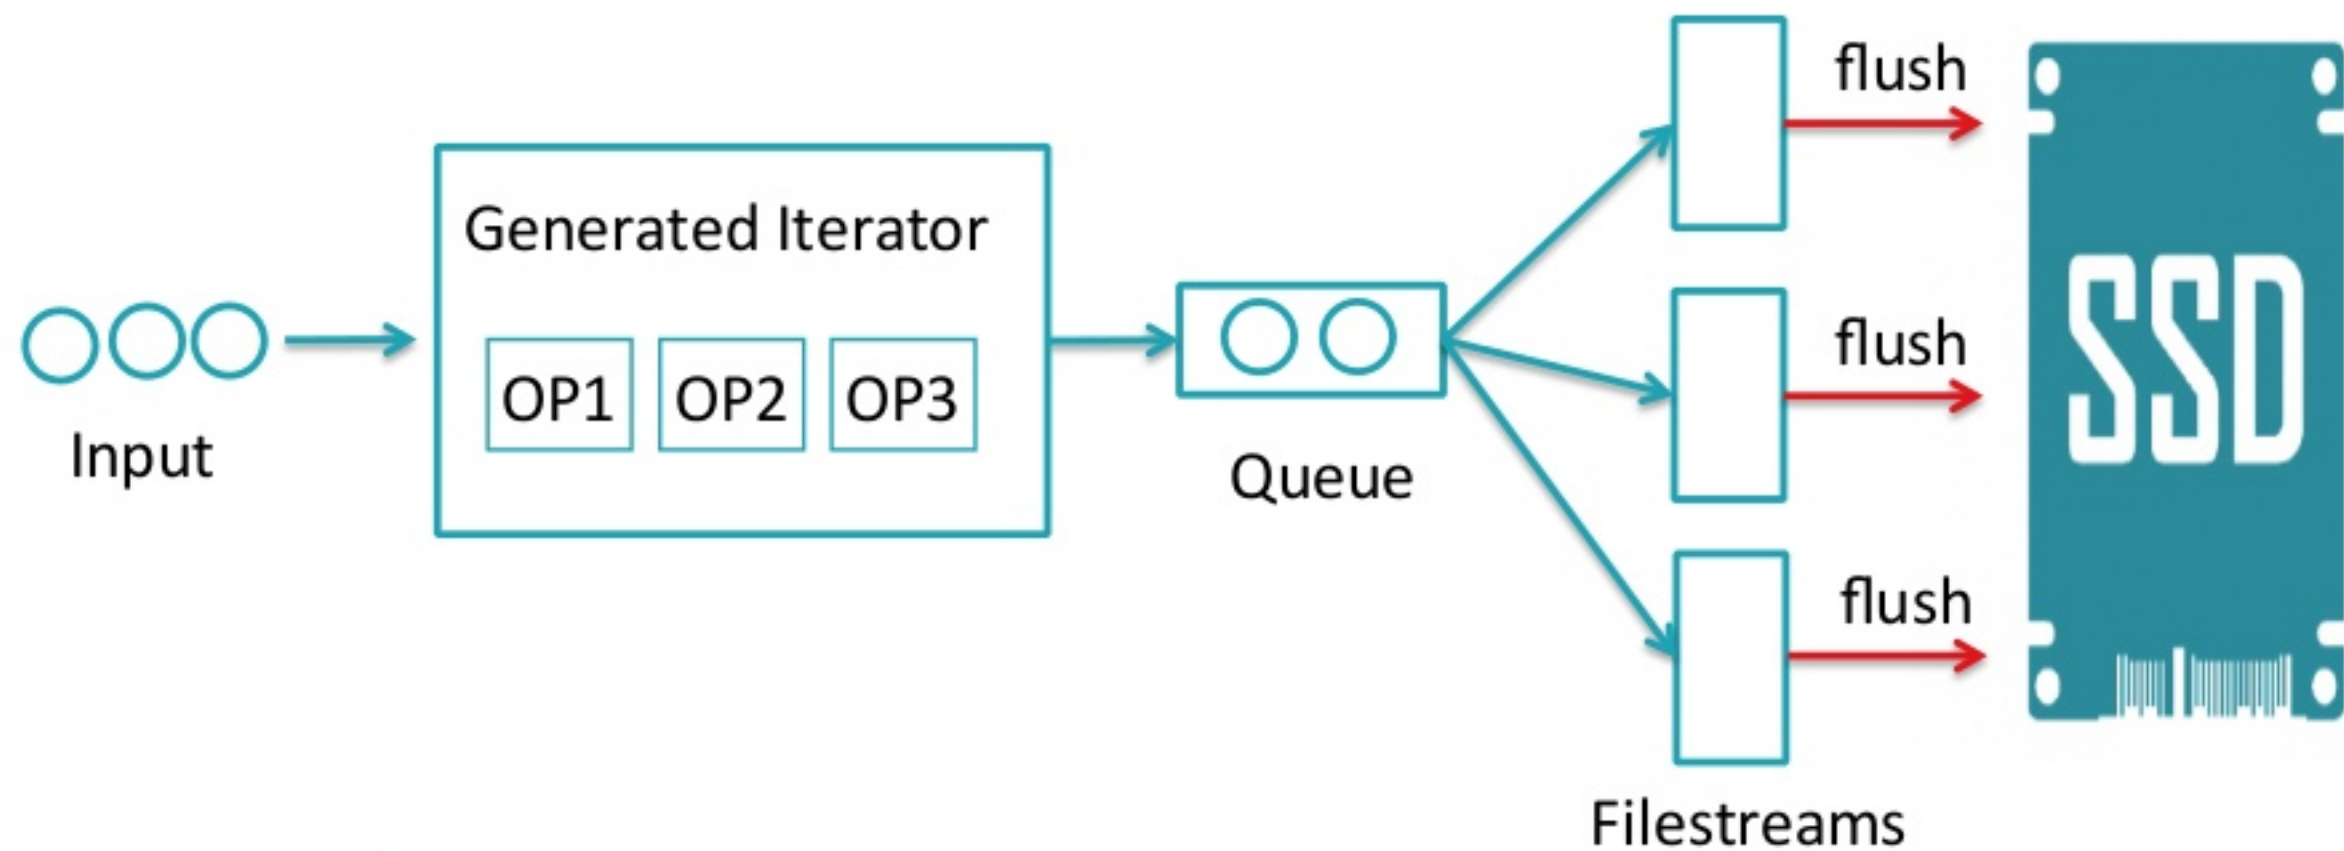
Reducers



Case 3



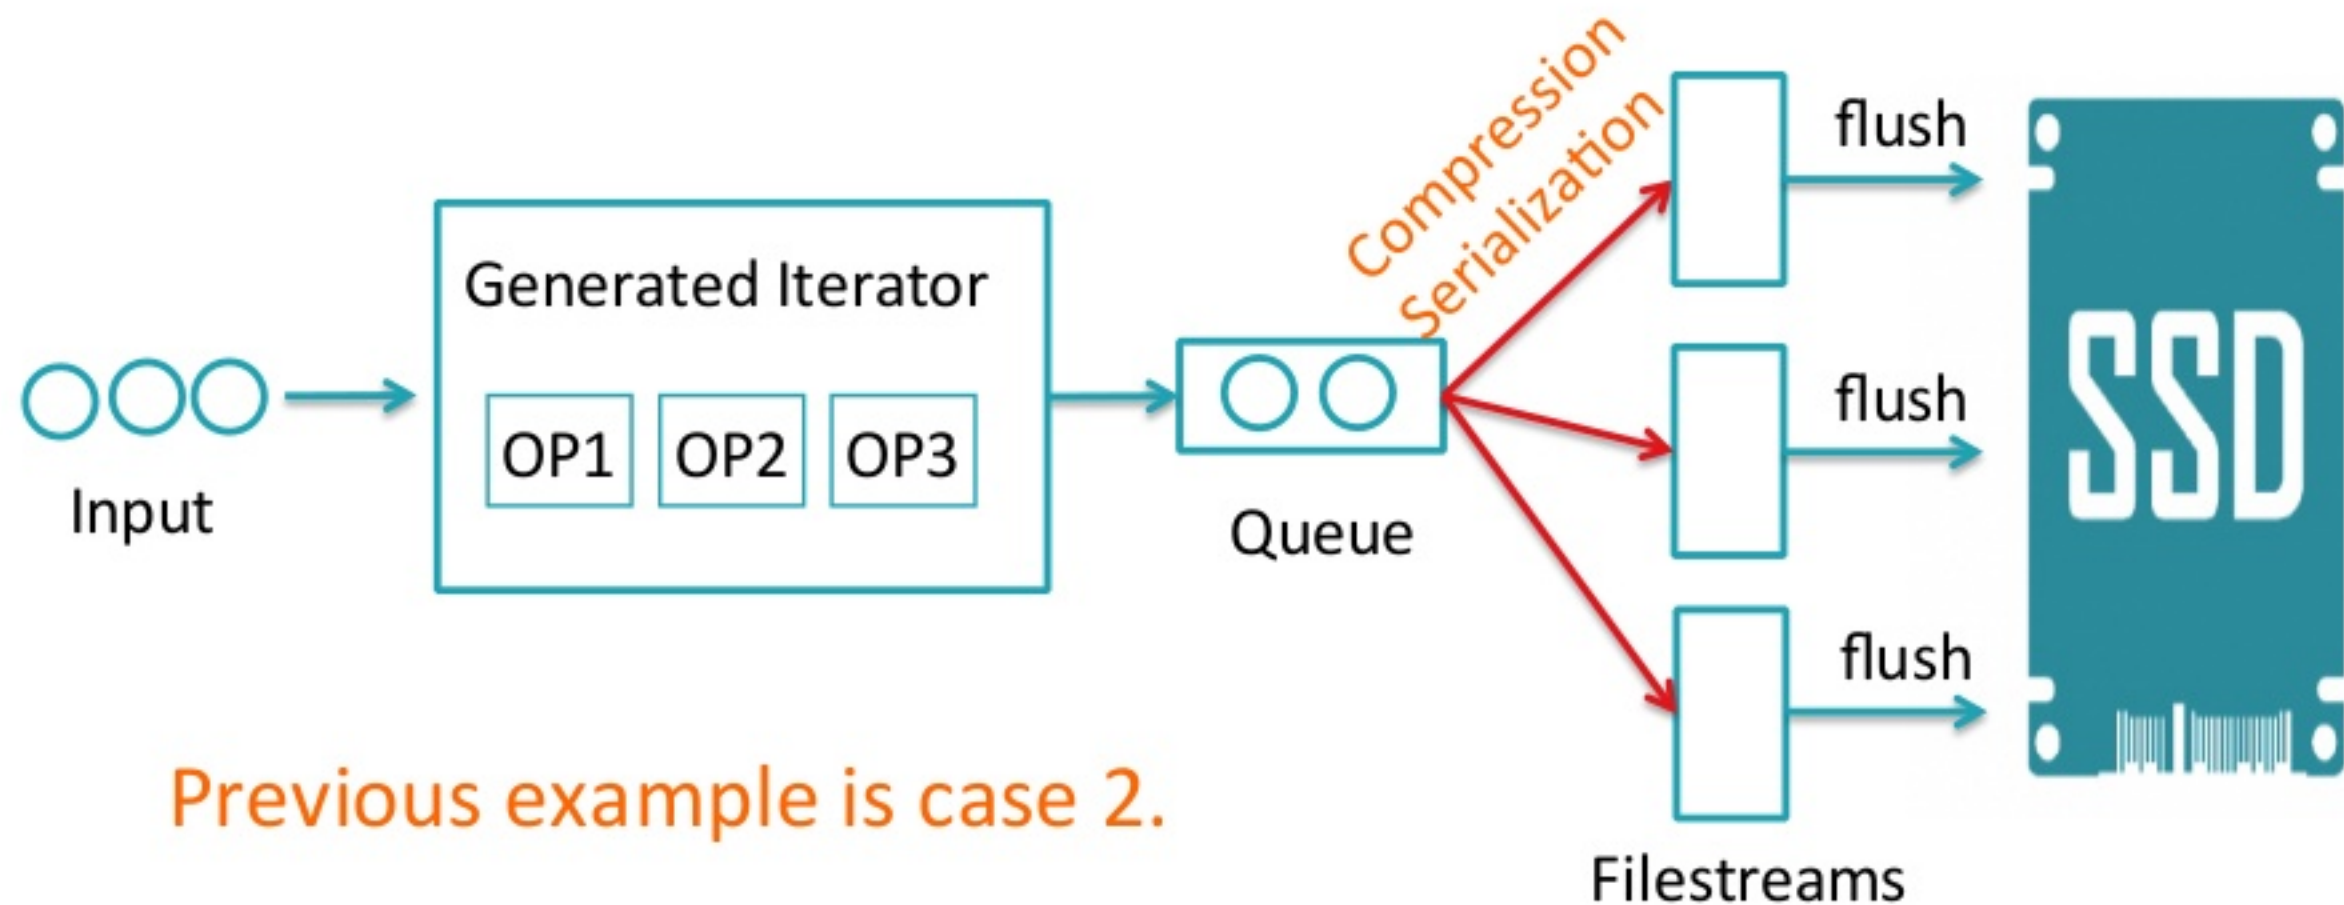
Case 3



Case 3: I/O contentions


Our first attempt should work in this case!

Can we improve case 2?



Attempt 2: get rid of ser/deser, etc


- Attempt 2: create $N \times M$ queues (N =mappers, M =reducers)
push corresponding records into queues



- No serialization
- No copy (data structure shared by both sides)


Attempt 2: get rid of ser/deser, etc

- Attempt 2: create $N \times M$ queues (N =mappers, M =reducers)
push corresponding records into queues




Attempt 2: get rid of ser/deser, etc

- Attempt 2: create NxM queues (N=mappers, M=reducers)
push corresponding records into queues



Attempt 3: avoiding GC

- Attempt 3: instead of queue, copy records to memory pages



Number of objects: \sim records \rightarrow \sim pages

NxM pages, or alternatively, one page per reducer

Spark SQL

- Unsafe Row (a buffer-backed row format):
- `row.pointTo(buffer)`


Instantaneous creation of unsafe rows by pointing to different offsets in the page




record1	recor..
..d2	record3
...	
...	
...	

Attempt 3: avoiding GC

- Attempt 3: copy records onto large memory pages




Consistent improvement with varying size



```
spark.range(N).repartition().selectExpr("sum(id)").collect()
```

Use N from 2^{20} to 2^{27}

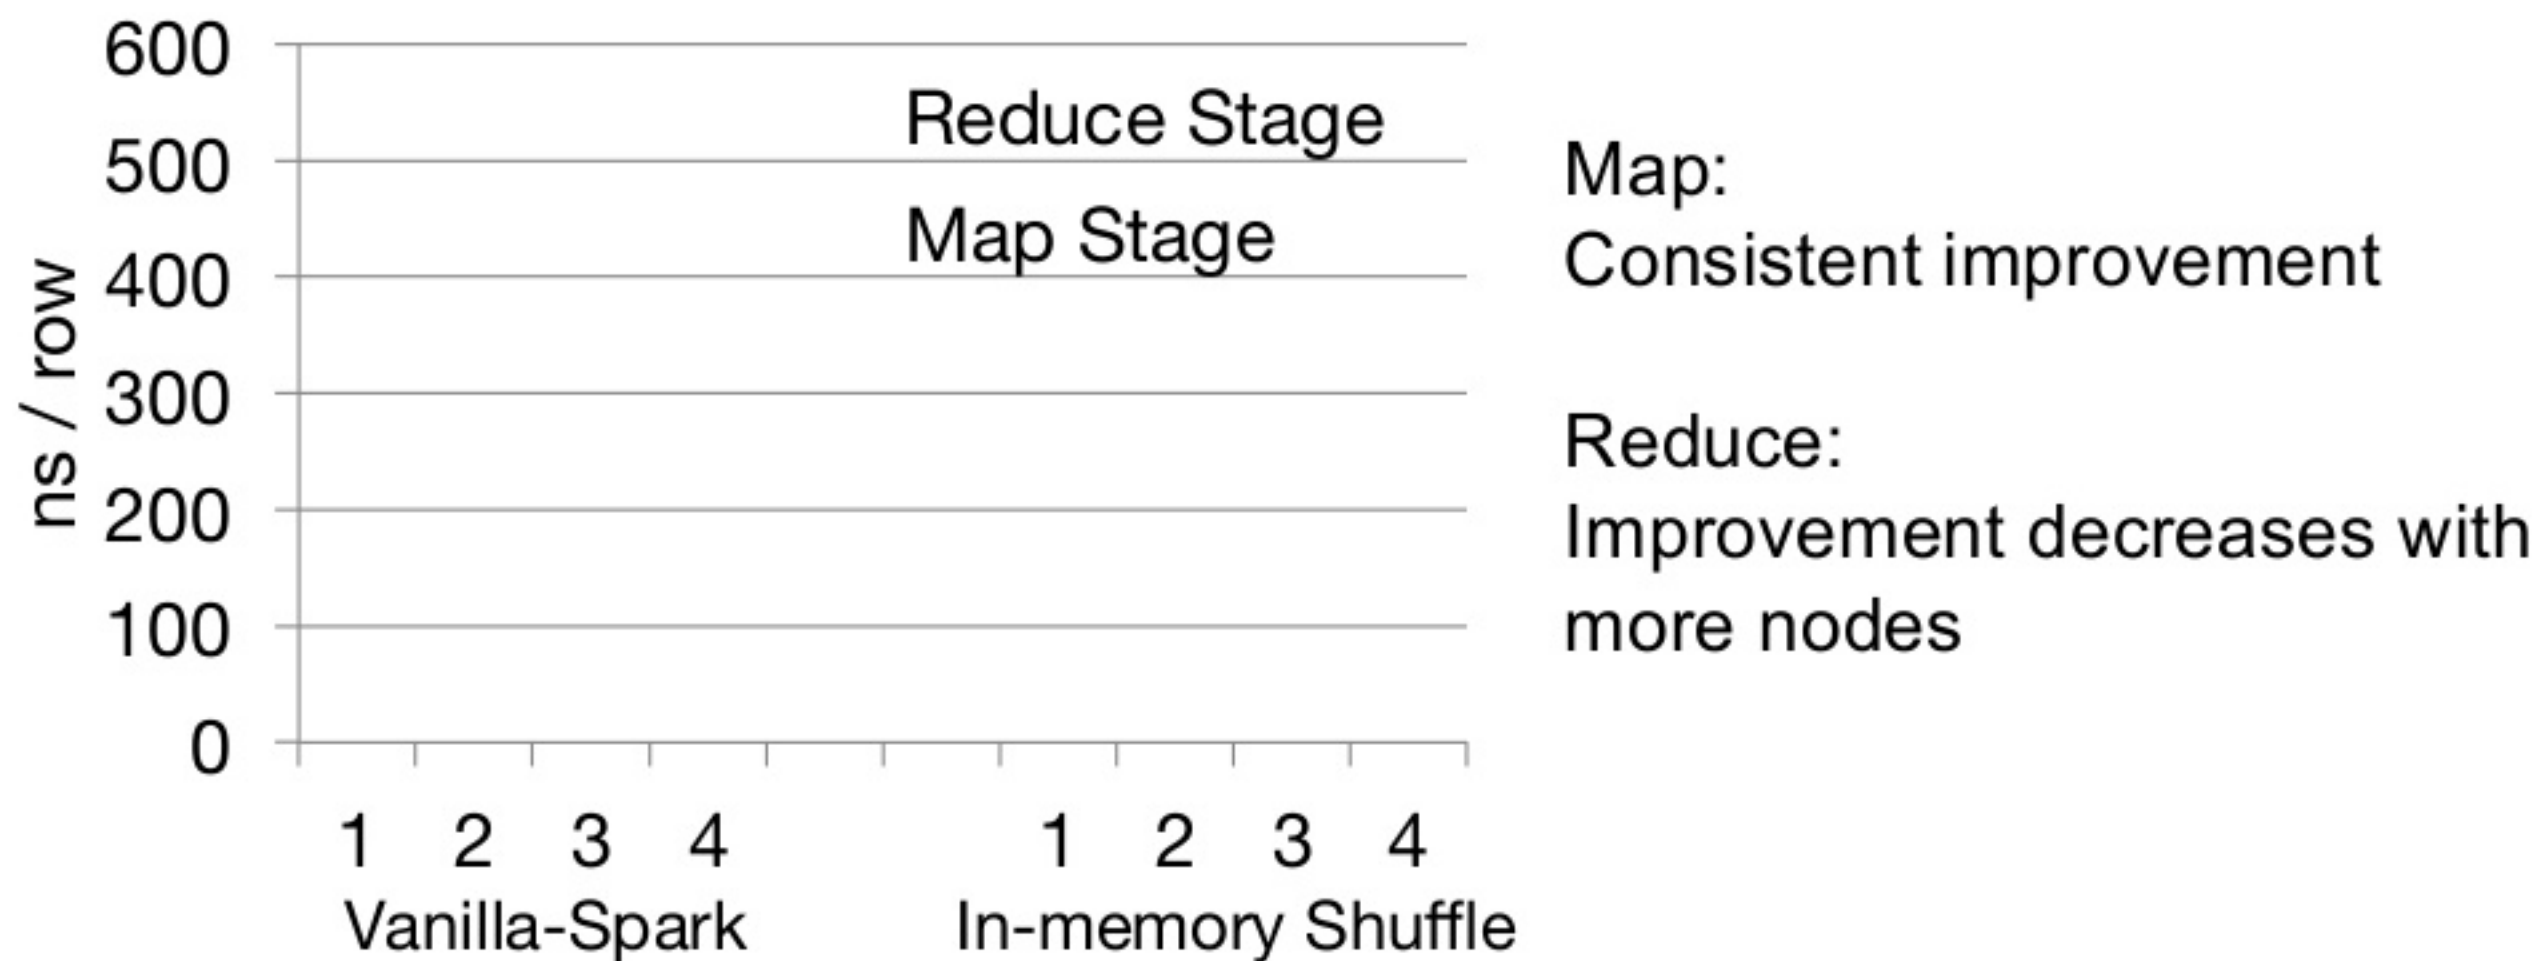
TPC-DS performance (single node)



Extending to multiple nodes

- Implementation
 - All data goes to memory
 - For remote transfer, copy from memory to network buffer
- A more memory-preserving way...
 - Local transfer goes to memory
 - Remote transfer goes to disk
 - Cons1: have to enforce stricter locality on reducers
 - Cons2: cannot avoid I/O contentions


Simple shuffle job



```
spark.range(N).repartition().selectExpr("sum(id)").collect()
```

TPC-DS performance (x1.xlarge32)

- SF=100
- Pick top 5 queries from single node experiment
- Best of 10 runs



Many other performance bottlenecks need investigation!

Summary

- Spark on many-core requires many architectural changes
- In-memory shuffle
 - How to improve shuffle performance with memory
 - 31% improvement over Spark
- On-going research
 - Identify other performance bottlenecks

Thank you

Qifan Pu

qifan@cs.berkeley.edu