


Online Learning with Structured Streaming

Ram Sriharsha, Vlad Feinberg

 @halfabrane

Spark Summit, Brussels
27 October 2016




What is online learning?

- Update model parameters on each data point
 - In batch setting get to see the entire dataset before update
- Cannot visit data points again
 - In batch setting, can iterate over data points as many times as we want!

An example: the perceptron

Goal: Find the best line separating positive
From negative examples on a plane

Update Rule: if $(y \neq \text{sign}(w \cdot x))$, $w \rightarrow w + y(w \cdot x)$



Why learn online?


- I want to adapt to changing patterns *quickly*
 - data distribution can change
 - e.g, distribution of features that affect learning might change over time
- I need to learn a good model *within resource + time* constraints (*large-scale learning*)
 - Time to a given accuracy might be faster for certain online algorithms

Online Classification Setting

- Pick a hypothesis
- For each labeled example (\mathbf{x}, y) :
 - Predict label \tilde{y} using hypothesis
 - Observe the loss $\mathcal{L}(y, \tilde{y})$ (and its gradient)
 - Learn from mistake and update hypothesis
- Goal: to make as few mistakes as possible in comparison to the *best* hypothesis in *hindsight*

An example: Online SGD

- Initialize weights \mathbf{w}
- Loss function \mathcal{L} is known.
- For each labeled example (\mathbf{x}, y) :
 - Perform update $\mathbf{w} \rightarrow \mathbf{w} - \eta \nabla \mathcal{L}(y, \mathbf{w} \cdot \mathbf{x})$
- For each new example \mathbf{x} :
 - Predict $\tilde{y} = \sigma(\mathbf{w} \cdot \mathbf{x})$ (σ is called link function)



Distributed Online Learning

- *Synchronous*
 - On each worker:
 - Load training data, compute gradients and update model, push model to driver
 - On some node:
 - Perform model merge
- *Asynchronous*
 - On each worker:
 - Load training data, compute gradients and push to server
 - On each server:
 - Aggregate the gradients, perform update step

Challenges

- Not all algorithms admit *efficient* online versions
- Lack of infrastructure
 - (Single machine) Vowpal Wabbit works great but hard to use from Scala, Java and other languages.
 - (Distributed) No implementation that is *fault tolerant, scalable, robust*
- Lack of framework in open source to provide extensible algorithms
 - Adagrad, normalized learning, L1 regularization, ...
 - Online SGD, FTRL, ...

Structured Streaming



Structured Streaming

1. One single API **DataFrame** for everything
 - Same API for machine learning, batch processing, graphX
 - Dataset is a typed version of DataFrame for Scala and Java
2. End-to-end exactly-once guarantees
 - The guarantees extend into the sources/sinks, e.g. MySQL, S3
3. Understands external event-time
 - Handling late arriving data
 - Support sessionization based on event-time

How does it work?


at any time, the output of the application is equivalent to executing a batch job on a prefix of the data

The Model

Input: data from source as an append-only table

Trigger: how frequently to check input for new data

Query: operations on input
usual map/filter/reduce
new window, session ops

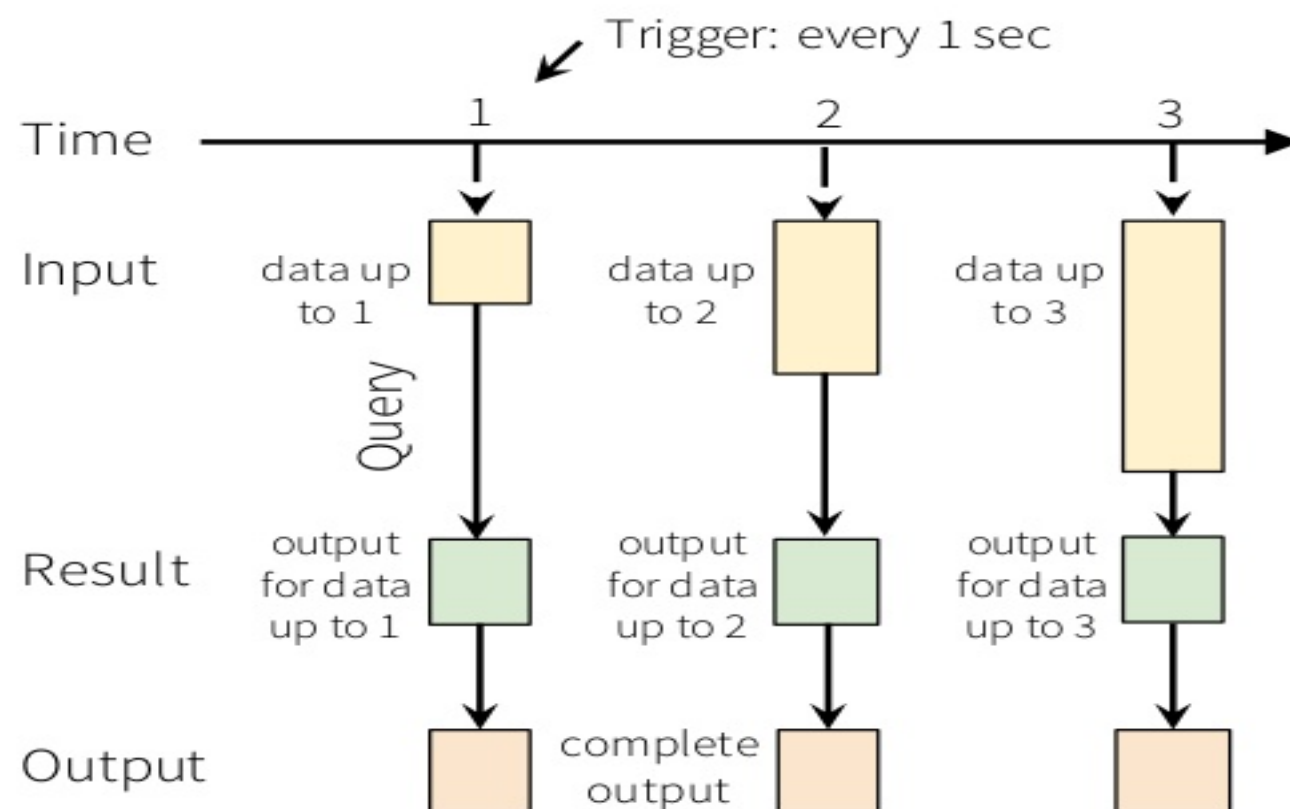


The Model

Result: final operated table
updated every trigger interval

Output: what part of result to write
to data sink after every trigger

Complete output: Write full result table every time



The Model

Result: final operated table
updated every trigger interval


Output: what part of result to write
to data sink after every trigger

Complete output: Write full result table every time

Delta output: Write only the rows that changed
in result from previous batch

Append output: Write only new rows

*Not all output modes are feasible with all queries




Streaming ML on Structured Streaming



Streaming ML on Structured Streaming

Input: append only table containing labeled examples

Query: Stateful aggregation query: picks up the last trained model, performs a distributed update + merge

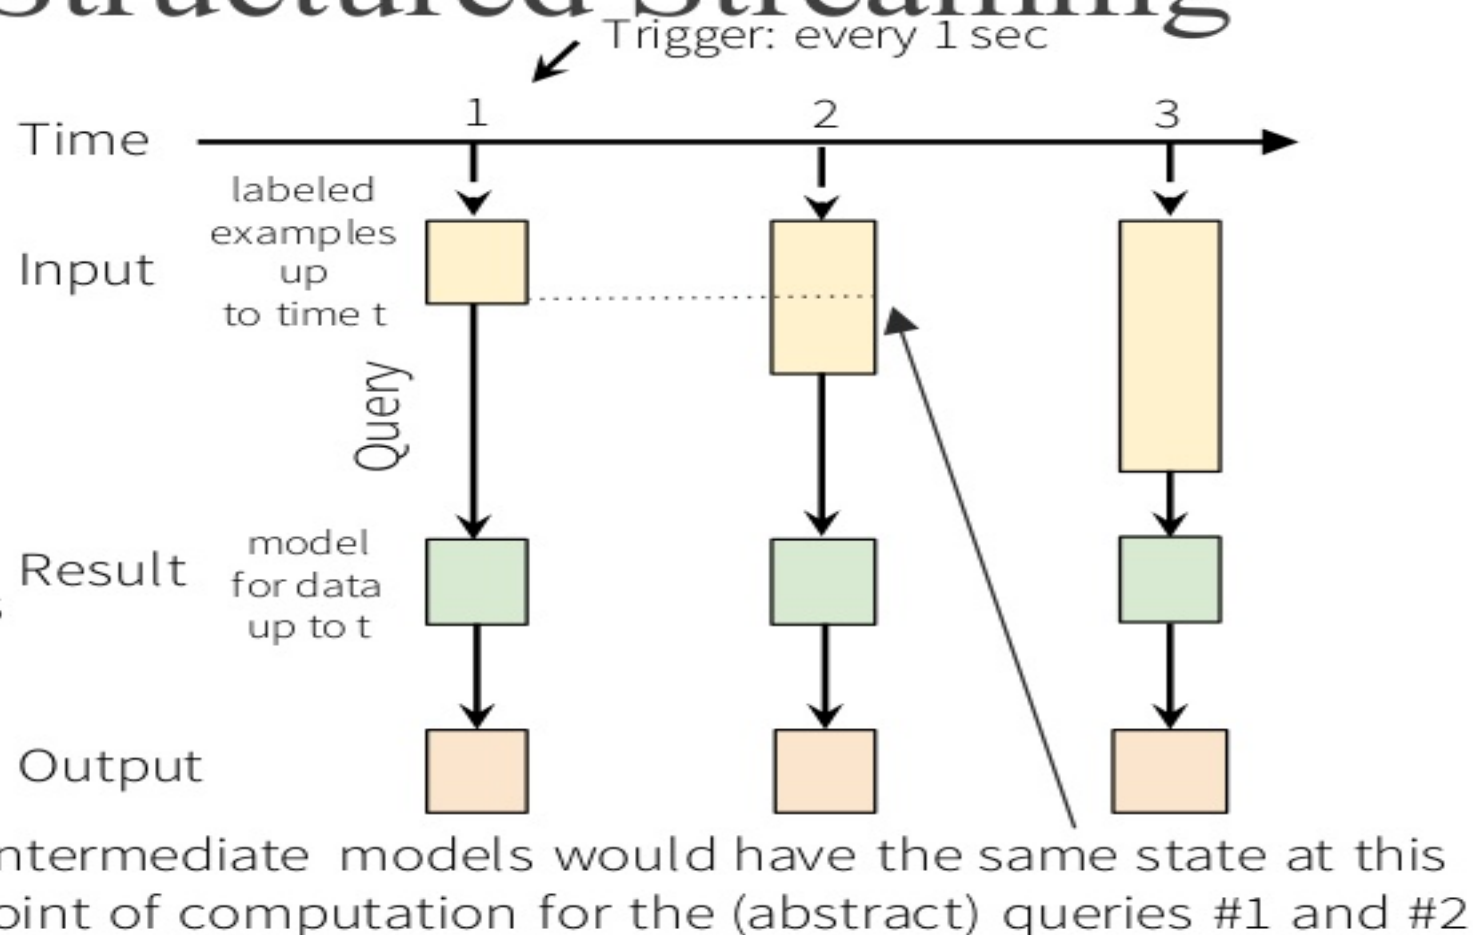


Streaming ML on Structured Streaming

Result: table of model parameters updated every trigger interval

Complete mode: table has one row, constantly being updated

Append mode (in the works): table has timestamp-keyed model, one row per trigger



Why is this hard?

- Need to update model, i.e.
 - `Update(previousModel, newDataPoint) = newModel`
- Typical aggregation is associative, commutative
 - e.g. `sum(P1: sum(sum(0, data[0]), data[1]), P2: sum(sum(0, data[2]), data[3]))`
- General model update violates associativity + commutativity!

Solution: Make Assumptions

- Result may be partition-dependent, but we don't care as long as we get some valid result.

average-models(
 P1: update(update(previous model, data[0]), data[1]),
 P2: update(update(previous model, data[2]), data[3]))

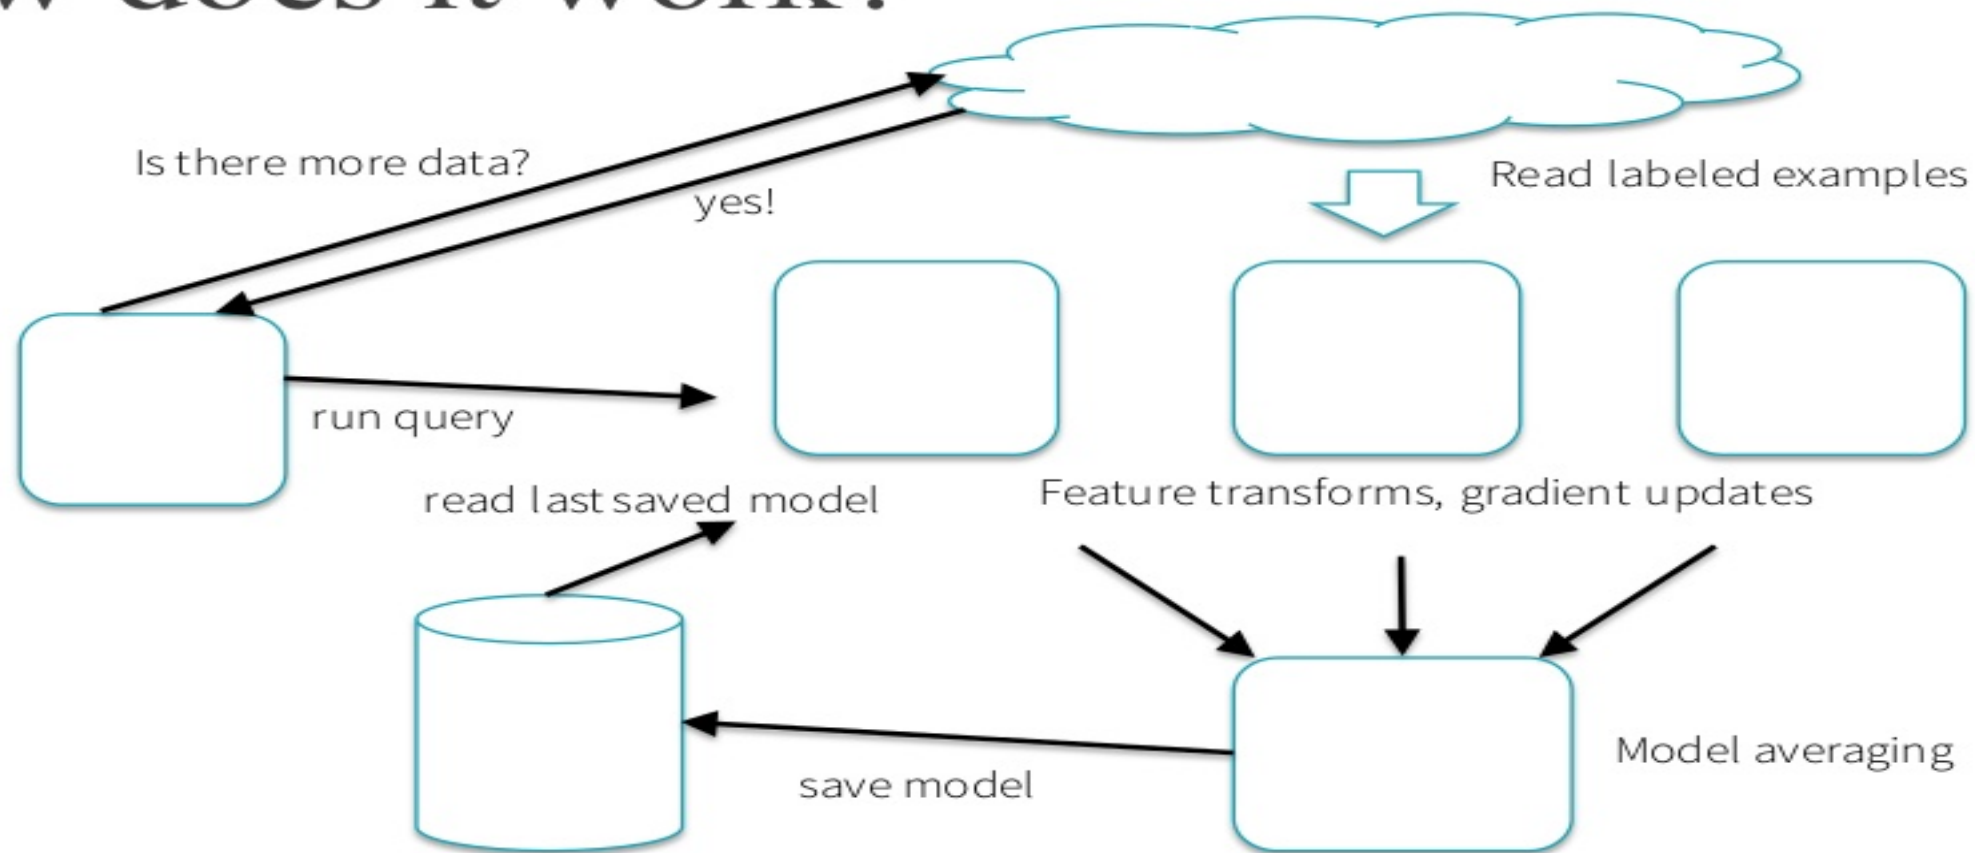
- Only partition-dependent if update and average don't commute - can still be deterministic otherwise!

Stateful Aggregator

- Within each partition
 - Initialize with previous state (instead of zero in regular aggregator)
 - For each item, update state
- Perform reduce step
- Output final state

Very general abstraction: works for sketches, online statistics (quantiles), online clustering ...

How does it work?



APIs



Spark Summit Brussels
27 October 2016



ML Estimator on Streams

- Interoperable with ML pipelines




Input: stream of labelled data
Output: stream of models, updated over time.

Batch Interoperability

- Seamless application on batch datasets

```
model = estimator.fit(batchDF)
```



Feature Creation

- Handle new features as they appear (ex., IPs in fraud detection)
 - Provide transformers, such as the HashingEncoder, that apply the hashing trick.
 - Encode arbitrary (possibly categorical data) without knowing cardinality ahead of time by using a high-dimensional sparse mapping.

API Goals

- Provide modern, regret-minimization-based online algorithms.
 - Online Logistic Regression
 - Adagrad
 - Online gradient descent
 - L2 regularization
- Input streams of any kind accepted.
- Streaming aware feature engineering

What's next?



Spark Summit Brussels
27 October 2016



What's next?

- More bells and whistles
 - Adaptive normalization
 - L1 regularization
- More algorithms
 - Online quantile estimation?
 - More general Sketches?
 - Online clustering?
- Scale testing and benchmarking

Demo



Spark Summit Brussels
27 October 2016

