

# Apache® Spark™ MLlib 2.0 Preview: Data Science and Production

Joseph K. Bradley


June 8, 2016



# Who am I?

- Apache Spark committer & PMC member
- Software Engineer @ Databricks
- Ph.D. in Machine Learning from Carnegie Mellon U.

# MLlib trajectory



# DataFrame-based API for MLlib

a.k.a. “Pipelines” API, with utilities for constructing ML Pipelines

In 2.0, the DataFrame-based API will become the primary API for MLlib.

- Voted by community
- `org.apache.spark.ml`, `pyspark.ml`

The RDD-based API will enter maintenance mode.

- Still maintained with bug fixes, but no new features
- `org.apache.spark.mllib`, `pyspark.mllib`

# Goals for MLlib in 2.0

## Major initiatives

- Generalized Linear Models
  - Python & R API expansion
  - ML persistence: saving & loading models & Pipelines
- 
- Exploratory analysis
- Production

## Also in 2.0:

- Sketching algorithms: <http://databricks.com/blog/2016/05/19>
- For details, see [SPARK-12626](#) roadmap JIRA + mailing list discussions.



MLlib for exploratory analysis


Generalized Linear Models (GLMs)

Python & R APIs

# Generalized Linear Models (GLMs)

Arguably *the* most important class of models for ML

- Logistic regression
  - Linear regression
  - Many other types of models
  - Model summary statistics
- In Spark 1.6 & earlier



# GLMs in 2.0

Model family	Supported link functions
Gaussian	Identity, Log, Inverse
Binomial	Logit, Probit, CLogLog
Poisson	Log, Identity, Sqrt
Gamma	Inverse, Identity, Log

Fixes for corner cases

- E.g., handle invalid labels gracefully

`GeneralizedLinearRegression`

- Max 4096 features
- Solved using Iteratively Reweighted Least Squares (IRLS)

`LinearRegression` &  
`LogisticRegression`

- Millions of features
- Solved using L-BFGS / OWL-QN



# Python & R APIs for MLlib

Goal: Expand ML APIs for critical languages for data science

## Python

- **Clustering algorithms**: Bisecting K-Means, Gaussian Mixtures, LDA
- **Meta-algorithms**: OneVsRest, TrainValidationSplit
- **GeneralizedLinearRegression**
- **Feature transformers**: ChiSqSelector, MaxAbsScaler, QuantileDiscretizer
- **Model inspection**: summaries for Logistic Regression, Linear Regression, GLMs

## R

- **Regression & classification**: Generalized Linear Regression, AFT survival regression
- **Clustering**: K-Means

MLlib in production

ML Persistence

Customizing Pipelines


# Why ML persistence?

Data Science

Prototype (Python/R)  
Create model

Software Engineering

Re-implement model for  
production (Java)  
Deploy model



# Why ML persistence?

## Data Science

Prototype (Python/R)  
Create Pipeline

- Extract raw features
- Transform features
- Select key features
- Fit multiple models
- Combine results to make prediction



## Software Engineering

Re-implement Pipeline for  
production (Java)  
Deploy Pipeline

- Extra implementation work
- Different code paths
- Synchronization overhead

# With ML persistence...

Data Science

Software Engineering

Prototype (Python/R)  
Create Pipeline

Load Pipeline (Scala/Java)  
`Model.load("s3n://...")`  
Deploy in production

Persist model or Pipeline:  
`model.save("s3n://...")`

```
graph TD; DS[Data Science] --> SE[Software Engineering]; DS --> Persist[Persist model or Pipeline: model.save("s3n://...")]; Persist --> SE;
```

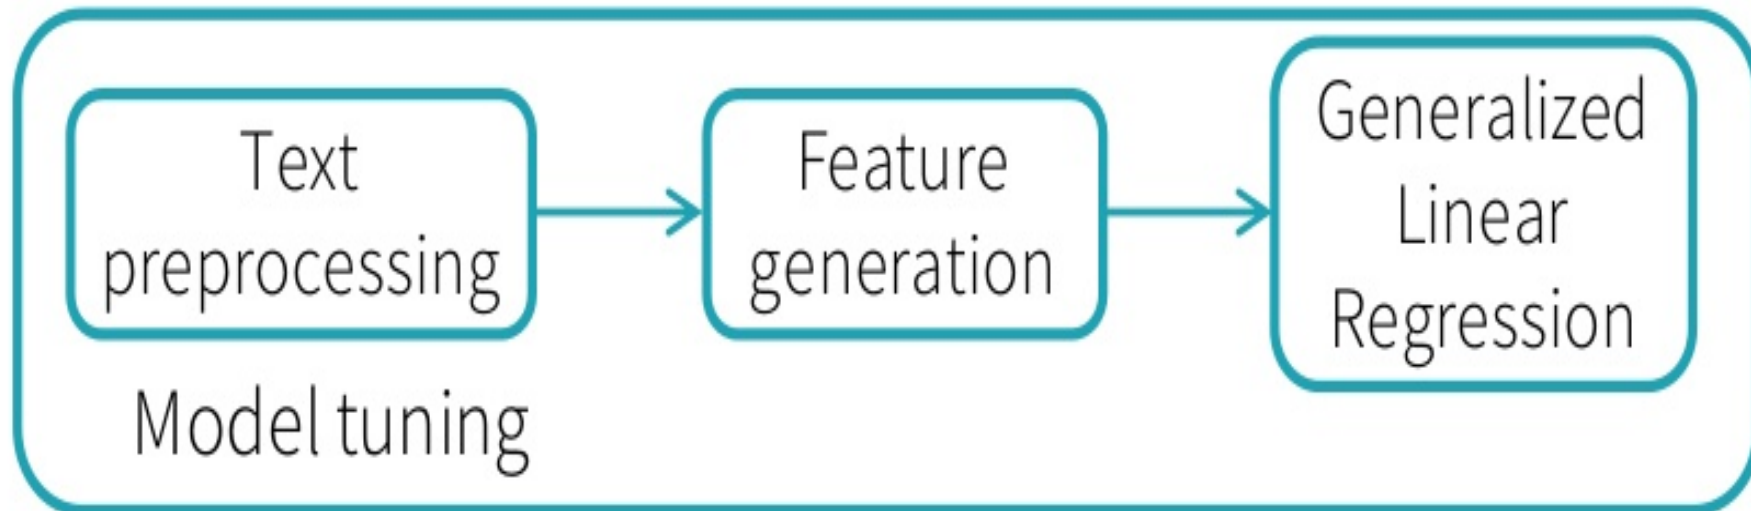
The diagram illustrates the workflow of ML persistence. It starts with 'Data Science' on the left, where a 'Prototype (Python/R)' is used to 'Create Pipeline'. An arrow points from this step to a central 'Persist model or Pipeline' step, which involves running `model.save("s3n://...")`. From the persistence step, an arrow points to the 'Software Engineering' side, where the pipeline is 'Load Pipeline (Scala/Java)' using `Model.load("s3n://...")` and then 'Deploy in production'.



# ML persistence status

	"recipe"	"result"
	Unfitted	Fitted
Model	✓	✓
Pipeline	✓	✓

Supported in MLlib's  
RDD-based API



# ML persistence status

Near-complete coverage in all Spark language APIs

- Scala & Java: complete
- Python: complete except for 2 algorithms
- R: complete for existing APIs

Single underlying implementation of models

Exchangeable data format

- JSON for metadata
- Parquet for model data (coefficients, etc.)

# Customizing ML Pipelines

MLlib 2.0 will include:

- 29 feature transformers (Tokenizer, Word2Vec, ...)
- 21 models (for classification, regression, clustering, ...)
- Model tuning & evaluation

But some applications require customized  
Transformers & Models.

# Options for customization

Extend abstractions

- Transformer
- Estimator & Model
- Evaluator

Existing use cases:

- Natural Language Processing (Snowball, Stanford NLP)
- Featurization libraries
- Many others!

UnaryTransformer: input → output

E.g.:

- Tokenizer: text → sequence of words
- Normalizer: vector → normalized vector

Simple API which provides:

- DataFrame-based API
- Built-in parameter getters, setters
- Distributed Row-wise transformation

# Persistence for customized algorithms

2 traits provide persistence for simple Transformers:

- `DefaultParamsWritable`
- `DefaultParamsReadable`

Simply mix in these traits to provide persistence in Scala/Java.

→ Saves & loads algorithm parameters

```
val myParam: Param[Double]
```

→ Does not save member data

```
val myCoefficients: Vector
```

For an example, see `UnaryTransformerExample.scala`  
in `spark/examples/`



# Recap

## Exploratory analysis

- Generalized Linear Regression
- Python & R APIs

## Production

- ML persistence
- Customizing Pipelines

Many thanks to the community  
for contributions & support!

# What's next?

Prioritized items on the 2.1 roadmap JIRA (SPARK-15581):

- Critical feature completeness for the DataFrame-based API
  - Multiclass logistic regression
  - Statistics
- Python API parity & R API expansion
- Scaling & speed tuning for key algorithms: trees & ensembles

## GraphFrames

- Release for Spark 2.0
- Speed improvements (join elimination, connected components)

# Get started

- Get involved via roadmap JIRA (SPARK-15581) + mailing lists

- Download notebook for this talk  
<http://dbricks.co/1UfvAH9>

Try out the Apache Spark 2.0 preview release:

<http://databricks.com/try>

- ML persistence blog post  
<http://databricks.com/blog/2016/05/31>

# *Thank you!*

Office hour @ 2:45pm today (Expo Hall)

Twitter: @jkbatcmu