



# Efficient state management with Spark 2.0 and scale-out databases

Jags Ramnarayan

Co-founder, CTO, SnappyData

Barzan Mozafari

Strategic Advisor @ Snappydata


Assistant Professor,

Univ. of Michigan, Ann Arbor



**SPARK SUMMIT 2016**  
DATA SCIENCE AND ENGINEERING AT SCALE  
JUNE 6-8, 2016 SAN FRANCISCO

# SnappyData's focus – DB on Spark



Funded by Pivotal, GE, GTD Capital

- New Spark-based open source project started by Pivotal GemFire founders+engineers
- Decades of in-memory data management experience
- Focus on real-time, operational analytics: Spark inside an OLTP+OLAP database

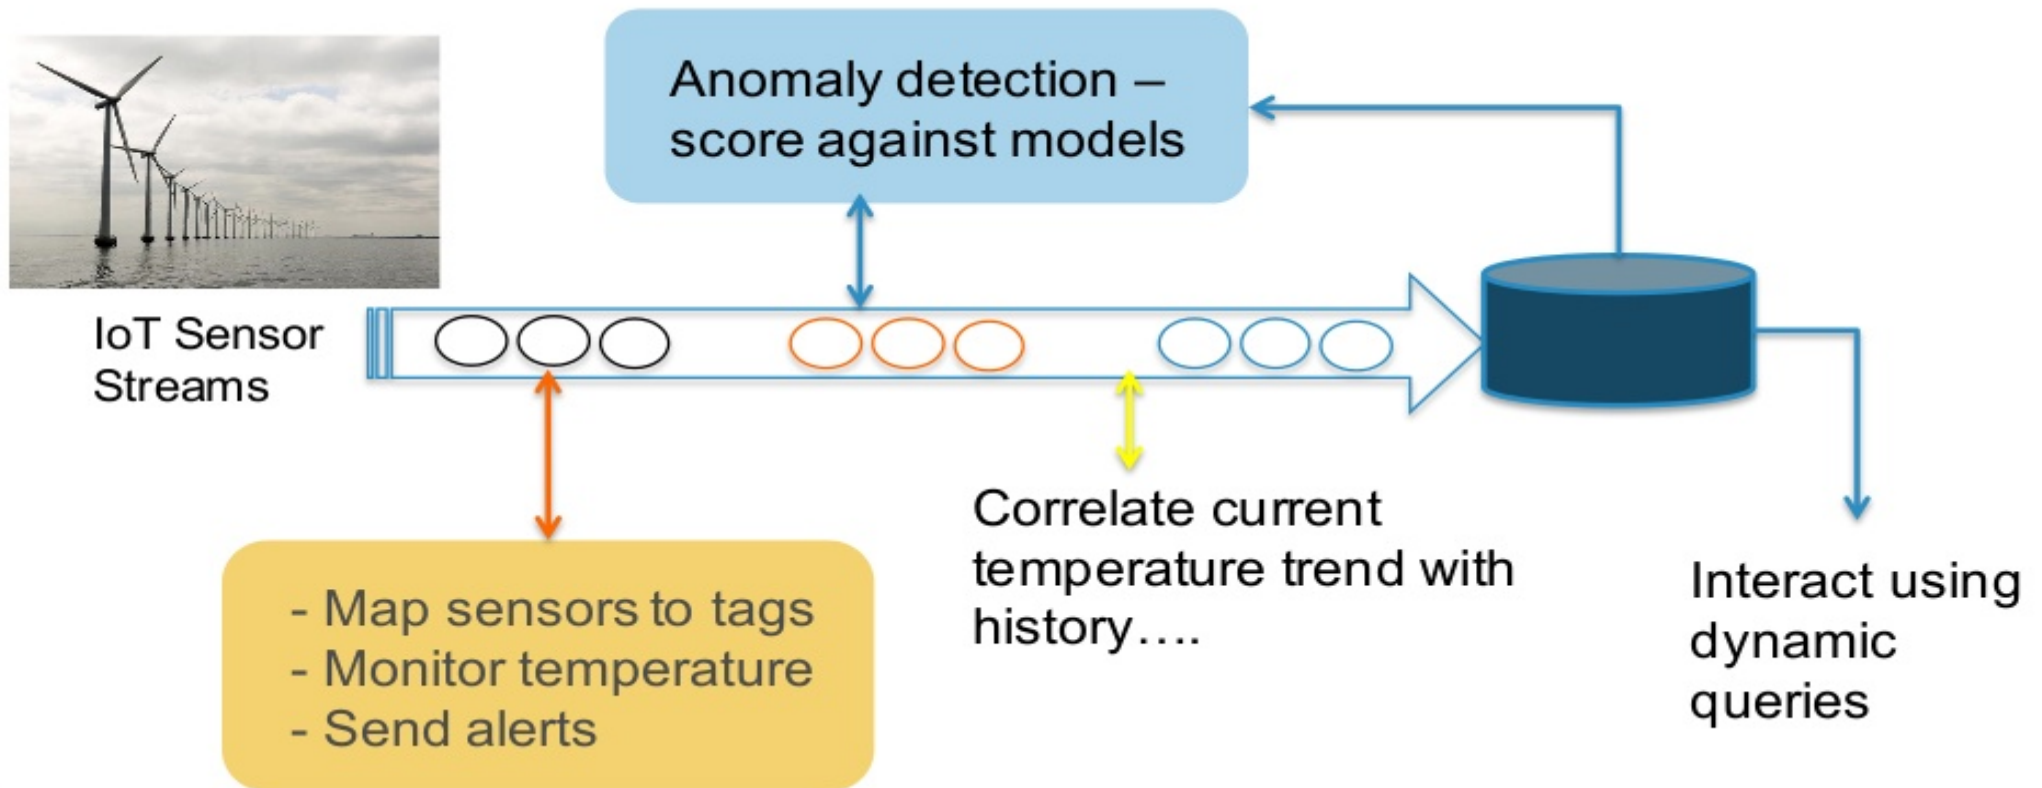


# Agenda

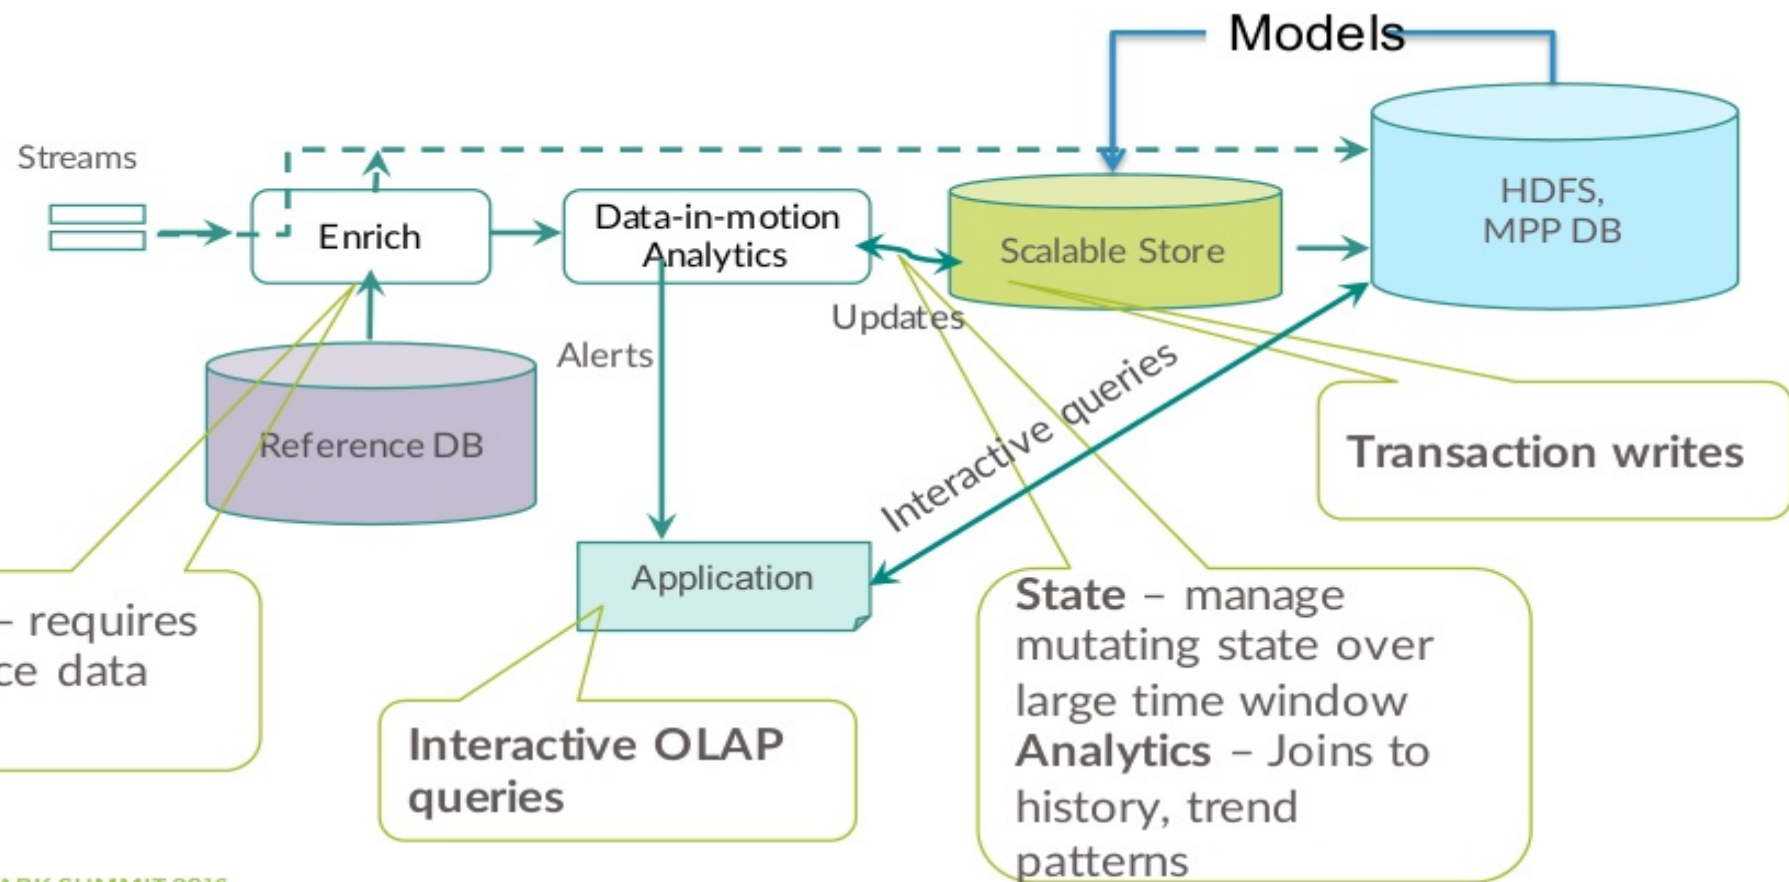
- Mixed workloads are important but complex today
- State management challenges with such workloads in spark Apps
- The SnappyData solution
- Approximate query processing
- How do we extend spark for real time, mixed workloads?
- Q&A



# Mixed Workloads are increasingly common



# Mixed workload Architecture is Complex (Lambda)



## Lambda Architecture is Complex

- Complexity: learn and master multiple products → data models, disparate APIs, configs
- Slower
- Wasted resources





# Can we simplify & optimize?

Perhaps a single clustered DB that can manage stream state, transactional data and run OLAP queries?



# Deeper Look into Managing State in Spark Applications



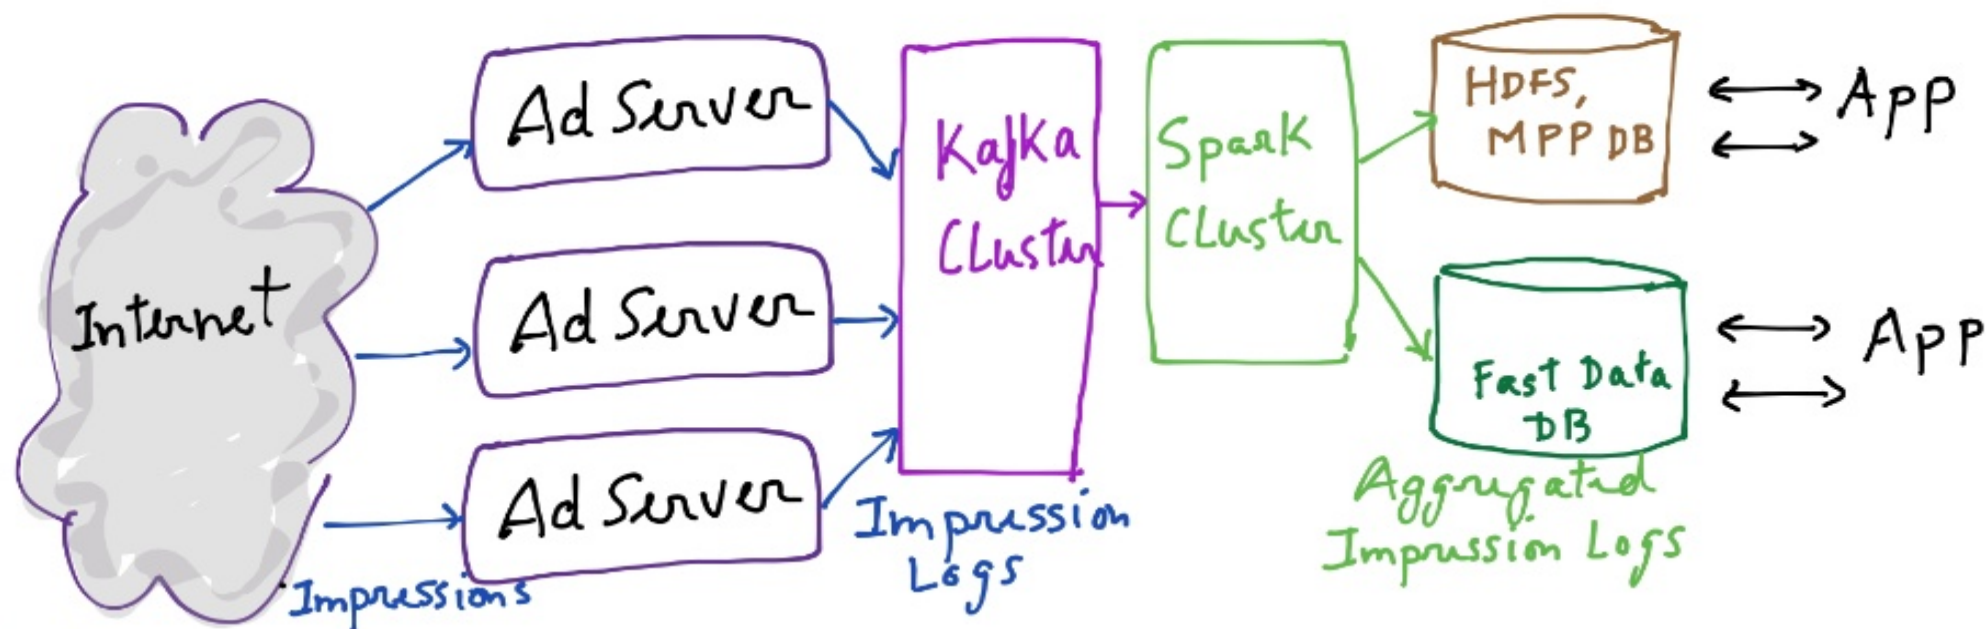
SPARK SUMMIT 2016

---



# Deeper Look – Ad Impression Analytics

Ad Network Architecture – Analyze log impressions in real time



# Ad Impression Analytics

To simplify, let's consider that impression logs are in this format:

timestamp		publisher	advertiser	website	geo	bid	cookie
2013-01-28 13:21:12		pub1	adv10	abc.com	NY	0.0001	1214
2013-01-28 13:21:13		pub1	adv10	abc.com	NY	0.0005	1214
2013-01-28 13:21:14		pub2	adv20	xyz.com	CA	0.0003	4321
2013-01-28 13:21:15		pub2	adv20	xyz.com	CA	0.0001	5675

Our goal is to aggregate these logs by publisher and geo, and compute the average bid, the number of impressions and the number of uniques by minute. So the aggregation will look something like:

timestamp		publisher	geo	avg_bid	imps	uniques
2013-01-28 13:21:00		pub1	NY	0.0003	256	104
2013-01-28 13:21:00		pub2	CA	0.0002	121	15
2013-01-28 13:22:00		pub1	NY	0.0001	190	98
2013-01-28 13:22:00		pub2	CA	0.0007	137	19



# Bottlenecks in the write path

```
val input :DataFrame= sqlContext.read  
  .options(kafkaOptions).format(..)  
  .stream("Kafka url")
```

Stream micro batches in parallel from Kafka to each Spark executor

```
val result :DataFrame = input  
  .where("geo != 'UnknownGeo'")  
  .groupBy(  
    window("event-time", "1 min"),  
    "Publisher", "geo")  
  .agg(avg("bid"), ....
```

Filter and collect event for 1 minute. Reduce to 1 event per Publisher, Geo every minute

Execute GROUP BY ... Expensive Spark Shuffle ...


```
val query = result.write.format("My Favorite NoSQLDB")  
  .outputMode("append")  
  .startStream("dest-path")
```

Shuffle again in DB cluster ... data format changes ... serialization costs



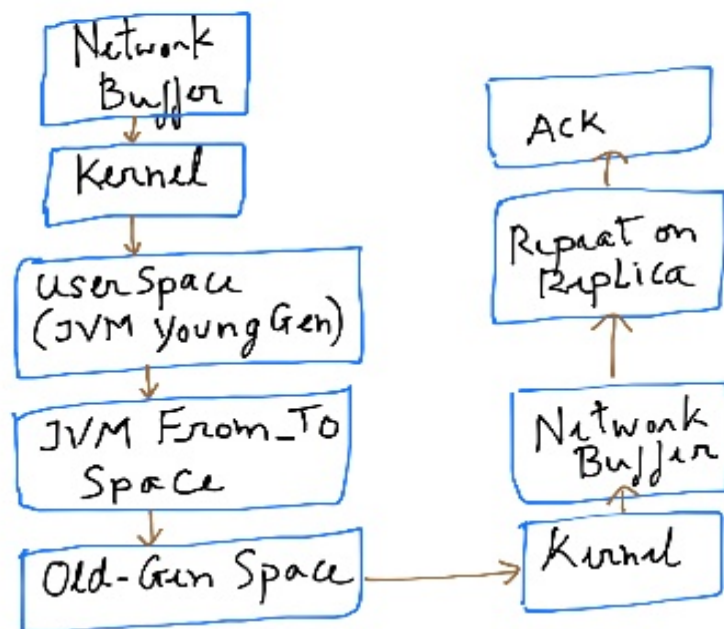
# Bottlenecks in the Write Path

## Shuffle Costs (Copying, Serialization)



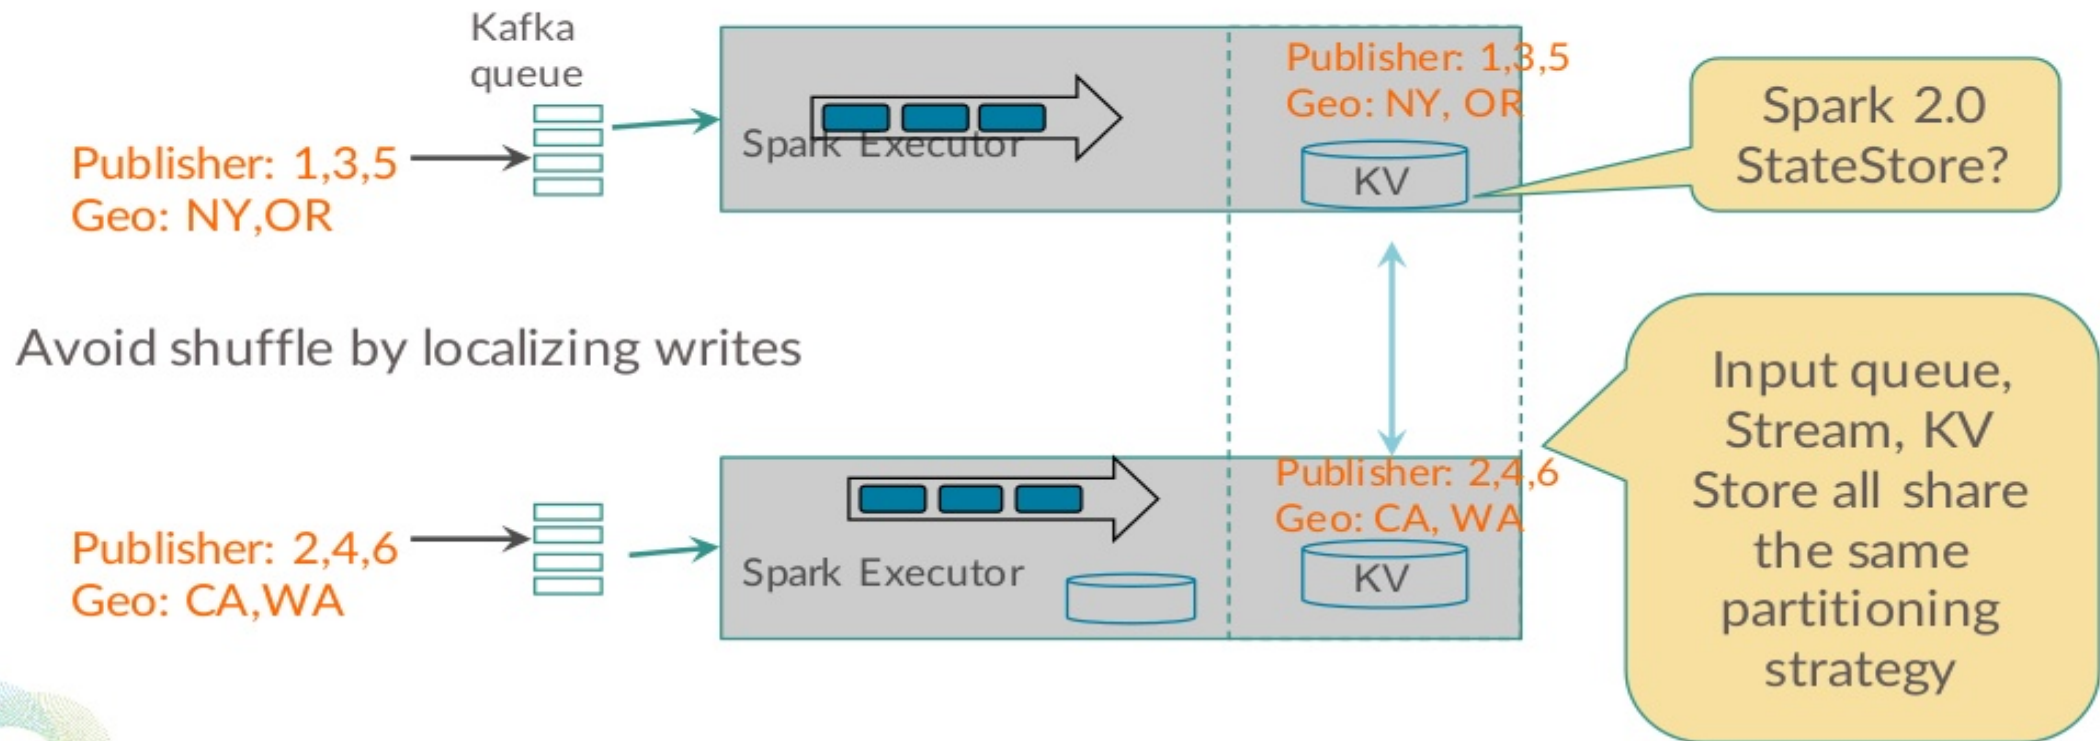
- Aggregations – GroupBy, MapReduce
- Joins with other streams, Reference data

## Excessive copying in Java based Scale out stores






# Avoid bottleneck - Localize process with State



# New State Management API in Spark 2.0 – KV Store

- Preserve state across streaming aggregations across multiple batches
- Fetch, store KV pairs
- Transactional
- Supports versioning (batch)
- Built in store that persists to HDFS





# Impedance mismatch with KV stores?

We want to run interactive “scan”-intensive queries:

- Find total *uniques* for Ads grouped on geography
- Impression trends for Advertisers(group by query)

Two alternatives: row-stores(all KV Stores) vs. column stores



Fast key-based lookup

But, too slow to run  
aggregations, scan based  
interactive queries

Emp_no	Dept_id	Hire_date	Emp_in	Emp_fn
1	1	2001-01-01	Smith	Bob
2	1	2002-02-01	Jones	Jim
3	1	2002-05-01	Young	Sue
4	2	2003-02-01	Stemle	Bill
5	2	1999-06-15	Aurora	Jack
6	3	2000-08-15	Jung	Laura

Fast scans, aggregations

Updates and random  
writes are very difficult

### Row-Oriented Database

1	1	2001-01-01	Smith	Bob
2	1	2002-02-01	Jones	Jim
3	1	2002-05-01	Young	Sue


Consume too  
much memory

### Column-Oriented Database

1	2	3	4	5
1	1	1	2	2
2001-01-01	2002-02-01	2002-02-01	2002-02-01	2002-02-01

# Why columnar storage in-memory?

## Hardware Changes: The Memory Wall



Source: MonetDB



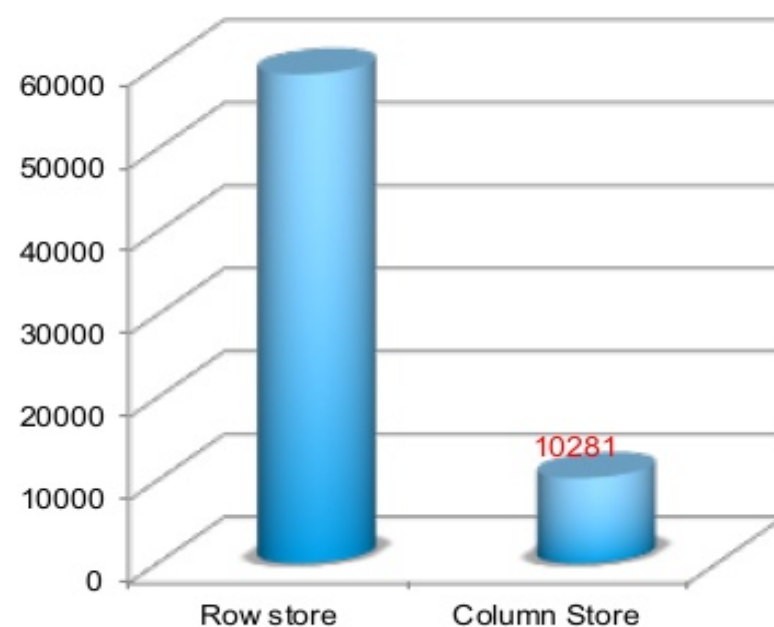
SPARK SUMMIT 2016

## But, are in-memory column-stores fast enough?

AWS c4.2xlarge ; 4 x (8 cores, 15GB)

Column Table: AdImpressions; 450million rows

```
select count(*) from adImpressions  
  
group by geo  
  
order by count desc  
  
limit 20;
```



Single User, 10+seconds. Is this Interactive speed?




SPARK SUMMIT 2016




# Use statistical techniques to shrink data!

- Most apps happy to tradeoff 1% accuracy for 200x speedup!
  - Can usually get a 99.9% accurate answer by only looking at a tiny fraction of data!
- Often can make perfectly accurate decisions with imperfect answers!
  - A/B Testing, visualization, ...
- The data itself is usually noisy
  - Processing entire data doesn't necessarily mean exact answers!
- Inference is probabilistic anyway



Original (2 billion points), 71 mins



VAS (1M), 3 secs



# Our Solution: SnappyData

Open Sourced @ <https://github.com/SnappyDataInc/snappydata>



SPARK SUMMIT 2016




# SnappyData: A New Approach



**A Single Unified Cluster:** OLTP + OLAP + Streaming  
for real-time analytics



Huge ecosystem



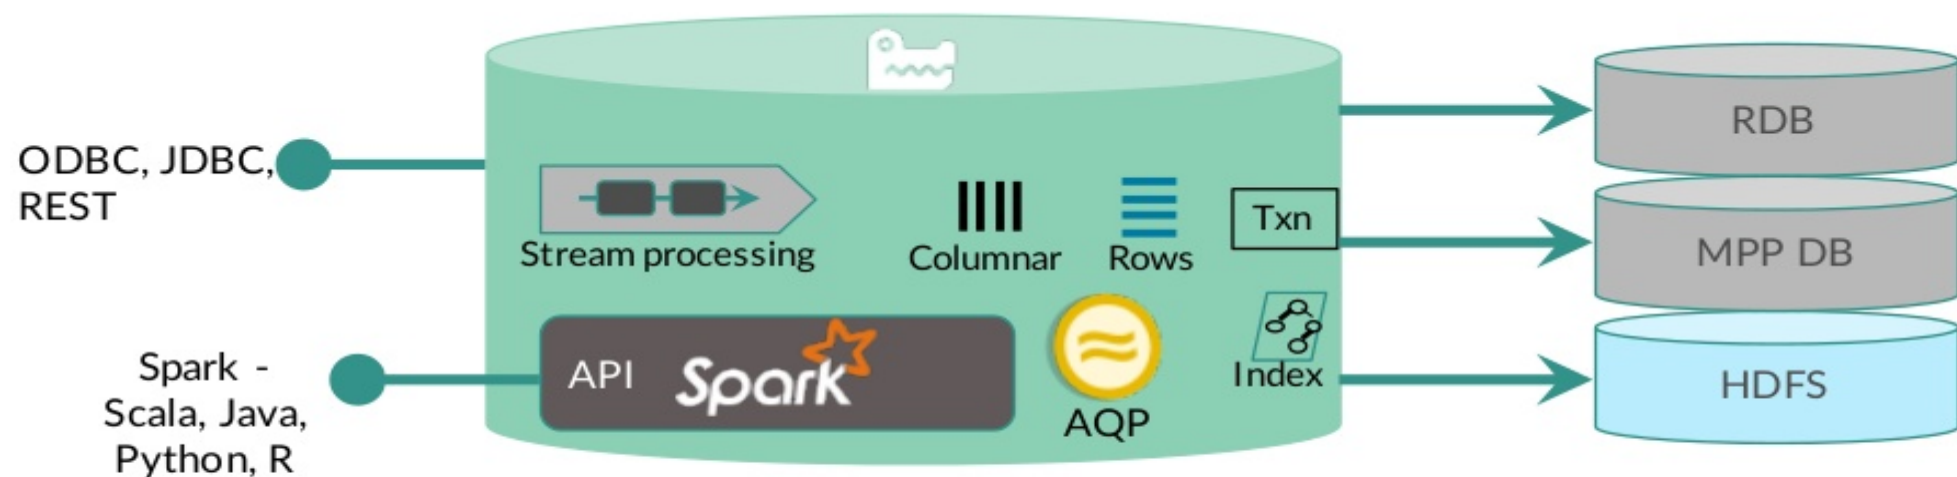
Matured over 13 years



SPARK SUMMIT 2016

Vision: Drastically reduce the cost and  
complexity in modern big data

# Unified In-memory DB for Streams, Txn, OLAP queries



First commercial product with  
Approximate Query Processing (AQP)

Real-time operational Analytics – TBs in memory



SPARK SUMMIT 2016

# Features

- **Deeply integrated database for Spark**
  - 100% compatible with Spark
  - Extensions for Transactions (updates), SQL stream processing
  - Extensions for High Availability
  - Approximate query processing for interactive OLAP
- **OLTP+OLAP Store**
  - Replicated and partitioned tables
  - Tables can be Row or Column oriented (in-memory & on-disk)
  - SQL extensions for compatibility with SQL Standard
    - create table, view, indexes, constraints, etc



## Approximate Query Processing Features

- **Uniform random sampling**
- **Stratified sampling**
  - Solutions exist for stored data (BlinkDB)
  - SnappyData works for infinite streams of data too
- Support for **synopses**
  - Top-K queries, heavy hitters, outliers, ...
- **Exponentially decaying windows** over time



## Interactive-Speed Analytic Queries – Exact or Approximate

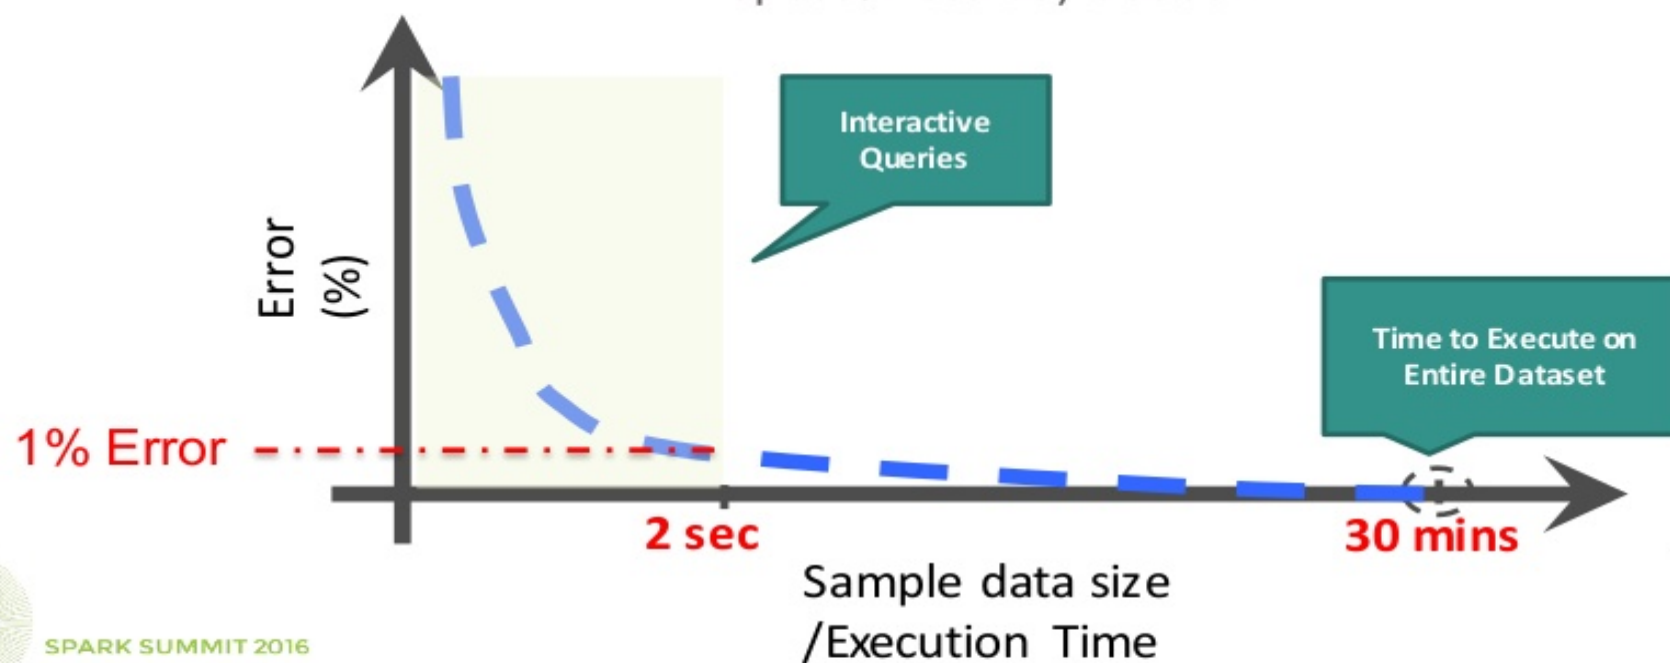
`Select avg(Bid), Advertiser from T1 group by Advertiser`

`Select avg(Bid), Advertiser from T1 group by Advertiser with error 0.1`

100 secs


2 secs

Speed/Accuracy tradeoff



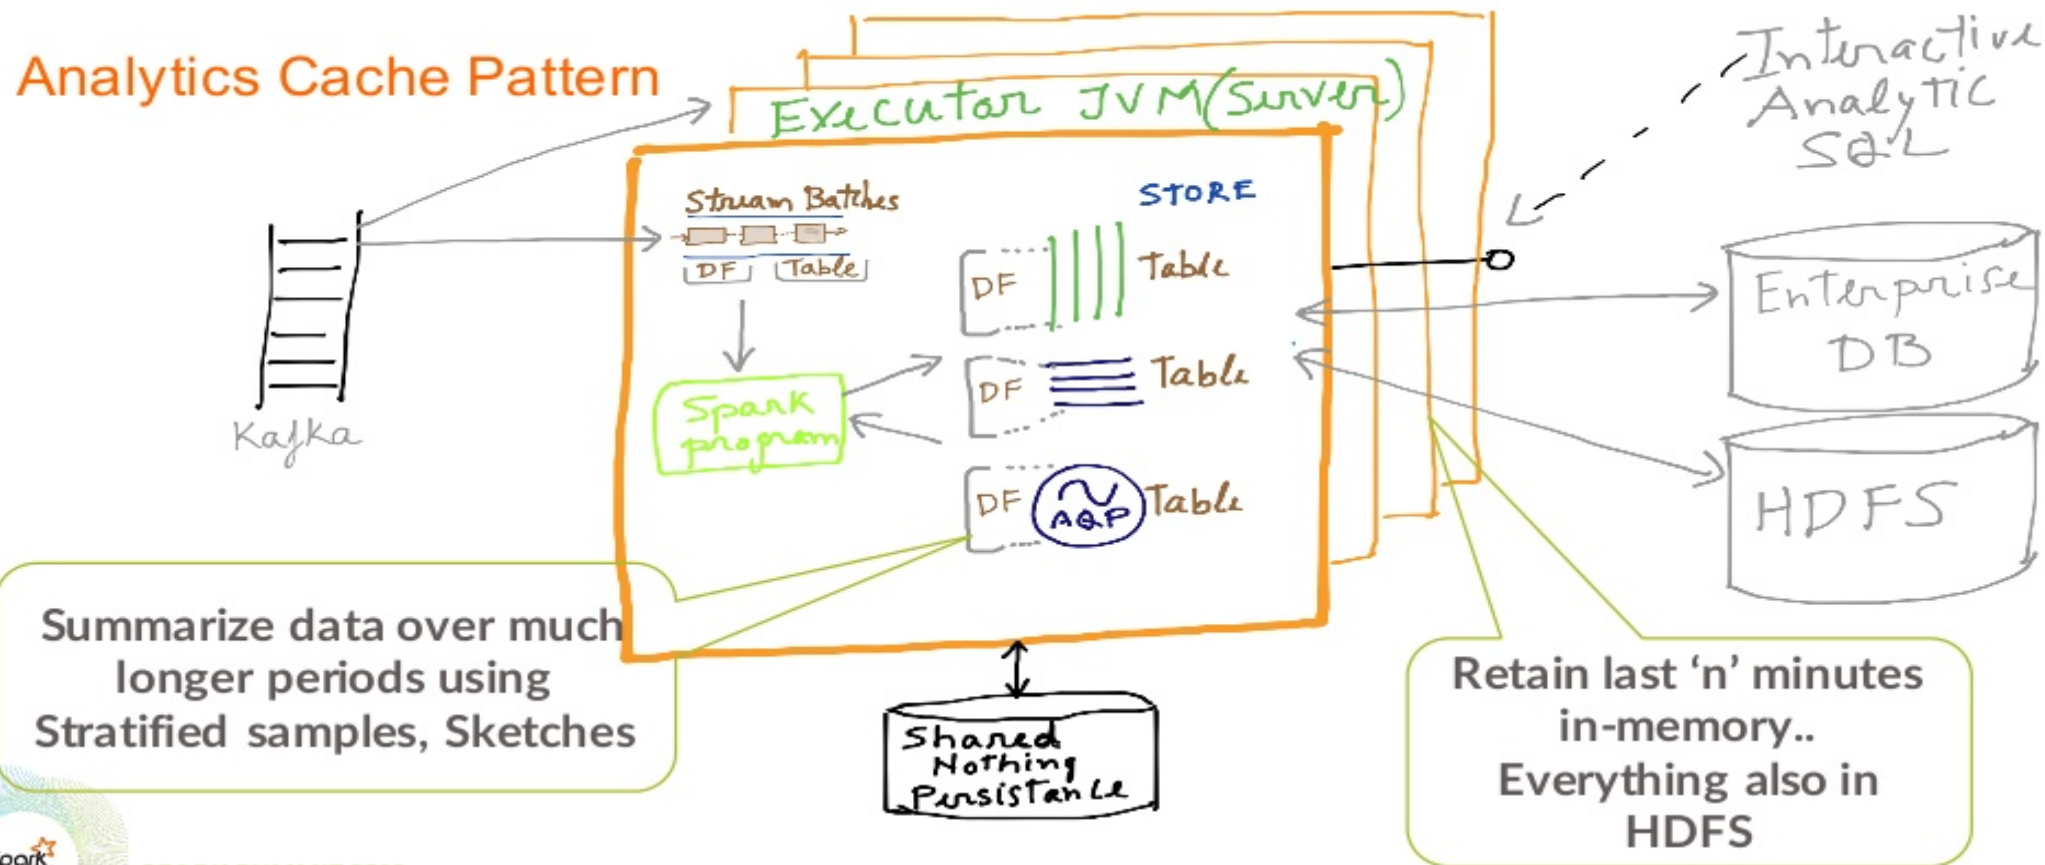
SPARK SUMMIT 2016

# Airline Ontime performance Analytics – Demo



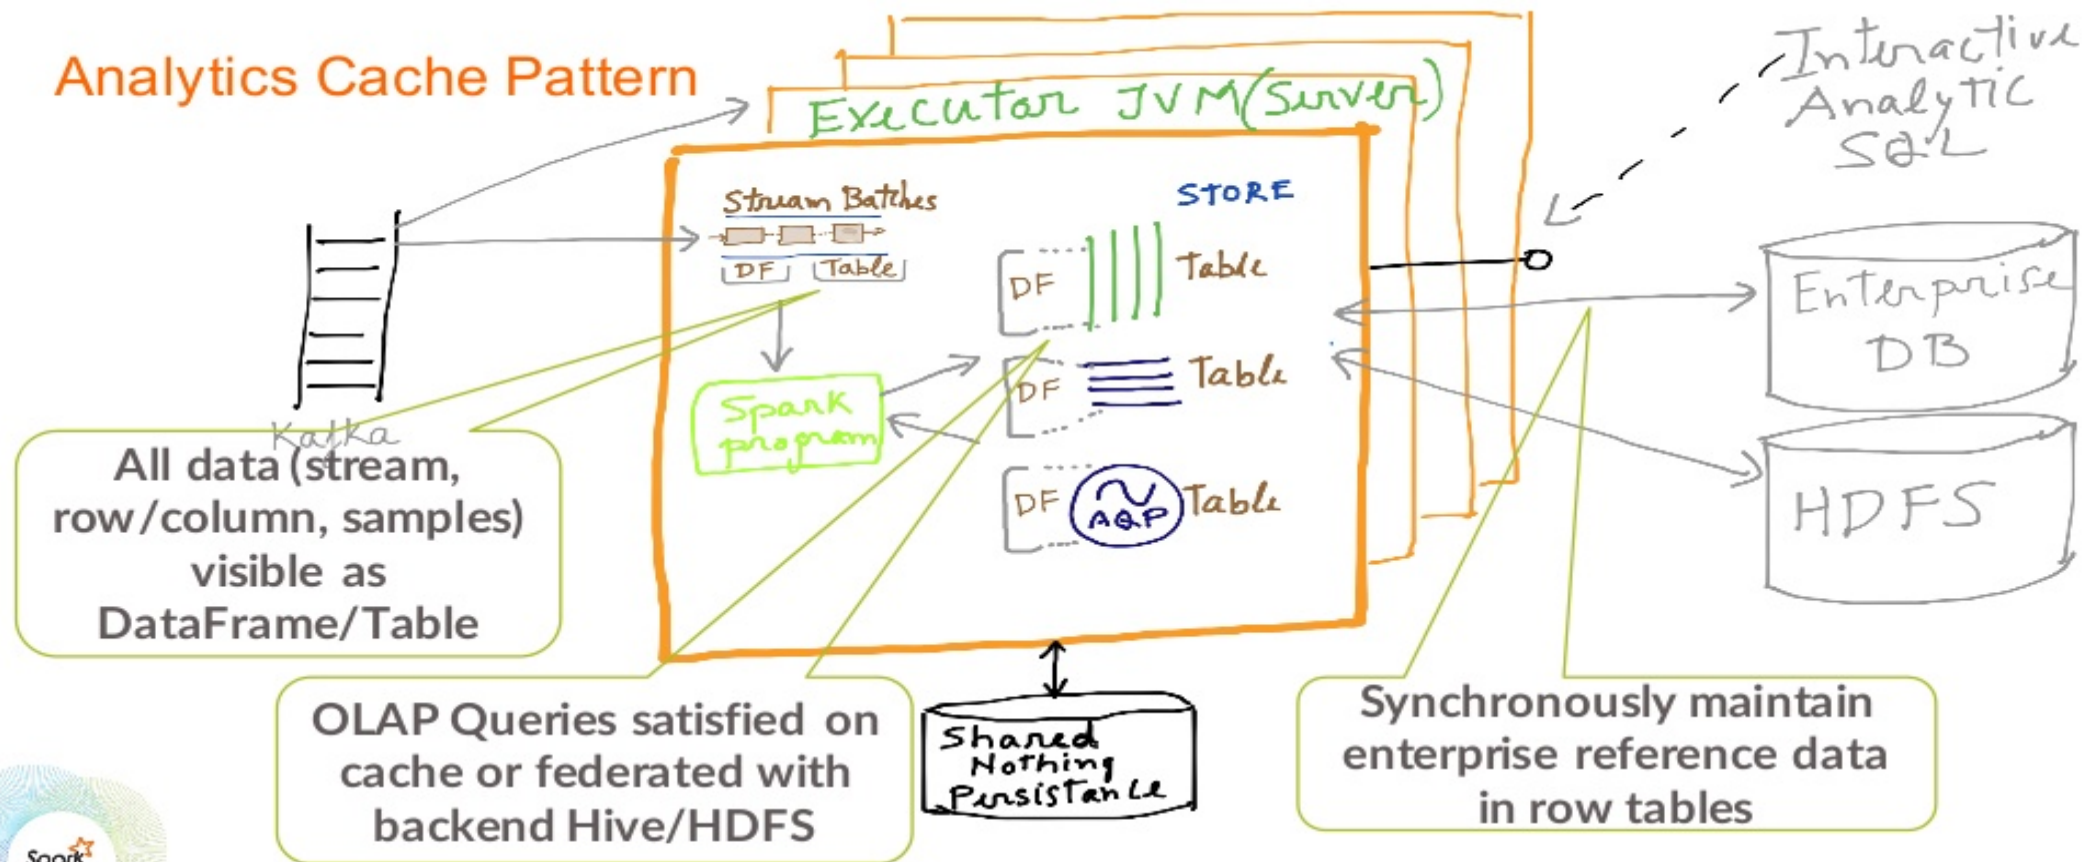


# Revisiting AdAnalytics – Spark with colocated store



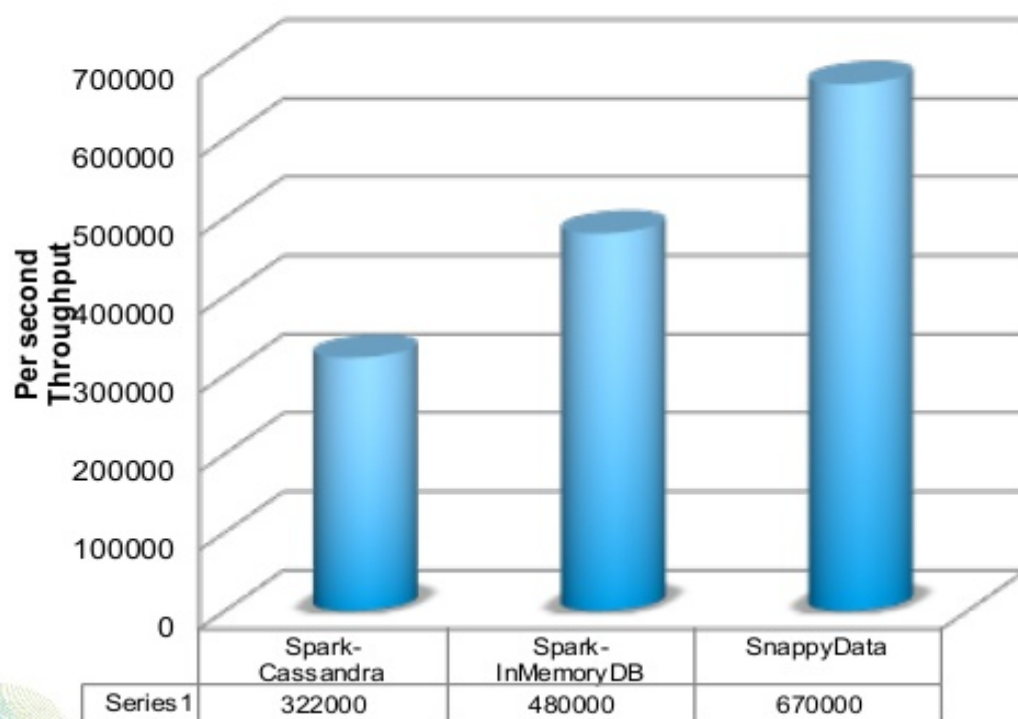
# Revisiting AdAnalytics – Spark with colocated store

## Analytics Cache Pattern



# Concurrent Ingest + Query Performance

**Stream ingestion rate**  
(On 4 nodes with cap on CPU to allow for queries)




- AWS 4 c4.2xlarge instances  
- 8 cores, 15GB mem
- Each node parallelly ingests stream from Kafka
- Parallel batch writes to store (32 partitions)
- Only few cores used for Stream writes as most of CPU reserved for OLAP queries



SPARK SUMMIT 2016

<https://github.com/SnappyDataInc/snappy-poc>

# Concurrent Ingest + Query Performance



Sample “scan” oriented OLAP query(Spark SQL) performance executed while ingesting data



SPARK SUMMIT 2016

<https://github.com/SnappyDataInc/snappy-poc>

# How SnappyData Extends Spark

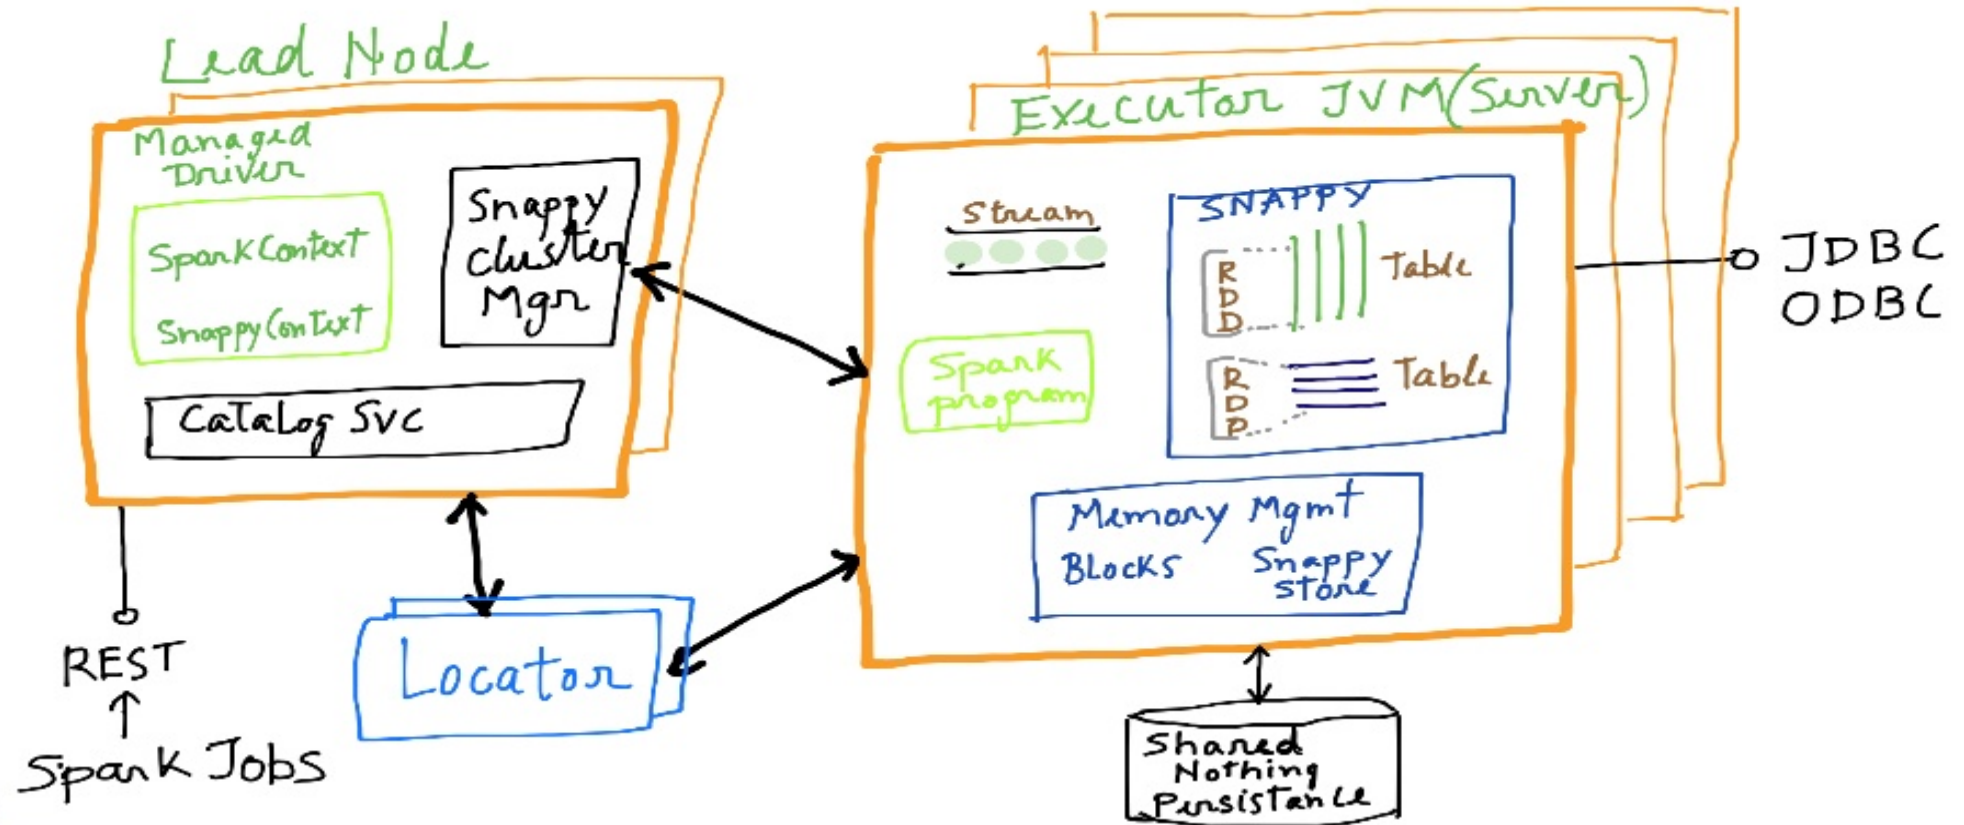


SPARK SUMMIT 2016

---

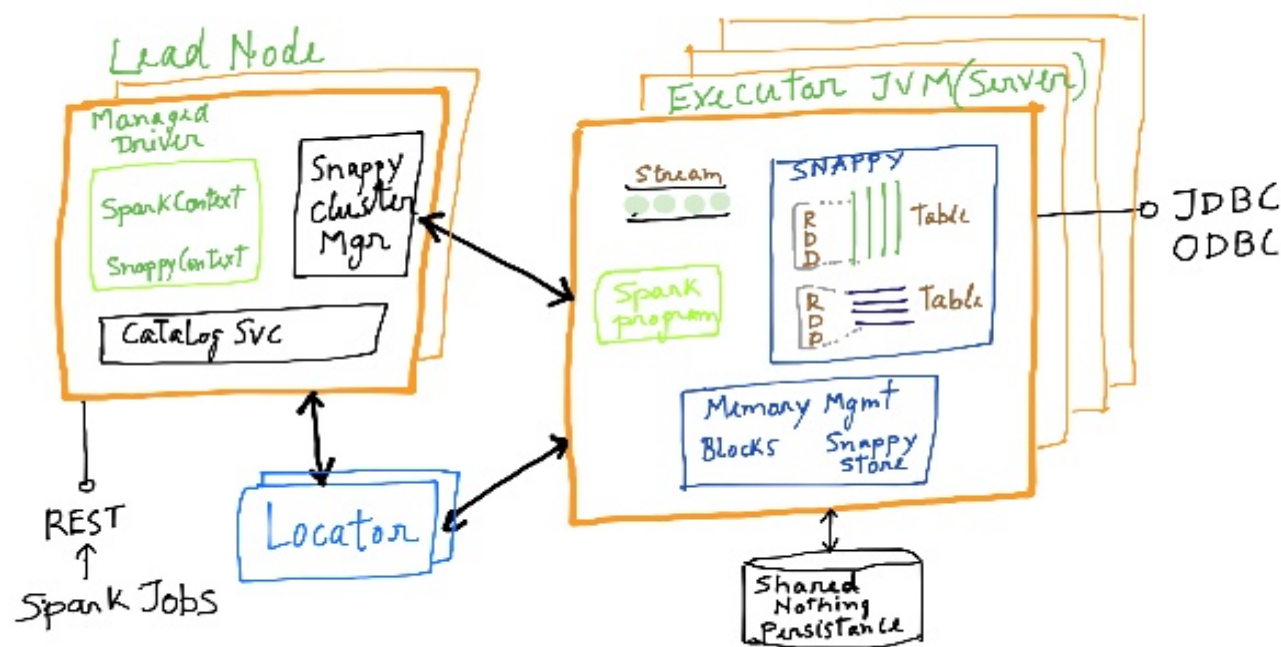


# Unified Cluster Architecture






# How do we extend Spark for Real Time?



- Spark Executors are long running. **Driver failure doesn't shutdown Executors**
- **Driver HA** – Drivers run “Managed” with standby secondary
- **Data HA** – Consensus based clustering integrated for eager replication

# How do we extend Spark for Real Time?



- By pass scheduler for low latency SQL
- **Deep integration with Spark Catalyst(SQL)** – collocation optimizations, indexing use, etc
- **Full SQL support** – Persistent Catalog, Transaction, DML



# Unified OLAP/OLTP/streaming with Spark

- Far fewer resources: TB problem becomes GB.
  - CPU contention drops
- Far less complex
  - single cluster for stream ingestion, continuous queries, interactive queries and machine learning
- Much faster
  - compressed data managed in distributed memory in columnar form reduces volume and is much more responsive



# SnappyData is Open Source

- [Ad Analytics example/benchmark -  
https://github.com/SnappyDataInc/snappy-poc](https://github.com/SnappyDataInc/snappy-poc)
- <https://github.com/SnappyDataInc/snappydata>



**Join SnappyData for pizza and drinks, tonight at *Tradition!***

JUNE 7, 2016 @ 7:30 PM

441 JONES ST, SAN FRANCISCO, CA  
(JUST AROUND THE CORNER)



# THANK YOU.

Drop by our booth to learn more.




**SPARK SUMMIT 2016**  
DATA SCIENCE AND ENGINEERING AT SCALE  
JUNE 6-8, 2016 SAN FRANCISCO




# Snappy Spark Cluster Deployment topologies

## Unified Cluster



- Snappy store and Spark Executor share the JVM memory
- Reference based access – zero copy

## Split Cluster



- SnappyStore is isolated but use the same COLUMN FORMAT AS SPARK for high throughput

# Simple API – Spark Compatible

- Access Table as DataFrame  
Catalog is automatically recovered

```
val impressionLogs: DataFrame = context.table(colTable)
val campaignRef: DataFrame = context.table(rowTable)
```

```
val parquetData: DataFrame = context.table(parquetTable)
<... Now use any of DataFrame APIs ... >
```

- Store RDD[T]/DataFrame can be stored in SnappyData tables
- Access from Remote SQL clients
- Additional API for updates, inserts, deletes

```
//Save a dataframe using the Snappy or spark context ...
context.createExternalTable("T1", "ROW", myDataFrame.schema,
props );
```

```
//save using DataFrame API
dataDF.write.format("ROW").mode(SaveMode.Append).options(prop
s).saveAsTable("T1");
```



# Extends Spark

```
CREATE [Temporary] TABLE [IF NOT EXISTS] table_name
(
    <column definition>
) USING 'JDBC | ROW | COLUMN'
OPTIONS (
    COLOCATE_WITH 'table_name',          // Default none
    PARTITION_BY 'PRIMARY KEY | column name', // will be a replicated table, by default
    REDUNDANCY '1',                      // Manage HA
    PERSISTENT "DISKSTORE_NAME ASYNCHRONOUS | SYNCHRONOUS",
    // Empty string will map to default disk store.
    OFFHEAP "true | false"
    EVICTION_BY "MEMSIZE 200 | COUNT 200 | HEAPPERCENT",
    ....
    [AS select_statement];
```



# Simple to Ingest Streams using SQL

