

Large Scale Multimedia Data Processing on Spark and Related Applications at Baidu

Quan Wang
Baidu USA



SPARK SUMMIT 2016
DATA SCIENCE AND ENGINEERING AT SCALE
JUNE 6-8, 2016 SAN FRANCISCO


Motivations

- Why Multimedia?
- Examples of large scale multimedia processing at Baidu:
 - HD map generation and simulation for self-driving cars
 - Image feature exaction and transform for CTR predictions



Application Example: Self-Driving Cars

- Maps for navigation, planning and localization:

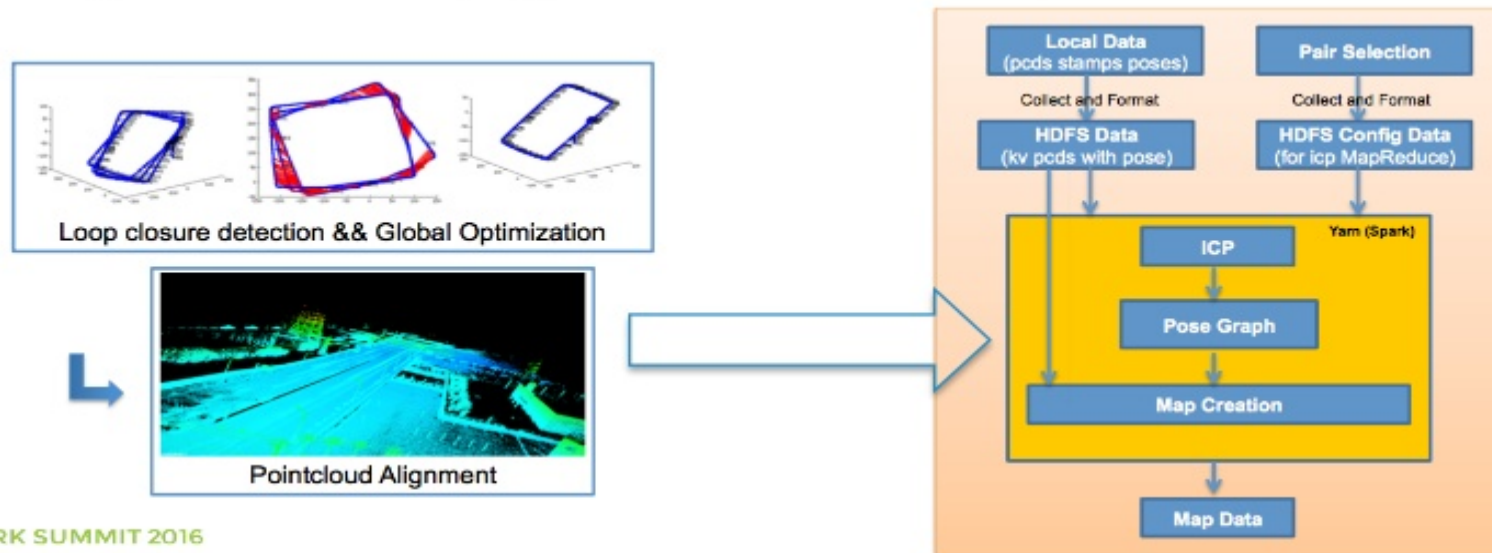


- HD Map generation pipeline:



Data challenge in Self-driving Cars Project

- Backend map generation and simulation:
 - Point clouds inputs of 40MB/s for one single 3D LiDAR sensor
 - Counting all 2D/3D sensors => TB level of data, per hour, per car
- Example: HD map generation on Spark



Another Example: Baidu Image Search

- Billions of images need feature extractions & transformations for deep learning applications:
 - Image recognition and classification
 - Ranking for best picture to show
 - CTR prediction



Challenges

- Core functions for feature extraction
- Efficient large scale distributed executions of feature extraction with multimedia input support
- Plug and play for any feature extraction executable never designed for Spark; Flexible and easy to use for platform users

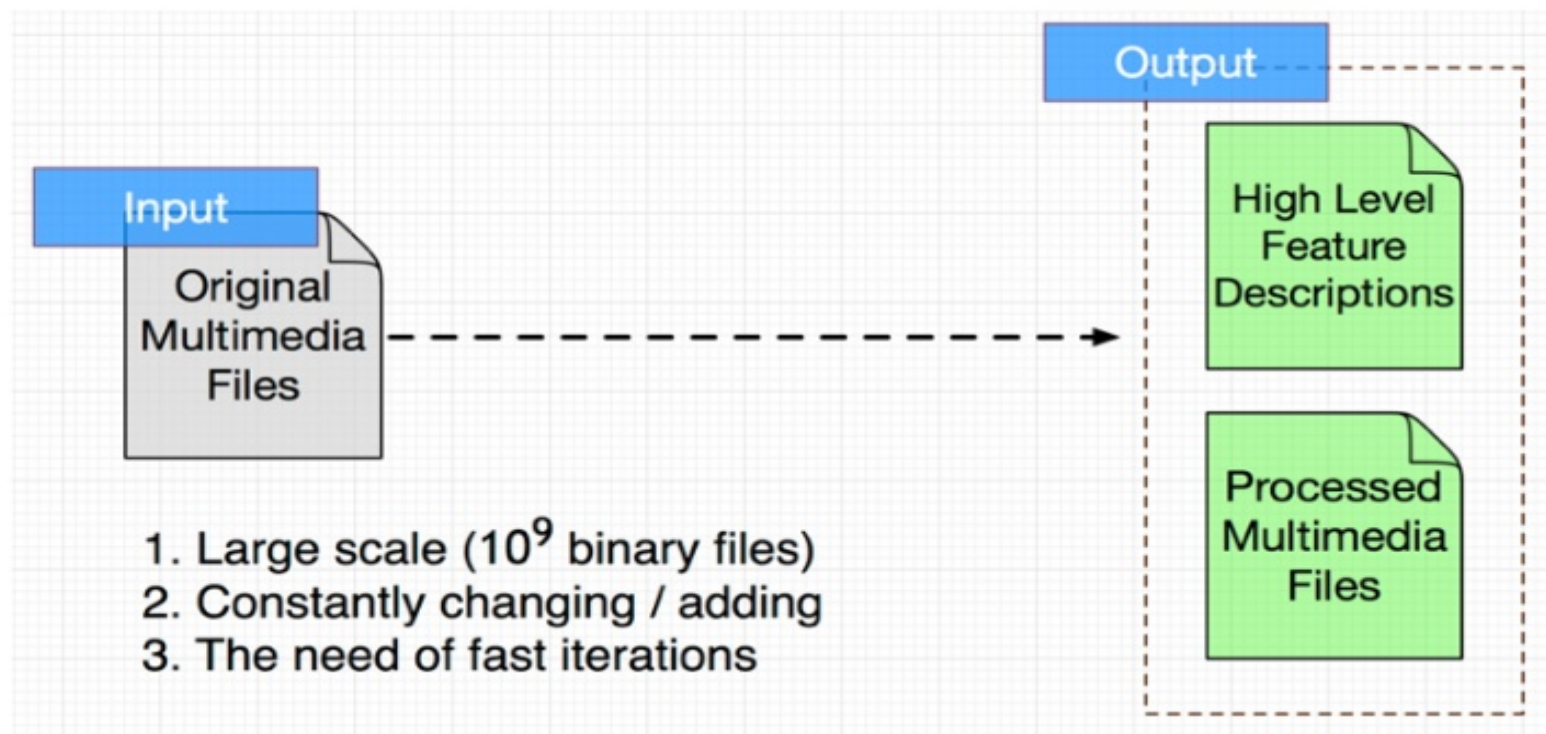


Feature Extraction Core Function

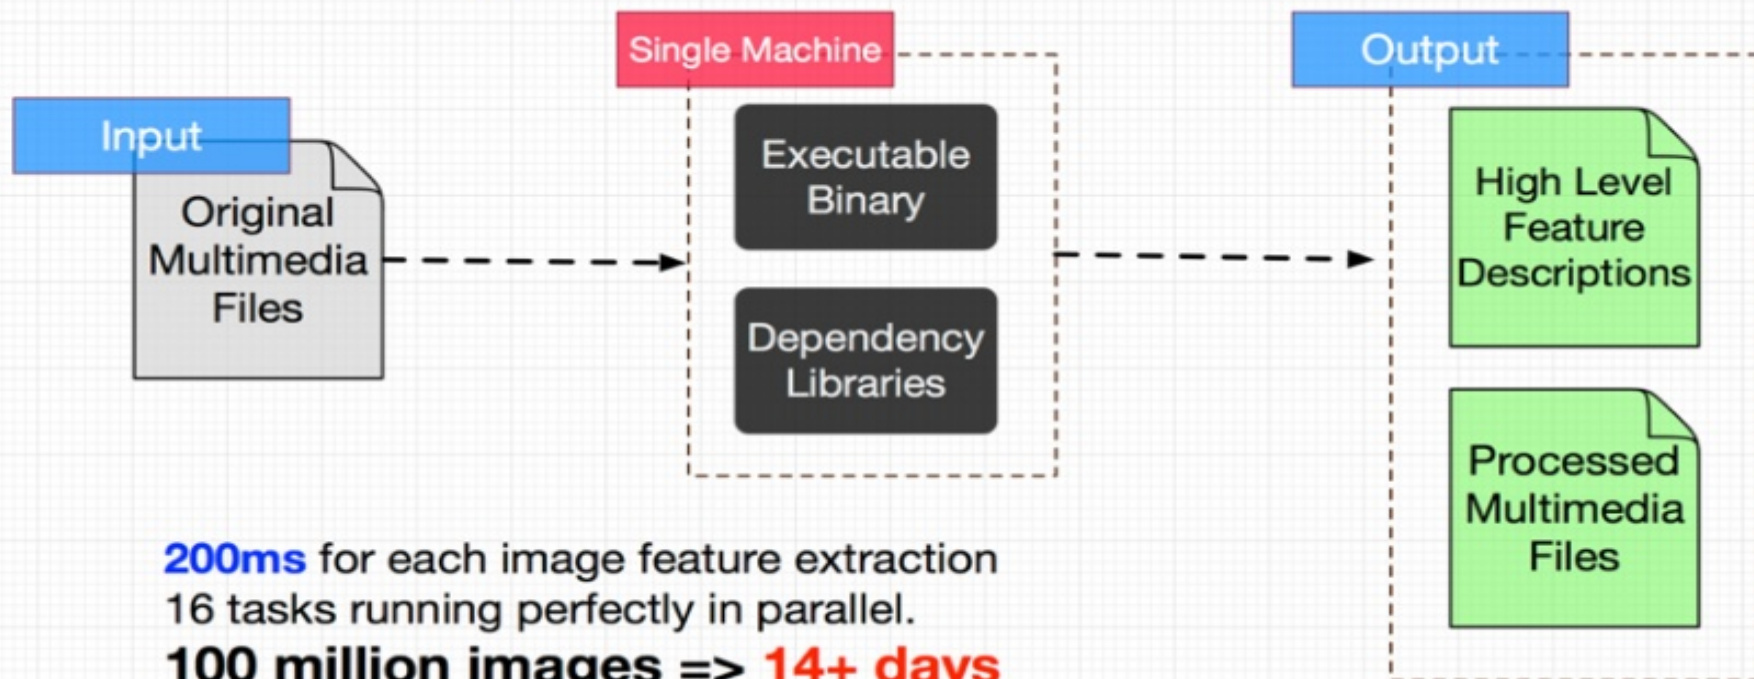
- Feature extraction C++ program depends on CDNN + OpenCV library
 - Compute the per pixel difference based on a pre-computed mean
 - Feed the difference values into a pre-computed CDNN model
 - Produce image features after multi-layers of computations
- Need streaming/pipe based function support



End to End Requirements



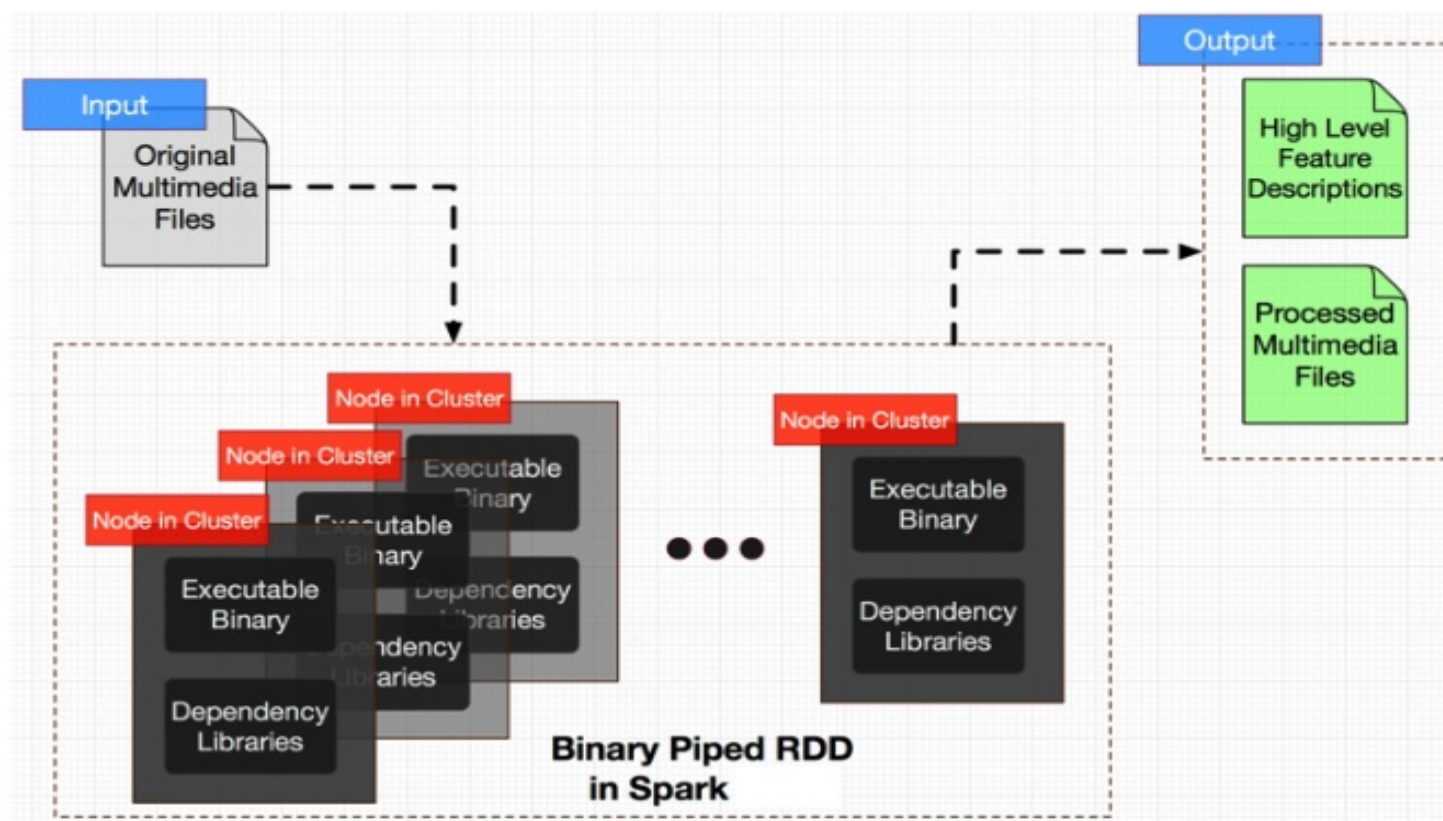
Single Node Execution



Distribute to 500 machines
How to manage?

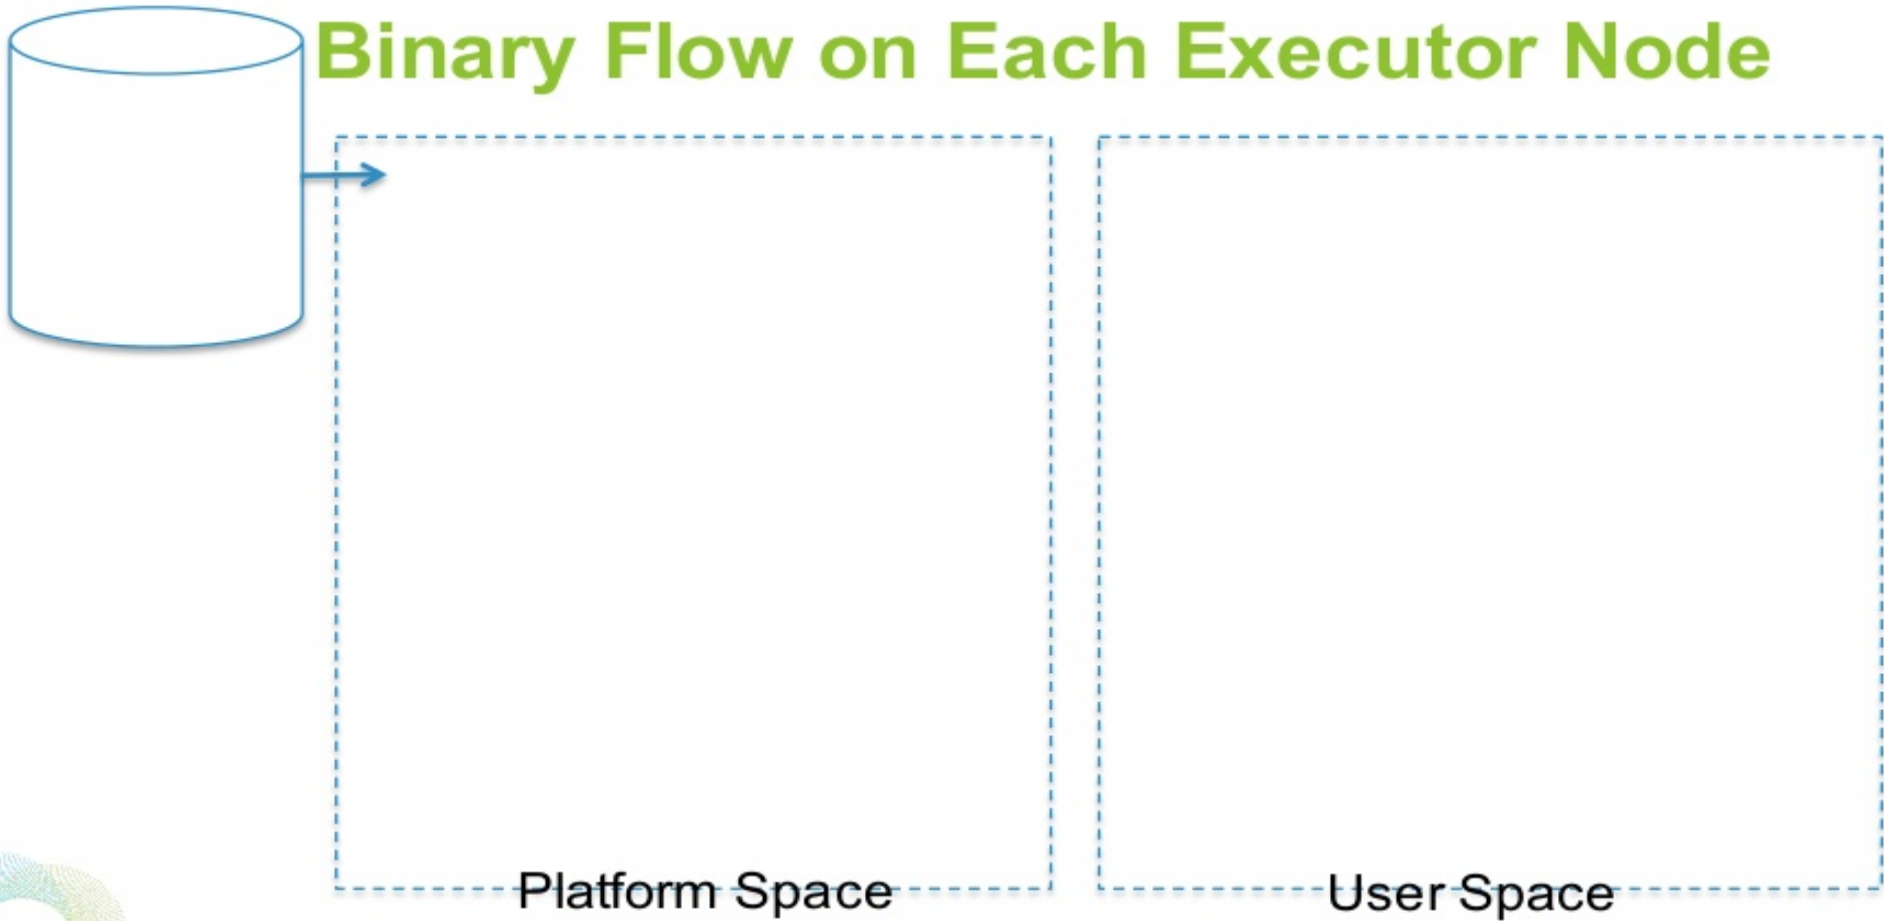


Distributed Execution

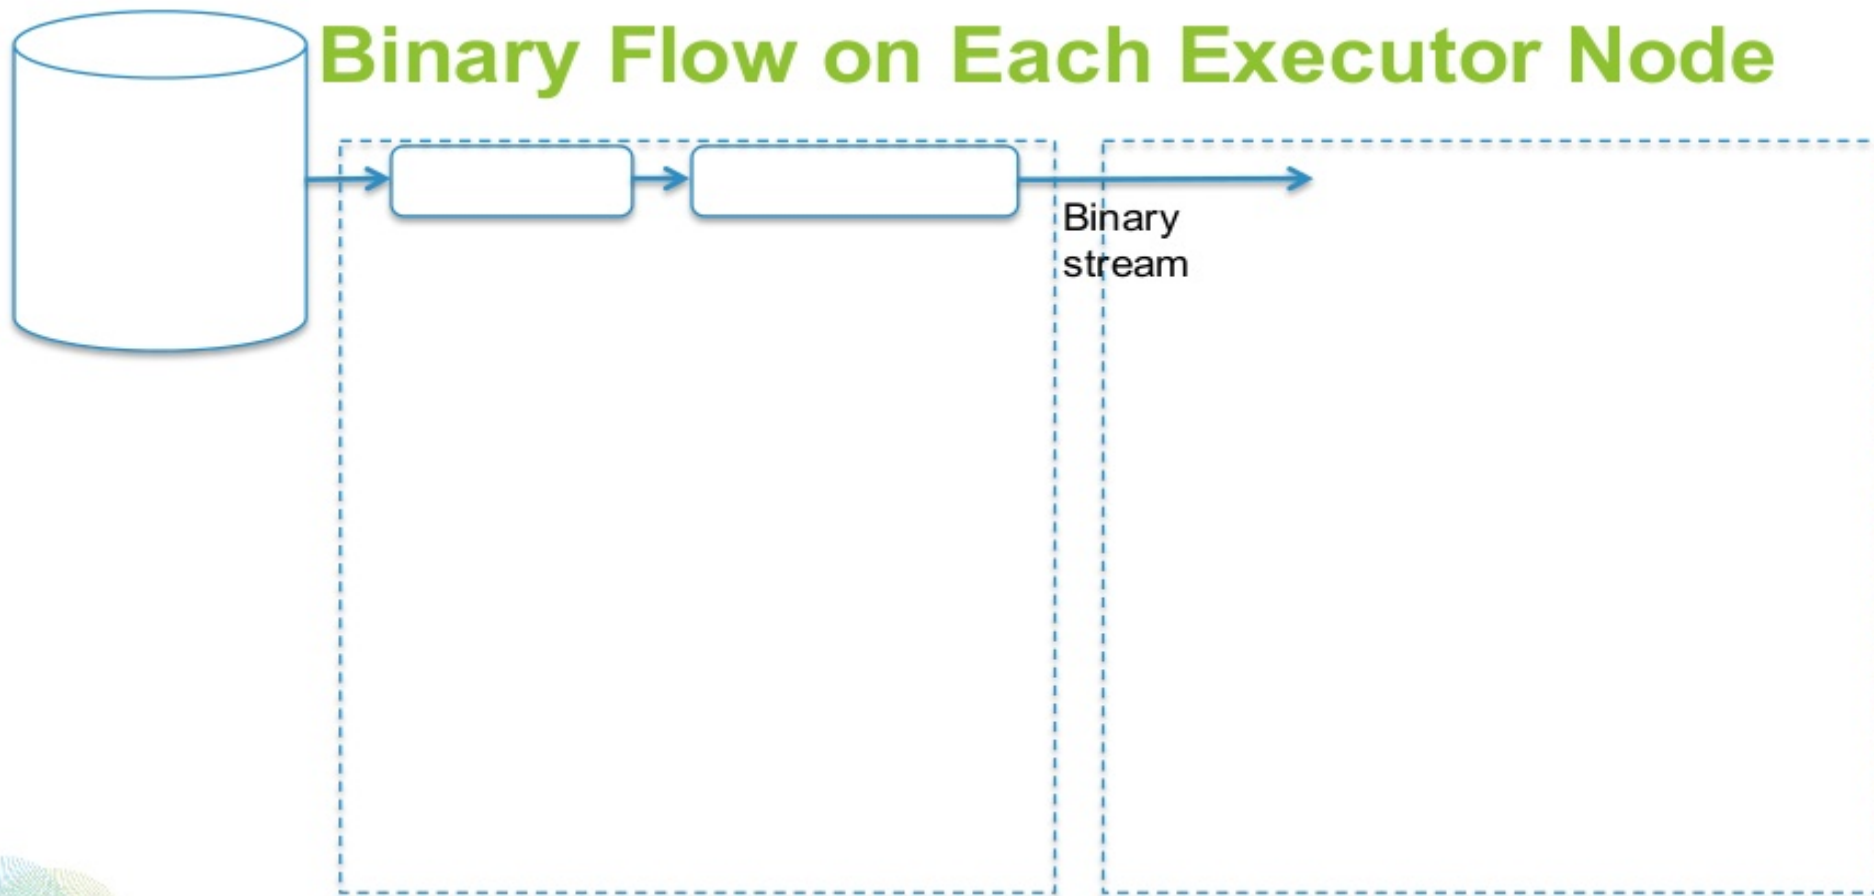


Technical Details:

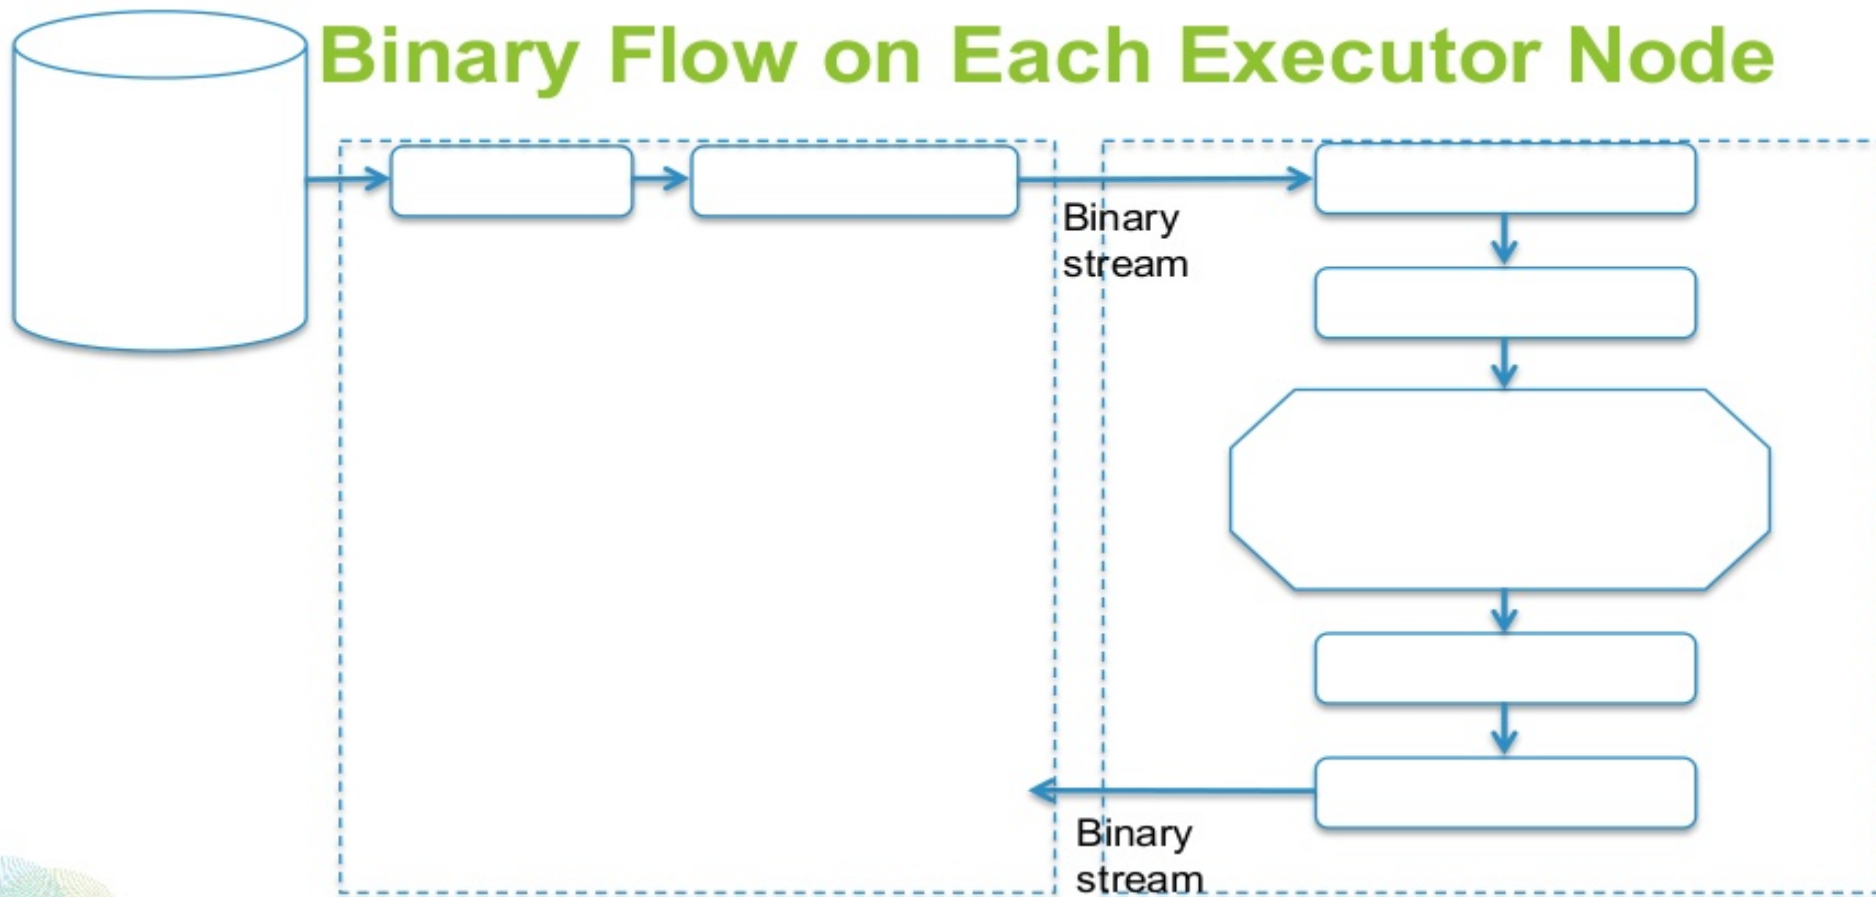
Binary Flow on Each Executor Node



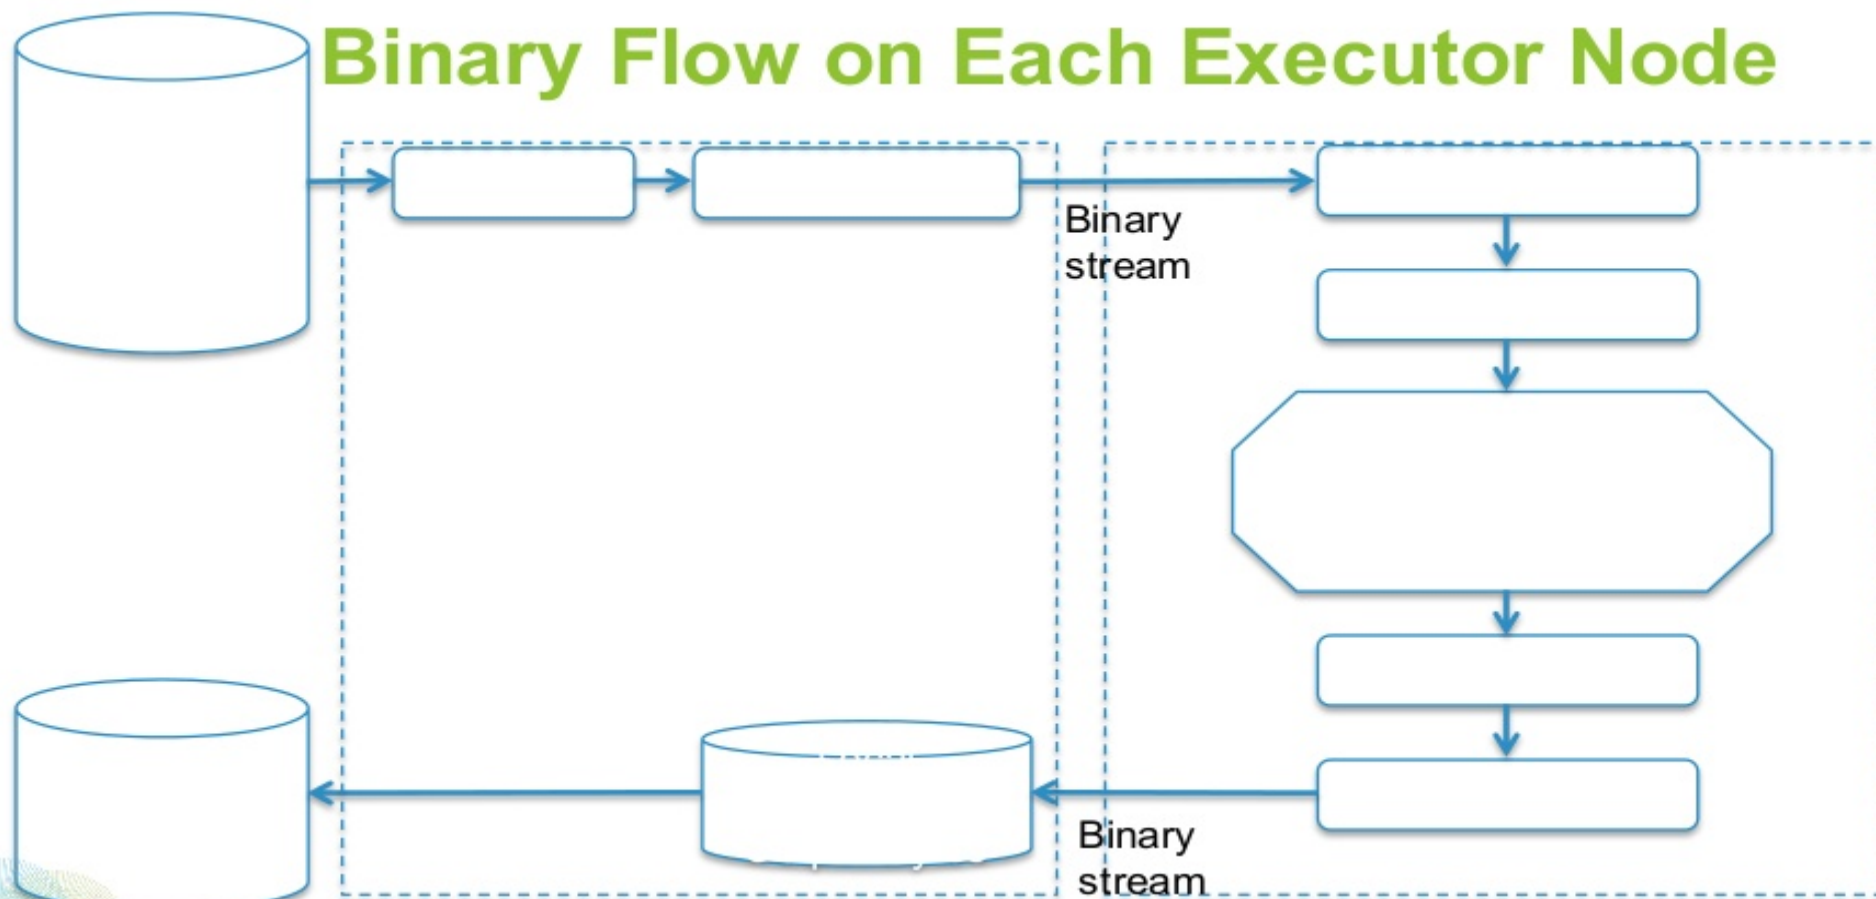
Technical Details: Binary Flow on Each Executor Node



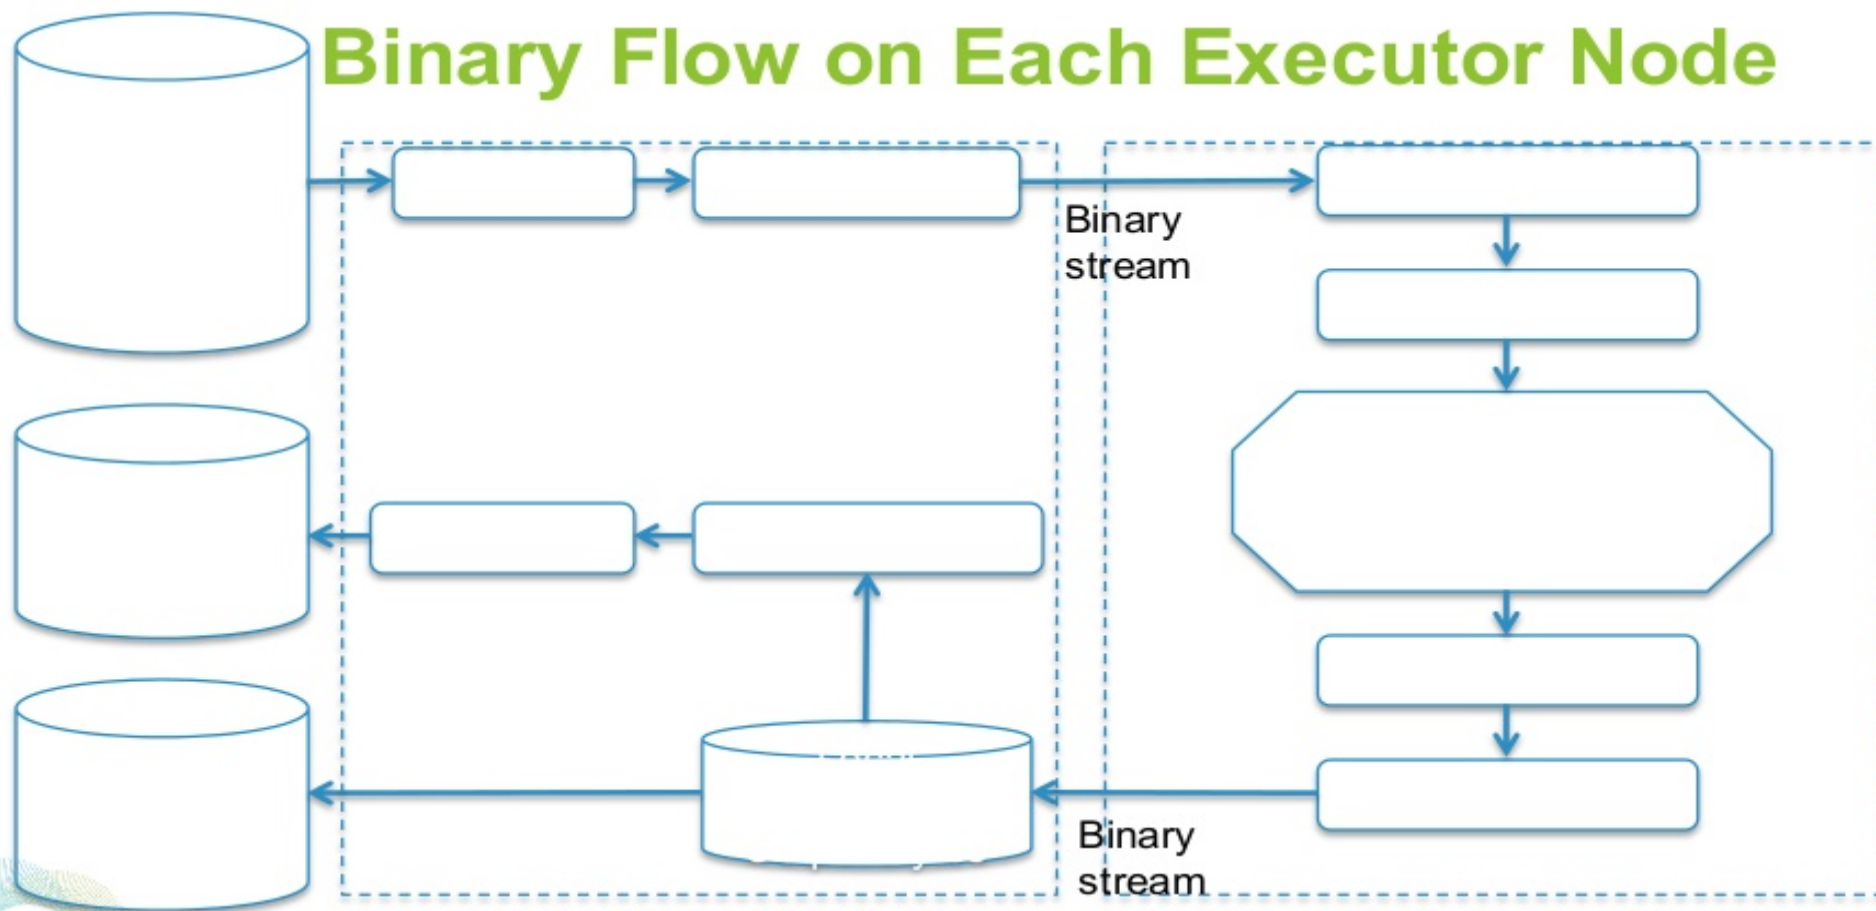
Technical Details: Binary Flow on Each Executor Node



Technical Details: Binary Flow on Each Executor Node



Technical Details: Binary Flow on Each Executor Node



Implementation Highlights

- All data serialization and encoding inside Spark RDD, thus linear scalable
- Flexible output format (feature vectors, ranking scores, processed binaries, etc.)
- Easy to plug in customized encoding/serialization functions directly into platform
- Support of passing spark internal information (e.g. partition id, task attempt id) into user program



Flexibility


```
/**
 * A Binary Piped RDD which pipes the contents of each parent partition to an external command in the format
 * of binary streams and returns the output as a collection of bytes.
 *
 * @param savePath Optional file system path to save the decoded and de-serialized output bytes,
 *                  only needed for the saveResults API
 * @param printPipeContext Optional plug in function to add a context header to the binary stream
 * @param printRDDElement Optional plug in function for customized serialization and encoding
 * @param saveRDDElement Optional plug in function for customized decoding and de-serialization
 */
class BinPipedRDD[T: ClassTag](
    prev: RDD[(String, PortableDataStream)],
    command: Seq[String],
    savePath: String,
    envVars: Map[String, String],
    printPipeContext: (Array[Byte] => Unit) => Unit,
    printRDDElement: ((String, PortableDataStream), Array[Byte] => Unit) => Unit,
    saveRDDElement: (Iterator[Byte]) => Unit,
    separateWorkingDir: Boolean)
    extends RDD[Byte](prev) {
```

Flexible output format (feature vectors, ranking scores, processed binaries, etc.)
Easy to plug in Customized encoding/serialization functions directly into Spark



Performance Results

Running on a Spark cluster, over 11k images inputs with archive function on, running feature extraction on each image



5.6 images per core per second, and linear scalable~!

Going to production: 100 million image files, 10000 cores cluster => less than 0.5 hour

In production: daily incremental 3 million images, 500 cores cluster => within 20 minutes



Conclusion

- General data intelligence and analysis
 - Binary input format + Pipe based bin/lib execution

Missing functionality in Spark/Hadoop

- Introduce the Binary Piped RDD for:
 - Platform level abstraction of input data format in their **original binary form**
 - Seamless streaming to and from **existing executable/libraries** for high level data analysis and understanding
 - **Linear scalability** with input data



THANK YOU.

Questions?



SPARK SUMMIT 2016
DATA SCIENCE AND ENGINEERING AT SCALE
JUNE 6-8, 2016 SAN FRANCISCO