


Tegra: Time-evolving Graph Processing on Commodity Clusters

Anand Iyer, Ion Stoica

Presented by Ankur Dave




Graphs are everywhere...



Social Networks

Graphs are everywhere...




Gnutella network subgraph


Graphs are everywhere...



Graphs are everywhere...



Metabolic network of a single cell organism




Tuberculosis

Plenty of interest in processing them

- Graph DBMS 25% of all enterprises by 2017¹
- Many open-source and research prototypes on distributed graph processing frameworks: Giraph, Pregel, GraphLab, Chaos, GraphX, ...

Real-world Graphs are Dynamic


A world map with a dark background, overlaid with numerous blue circular markers of varying sizes. These markers represent earthquake occurrences and their density patterns. The markers are most concentrated in the western United States, the western Pacific, and along the Mid-Atlantic Ridge. Other notable clusters are visible in Japan, the Philippines, and the Indian subcontinent. The map includes labels for various countries and regions.

Many interesting business and research insights possible by processing them...


USGS Earthquake
Occurrence
Density Patterns

...but little work on incrementally updating
computation results or on window-based operations

A Motivating Example...



A Motivating Example...



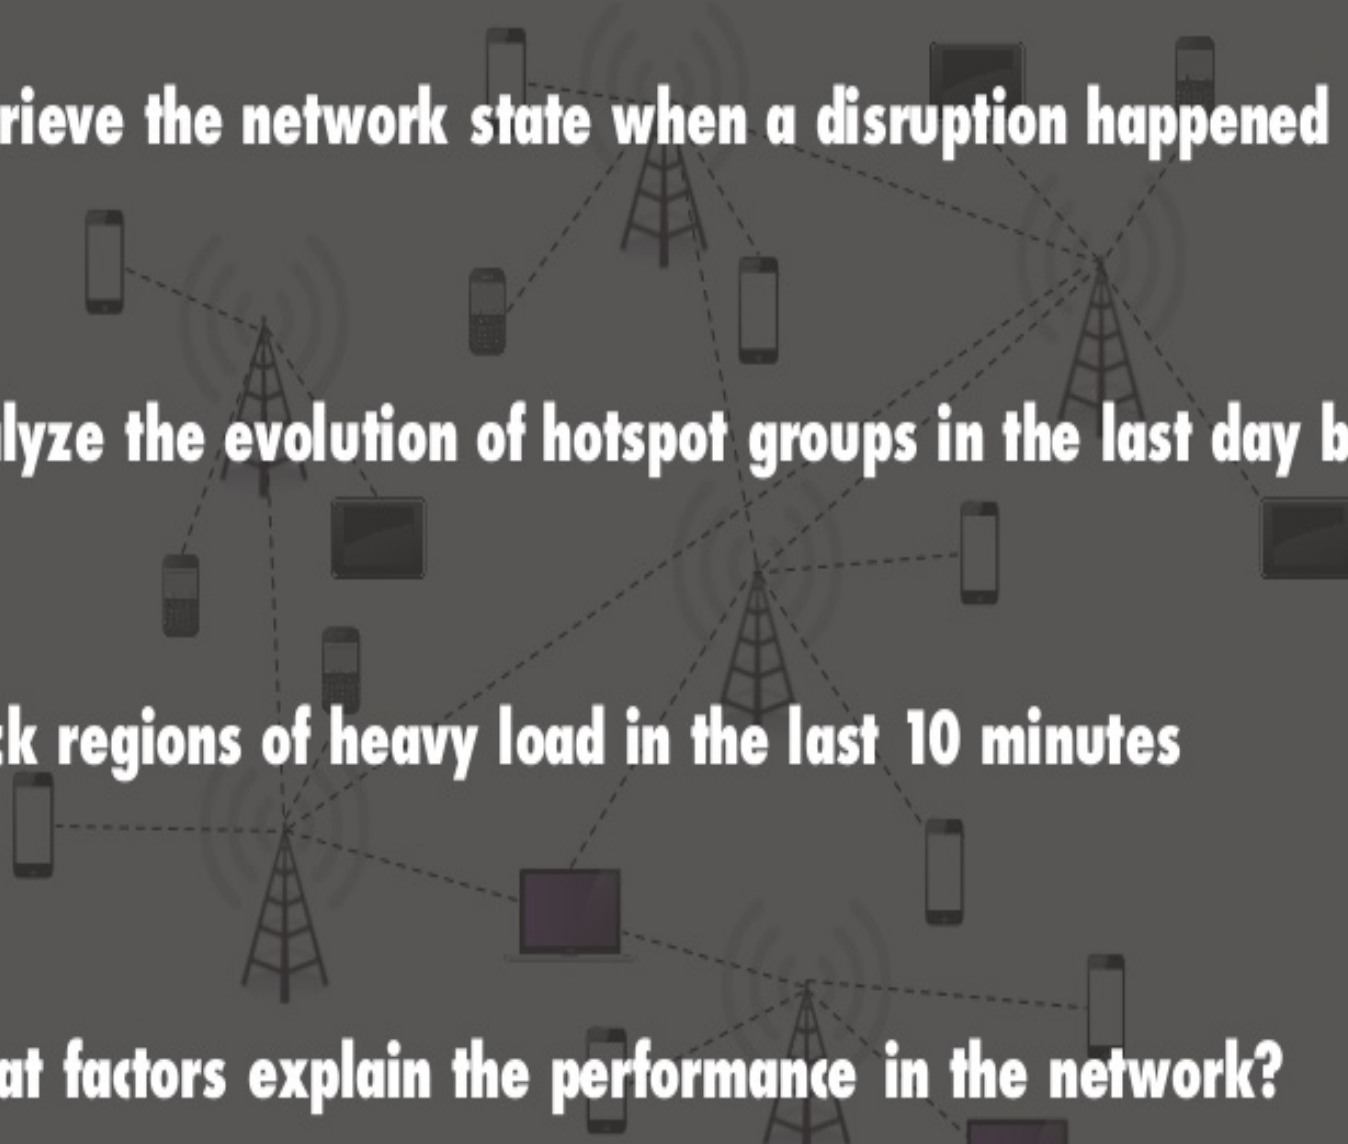
A Motivating Example...

Retrieve the network state when a disruption happened

Analyze the evolution of hotspot groups in the last day by hour

Track regions of heavy load in the last 10 minutes

What factors explain the performance in the network?



Tegra


How do we perform efficient computations on time-evolving, dynamically changing graphs?

Goals


- **Create and manage time-evolving graphs**
 - *Retrieve the network state when a disruption happened*
- **Temporal analytics on windows**
 - *Analyze the evolution of hotspot groups in the last day by hour*
- **Sliding window computations**
 - *Track regions of heavy load in the last 10 minutes interval*
- **Mix graph and data parallel computing**
 - *What factors explain the performance in the network*

Tegra


Graph Snapshot Index




Timelapse Abstraction



Lightweight Incremental Computations




Pause-Shift-Resume Computations



Goals

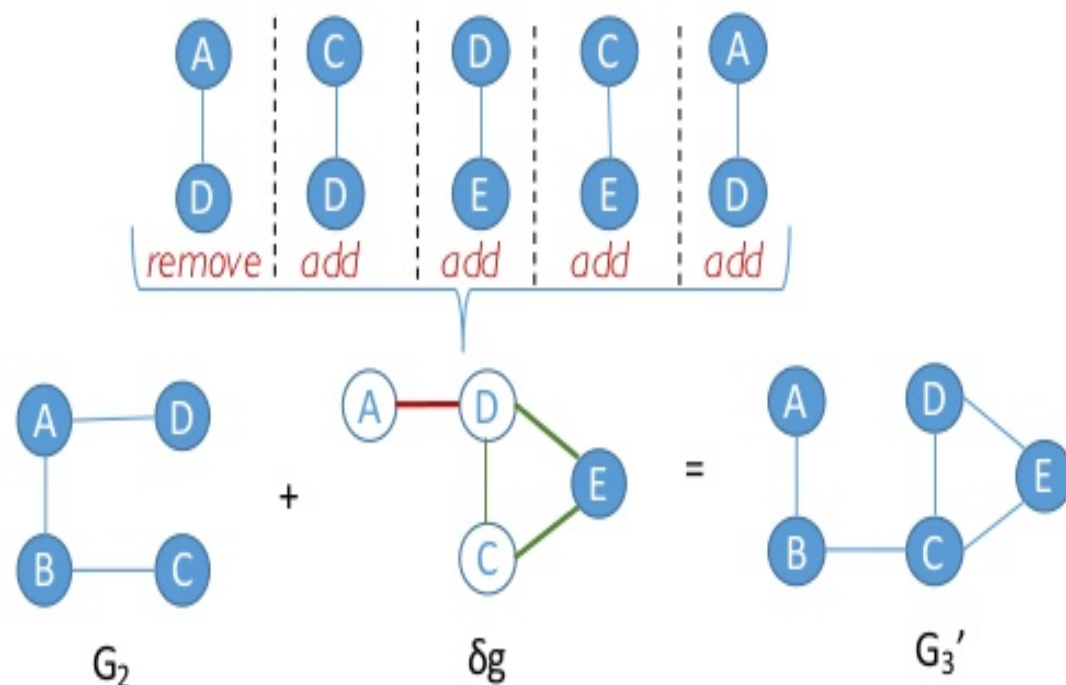
- Create and manage time-evolving graphs using Graph Snapshot Index
 - *Retrieve the network state when a disruption happened*
- Temporal analytics on windows
 - *Analyze the evolution of hotspot groups in the last day by hour*
- Sliding window computations
 - *Track regions of heavy load in the last 10 minutes interval*
- Mix graph and data parallel computing
 - *What factors explain the performance in the network*

Representing Evolving Graphs





Graph Composition

Updating graphs depend on application semantics




Maintaining Multiple Snapshots

Store entire
snapshots



+ Efficient retrieval
- Storage overhead


Store only deltas



+ Efficient storage
- Retrieval overhead

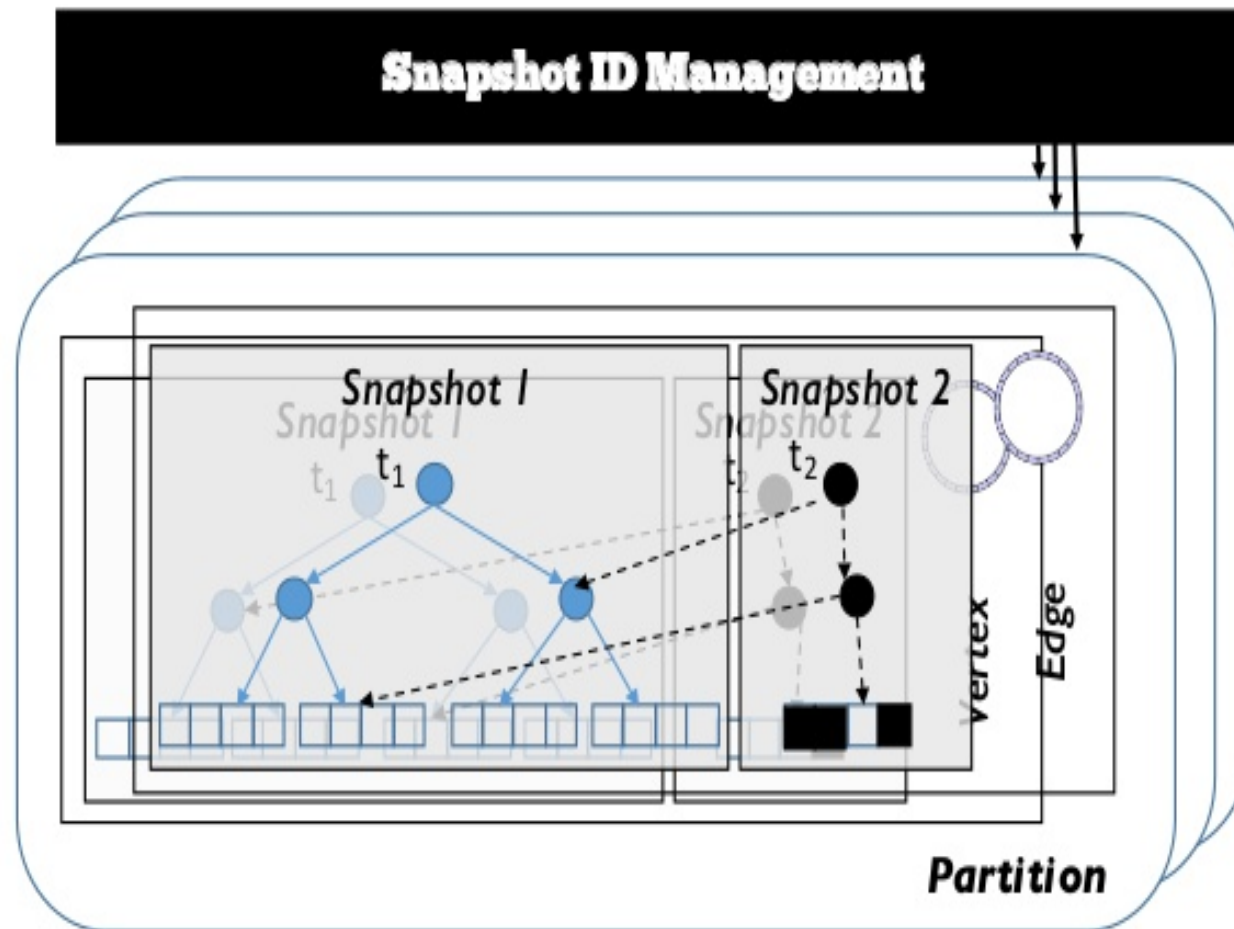
Maintaining Multiple Snapshots

Use a structure-sharing persistent data-structure.




Persistent Adaptive Radix Tree (PART)
is one solution available for Spark.

Graph Snapshot Index



Graph Windowing




Goals

- Create and manage time-evolving graphs using Distributed Graph Snapshot Index
 - *Retrieve the network state when a disruption happened*
- **Temporal analytics on windows using Timelapse Abstraction**
 - *Analyze the evolution of hotspot groups in the last day by hour*
- Sliding window computations
 - *Track regions of heavy load in the last 10 minutes interval*
- Mix graph and data parallel computing
 - *What factors explain the performance in the network*

Graph Parallel Computation

Many approaches

- Vertex centric, edge centric, subgraph centric, ...




Timelapse

Operations on windows of snapshots result in redundant computation

Timelapse

Instead, expose the temporal evolution of a node




- Significant reduction in messages exchanged between graph nodes
- Avoids redundant computations

Timelapse API


```
class Graph[V, E] {  
  // Collection views  
  def vertices(sid: Int): Collection[(Id, V)]  
  def edges(sid: Int): Collection[(Id, Id, E)]  
  def triplets(sid: Int): Collection[Triplet]  
  // Graph-parallel computation  
  def mrTriplets(f: (Triplet) => M,  
    sum: (M, M) => M,  
    sids: Array[Int]): Collection[(Int, Id, M)]  
  // Convenience functions  
  def mapV(f: (Id, V) => V,  
    sids: Array[Int]): Graph[V, E]  
  def mapE(f: (Id, Id, E) => E,  
    sids: Array[Int]): Graph[V, E]  
  def leftJoinV(v: Collection[(Id, V)],  
    f: (Id, V, V) => V,  
    sids: Array[Int]): Graph[V, E]  
  def leftJoinE(e: Collection[(Id, Id, E)],  
    f: (Id, Id, E, E) => E,  
    sids: Array[Int]): Graph[V, E]  
  def subgraph(vPred: (Id, V) => Boolean,  
    ePred: (Triplet) => Boolean,  
    sids: Array[Int]): Graph[V, E]  
  def reverse(sids: Array[Int]): Graph[V, E]  
}
```

Temporal Operations

Bulk Transformations



Bulk Iterative Computations



How did the hotspots change over this window?


Evolution analysis with no state keeping requirements

Goals


- Create and manage time-evolving graphs using Distributed Graph Snapshot Index
 - *Retrieve the network state when a disruption happened*
- Temporal analytics on windows using Timelapse Abstraction
 - *Analyze the evolution of hotspot groups in the last day by hour*
- **Sliding window computations**
 - *Track regions of heavy load in the last 10 minutes interval*
- Mix graph and data parallel computing
 - *What factors explain the performance in the network*

Incremental Computation

- If results from a previous snapshot is available, how can we reuse them?
- Three approaches in the past:
 - Restart the algorithm
 - Redundant computations
 - Memoization (GraphInc¹)
 - Too much state
 - Operator-wise state (Naiad^{2,3})
 - Too much overhead




Incremental Computation



- Can keep state as an efficient time-evolving graph
- Not limited to vertex-centric computations

Incremental Computation

- Some iterative graph algorithms are robust to graph changes
 - Allow them to proceed without keeping any state




Goals

- Create and manage time-evolving graphs using Distributed Graph Snapshot Index
 - *Retrieve the network state when a disruption happened*
- Temporal analytics on windows using Timelapse Abstraction
 - *Analyze the evolution of hotspot groups in the last day by hour*
- Sliding window computations
 - *Track regions of heavy load in the last 10 minutes interval*
- **Mix graph and data parallel computing**
 - *What factors explain the performance in the network*

Implementation & Evaluation

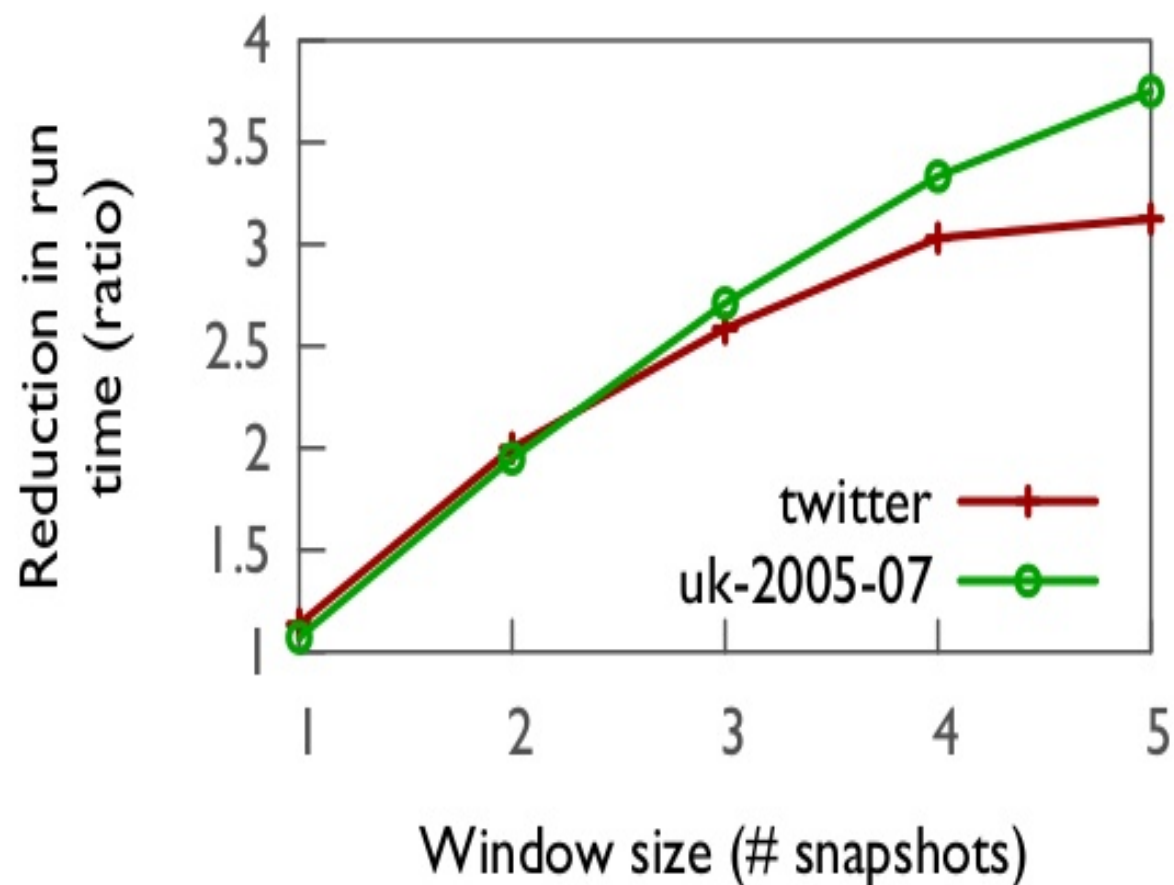
- Implemented as a major enhancement to GraphX
- Evaluated on two open source real-world graphs
 - **Twitter:** 41,652,230 vertices, 1,468,365,182 edges
 - **uk-2007:** 105,896,555 vertices, 3,738,733,648 edges

Preliminary Evaluation



Tegra can pack more snapshots in memory

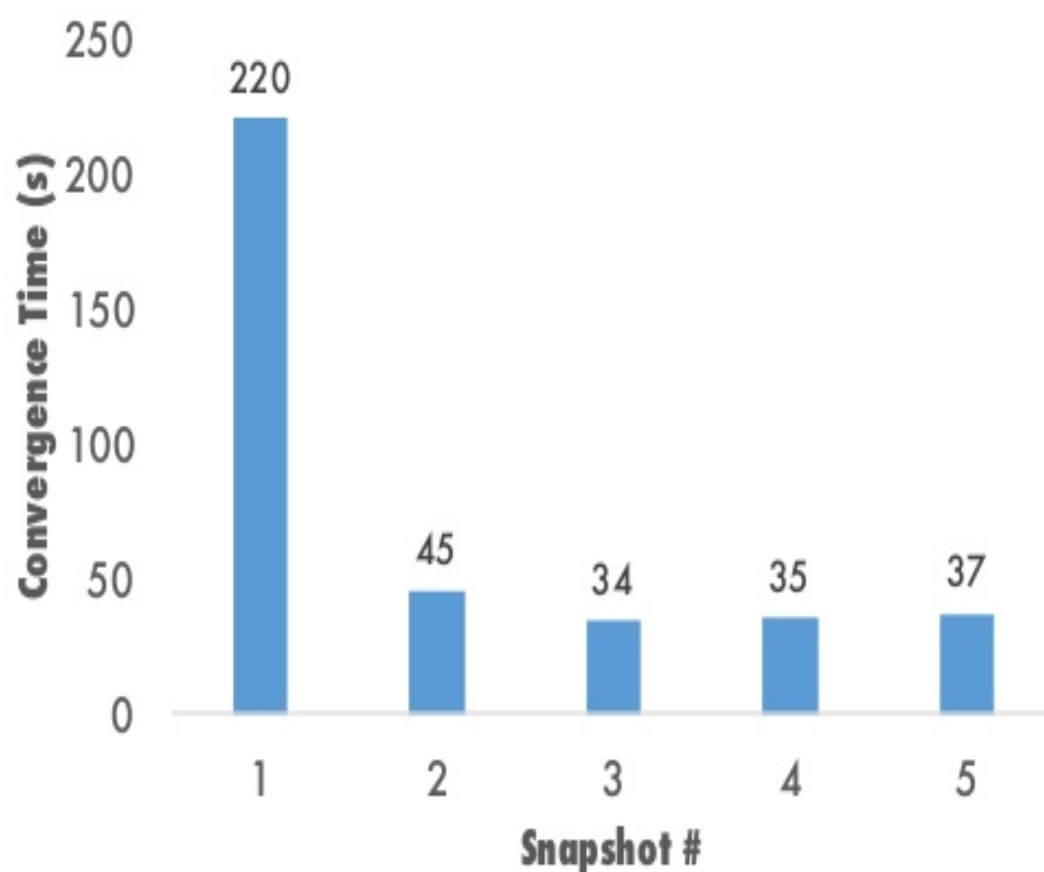
Preliminary Evaluation



Timelapse results in run time reduction

Preliminary Evaluation

Effectiveness of incremental computation



Summary

- Processing time-evolving graph efficiently can be useful.
- Efficient storage of multiple snapshots and reducing communication between graph nodes key to evolving graph analysis.

Ongoing Work

- Expand timelapse and its incremental computation model to other graph-parallel paradigms
 - Other interesting graph algorithms:
 - Fraud detection/prediction/incremental pattern matching
- Add graph querying support
 - Graph queries and analytics in a single system
- **Stay tuned for code release!**