

# Conviva Unified Framework (CUF) for Real Time, Near Real Time and Offline Analysis of Video Streaming With Spark and Databricks

Jibin Zhan  
CONVIVA®



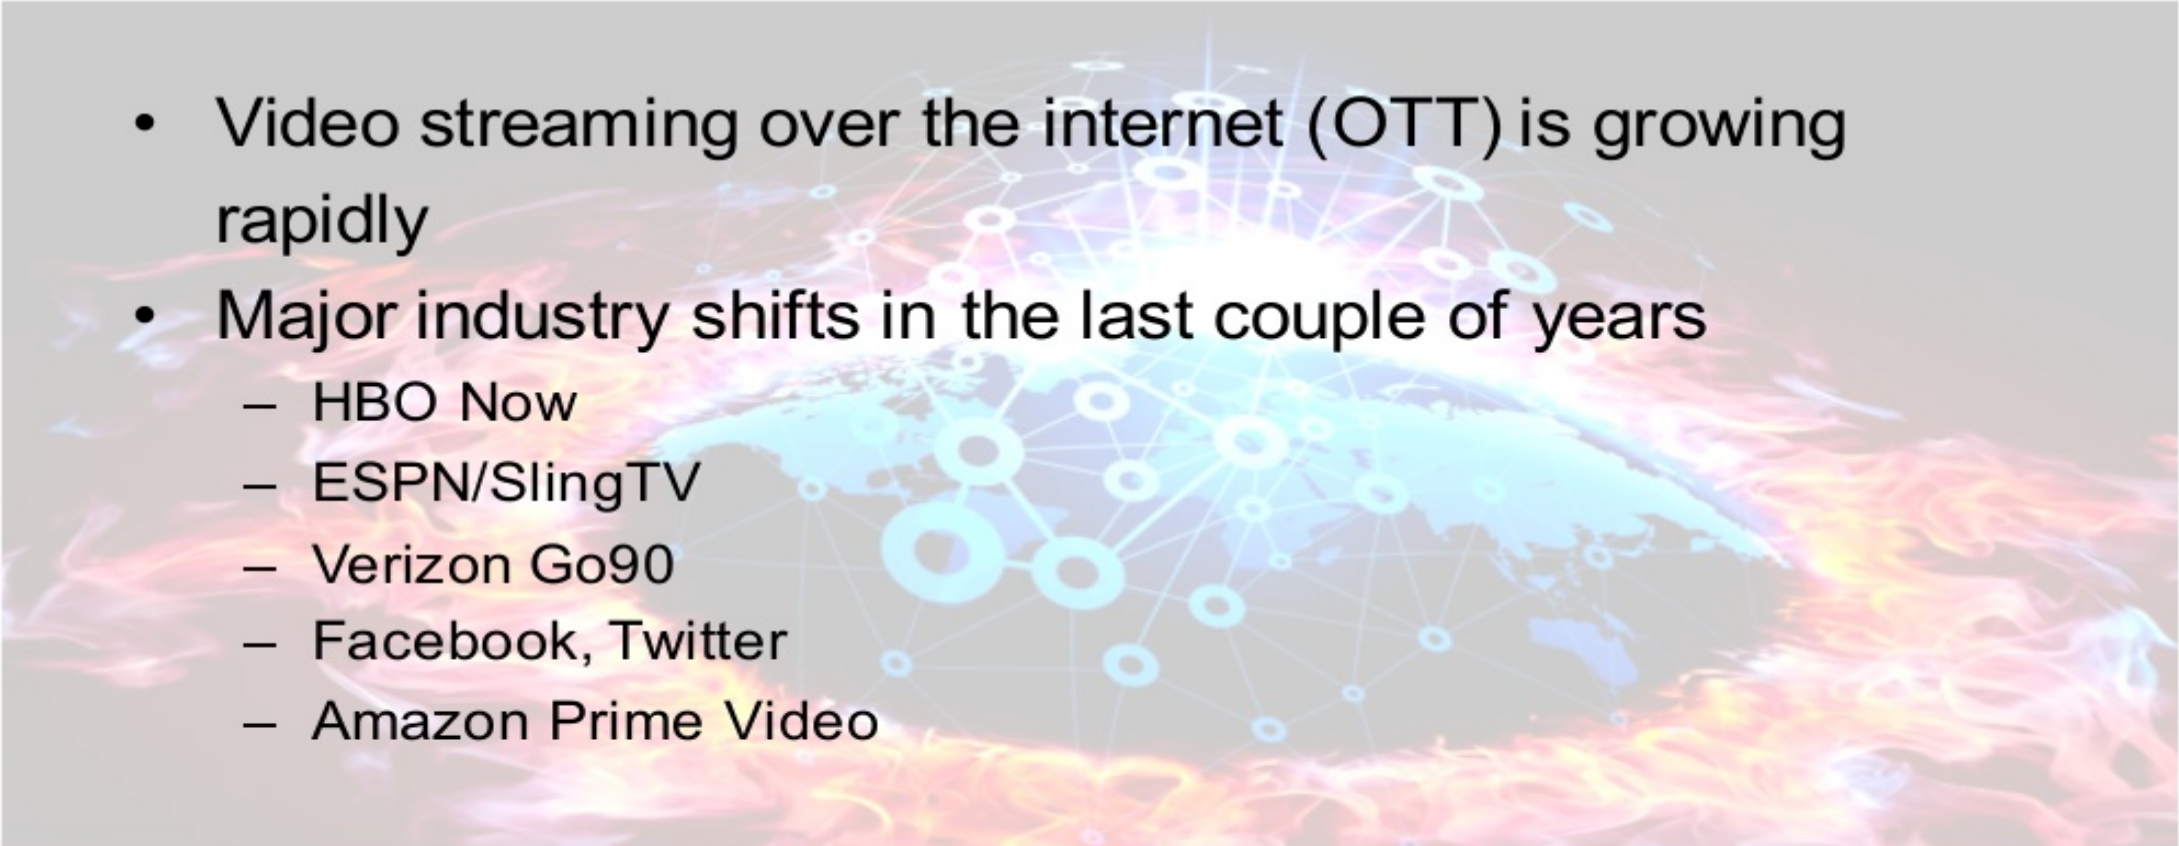
SPARK SUMMIT 2016  
DATA SCIENCE AND ENGINEERING AT SCALE  
JUNE 6-8, 2016 SAN FRANCISCO

# CONVIVA<sup>®</sup>

is a video **experience**  
management platform  
that maximizes viewer  
**engagement**

Unleashing the Power of OTT



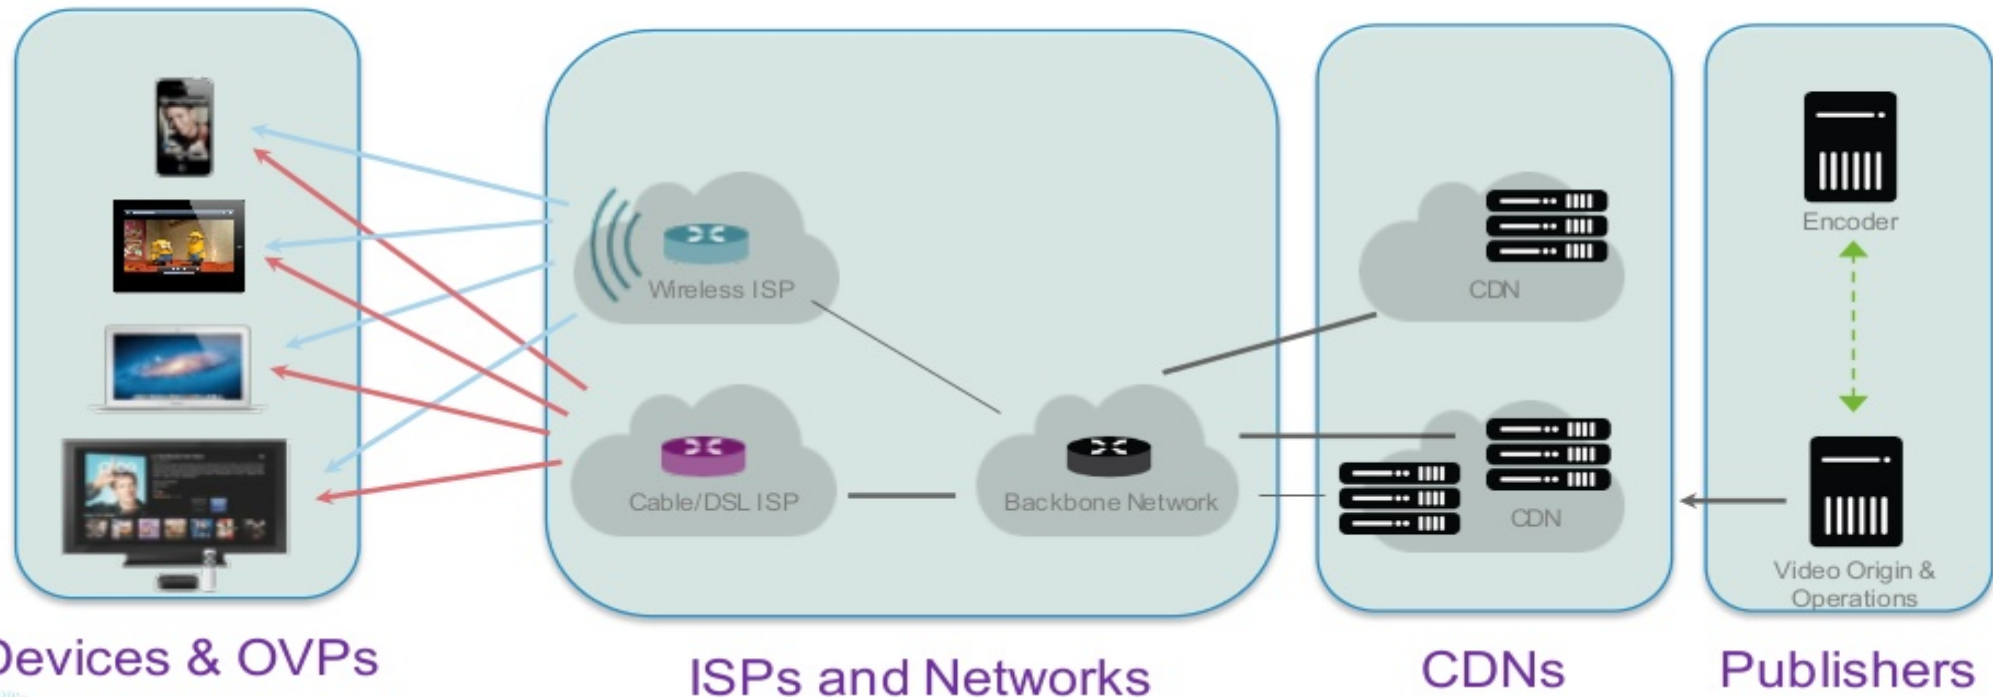
- 
- Video streaming over the internet (OTT) is growing rapidly
  - Major industry shifts in the last couple of years
    - HBO Now
    - ESPN/SlingTV
    - Verizon Go90
    - Facebook, Twitter
    - Amazon Prime Video

**Online Video – A Hugely Important Application  
“Big Bang” Moment is Unfolding – Right Now**



# Internet Video Streaming is Hard


Many parties, many paths but no E2E owner



# QoE is Critical to Engagement

## For both – Video and Advertisement business

HOW LIKELY ARE YOU TO WATCH  
FROM THAT SAME PROVIDER AGAIN?



33.6%  
VERY  
UNLIKELY

24.8%  
UNLIKELY

24.6%  
UNCHANGED

**58.4%**  
CHURN RISK

Source: Conviva, "2015 Consumer Survey Report"



SPARK SUMMIT 2016

## Viewers are expecting TV like quality and better

A background image showing two hands, one from the top right and one from the bottom left, reaching towards each other in a gesture of connection or support.

# Experience Matters!!

Must solve for EXPERIENCE and ENGAGEMENT

Success is more than just great content...  
Experience impacts engagement

Competition for eyeballs increasing...  
Internet of Content > Traditional TV viewing

TV revenues are up for grabs...  
Internet offers SVOD, AVOD, PPV &  
“Unbundled choices”

Publishers and  
Service Providers  
cannot lose touch  
with viewers’  
experience

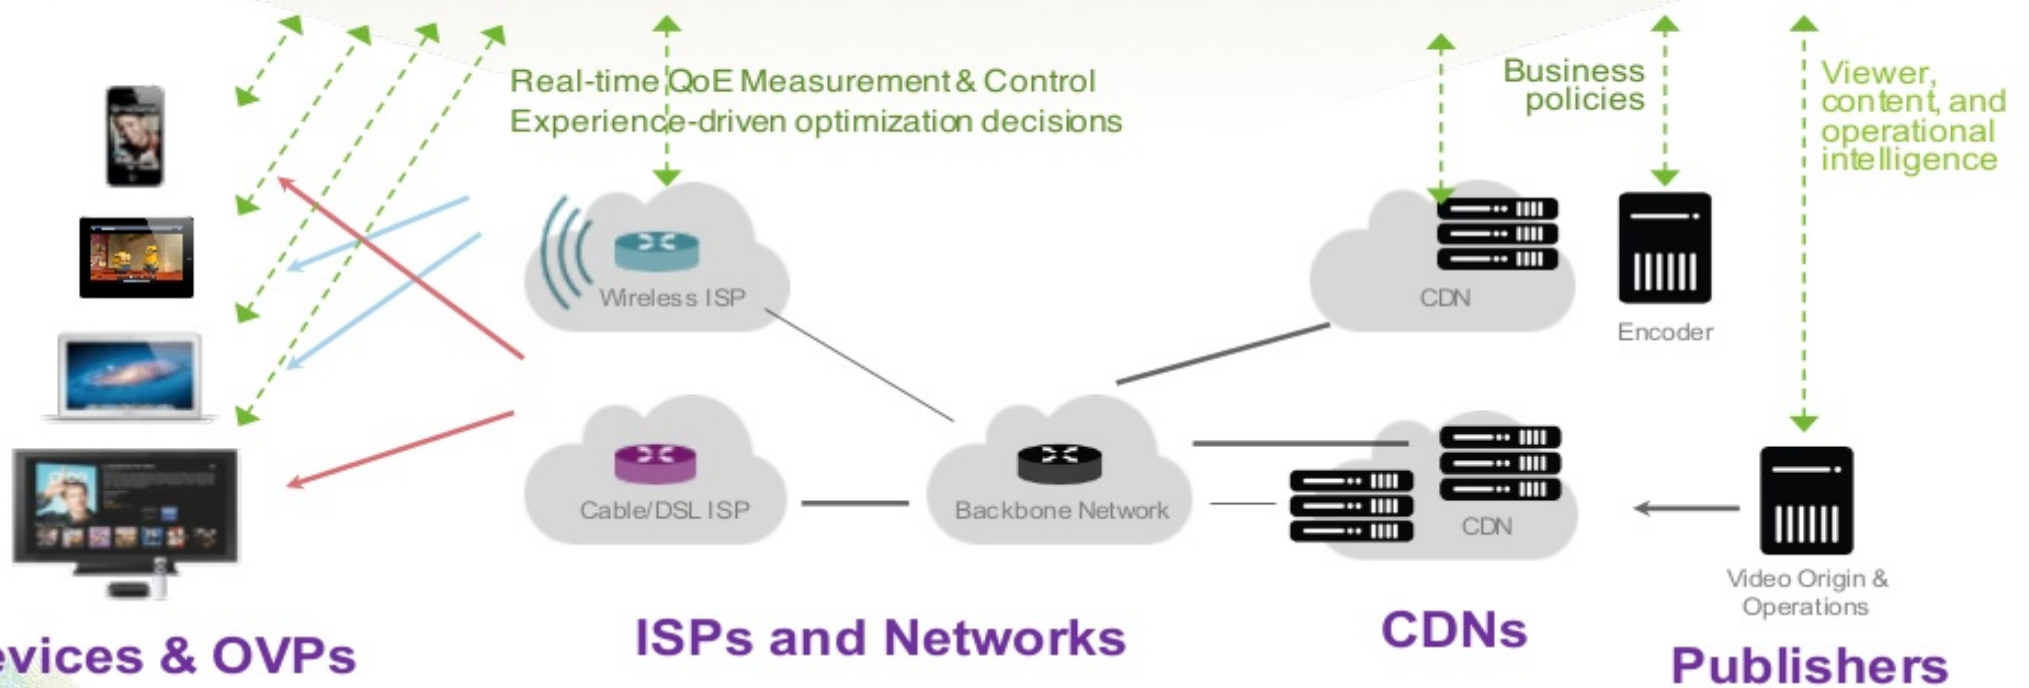
OR ELSE All bets are off!



SPARK SUMMIT 2016

CIVA

# CONVIVA EXPERIENCE MANAGEMENT PLATFORM



Devices & OVPs


ISPs and Networks

CDNs

Publishers



# Granular Dataset Indexed by Rich Metadata



SPARK SUMMIT 2016



# Scale of Deployment

AVOD													
													
													
SVOD													
													
Infra													



SPARK SUMMIT 2016

# Scale of Deployment



SPARK SUMMIT 2016

# Use Cases requiring 3 Stacks



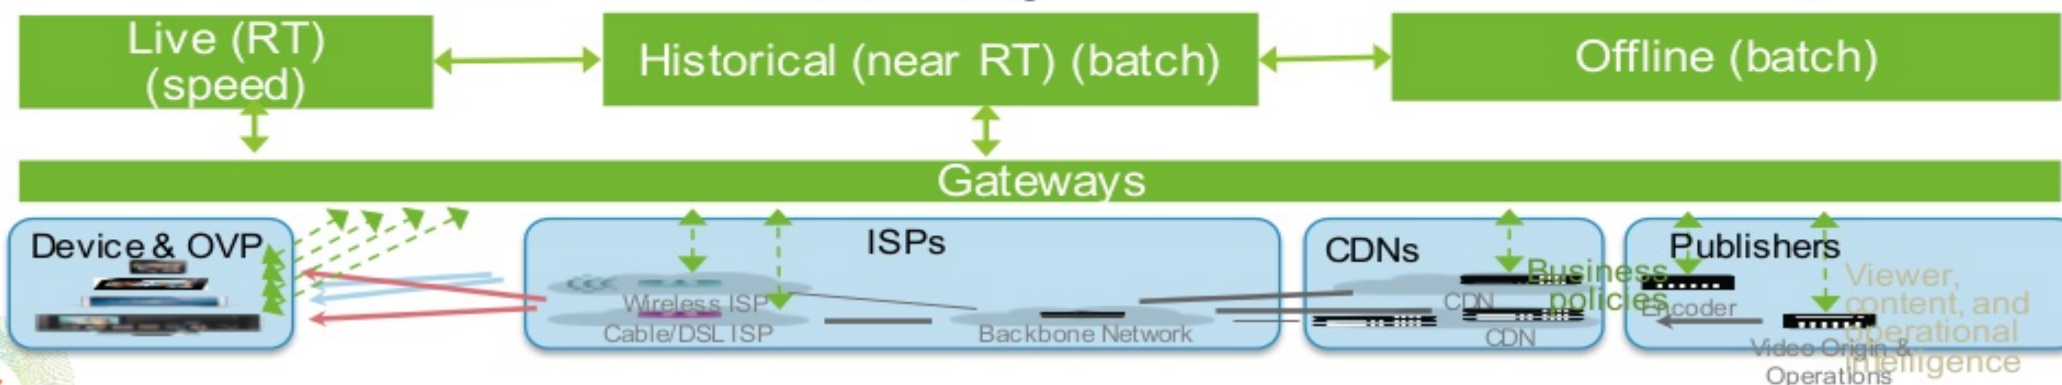
- Real time metrics
- Real time alerts
- Real time optimization




- Near real time metrics
- Historical trending
- Benchmarking



- In depth & ad hoc analysis
- Data exploration
- ML model training



## Old Architecture



- RT and near RT stacks get input from Kafka independently
- RT and near RT run independently (except some RT results saved to HDFS for some near RT calculation)
- Offline gets data from near RT Hadoop, with additional calculation specifics to offline analysis.
- Hive/Spark/R/NeuralNet etc. are used for various offline tasks






## Major Issues with old stack

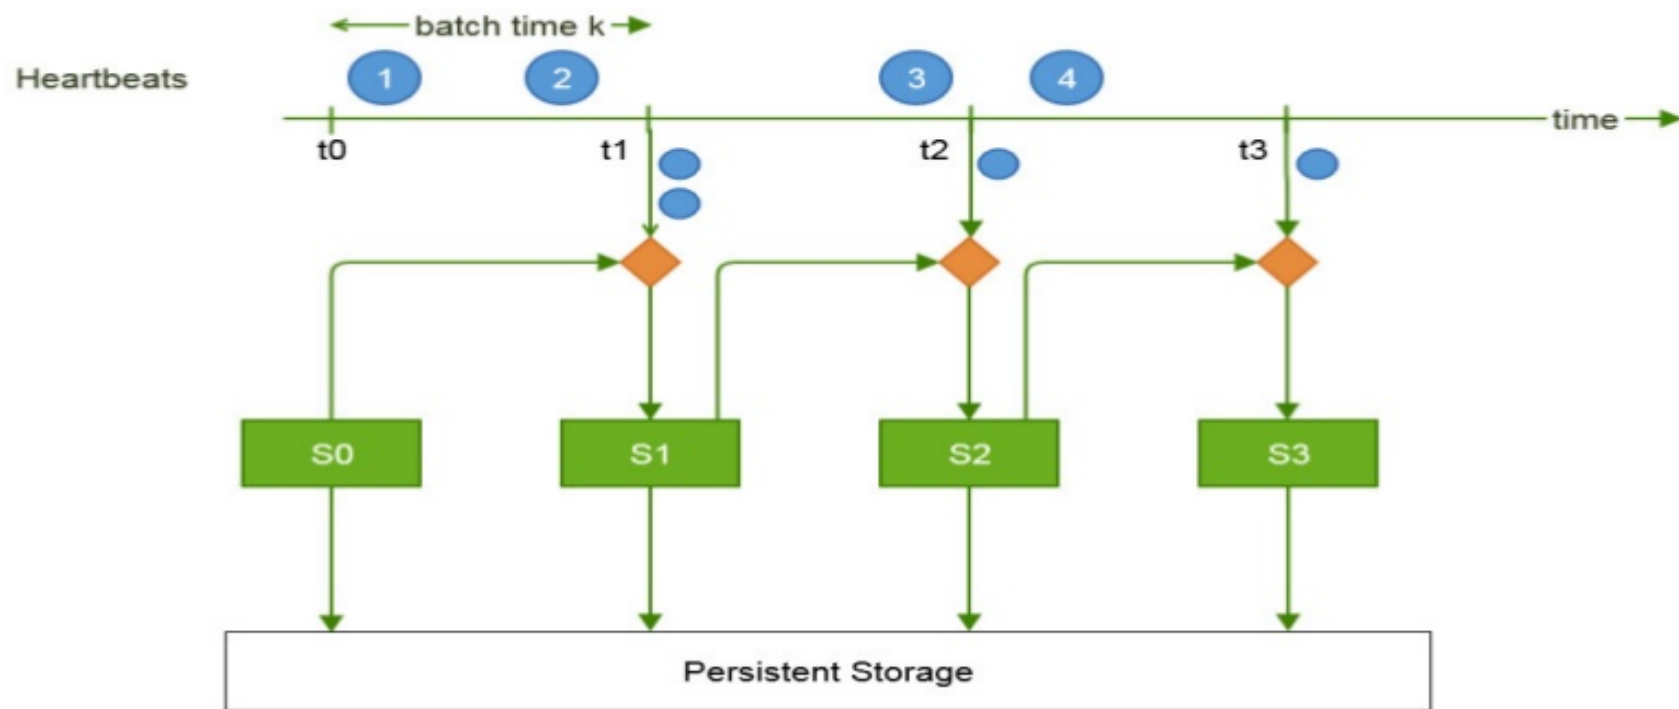
- **Code discrepancy among all 3 separate stacks**
  - RT: Pure updating model vs near RT: batch model
  - Offline: separate Hive layer; can have different calculation logic scattered in hive queries. (some standard UDFs/UDAFs help to certain extend)
- **A very complex and vulnerable RT stack**
  - Tricky thread locking
  - Mutable objects
  - Fixed data flow, specific delicate data partition, load balance.
- **Metric discrepancies cross all 3 stacks**
- **Different stacks also incur a lot of overhead of development, deployment**



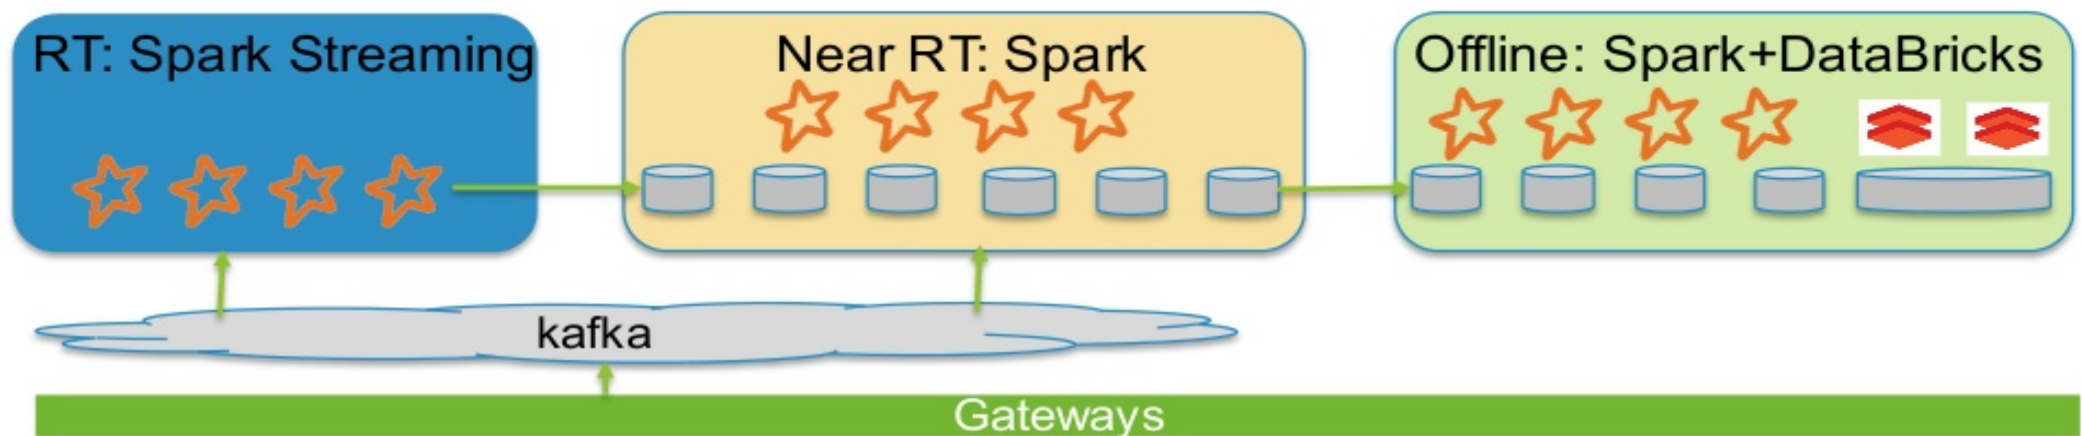
# Proprietary Real Time Stack



# Hadoop Batch Based Near RT Stack



# New Architecture

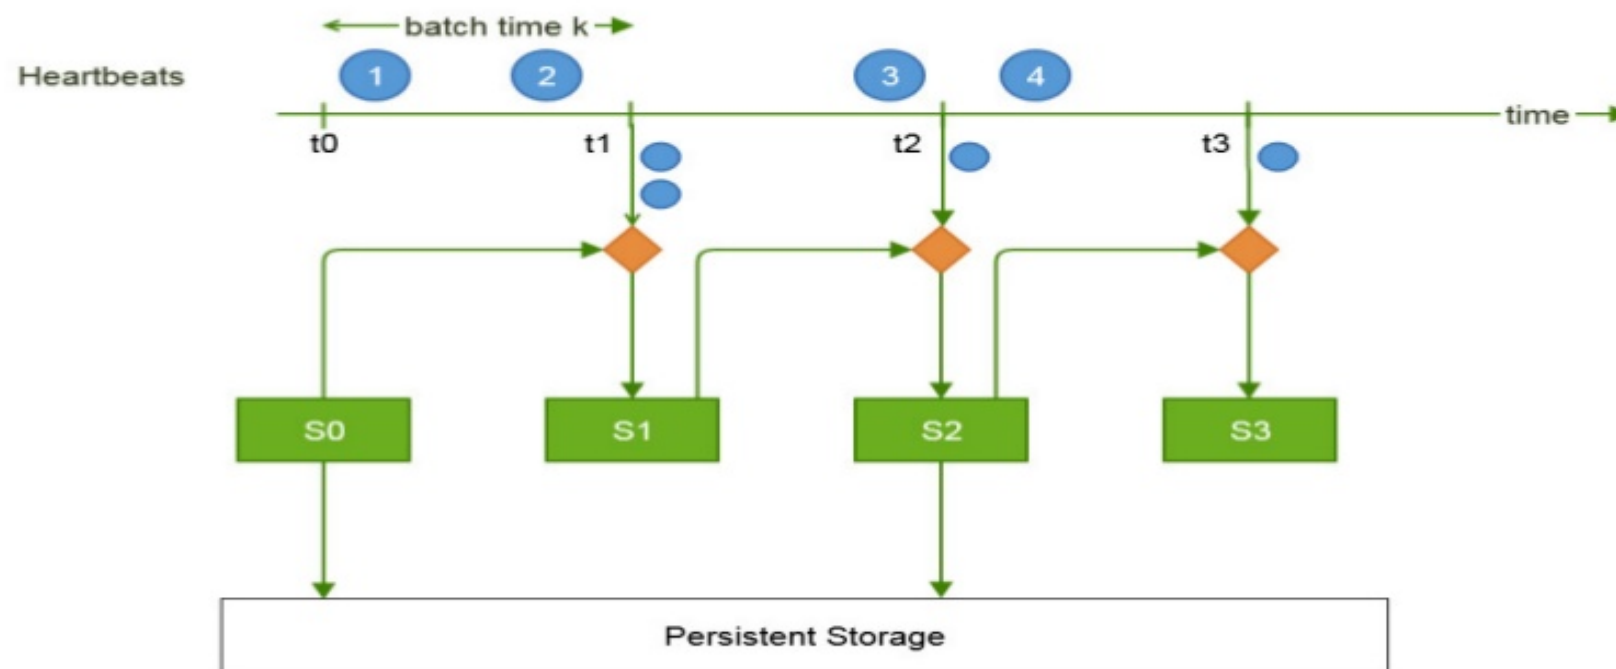


- All Converging to Spark based Technology.
- Max. sharing of code cross all 3 stacks
- Offline: with better cluster management (Databricks), running over many on-demand clusters instead of one pre-built





# Unified Stack



## Unified Stack High Level Code

```
val rawlogs = (DStream from rawlogs)

val ssMapped = rawlogs.transform {
  (map function)
}

val ssReduced = ssMapped.updateStateByKey{
  (reduce/update function)
}

(every n batches)
saveToStorage(ssReduced)
```



## updateStateByKey, mapWithState

- Acts as the 'reduce' phase of the computation
- Helps maintain the evolving state shown earlier
- The performance of `updateStateByKey` is proportional to the size of the state instead of the size of the batch data
- In Spark 1.6, will be replaced in our workflow by `mapWithState`, which only updates as needed



# Deployment

- RT portion In production environment for ~5 months
- Backward compatible migration first, major improvements later
- Performance tuning is important
- For RT: `Checkpointing`, reliability vs performance





## Importance of the Offline Stack

- For data centric companies, most important innovations are happening with exploring and learning through the big data
- Speed and efficiency of offline data analysis and learning is the key to the success
- Data and Insights accessible to many internal organizations besides data scientists (customer support, SE, PM,...) is extremely important for the overall success



## What's Important to Data Scientists

- Efficient access to all the data
- Can work independently with all the resources needed.
- Can work with the teams (internally and with other teams)
- Interactivity for data exploration
- Easy to use, powerful data visualization
- Machine learning tools

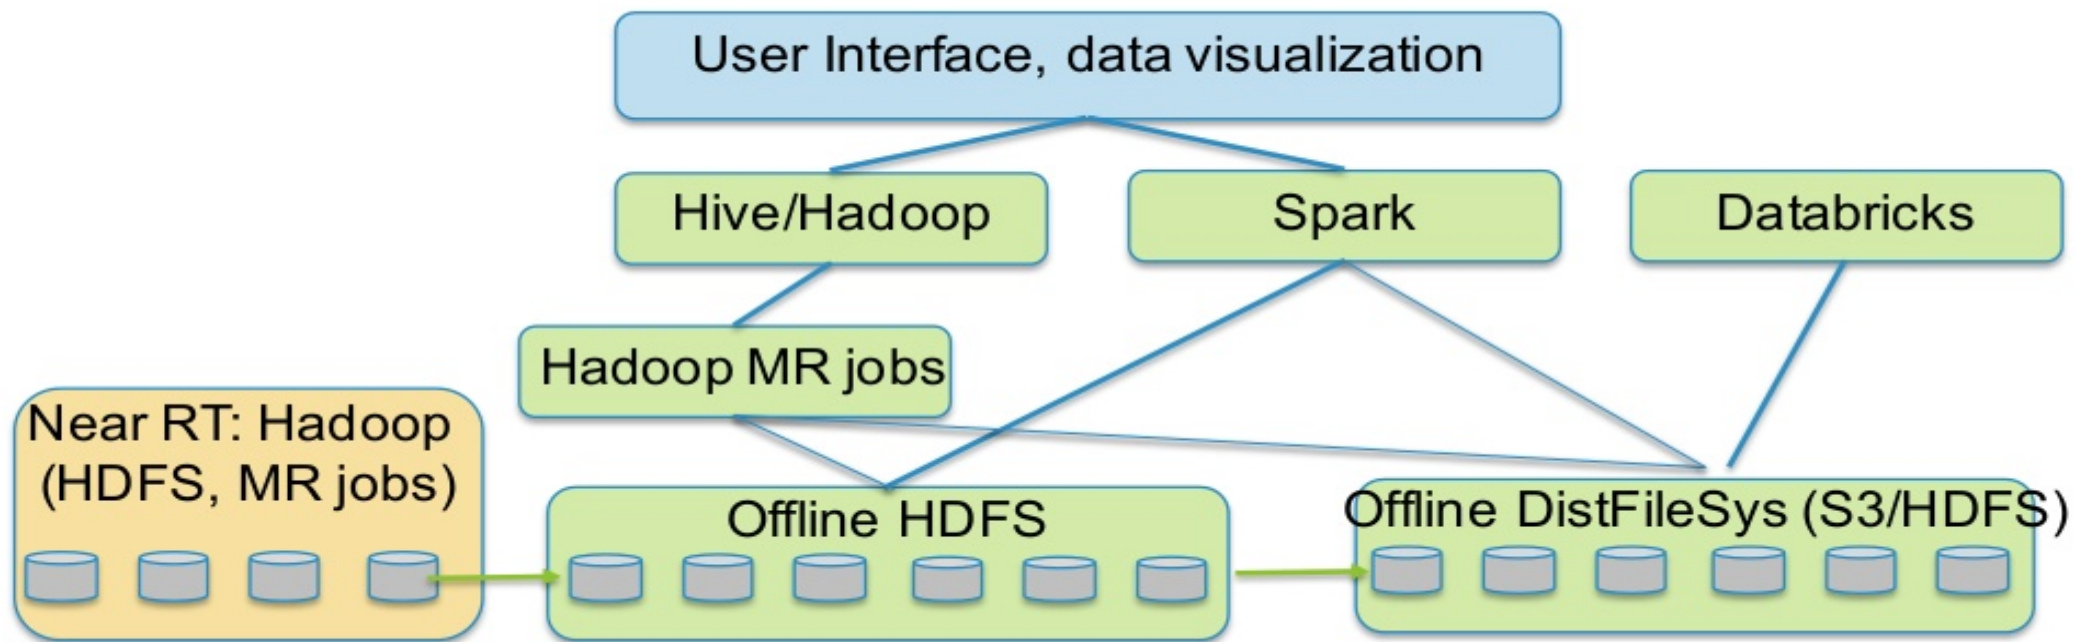


## What's Important to Data Scientists

- Re-use of existing production logic/code when applicable
- Easy transfer of work into production
- Integrated environments with engineering discipline
  - Code management and version control
  - Design and code reviews

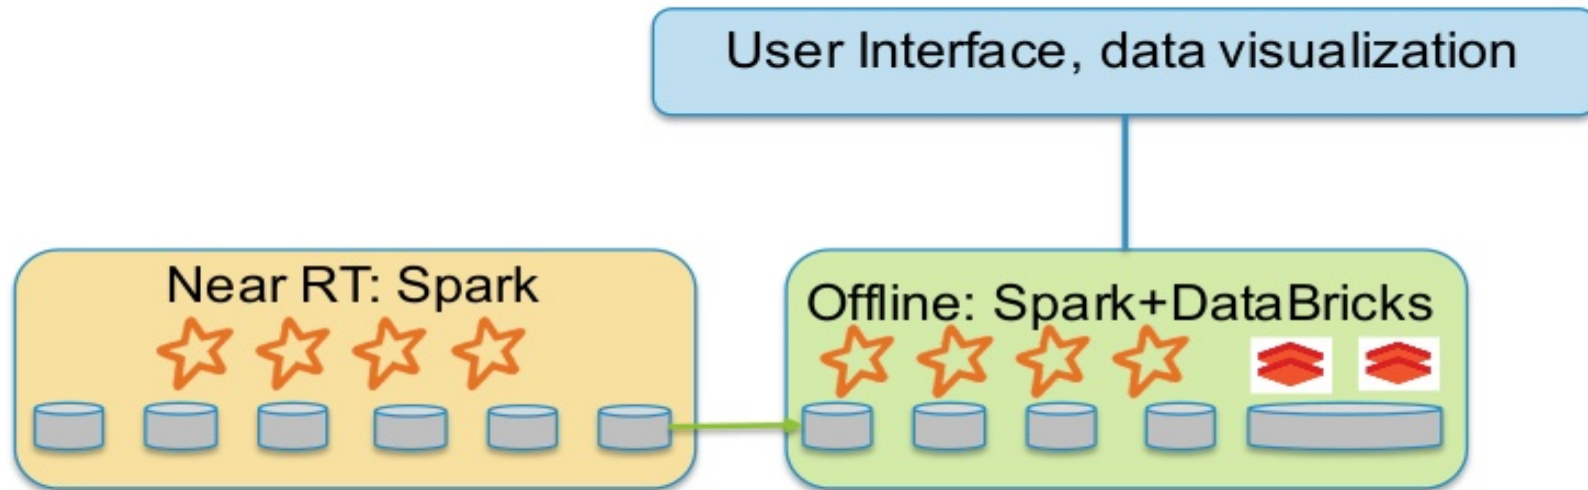


# Old Architecture





# New Architecture



# Benefits of Databricks

- Cluster management:
  - Instead of one shared cluster, everyone can launch/manage his/her own clusters
- Interactive environments
  - Notebook environment is very convenient and powerful
- Easy to share/working together
  - Sharing notebooks are easy (with permission control)
- Data visualization
  - Good visualization tools: matplotlib, ggplot inside R
- Reasonably good machine learning tools
  - MLlib, R, other packages (H2O)



## Benefits of Databricks

- Same code can potentially be moved to other stacks and production
- New features built faster here:
  - Harder to change production environment
  - New features developed, tested and deployed faster
- Huge efficiency gain for the data science team
- Production issue debugging also using Databricks with big efficiency gain



## ML Example: Long Term Retention

- Long Term Retention Analysis
  - Months/years of data: many billions of views, many millions of subscribers per publisher/service provider.
  - Determining appropriate time interval for subscriber history and for subscriber abandonment.
  - Finding best features for predictive model.
  - Handling categorical features with too many possible values.



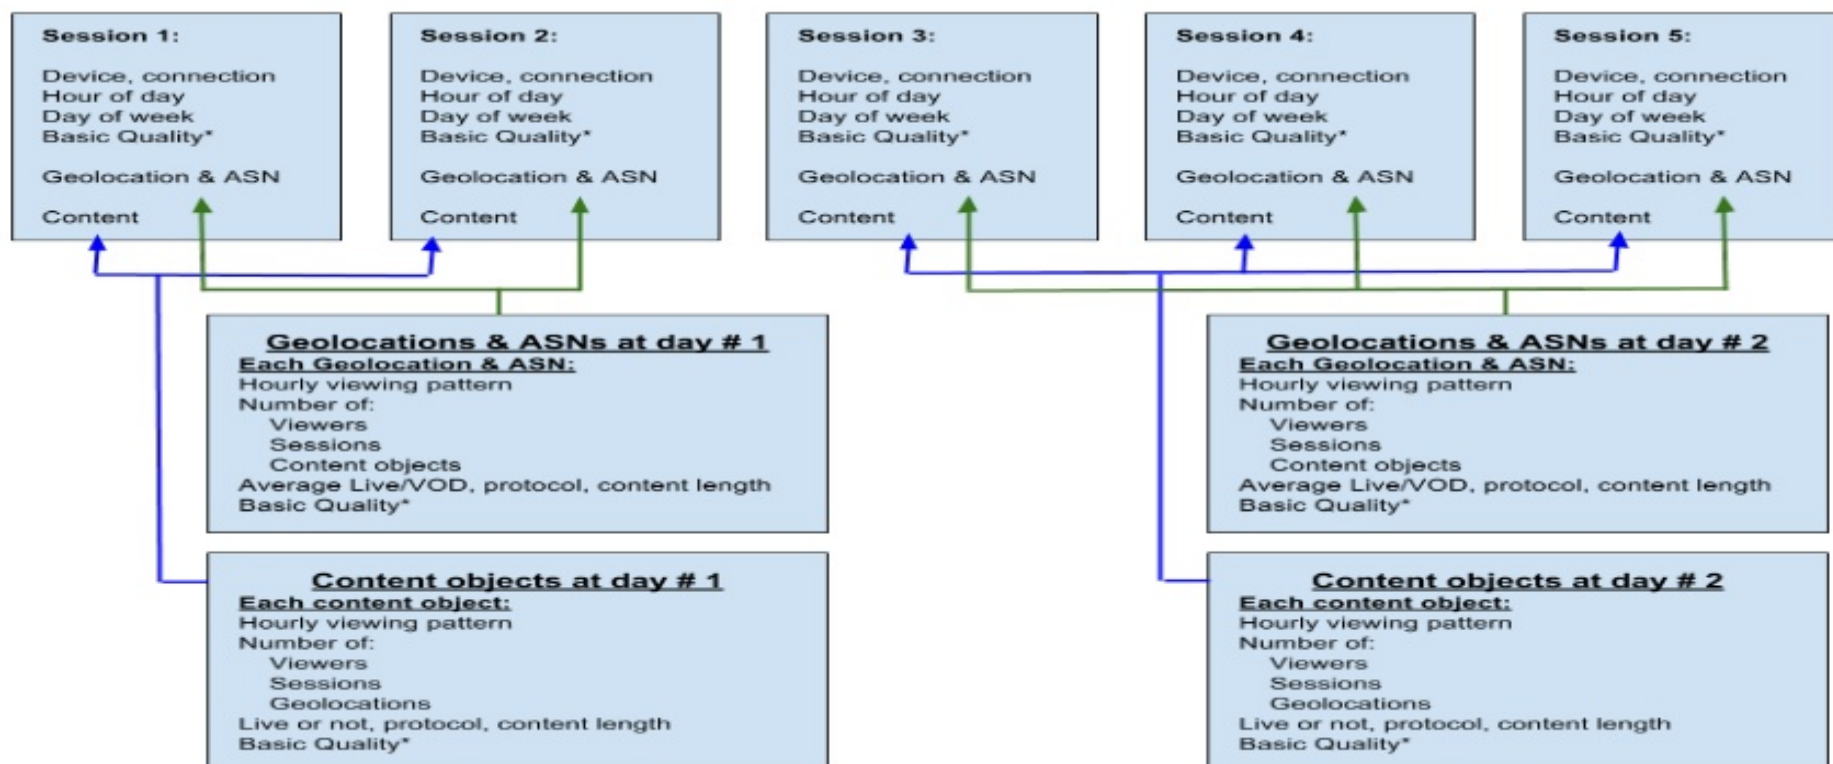
# Characterization of Categorical Features

- One-hot encoding:
  - Some categorical features (e.g. Device) with dense limited values
- Some features have too many sparse categorical values
  - City & ASN
  - Video Content
- Aggregated features of many months of subscriber behavior:
  - All content that the subscriber watched
  - all Geolocations from which the viewer watched





# Geolocation & Day



(\*) **Basic Quality:** Failed/not, Start-time, Buffering, Bit-rate



SPARK SUMMIT 2016

# Work Flow Inside Databricks

- Create dataframes with features for each geo x day, content x day
- For each subscriber history, for each video session, replace geo and content with features of geo x day, content x day for that day
- Aggregate each subscriber history to obtain final features
- All done inside Databricks environment. Highly iterative process: especially related to feature design and extractions (many iterations)
- Use Spark MLlib, various models, such as Gradient Boosted Tree Regression
- User visualization inside Databricks

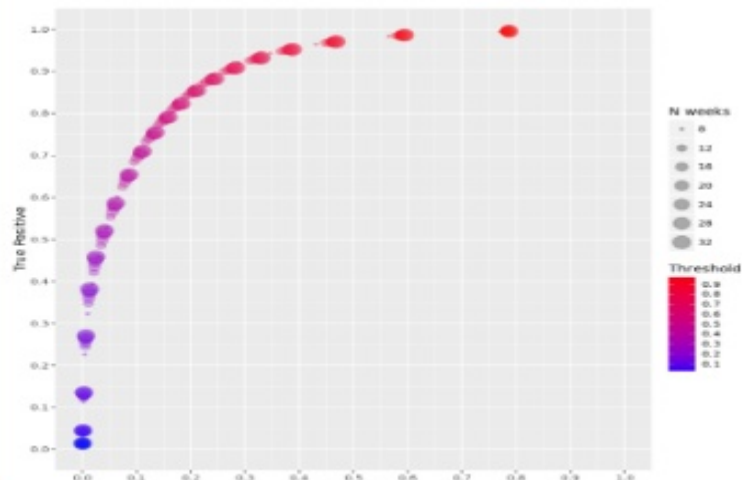


# Sample Results

## Common ROC (churn) representation

```
> ggplot(dfrROC, aes(x=dfrROC$FalsePositiveChurn, y=dfrROC$Selectivity)) +  
  geom_point(alpha=0.3, aes(size=dfrROC$N weeks, color=dfrROC$Threshold)) +  
  scale_size_continuous(name="N weeks", breaks=c(8,12,16,20,24,28,32), labels=c("8","12","16","20","24","28","32")) +  
  scale_color_gradient(name="Threshold", low="blue", high="red", breaks=c(0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9), labels=c("0.1","0.2","0.3","0.4","0.5","0.6","0.7","0.8","0.9")) +  
  ggtitle(expression(atop("ROC. Model is trained on 8 weeks history to predict next 8 weeks inactivity",  
    atop("The model is applied 8 weeks later. The prediction is compared with N weeks inactivity."), ""))) +  
  labs(x="False Positive", y="True Positive") +  
  coord_cartesian(xlim=c(0,1), ylim=c(0,1)) +  
  scale_x_continuous(breaks=c(0,0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9,1)) + scale_y_continuous(breaks=c(0,0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9,1))
```

ROC. Model is trained on 8 weeks history to predict next 8 weeks inactivity  
The model is applied 8 weeks later. The prediction is compared with N weeks inactivity.



SPARK SUMMIT 2016

## Much More Work Ahead

- Improve the real time performance, trading off latency vs metrics accuracy/failure handling.
- <100ms real time processing and response still need more work.
- Making sure modular design for all the logics so that it can be shared cross all 3 stacks wherever possible
- More exciting and challenge works ...



# We Are Hiring

<http://www.conviva.com/our-team/careers/>



SPARK SUMMIT 2016

---



# THANK YOU.

[jibin@conviva.com](mailto:jibin@conviva.com)



**SPARK SUMMIT 2016**  
DATA SCIENCE AND ENGINEERING AT SCALE  
JUNE 6-8, 2016 SAN FRANCISCO