

Vertica and Spark: Connecting Computation and Data

Rui Liu and Edward Ma

Hewlett Packard Enterprise Vertica Advanced R&D Labs



SPARK SUMMIT 2016
DATA SCIENCE AND ENGINEERING AT SCALE
JUNE 6-8, 2016 SAN FRANCISCO

Overview



Computation and data



Vertica
Analytics Platform

- The Vertica-Spark connector connects both data and computation between Vertica and Spark
 - VerticaRDD and Vertica Data Source APIs
 - Data-locality optimization
 - Computation pushdown
 - Save data from Spark to Vertica



SPARK SUMMIT 2016


HPE Vertica

Vertica
Analytics Platform


- An advanced SQL analytics platform
- Shared-nothing MPP architecture
- Column-oriented storage organization
- Standard SQL interface with many analytics capabilities built-in
- Extensible via User-Defined Functions



Connected pipeline



Connected pipeline




- Connects the computation pipelines of Spark and Vertica in an optimized way
- Utilizes parallel channels (parallel queries) for data movement
- Leverages data-locality optimization on Vertica
- Ensures computation push-down into Vertica as appropriate (filters and projections)



Vertica data segmentation on hash ring

```
CREATE TABLE t ...  
SEGMENTED BY HASH(id) ...
```

Id	Name	Age	email	...



- node0
- node1
- node2
- node3



Spark locality-aware partitions

Locality-aware partition:

- Uses Vertica data-distribution information
- Generates locality-aware queries

```
SELECT id, name FROM T
where
  0 < 0xffffffff & hash(id)
AND
  0xffffffff & hash(id) < 278985;
```

■ seg 0

■ seg 1

■ seg 2


■ seg 3

— Spark partition




Locality-aware query

Non locality-aware query




Locality-aware query



Vertica query pruning

- For a locality-aware query, the execution can be pruned
- Query only executes on nodes that contain the data




Arbitrary number of partitions

- seg 0
- seg 1
- seg 2
- seg 3

— Spark partition



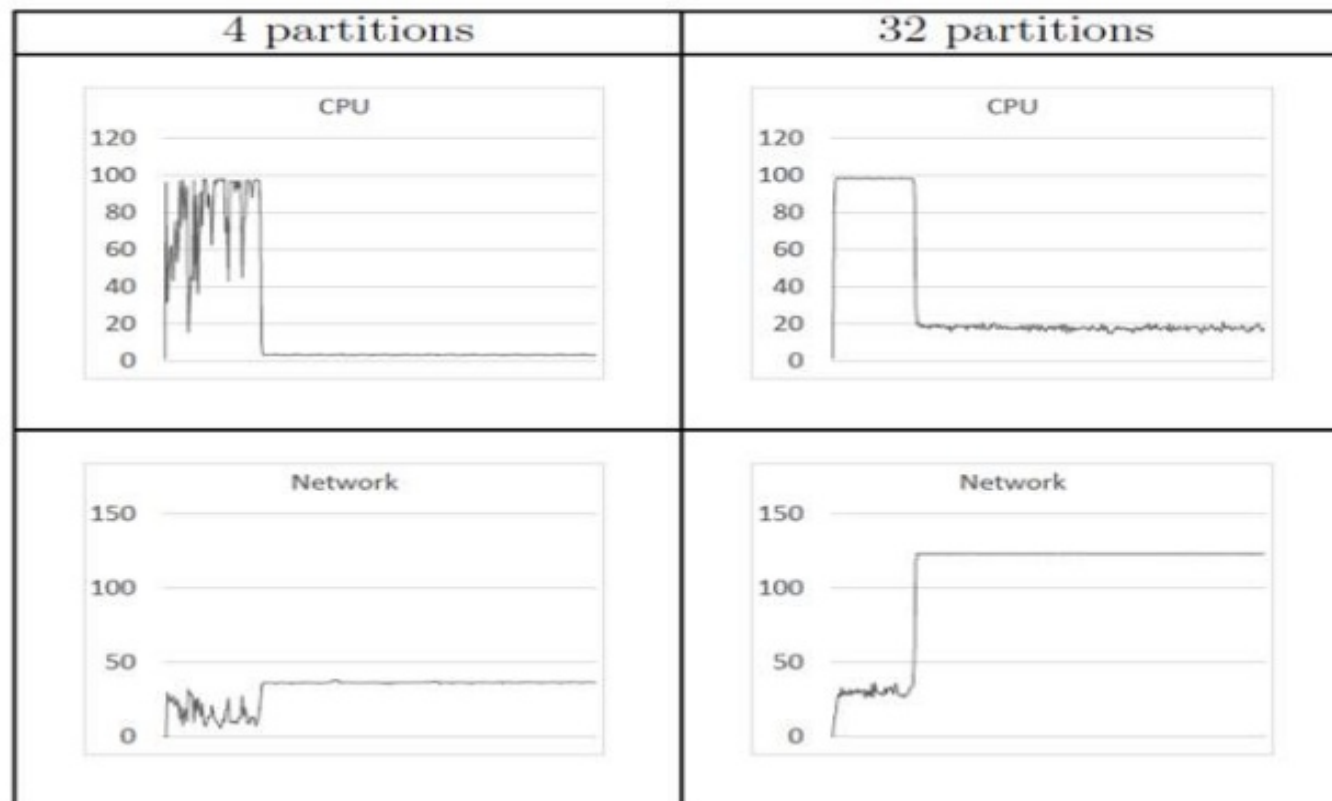
Spark partitions over unsegmented table replicas



} Spark partition



Performance characteristics




Scalability of throughput




Computation push-down

- Implemented:
 - Filters
 - Projections
 - Count(*)
- Works in progress:
 - Joins
 - Aggregations
- Future work:
 - User-Defined Functions



Computation push-down schemes


- Per-tuple computation (filtering, projection)
 - Always pushed down
- Joins & aggregations
 - Locality-aware: when join and aggregation keys are co-segmented in Vertica
 - Single query: run single query with join/aggregation inside Vertica



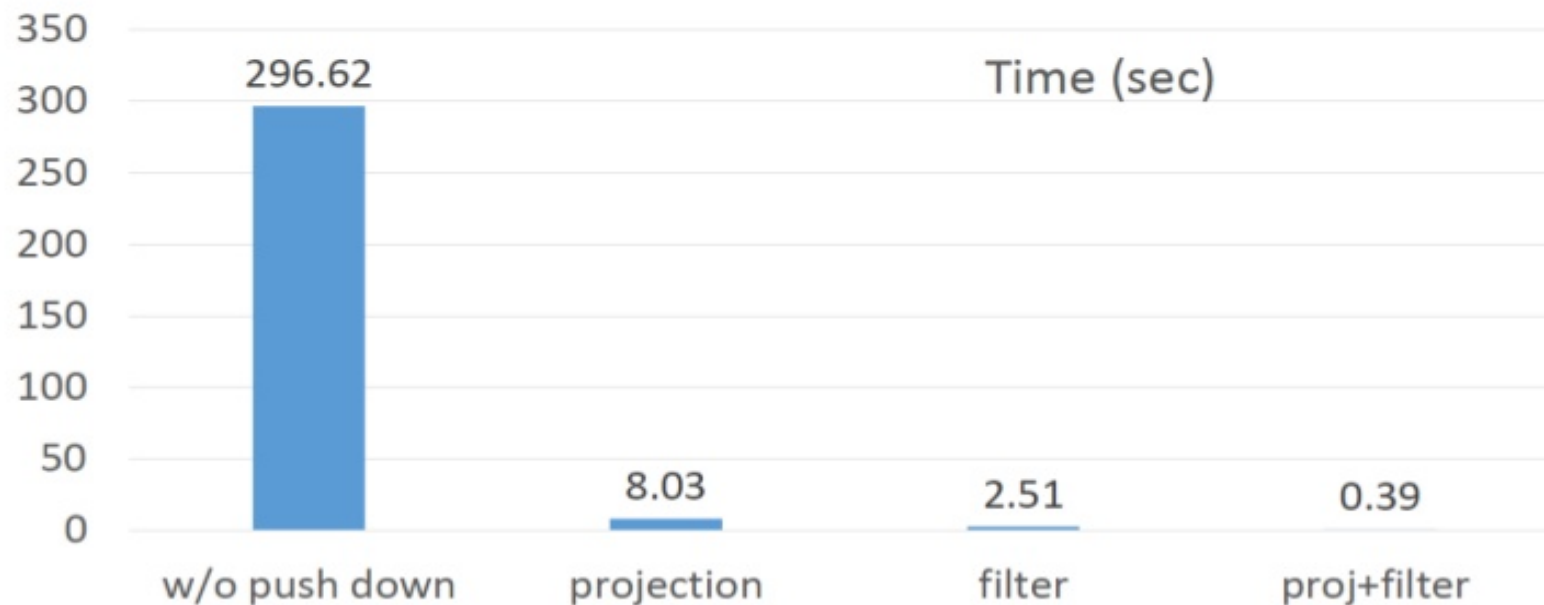
Original Query: `select c,d,f from A JOIN B on A.foo = B.foo`

Locality-aware
push-down
query:

```
select c,d,f from  
A JOIN B on A.foo = B.foo  
where  
0 < 0xffffffff & hash(foo) AND  
0xffffffff & hash(foo) < 278985 ...
```



Performance Impact of Pushdowns



Vertica Analytics

- Geospatial analytics
- Sentiment analysis
- Sessionization of event streams
- Time series pattern matching
- ... (Many more)



The Other Direction



Spark as an ETL Engine for Vertica



The Other Direction



Spark as an ETL Engine for Vertica

Objective: Want reliable and fast bulk loading



Two Possible Approaches

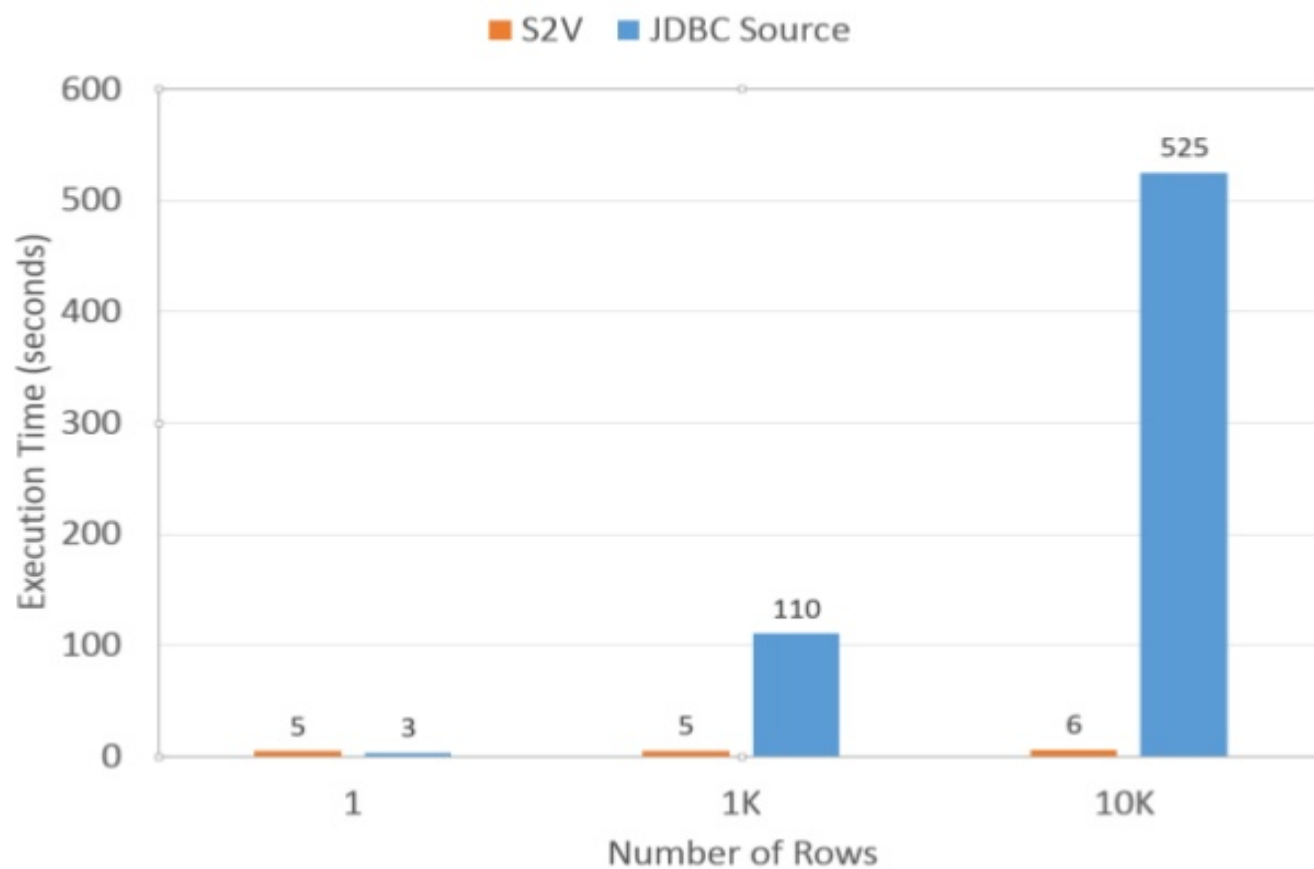
Direct



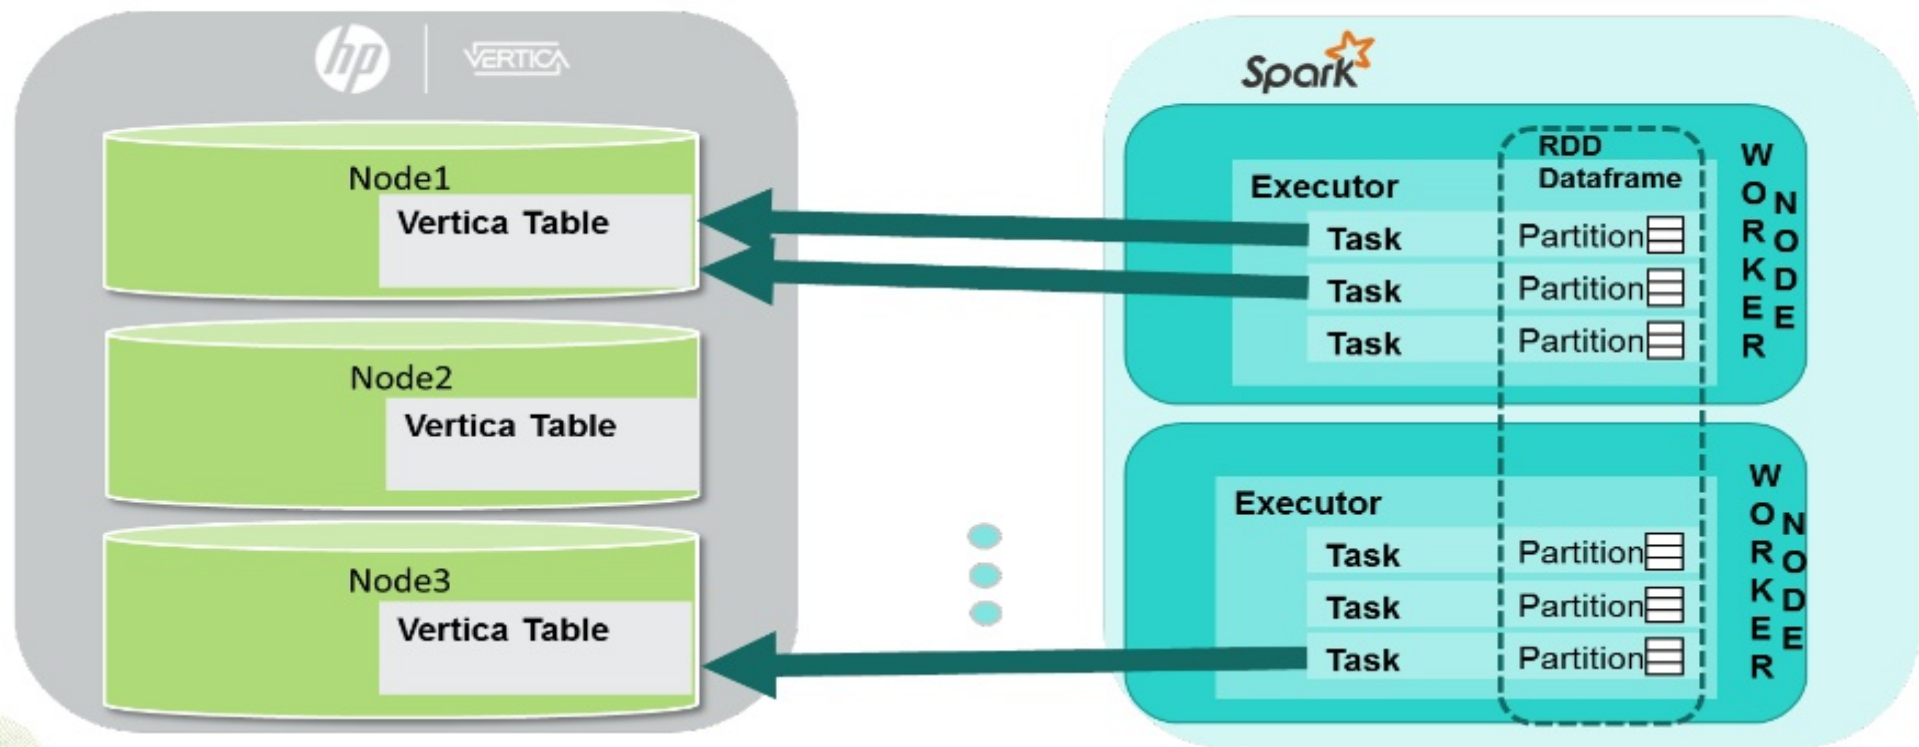
Indirect (Two-Stage)



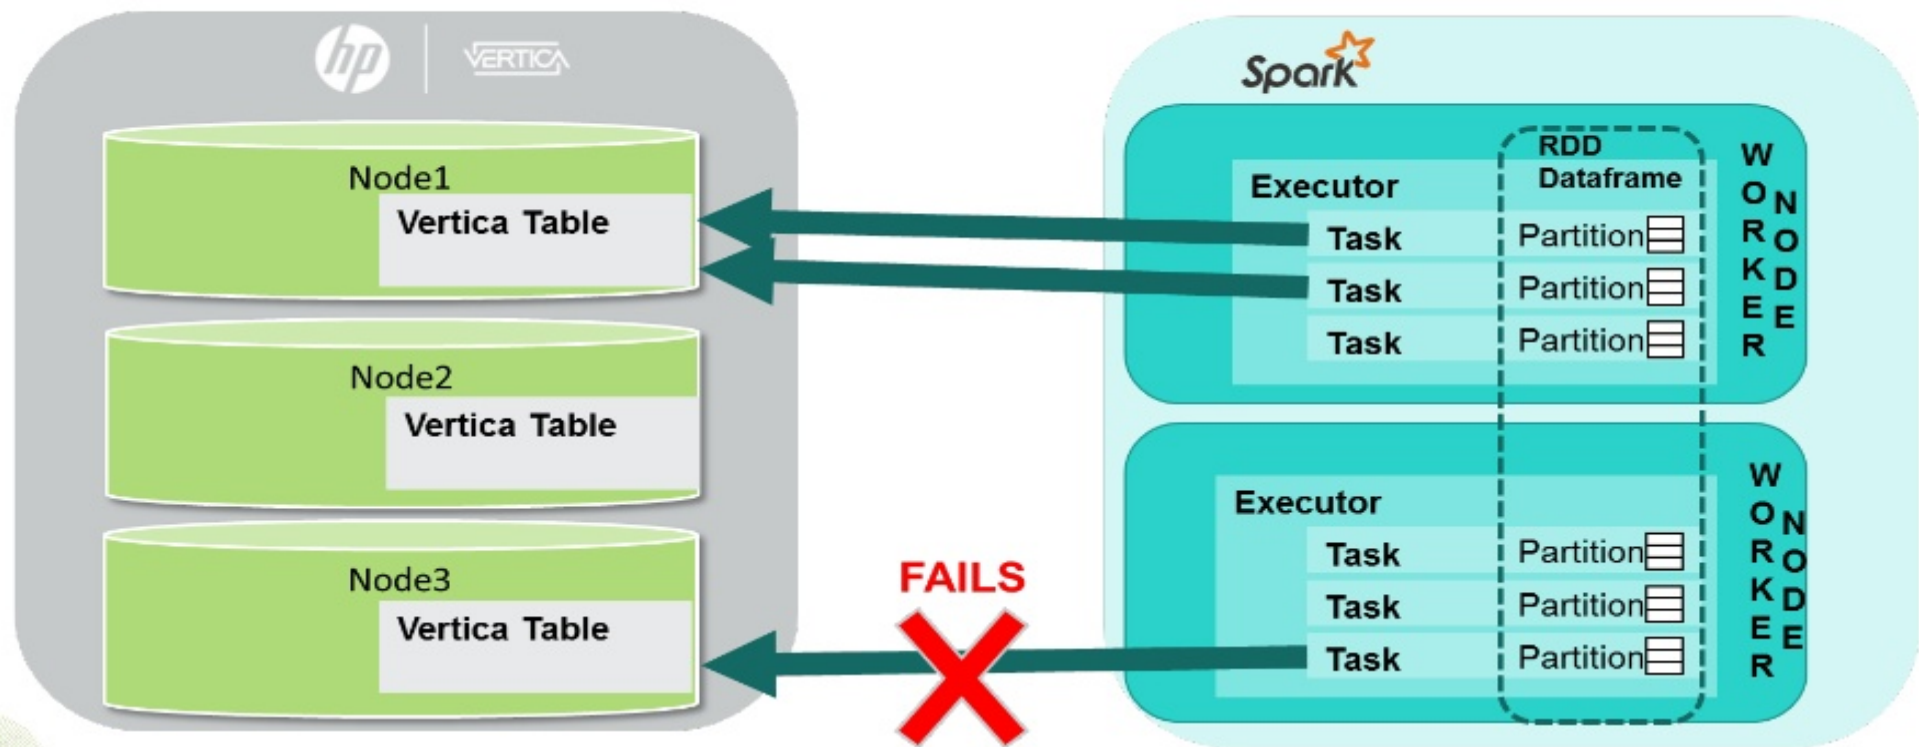
Comparison to JDBC Source



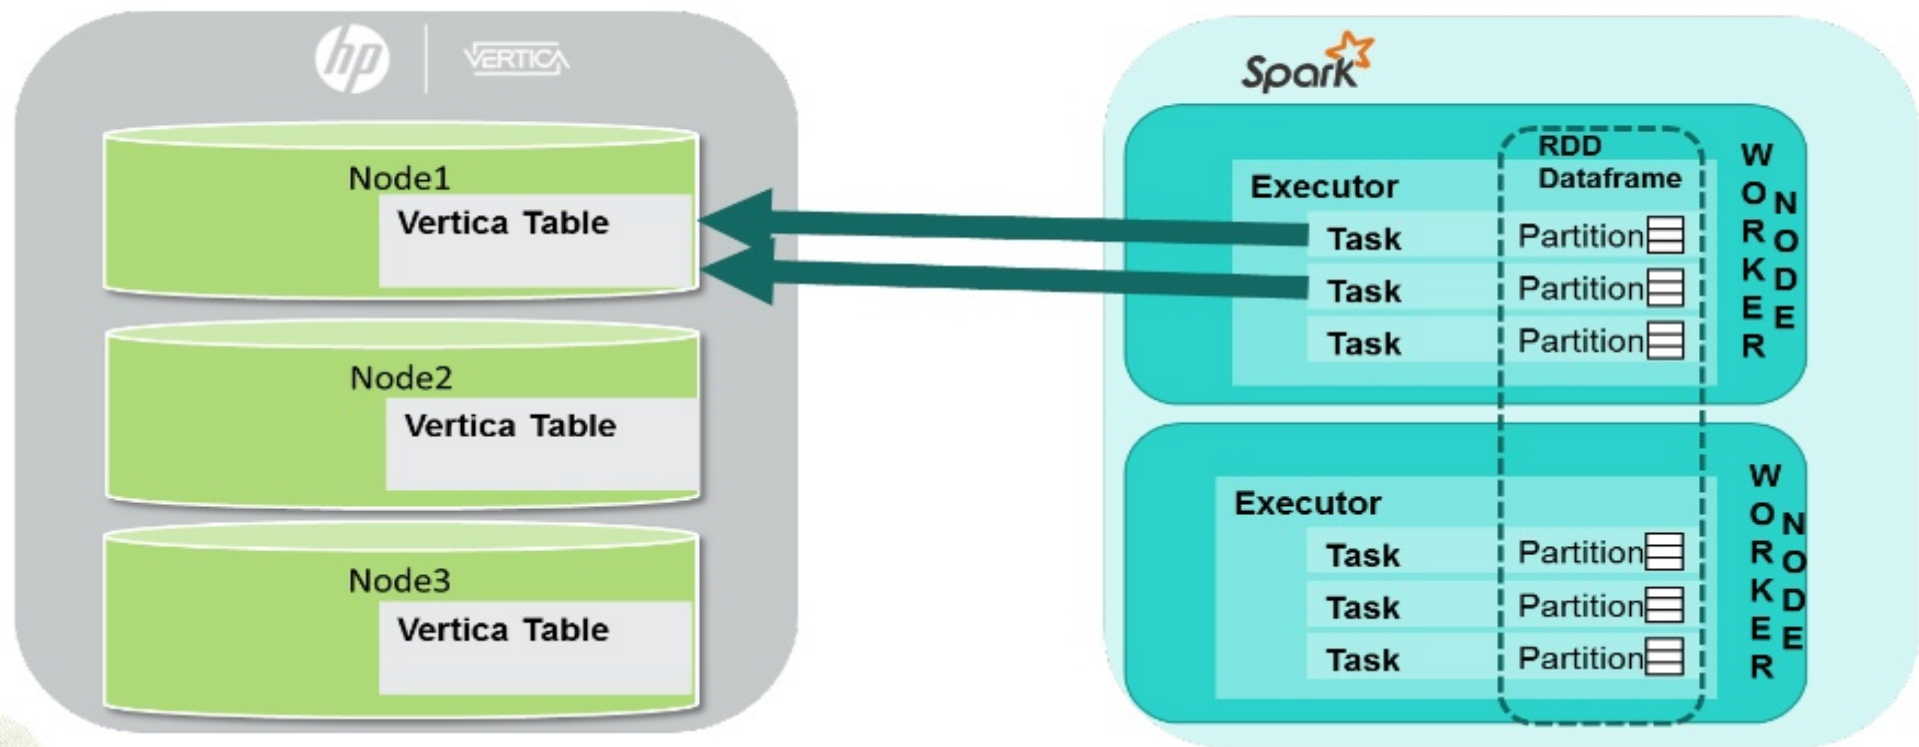
What could happen?



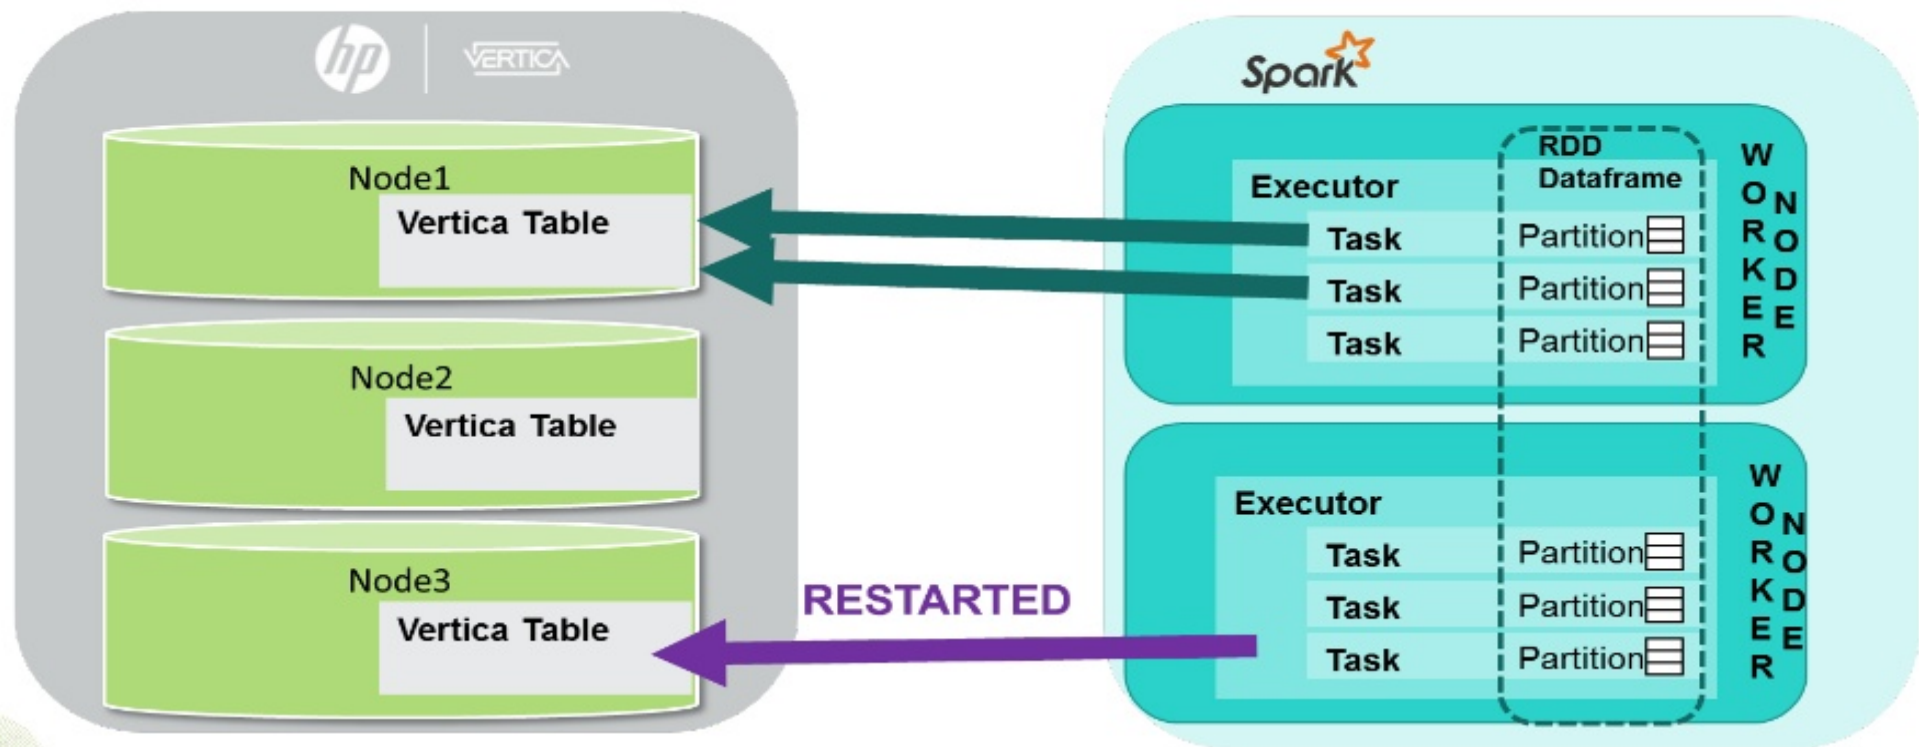
What could happen?



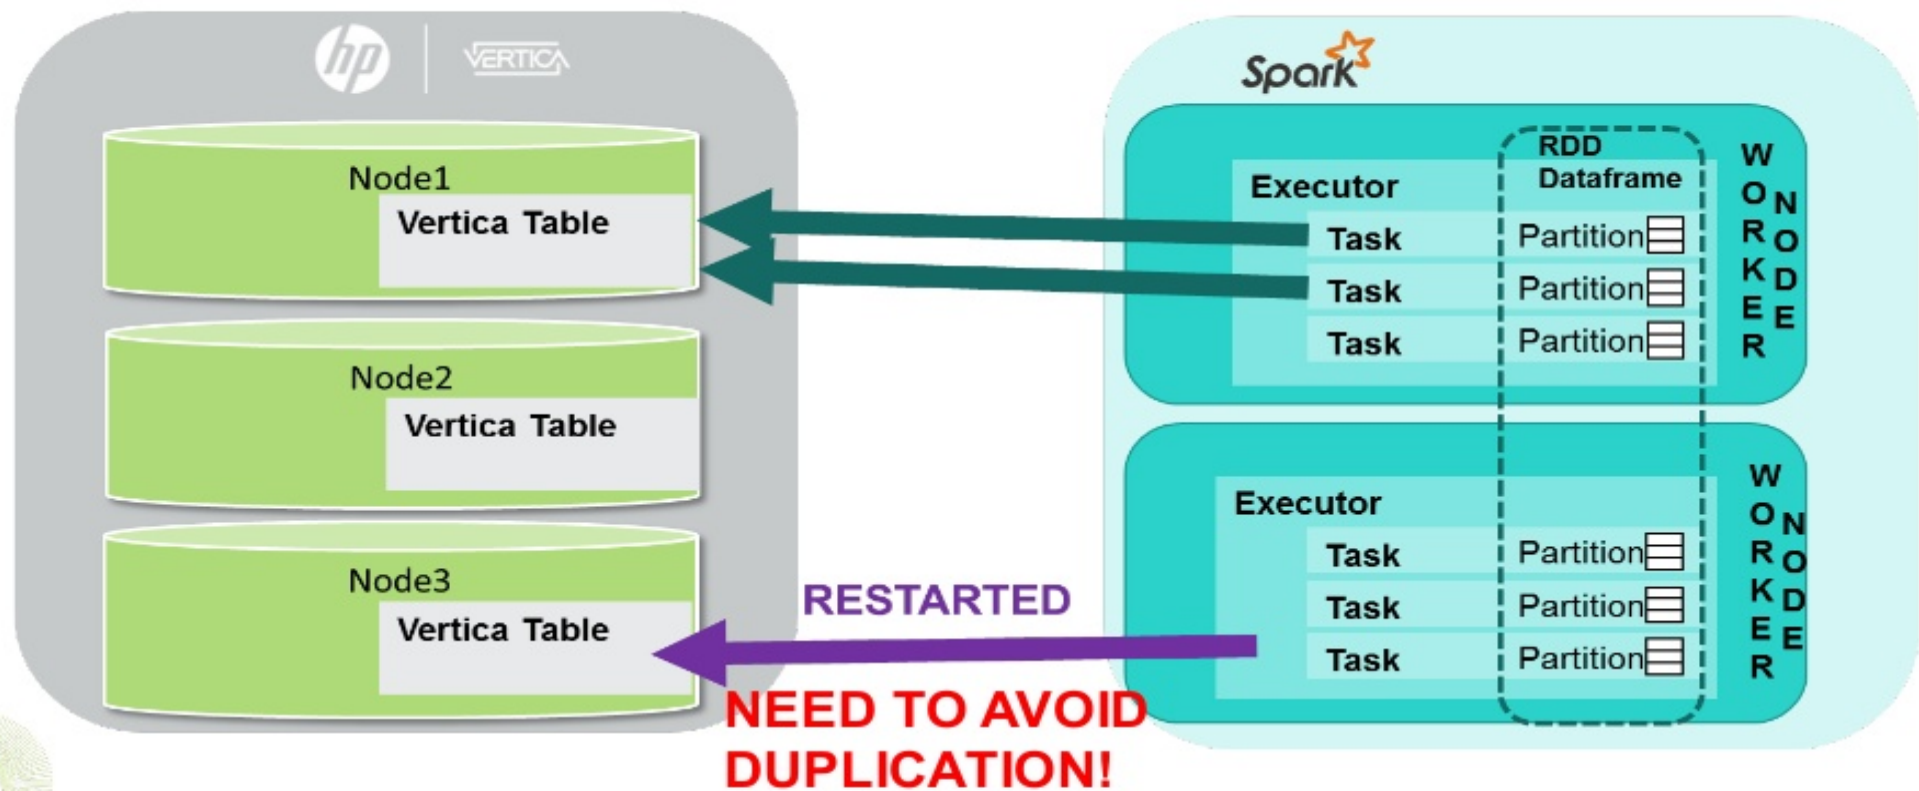
What could happen?



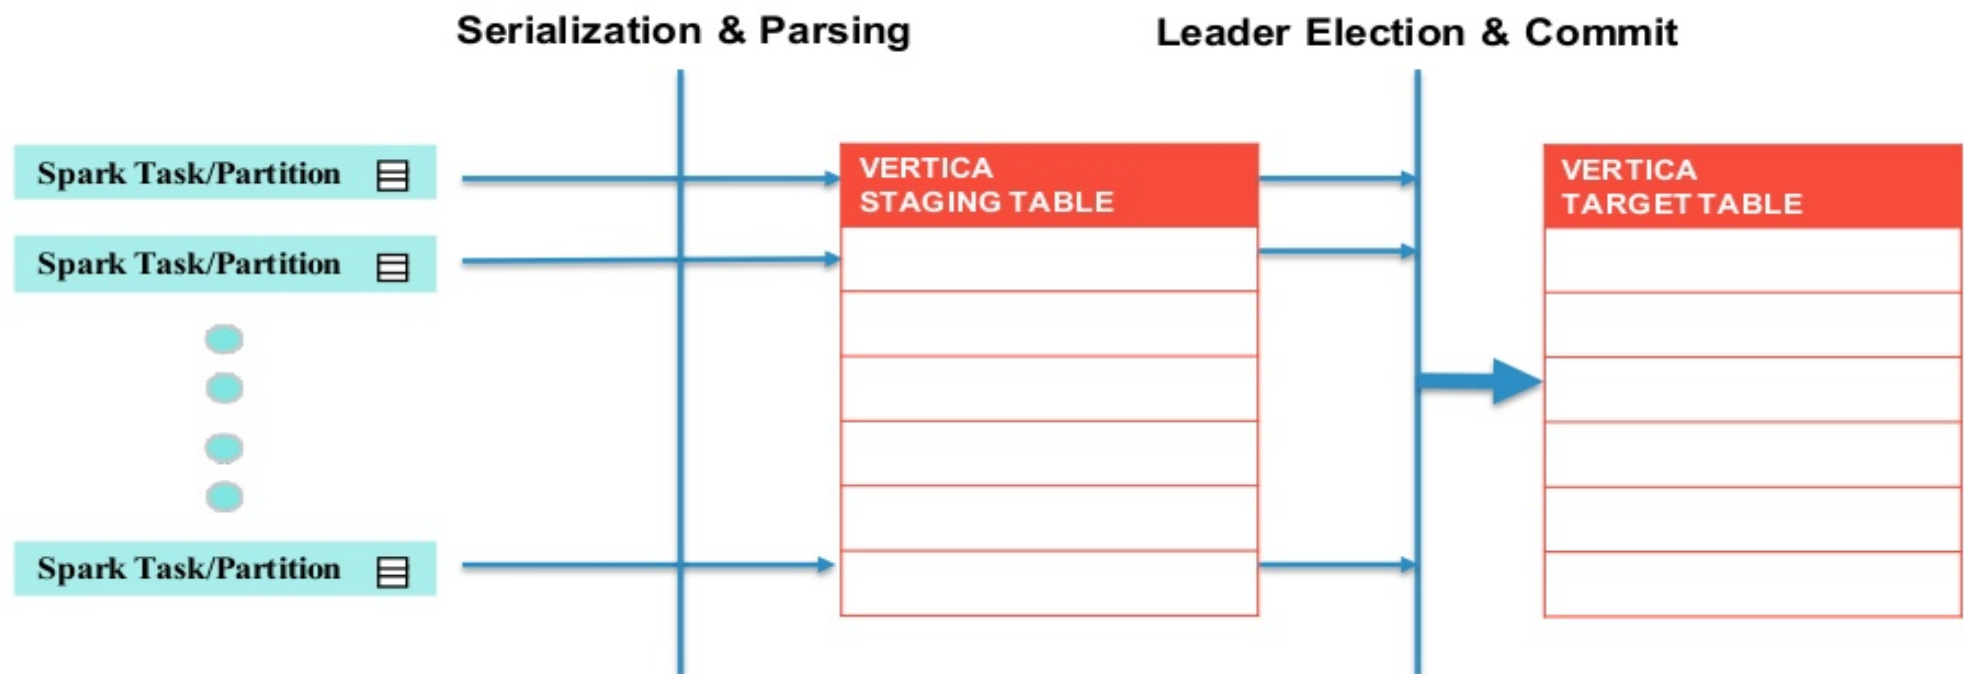
What could happen?



What could happen?



Big-Picture (Direct)




Exactly-once Semantics

Never results in partial or duplicate table loads

TASK_STATUS

STAGING_TABLE

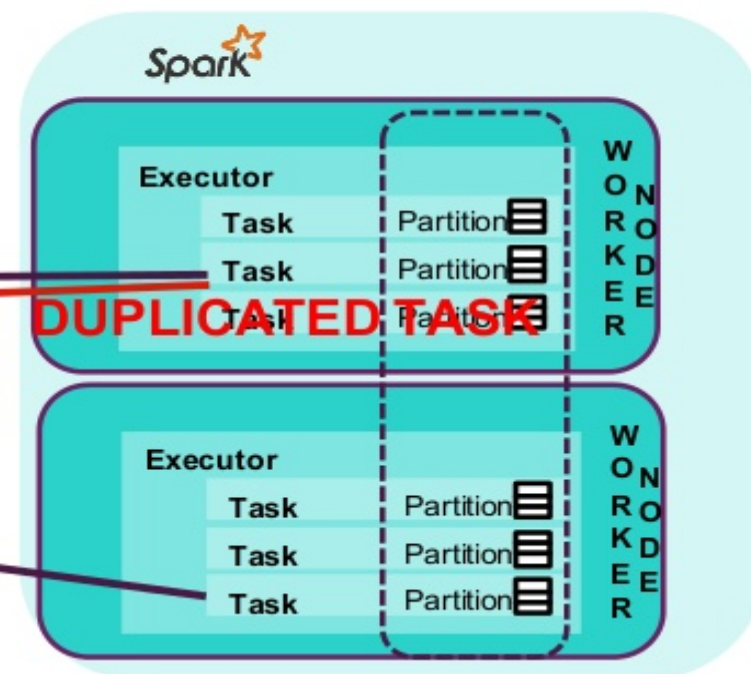


Exactly-once Semantics

Never results in partial or duplicate table loads

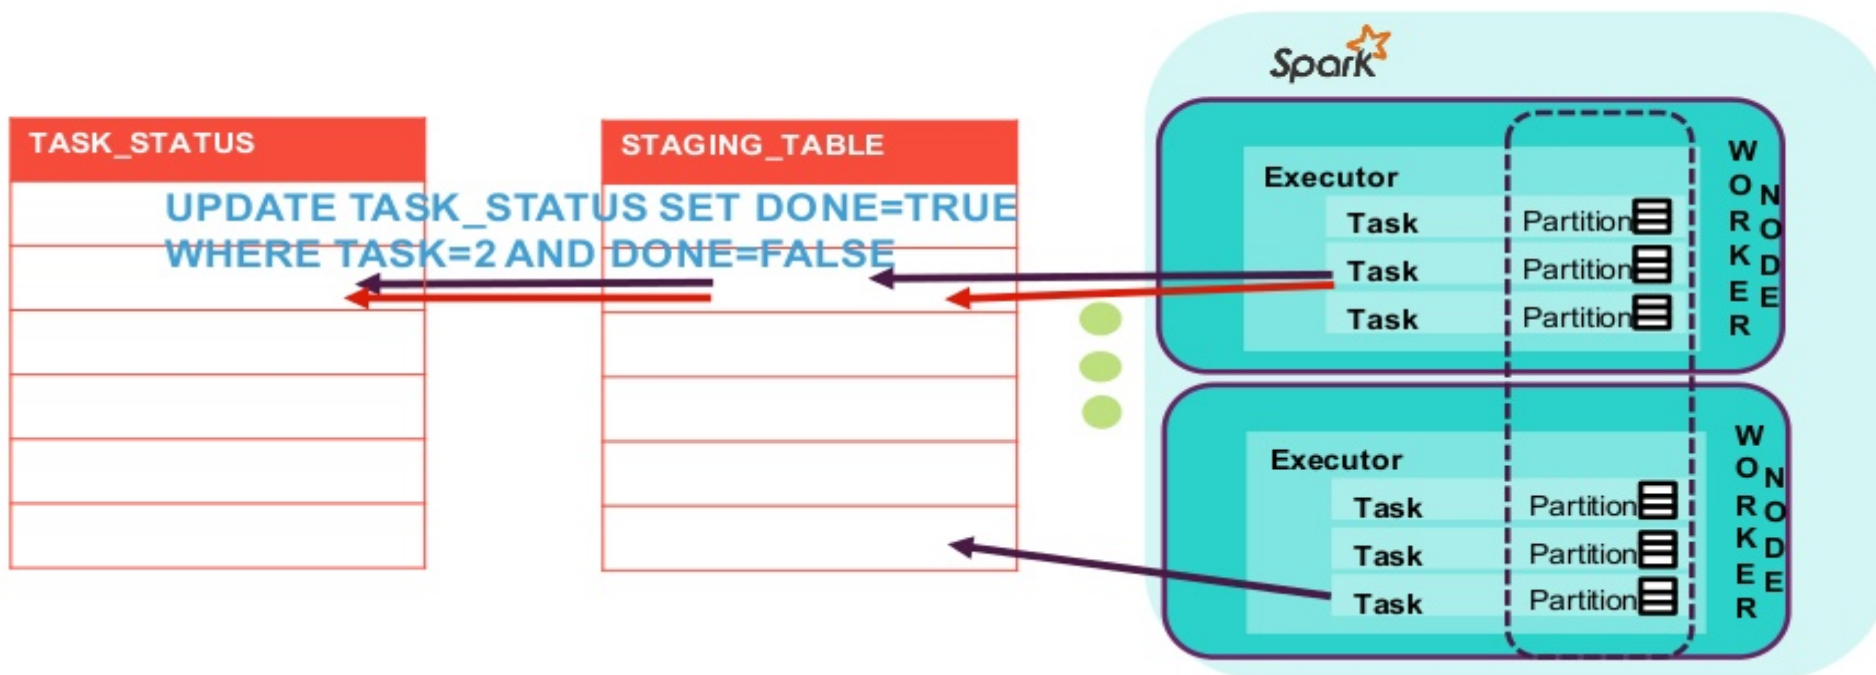
TASK_STATUS

STAGING_TABLE



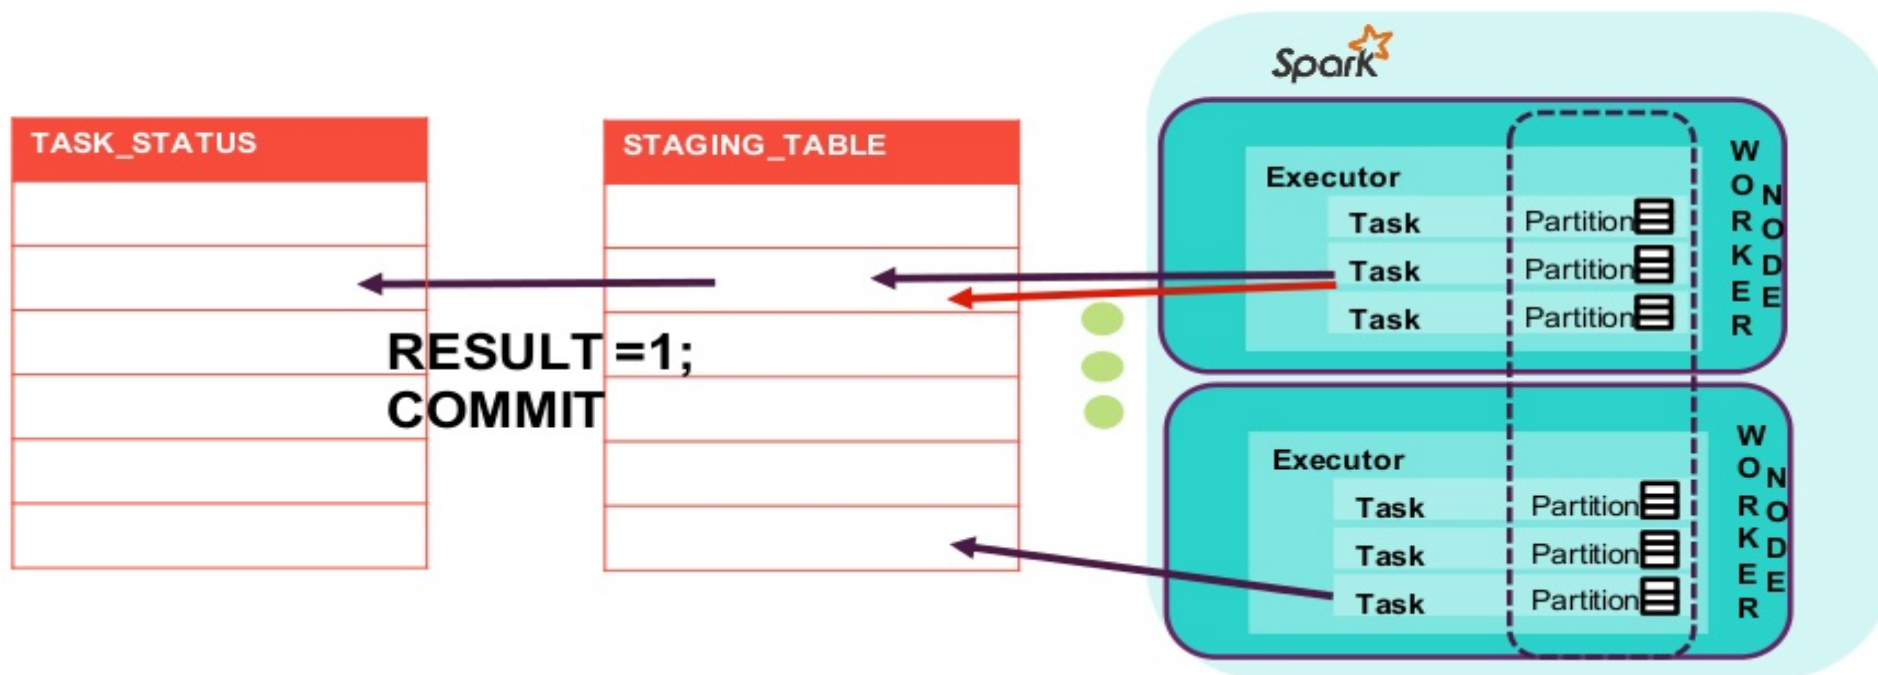
Exactly-once Semantics

Never results in partial or duplicate table loads



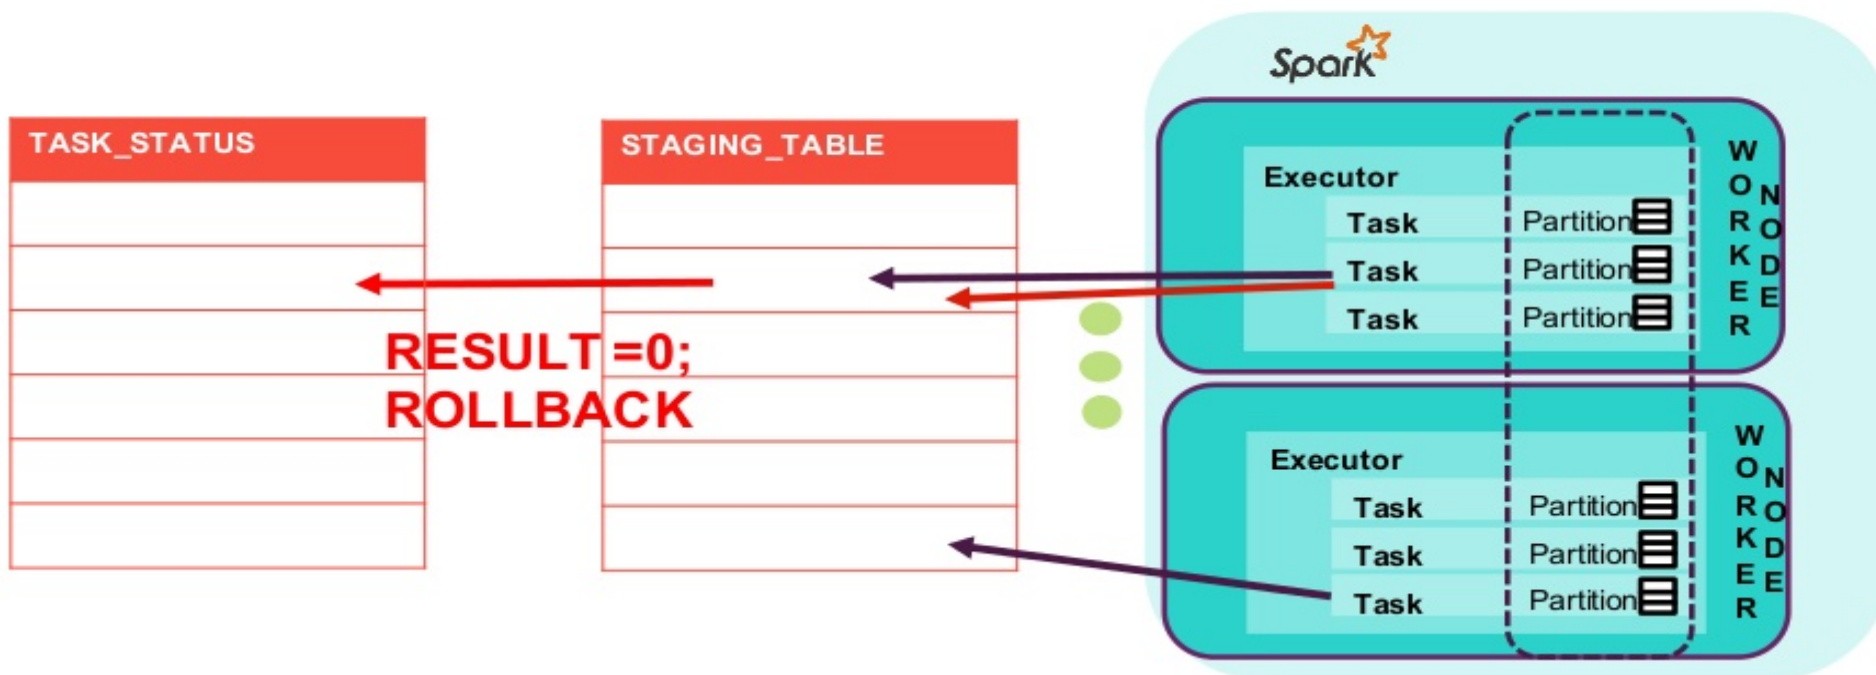
Exactly-once Semantics

Never results in partial or duplicate table loads



Exactly-once Semantics


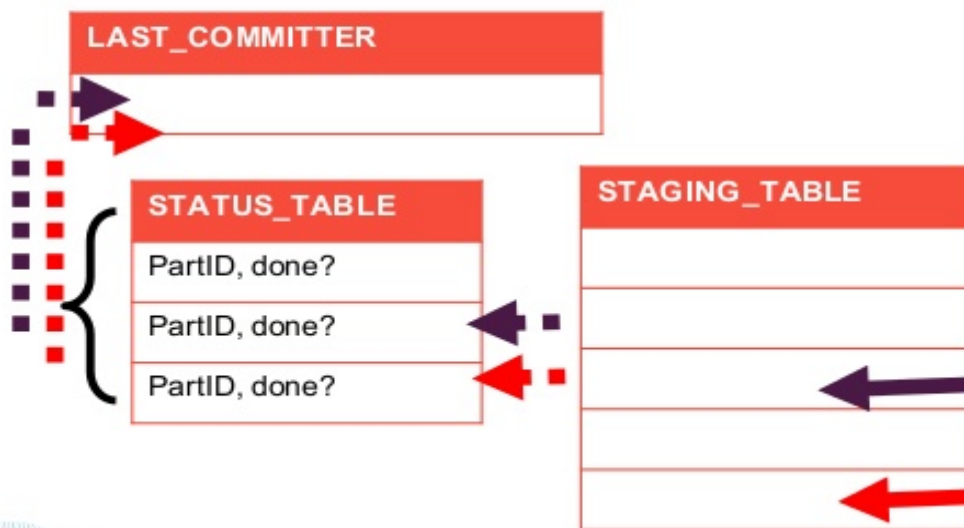
Never results in partial or duplicate table loads



Leader Election

Find Out Which Task Won The Race

UPDATE LAST_COMMITTER SET
ID={MY_ID} WHERE ID = -1


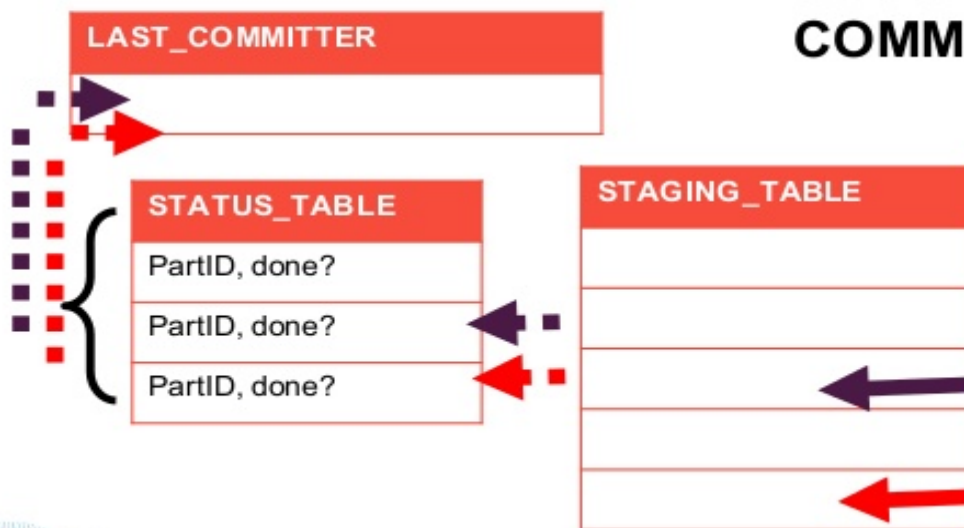


Leader Election

Find Out Which Task Won The Race

UPDATE LAST_COMMITTER SET
ID={MY_ID} WHERE ID = -1

SUCCESS ->
COMMITTER

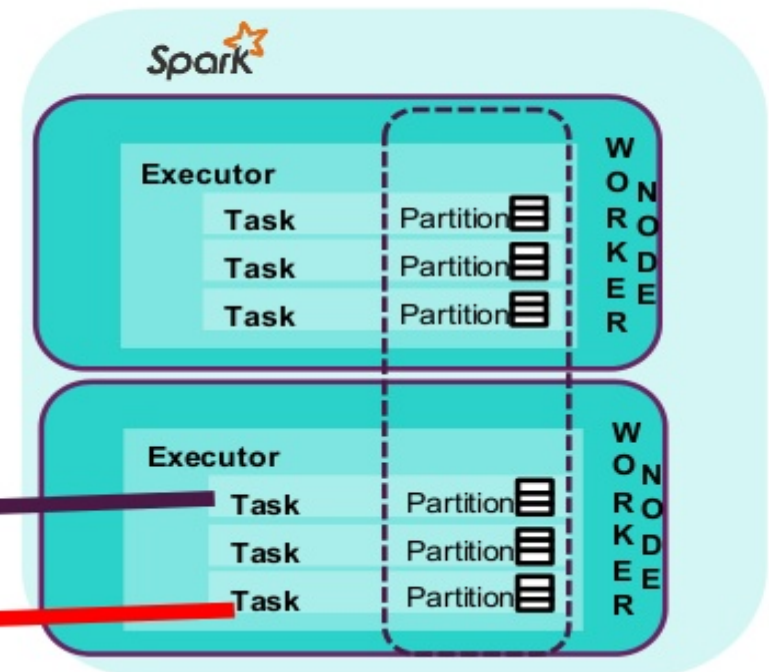
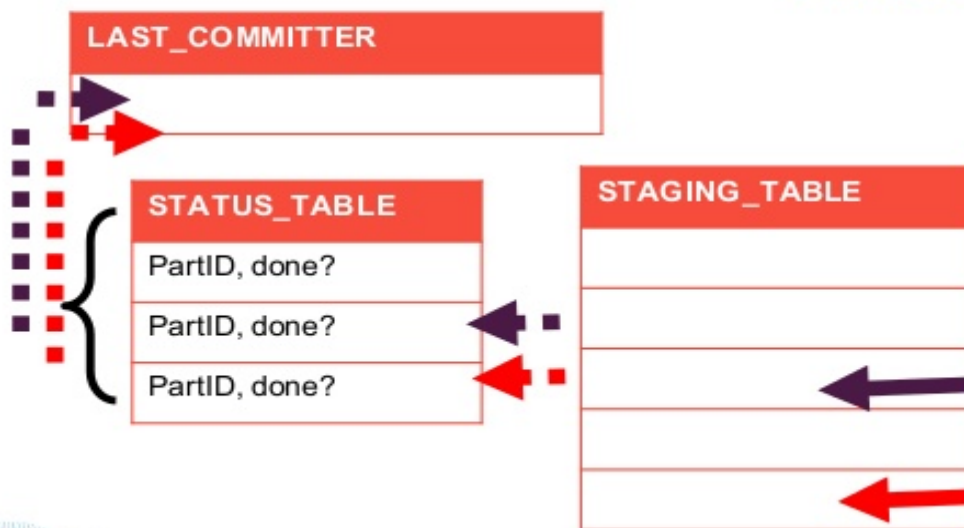


Leader Election

Find Out Which Task Won The Race

UPDATE LAST_COMMITTER SET
ID={MY_ID} WHERE ID = -1

FAIL -> DONE



Move to Target Table



Direct: Pros & Cons

- + Requires no external system beyond Vertica and Spark
- + Less I/O: Does not involve an additional write stage
- Exactly-once semantics are non-trivial to ensure




Big-Picture (Two-Stage)




SPARK SUMMIT 2016


Big-Picture (Two-Stage)



Big-Picture (Two-Stage)



Big-Picture (Two-Stage)




COPY target_table FROM 'hdfs://...' ON ANY NODE ORC;



SPARK SUMMIT 2016

Big-Picture (Two-Stage)



`COPY target_table FROM 'hdfs://...' ON ANY NODE ORC;`



SPARK SUMMIT 2016

Two-Stage: Pros & Cons

- + Utilizes Vertica's built-in parallel COPY command for bulk loads, and may allow for better scalability
- + Uses the optimized ORC reader for parsing (developed in collaboration with Hortonworks).
- + Simpler code logic, since only one query is involved
- More I/O: Requires an additional writing stage.
- Requires an additional dependency on intermediate storage system.



Conclusions

- **Vertica-to-Spark**

- Utilizes hash-range queries to minimize intra-Vertica data movement
- Node-pruning eliminates unneeded distributed query processing in Vertica
- Query pushdown to reduce data transfer and to take advantage of Vertica native analytics.

- **Spark-to-Vertica**

- Direct connector approach efficiently and reliably coordinates transfer with parallel Spark tasks
- A two-stage approach with an intermediary staging storage (HDFS) utilizes a single Vertica query to load and may be more efficient on resource consumption on Vertica side



Learn More About – and Try! - HPE Vertica

- Community Edition
 - Free Download 1TB, 3 nodes
 - my.vertica.com/
- Spark Connector – Beta Program
 - <https://saas.hpe.com/marketplace/big-data/hpe-vertica-connector-apache-spark>



SPARK SUMMIT 2016

THANK YOU.

Rui Liu (r.liu@hpe.com)

Edward Ma (ema@hpe.com)



SPARK SUMMIT 2016
DATA SCIENCE AND ENGINEERING AT SCALE
JUNE 6-8, 2016 SAN FRANCISCO