



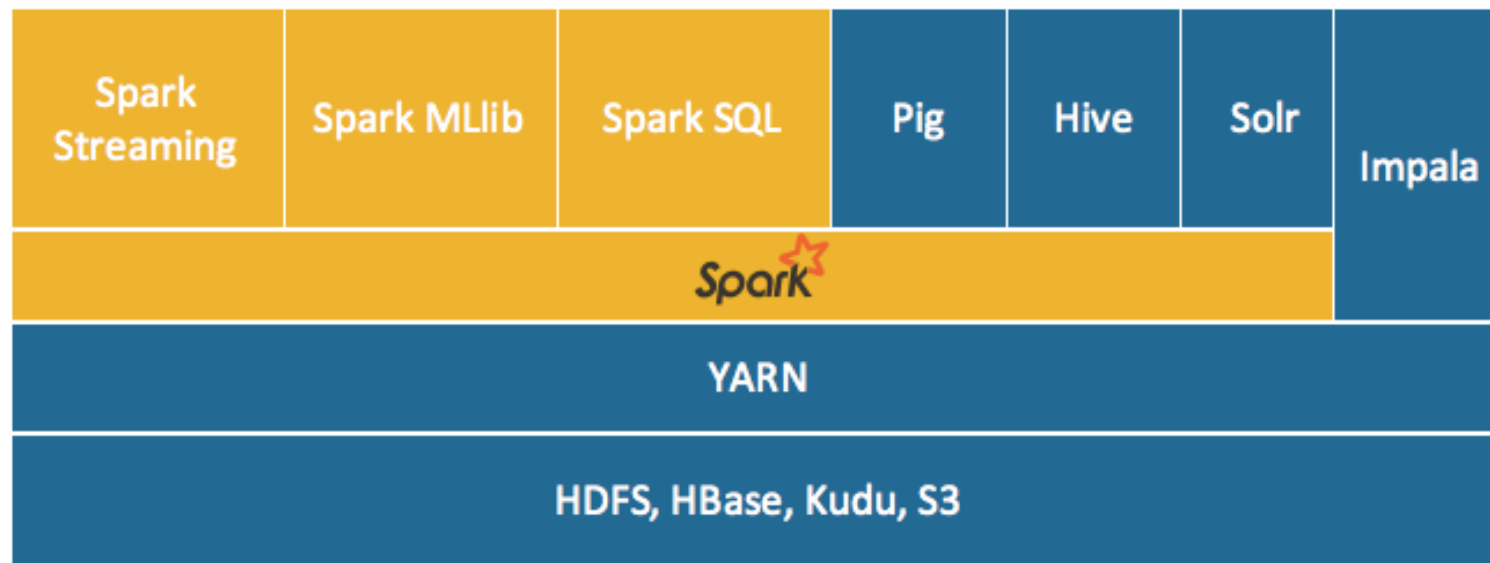
From MapReduce to Apache Spark: An Ecosystem Evolves

Doug Cutting (@cutting)


Chief Architect / Apache Hadoop Co-founder




The Merging Ecosystems



The New Use Cases are Amazing



Genomics
(Broad Institute)



Physics
(CERN)



Healthcare Delivery
(CHOA)

New Contribution: *Livy*, an Open Source REST Service for Spark (Alpha)

- For submitting Spark jobs from any web/mobile app, with no Spark client install — enabling new architectures and use cases
- Provides multi-tenancy and fault tolerance to support multiple users reliably
- Works in standalone mode, with YARN, or with Mesos
- No code changes needed
- Apache License. Initial contributors include employees of Cloudera, Intel, and Microsoft + more wanted!



Bring Your Questions About Livy to the Cloudera Booth



cloudera
Thank you

@cutting

cloudera.com/spark