



Massive Simulations In Spark: Distributed Monte Carlo For Global Health Forecasts

Kyle Foreman, PhD
07 June 2016

Simulations in Spark

- **Context**
- Motivation
- SimBuilder
- Backends
- Benchmarks
- Discussion

What is IHME?


- Independent research center at the University of Washington
- Core funding by Bill & Melinda Gates Foundation and State of Washington
- ~300 faculty, researchers, students and staff
- Providing rigorous, scientific measurement
 - What are the world's major health problems?
 - How well is society addressing these problems?
 - How should we best dedicate resources to improving health?

Goal:
improve health
by providing the best
information
on population health



Global Burden of Disease

- A systematic scientific effort to quantify the comparative magnitude of health loss due to diseases, injuries and risk factors
 - 188 countries
 - 300+ diseases
 - 1990 – 2015+
 - 50+ risk factors

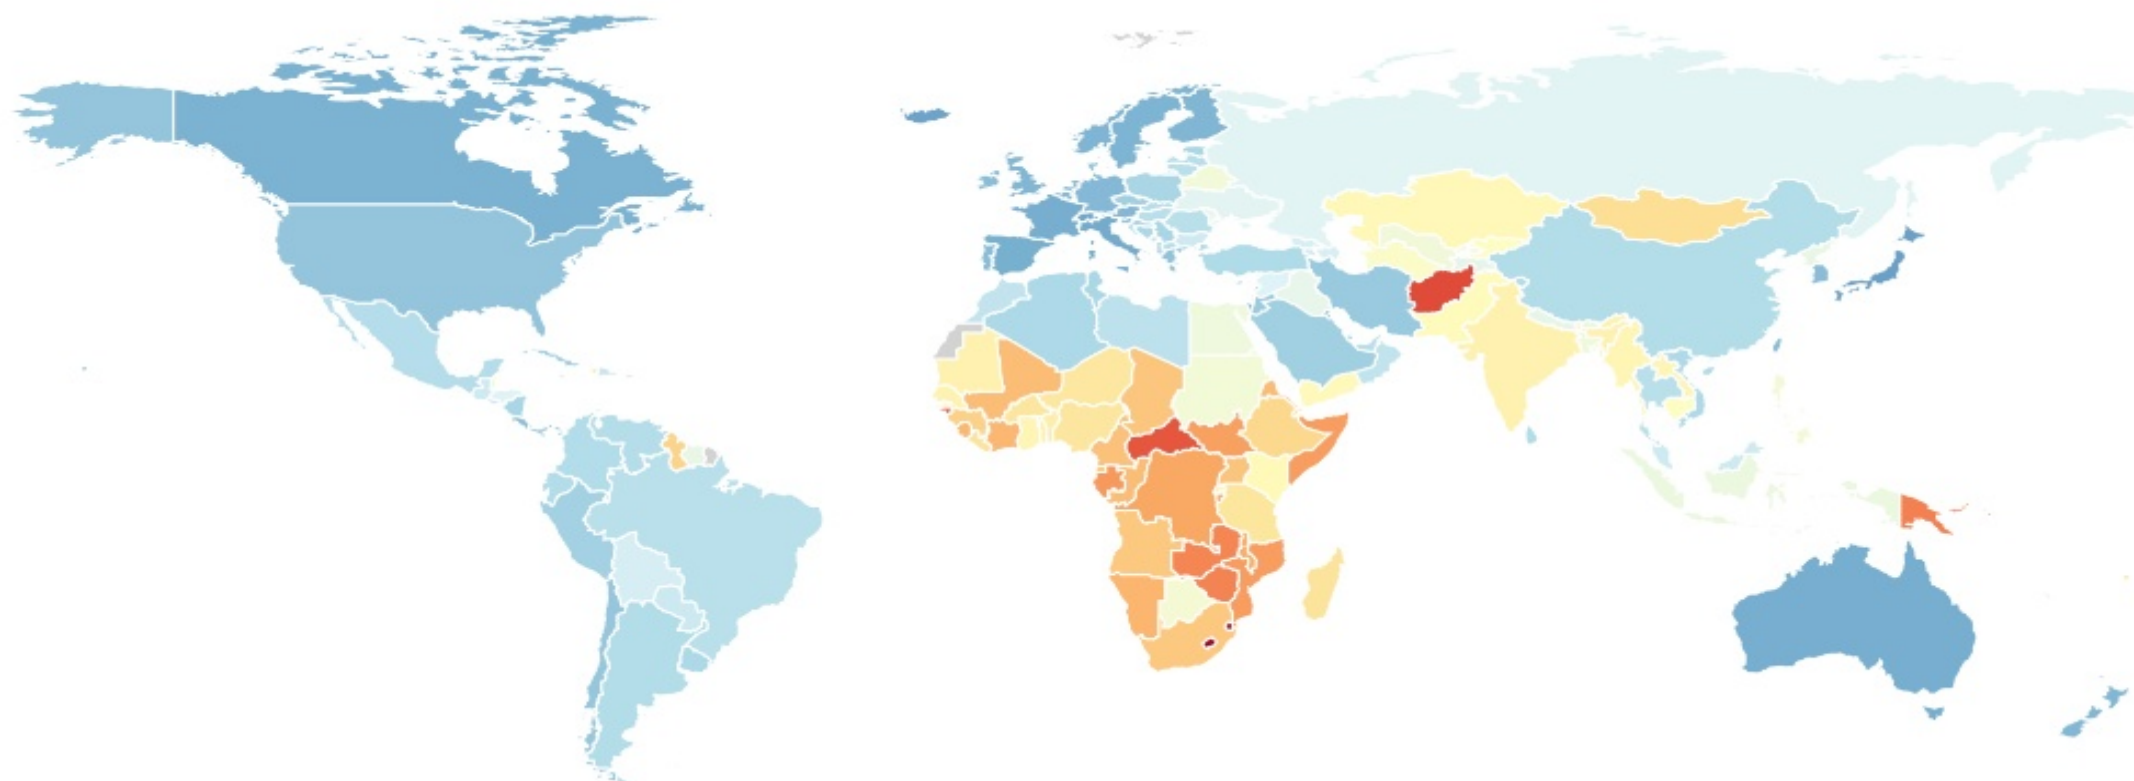


GBD Collaborative Model

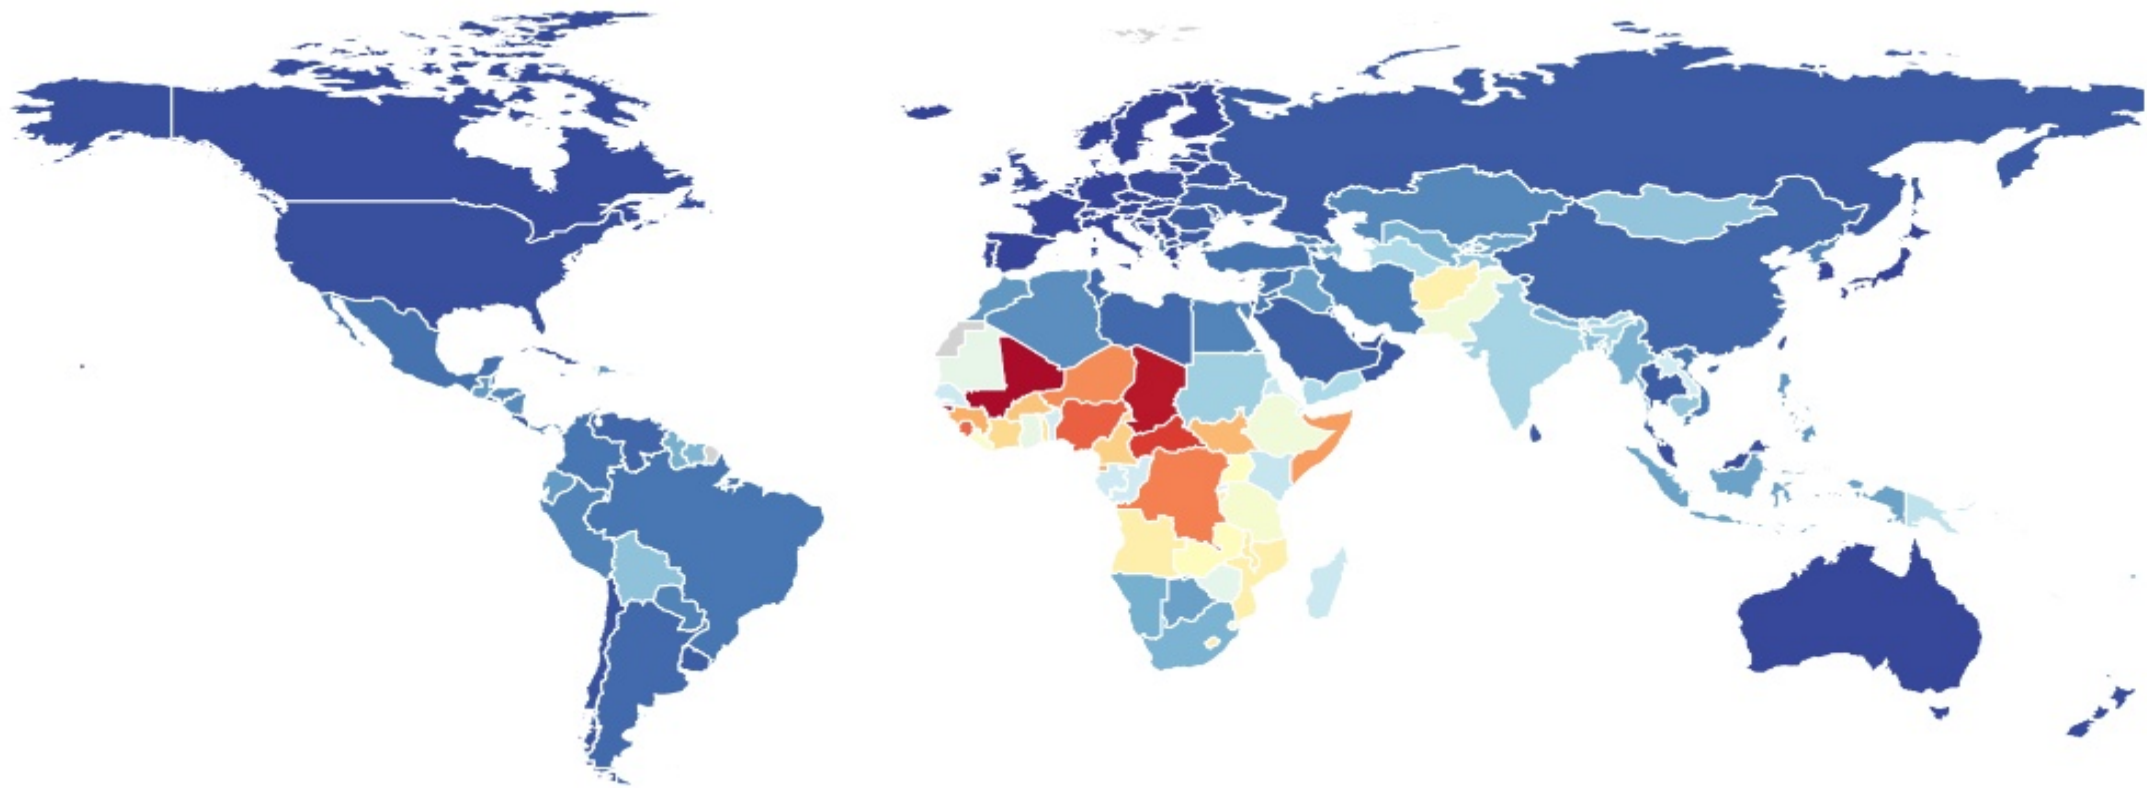


1,100 experts in 106 countries


Death Rate in 2013 (age-standardized)



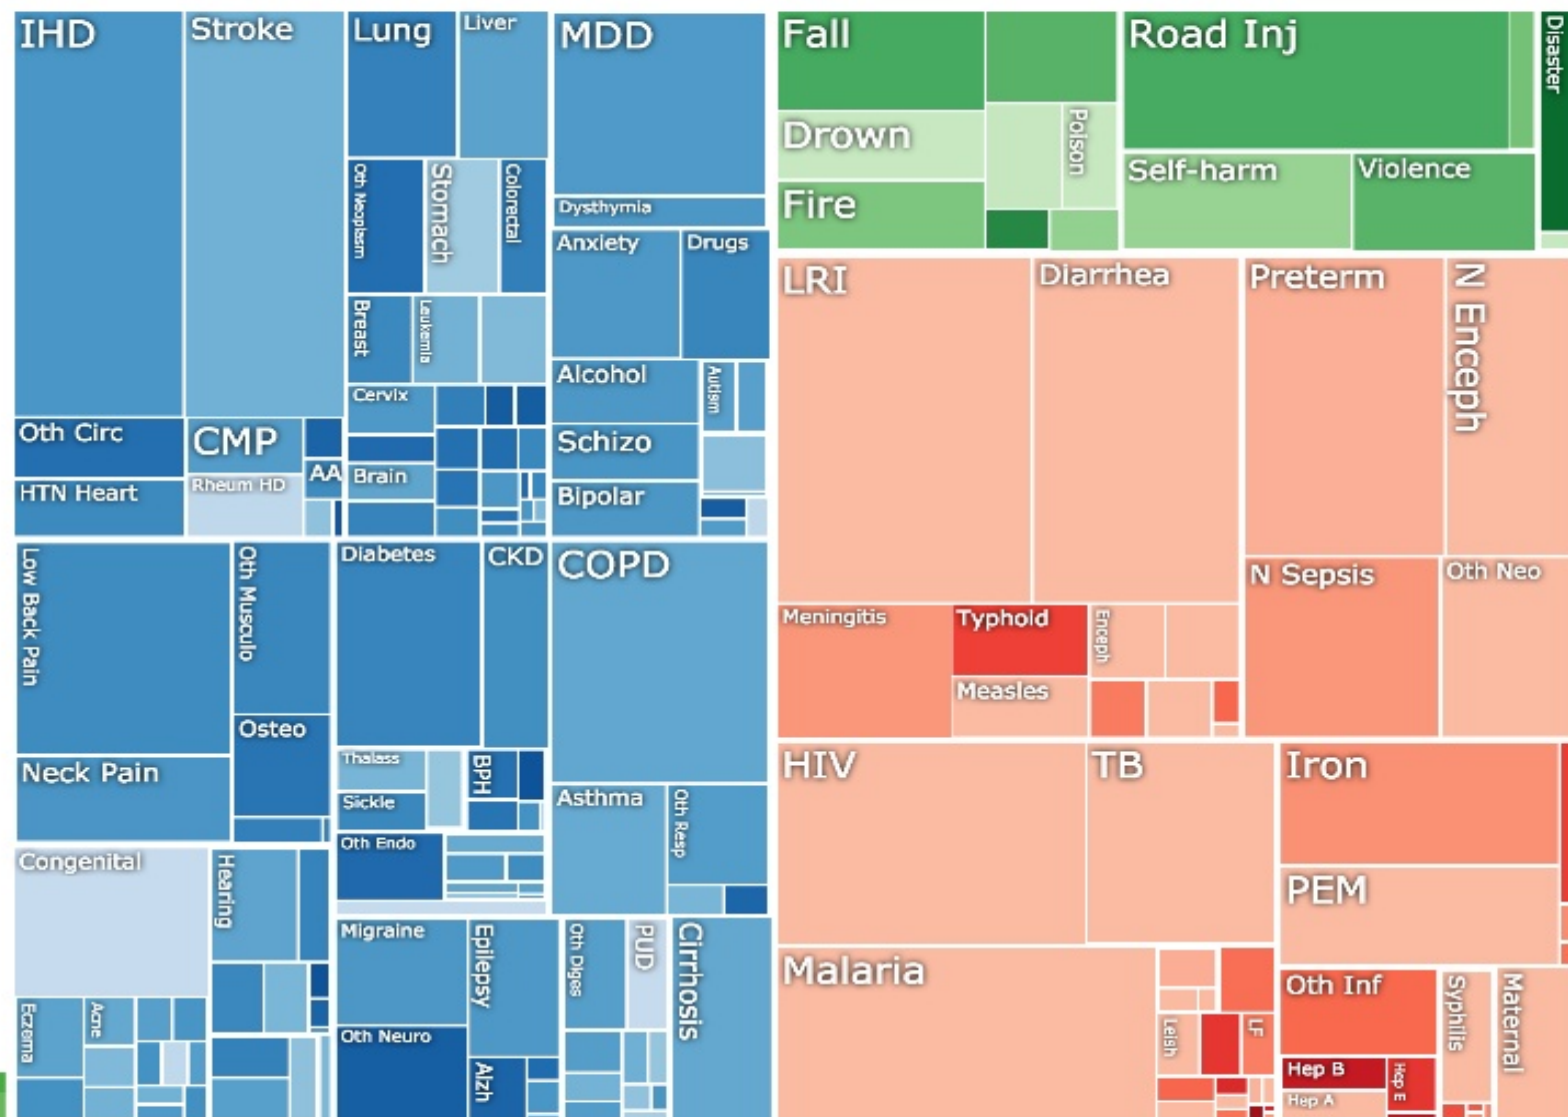
Childhood Death Rate in 2013



DALYs in High Income Countries (2010)



DALYs in Low & Middle Income Countries (2010)



Simulations in Spark

- Context
- **Motivation**
- SimBuilder
- Backends
- Benchmarks
- Discussion

Forecasting the Burden of Disease

1. Generate a baseline scenario projecting the GBD 25 years into the future
 - Mortality, morbidity, population, and risk factors
 - Every country, cause, sex, age
2. Create a comprehensive simulation framework to assess alternative scenarios
 - Capture the complex interrelationships between risk factors, interventions, and diseases to explore the effects of changes to the system
3. Build a modular and flexible platform that can incorporate new models to answer detailed “what if?” questions
 - E.g. introduction of a new vaccine, effects of global warming, scale up of ART coverage, risk of global pandemics, etc

Simulations

- Takes into account interdependencies between risk factors, diseases, etc.
 - Use draws of model parameters to advance from t to $t+1$, allowing for forecasts of other quantities (e.g. risk factors and mortality) to interact
- Modular structure allows for detailed “what ifs”
 - E.g. scale up of coverage of a new intervention
- Allows us to incorporate alternative models to test sensitivity to model choice and specification

Basic Simulation Process

1. Fit statistical models (**PyMB**) capturing most important relationships as statistical parameters
2. Generate many draws from those parameters, reflecting their uncertainty and correlation structure
3. Run Monte Carlo simulations to update each quantity over time, taking into account its dependencies

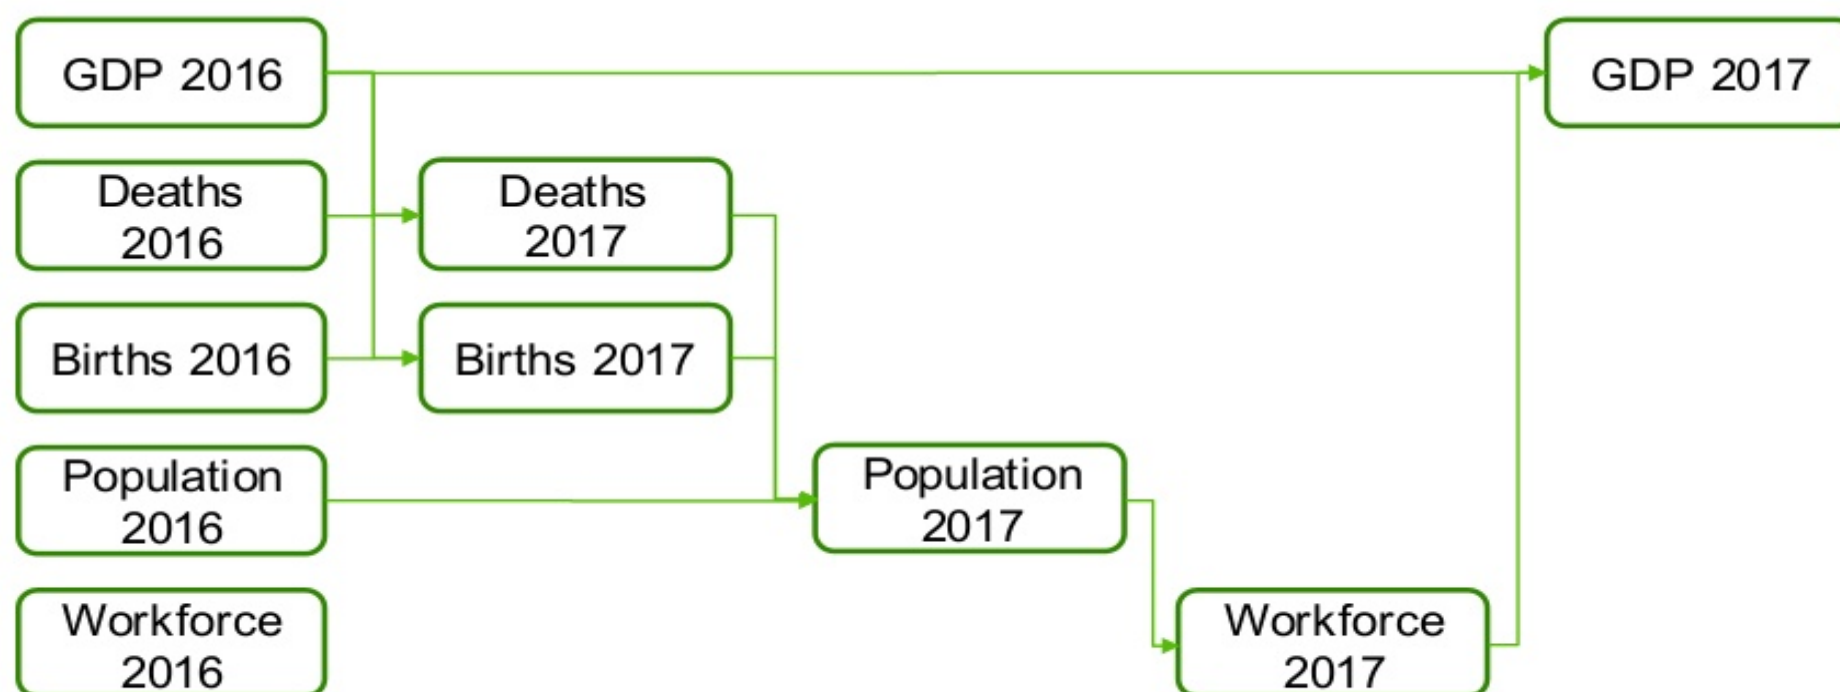
Simulations in Spark

- Context
- Motivation
- **SimBuilder**
- Backends
- Benchmarks
- Discussion

SimBuilder

- Directed Acyclic Graph (DAG) construction and validation
 - **yaml** and **sympy** specification
 - **curses** interface
 - **graphviz** visualization
- Flexible execution backends
 - **pandas**
 - **pyspark**

Example Global Health DAG



Creating a DAG with SimBuilder

1. Create a YAML for the overall simulation
 - Specify child models and the dimensions of the simulation
2. Build separate YAML files for each child model
 - Dimensions along which the model varies
 - Location of historical data
 - Expression for how to generate the quantity in $t+1$
3. Run SimBuilder to validate and explore DAG

Creating a DAG with SimBuilder

sim.yaml

name: gdp_pop

models:

- **name:** gdp_per_capita
- **name:** workforce
- **name:** mortality
- **name:** gdp
- **name:** pop

global_parameters:

- **name:** loc
set: List
values: ["USA", "MEX", "CAN", ...,]
- **name:** t
set: Range
start: 2014
end: 2040
- **name:** draw
set: Range
start: 0
end: 999

Creating a DAG with SimBuilder

pop.yaml

```
name: pop
version: 0.1
variables: [loc,t]
history:
  - type: csv
    path: pop.csv
```

Creating a DAG with SimBuilder

gdp_per_capita.yaml

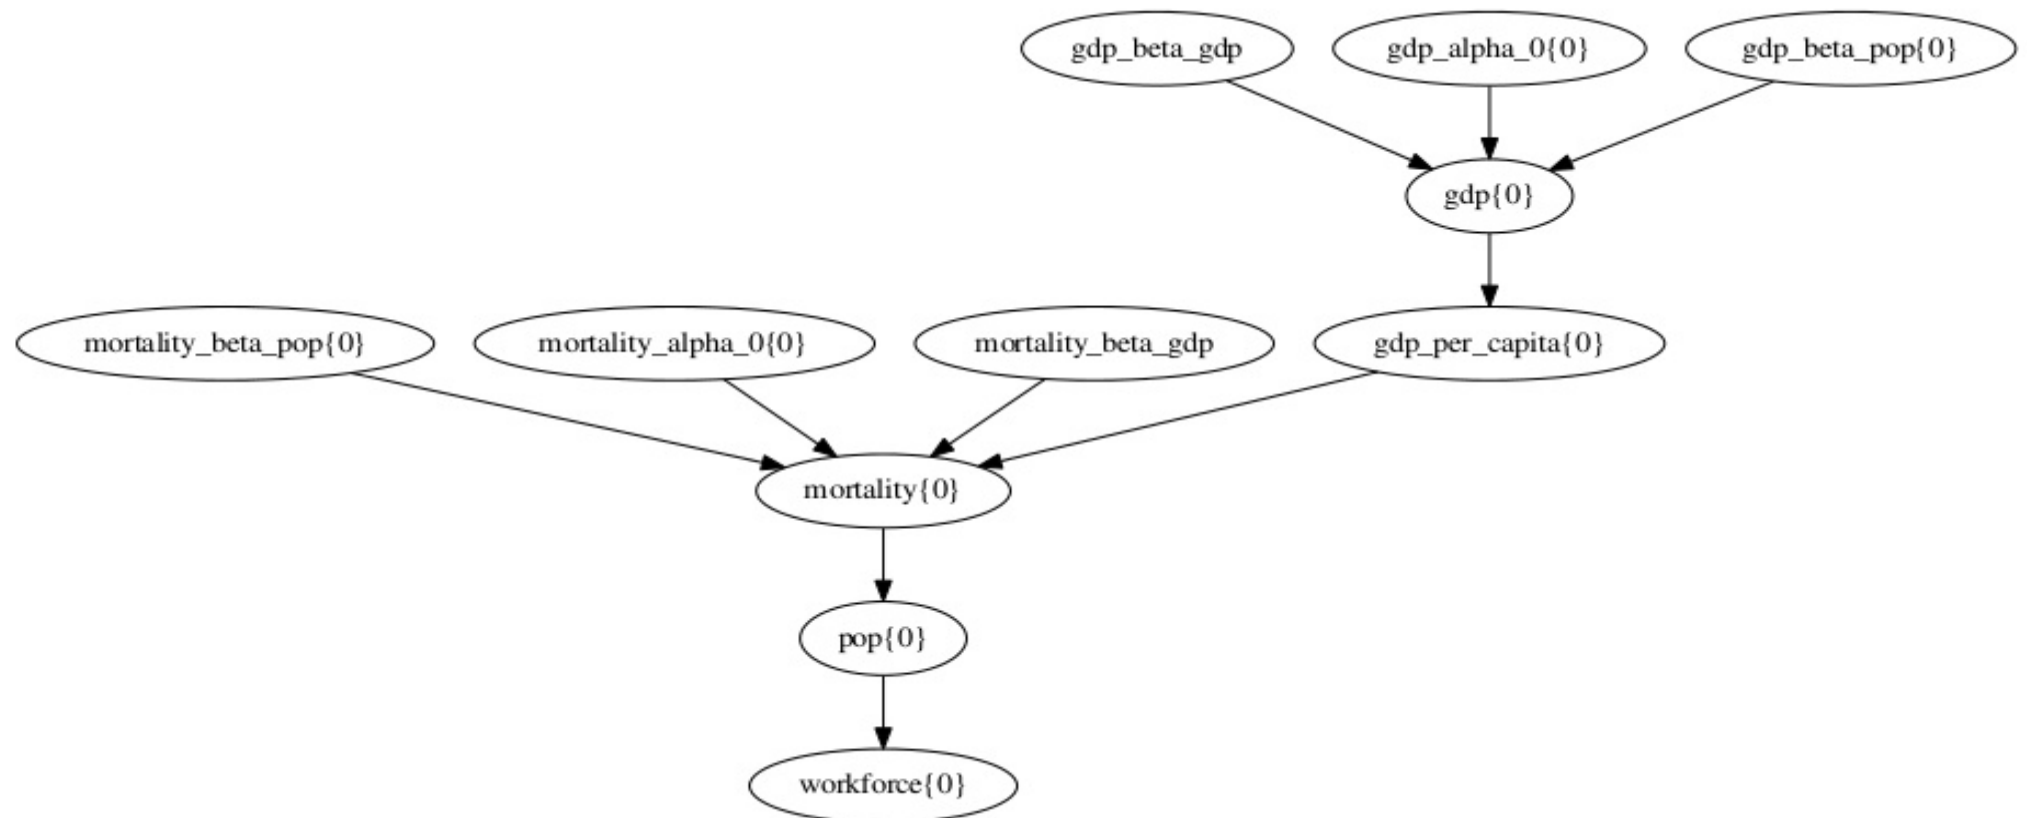
```
name: gdp_per_capita
version: 0.1
expr: gdp(draw, t, loc)/pop(draw, t-1, loc)
variables: [draw, t, loc]
history:
  - type: csv
    path: gdp_pc_w_draw.csv
```

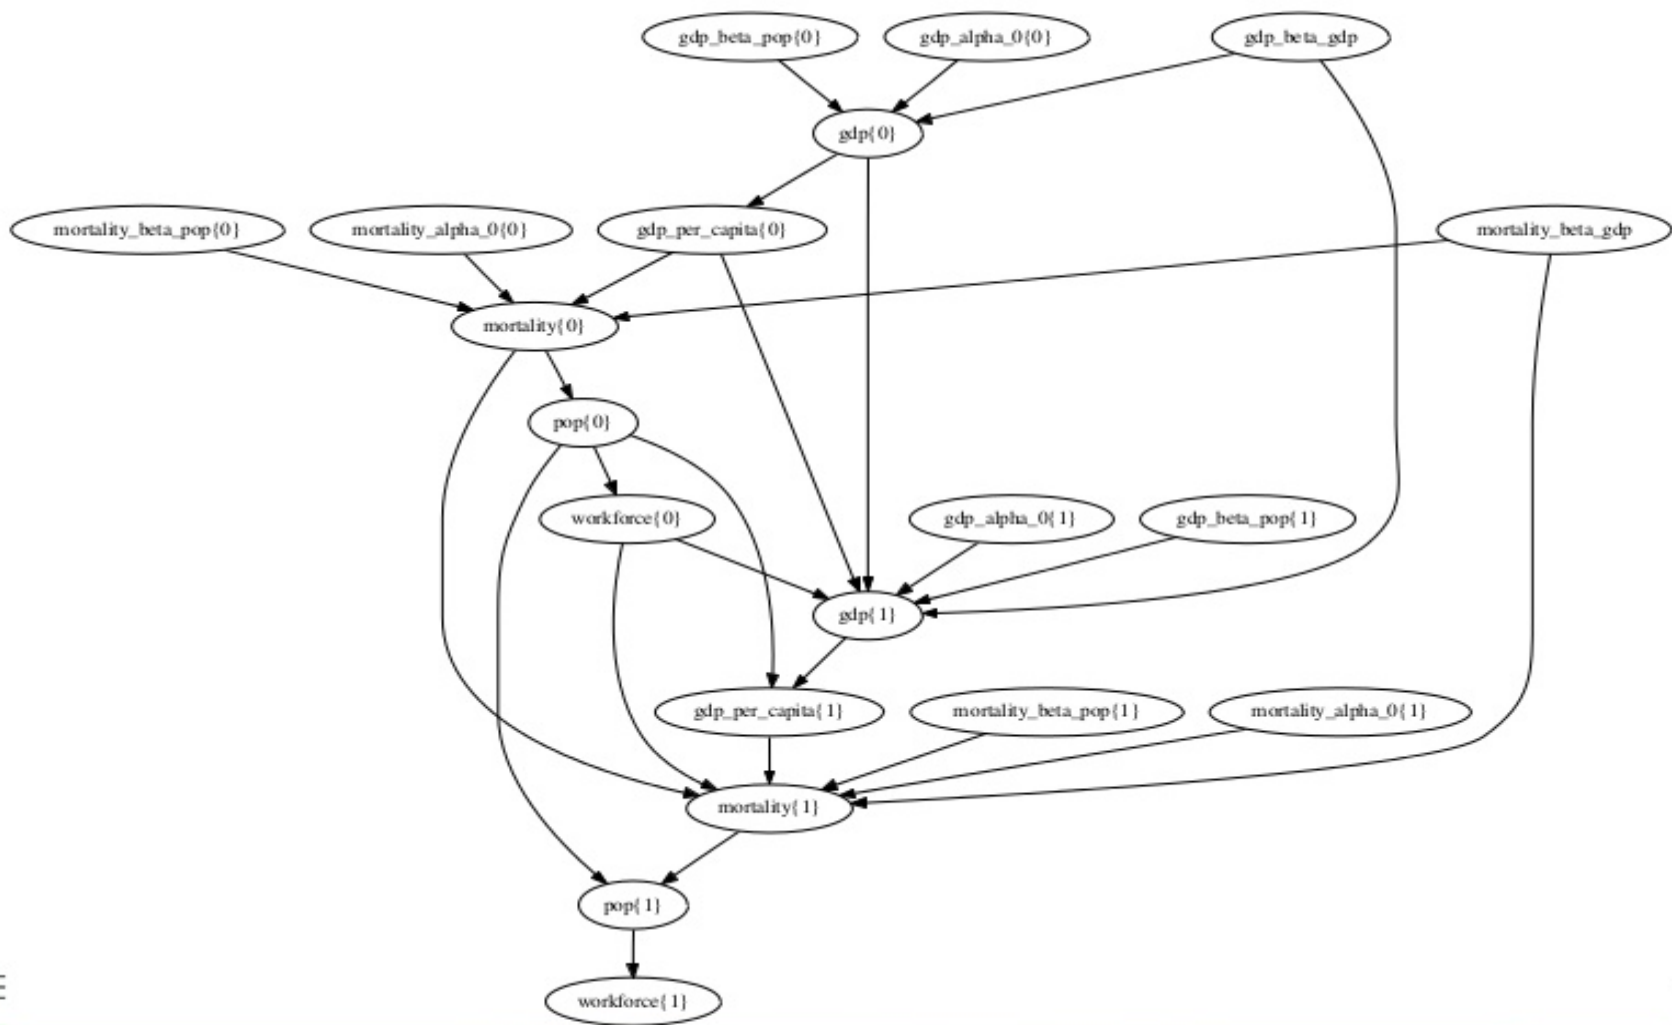
Creating a DAG with SimBuilder


gdp.yaml

```
name: gdp
version: 0.1
expr: gdp(d,loc,t-1) * exp(
    alpha_0(loc, t, draw) +
    beta_gdp(draw) * log(gdp_per_capita(draw, loc, t-1)) +
    beta_pop(draw, t) * workforce(loc, t)
)
variables: [draw, loc, t]
history:
  - type: csv
    path: gdp_w_draw.csv
model_parameters:
  - name: alpha_0
    variables: [draw, loc, t]
    history:
      - type: csv
        path: gdp/alpha_0.csv
  - name: beta_gdp
    variables: [d]
    history:
      - type: csv
        path: gdp/beta_gdp.csv
  - name: beta_pop
    variables: [d]
    history:
      - type: csv
        path: gdp/beta_pop.csv
```

DAG







Simulations in Spark

- Context
- Motivation
- SimBuilder
- **Backends**
- Benchmarks
- Discussion

SimBuilder Backends

- SimBuilder traverses the DAG
 - Backends plugin via API
- Backends provide
 - Data loading strategy
 - Methods to combine data from different nodes to calculate new nodes
 - Method to save computed data

Local Pandas Backend

- Useful for debugging and small simulations
- Master process maintains list of Model objects:
 - **pandas** DataFrames of all model data
 - Indexed by global variables
 - **sympy** expression for how to calculate $t+1$
 - Vectorized when possible, fallback to row-wise apply
- Simulating over time traverses DAG to join necessary DataFrames and compute results to fill in new columns

Spark DataFrame Backend

- Similar to local backend, swapping in Spark's DataFrame for Pandas'
- Spark context maintains list of Model objects:
 - **pyspark** DataFrames of all model data
 - Columns for each global variable and
 - **sympy** expression for how to calculate $t+1$
 - Calculated using row-wise apply
- Joins DataFrames where necessary

Spark + Numpy Backend

- Uses Spark to distribute and schedule the execution of vectorized **numpy** operations
- One RDD for each node and year:
 - Containing an ndarray and index vector for each axis
 - Can optionally partition the data array into key-value pairs along one or more axes (similar to **bolt**)
- To calculate a new node, all the parent RDD's are unioned and then reduced into a new child RDD

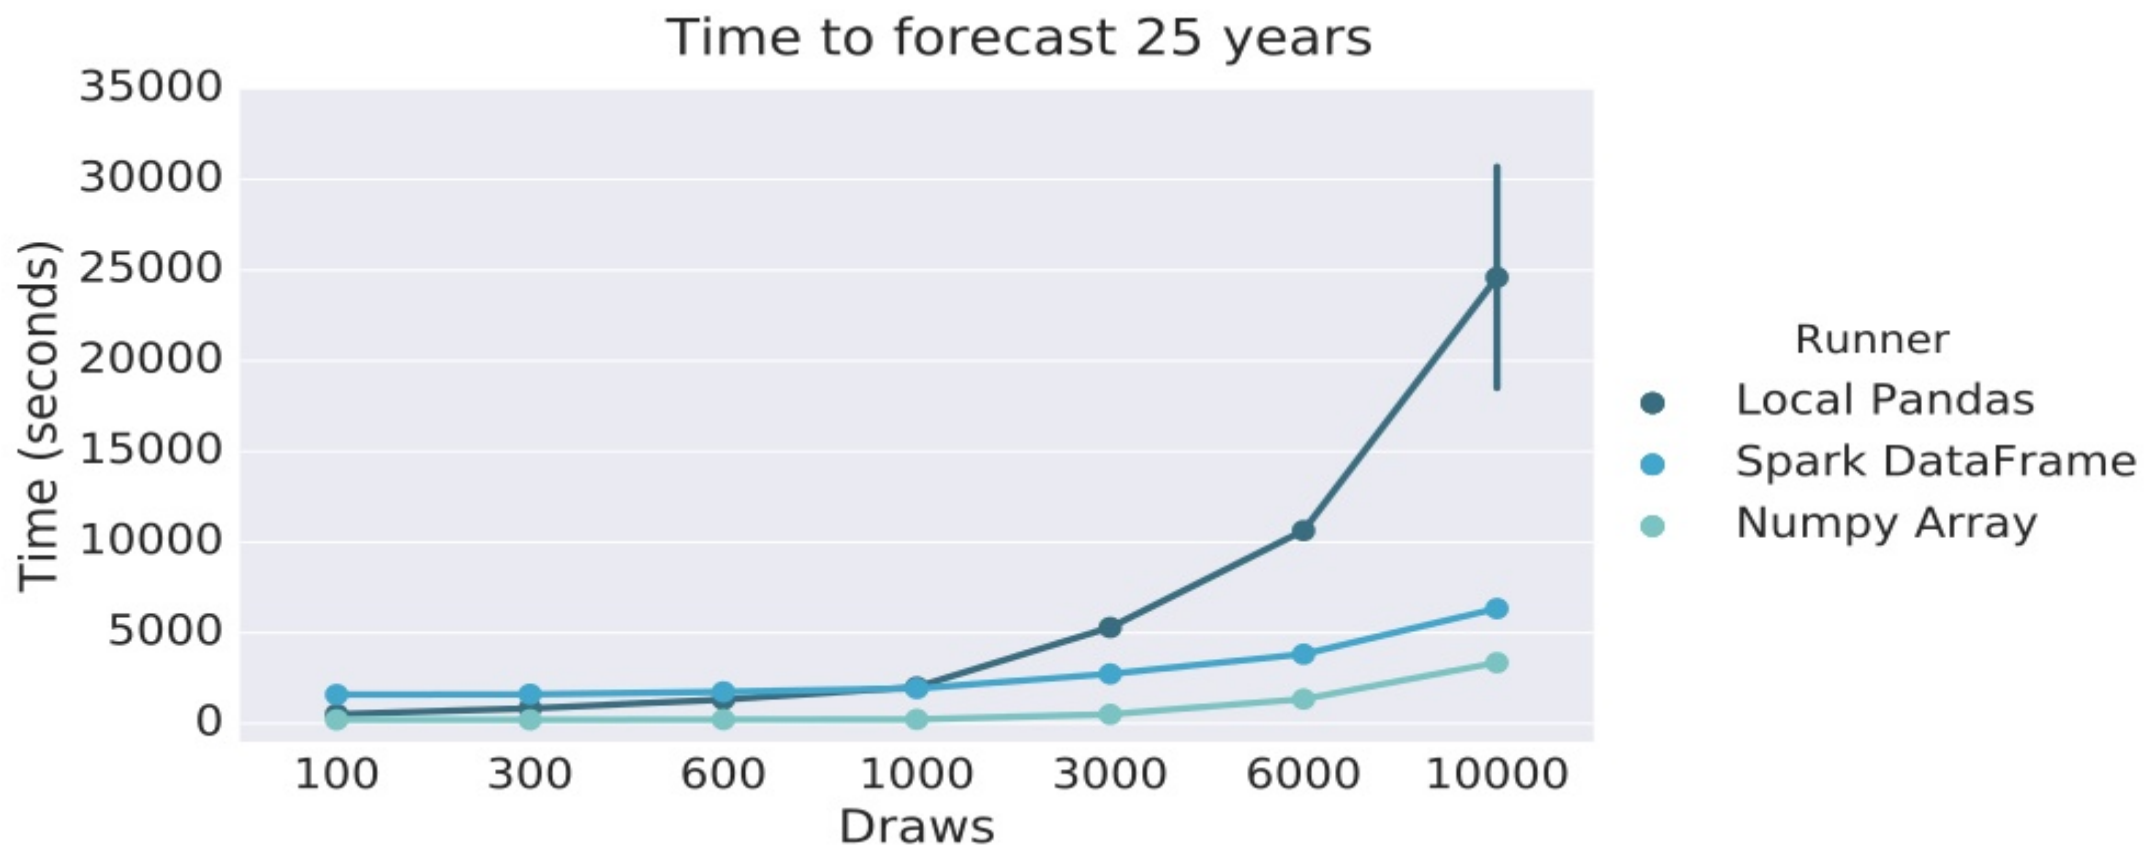
Simulations in Spark

- Context
- Motivation
- SimBuilder
- Backends
- **Benchmarks**
- Discussion

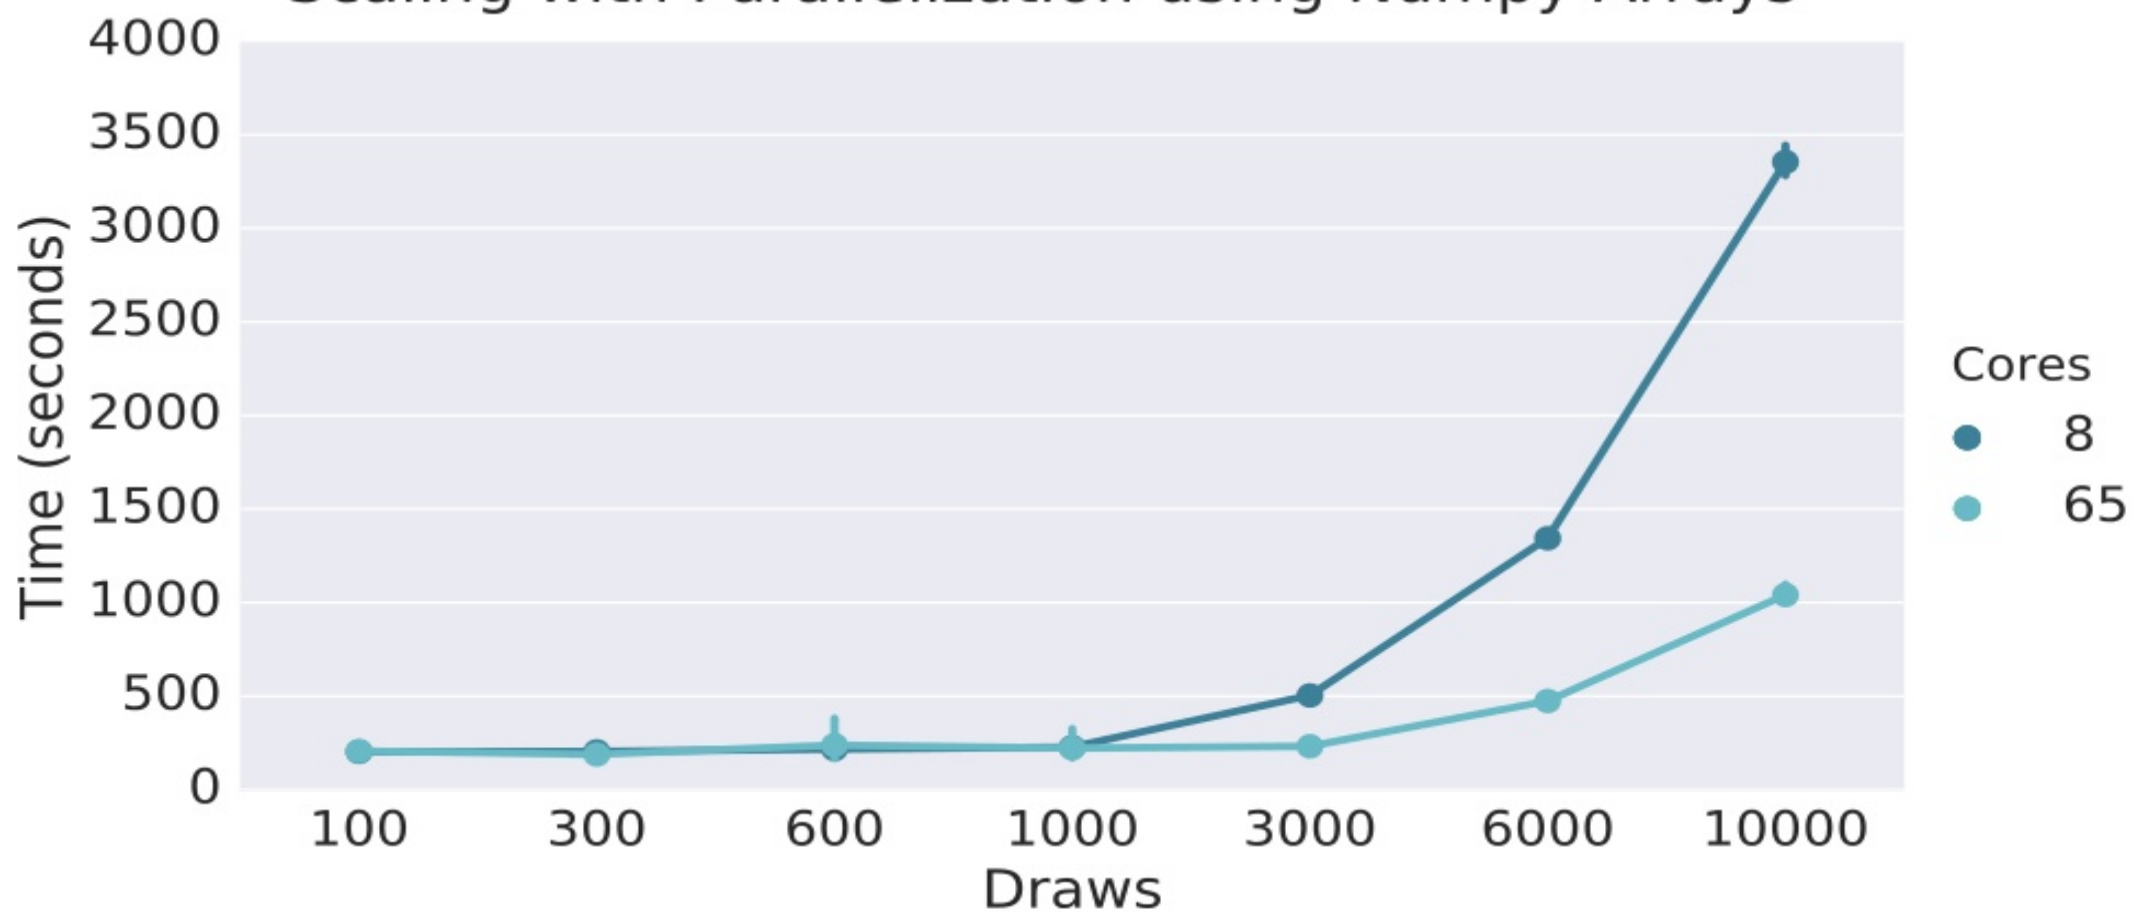
Benchmarking

- Single executor (8 cores, 32GB)
- Synthetic data
 - 65 nodes in DAG
 - 3 countries x 20 ages x 2 sexes x N draws
 - Forecasted forward 25 years

Benchmarks



Scaling with Parallelization using Numpy Arrays



Simulations in Spark

- Context
- Motivation
- SimBuilder
- Backends
- Benchmarks
- **Discussion**

Limitations of Spark DataFrame

- Majority of execution time is spent joining the DataFrame and aligning the dimension columns
- These times were achieved after careful tuning of partitions
 - Perhaps custom partitioners could reduce runtime further

Taking Advantage of Numpy

- For multidimensional datasets of consistent shape and size, simply aligning indices can eliminate the overhead of joins
 - Including generalizations to axes of size 1 to simplify coding
- Numpy's vectorized operations are highly tuned and scale well
 - Including multithreading when e.g. compiled against **MKL**

Future Development

- Experiment with better partitioning of Numpy RDDs
 - Perhaps use **bolt**?
- Improve partial graph execution
 - Executing the entire DAG at once often crashes the driver
 - Executing each year's partial DAG in sequence results in idle CPU time towards the end of that DAG
- Investigate more efficient Spark DataFrame joining methods for Panel data

Team



Kyle Heuton
krheuton@uw.edu



Chun-Wei Yuan
cwyuan@uw.edu



Kyle Foreman
kfor@uw.edu

healthdata.org