

Solr as a Spark SQL Datasource

Kiran Chitturi

Lucidworks



SPARK SUMMIT 2016
DATA SCIENCE AND ENGINEERING AT SCALE



The standard for
enterprise search.



Solr as a Spark SQL Data Source

- Read/write data from/to Solr as DataFrame
- Use Solr Schema API to access field-level metadata
- Push predicates down into Solr query constructs, e.g. fq clause
- Shard partitioning, intra-shard splitting, streaming results

```
// Connect to Solr
val opts = Map("zkhost" -> "localhost:9983", "collection" -> "nyc_trips")
val solrDF = sqlContext.read.format("solr").options(opts).load

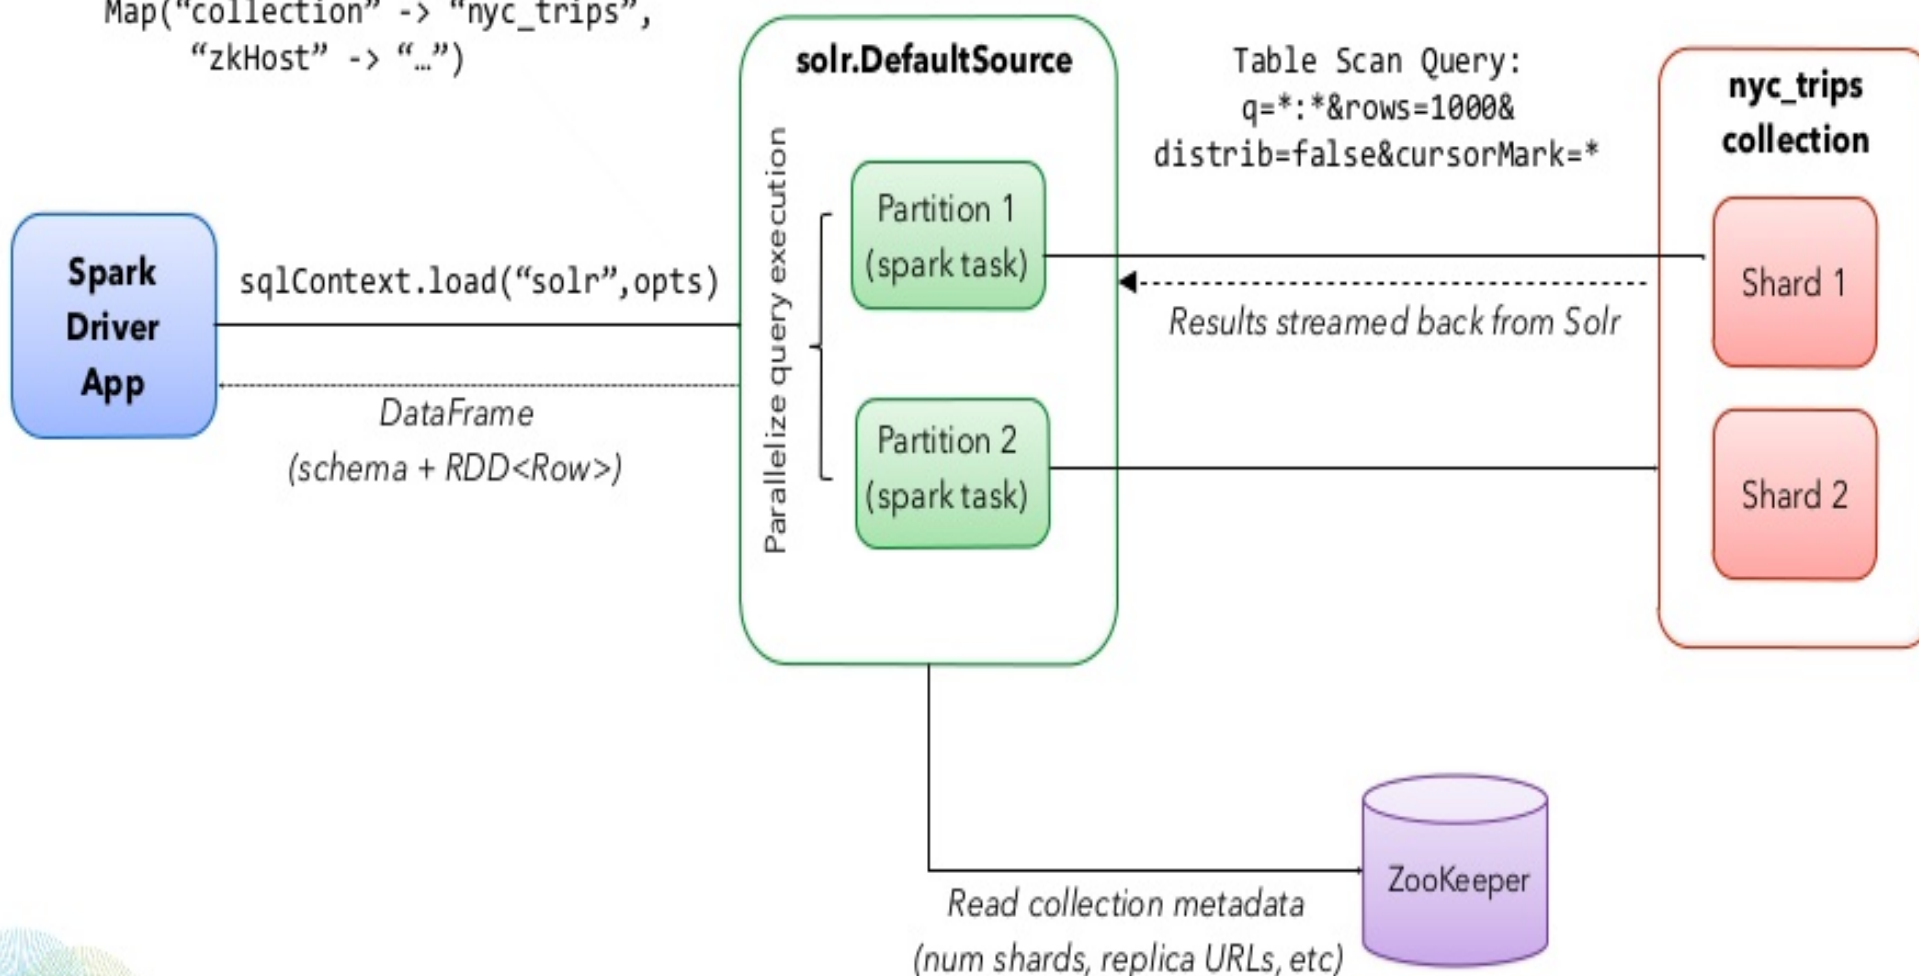
// Register DF as temp table
solrDF.registerTempTable("trips")

// Perform SQL queries
sqlContext.sql("SELECT avg(tip_amount), avg(fare_amount) FROM trips").show()
```



SolrRDD: Reading data from Solr into Spark

```
Map("collection" -> "nyc_trips",  
    "zkHost" -> "...")
```



SolrRDD: Reading data from Solr into Spark

```
Map("collection" -> "nyc_trips",  
    "zkHost" -> "...")
```

Table Scan Query:

```
q=*&rows=1000&  
distrib=false&cursorMark=*  
&split_field:[aa TO ba}
```

solr.DefaultSource

Parallelize query execution

Partition 1
(spark task)

Partition 2
(spark task)

Partition 3
(spark task)

Partition 4
(spark task)

&split_field:[ba TO bz]

&split_field:[aa TO ac}

&split_field:[ac TO ae]

**nyc_trips
collection**

Shard 1

Shard 2

ZooKeeper

Read collection metadata

sqlContext.load("solr",opts)

DataFrame
(schema + RDD<Row>)



Solr Streaming API for fast reads

- Contributed to spark-solr by Bloomberg team
- Extremely fast “table scans” over large result sets in Solr
- Relies on a column-oriented data structure in Lucene: docValues
- DocValues help speed up faceting and sorting too!
- Push SQL predicates down into Solr’s Parallel SQL engine available in Solr 6.x (Coming soon)



Data-locality Hint

- **SolrRDD** extends **RDD**[SolrDocument] (written in Scala)

- Give hint to Spark task scheduler about where data lives

```
override def getPreferredLocations(split: Partition): Seq[String] = {  
  
  // return preferred hostname for a Solr partition  
}
```

- Useful when Spark executor and Solr replicas live on same physical host, as we do in Fusion
- Query to a shard has a “preferred” replica; can fallback to other replicas if the preferred goes down (will be in 2.1)



Writing to Solr (aka indexing)

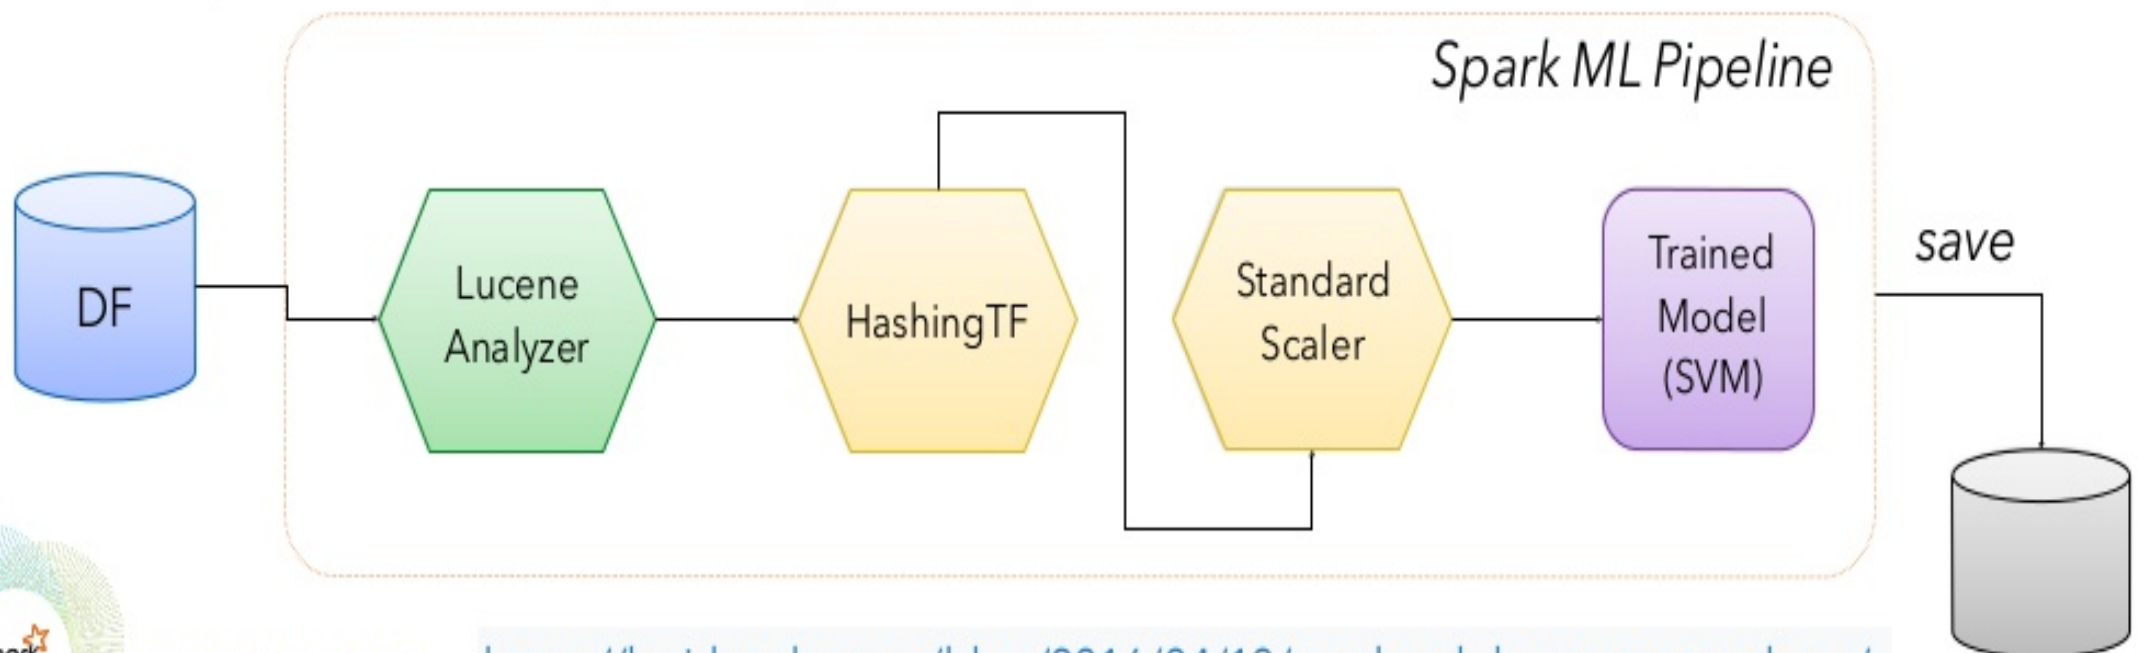
- Cloud-aware client sends updates to shard leaders in parallel
- Solr Schema API used to create fields on-the-fly using the DataFrame schema
- Better parallelism than traditional approaches like Solr DIH

```
val dbOpts = Map(  
  "url" -> "jdbc:postgresql:mydb",  
  "dbtable" -> "schema.table",  
  "partitionColumn" -> "foo",  
  "numPartitions" -> "10")  
  
val jdbcDF = sqlContext.read.format("jdbc").options(dbOpts).load  
  
val solrOpts = Map("zkhost" -> "localhost:9983", "collection" -> "mycoll")  
  
jdbcDF.write.format("solr").options(solrOpts).mode(SaveMode.Overwrite).save
```



Solr / Lucene Analyzers for Spark ML Pipelines

- Spark ML Pipeline provides nice API for defining stages to train / predict ML models
- Crazy idea ~ use battle-hardened Lucene for text analysis in Spark
- Pipelines support import/export
- Can try different text analysis techniques during cross-validation



Lucene Text Analysis

- Use Lucene text analyzers for text analysis
- Wide selection of tokenizers, filters, Character filters in Lucene
- JSON declared schema analysis definition

```
{ "analyzers": [{ "name": "...",  
                  "charFilters": [{ "type": "...", ... }, ... ],  
                  "tokenizer": { "type": "...", ... },  
                  "filters": [{ "type": "...", ... } ... ] } ],  
  { "name": "...",  
    "charFilters": [{ "type": "...", ... }, ... ],  
    "tokenizer": { "type": "...", ... },  
    "filters": [{ "type": "...", ... }, ... ] } ],  
  "fields": [{ "name": "...", "analyzer": "...",  
               { "regex": ".*", "analyzer": "...", ... } ] }
```



Lucene Text Analysis (example)

- Word count example of spark-solr README file

```
import com.lucidworks.spark.analysis.LuceneTextAnalyzer
val schema = """{ "analyzers": [{ "name": "StdTokLower",
                                |                               "tokenizer": { "type": "standard" },
                                |                               "filters": [{ "type": "lowercase" }] }],
                  | "fields": [{ "regex": ".+", "analyzer": "StdTokLower" }] }
                """.stripMargin
val analyzer = new LuceneTextAnalyzer(schema)
val file = sc.textFile("README.adoc")
val counts = file.flatMap(line => analyzer.analyze("anything", line))
                  .map(word => (word, 1))
                  .reduceByKey(_ + _)
                  .sortBy(_._2, false) // descending sort by count
println(counts.take(10).map(t => s"${t._1}(${t._2})").mkString(", "))
```

```
the(158), to(103), solr(86), spark(77), a(72), in(44), you(44), of(40), for(35), from(34)
```



Lucene Text Analysis (example)

- Word count example of spark-solr README file with stop word filtering

```
import com.lucidworks.spark.analysis.LuceneTextAnalyzer
val schema = """{ "analyzers": [{ "name": "StdTokLower",
    |                               "tokenizer": { "type": "standard" },
    |                               "filters": [{ "type": "lowercase" }] },
    |   { "name": "StdTokLowerStop",
    |       "tokenizer": { "type": "standard" },
    |       "filters": [{ "type": "lowercase" },
    |                   { "type": "stop" }] }],
    |   "fields": [{ "name": "all_tokens", "analyzer": "StdTokLower" },
    |               { "name": "no_stopwords", "analyzer": "StdTokLowerStop" } ]}
    """.stripMargin
val analyzer = new LuceneTextAnalyzer(schema)
val file = sc.textFile("README.adoc")
val counts = file.flatMap(line => analyzer.analyze("no_stopwords", line))
    .map(word => (word, 1))
    .reduceByKey(_ + _)
    .sortBy(_._2, false)
println(counts.take(10).map(t => s"${t._1} (${t._2})").mkString(", "))
```

```
solr(86), spark(77), you(44), from(34), source(32), query(25), option(25), collection(24), data(20), can(19)
```



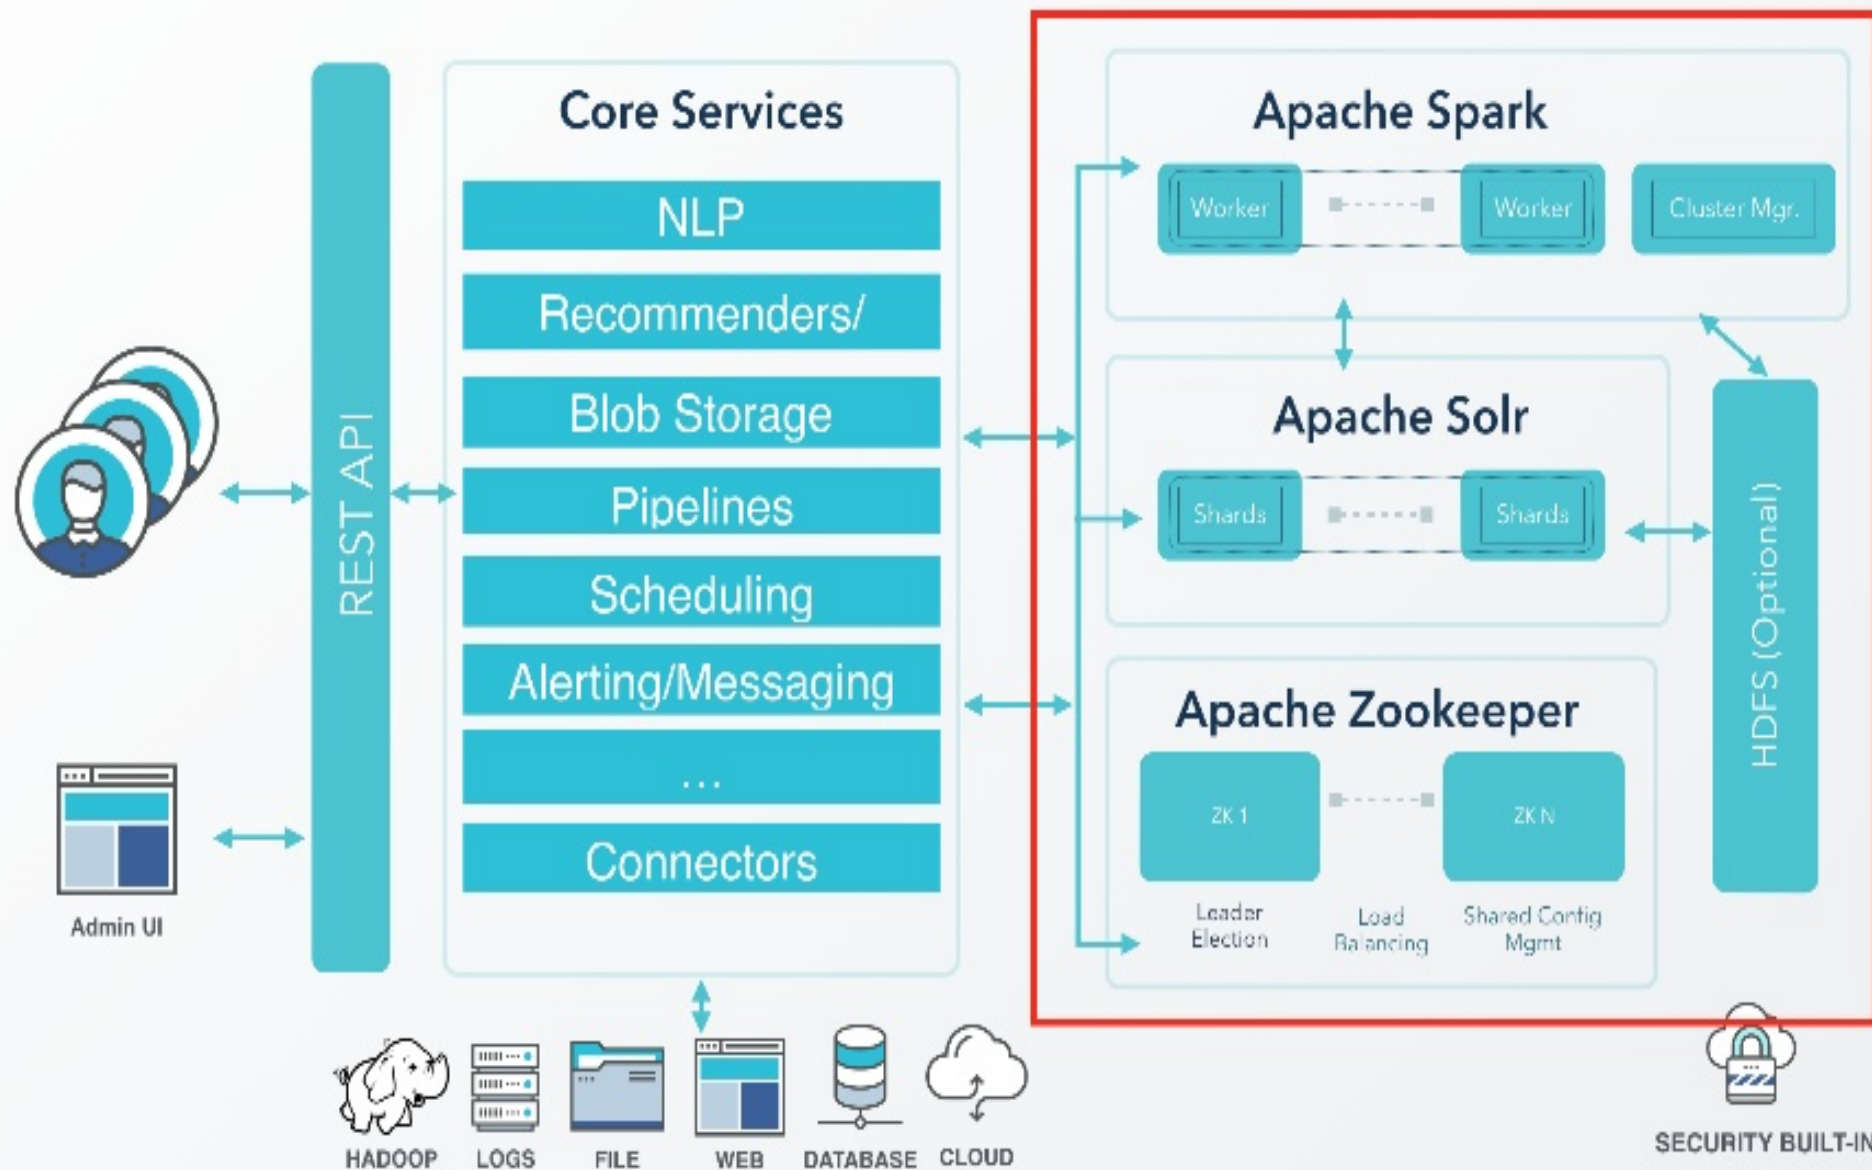
Lucene ML transformer (example)

- LuceneTextAnalyzerTransformer as an ML pipeline transformer
- Supports multi-valued input columns

```
val WhitespaceTokSchema =
  """{ "analyzers": [{ "name": "ws_tok", "tokenizer": { "type": "whitespace" } }],
    | "fields": [{ "regex": ".*", "analyzer": "ws_tok" }] }""".stripMargin
val StdTokLowerSchema =
  """{ "analyzers": [{ "name": "std_tok_lower", "tokenizer": { "type": "standard" },
    | "filters": [{ "type": "lowercase" } ] },
    | "fields": [{ "regex": ".*", "analyzer": "std_tok_lower" }] }""".stripMargin
[...]
val analyzer = new LuceneTextAnalyzerTransformer().setInputCols(contentFields).setOutputCol(WordsCol)
[...]
val pipeline = new Pipeline().setStages(Array(labelIndexer, analyzer, hashingTF, estimatorStage,
labelConverter))
```



Fusion Architecture



Fusion & Spark

- spark-solr 2.0.1 released, built into Fusion 2.4
- Users leave evidence of their needs & experience as they use your app
- Fusion closes the feedback loop to improve results based on user “signals”
- Train and serve ML Pipeline and mllib based Machine Learning models
- Run custom Scala “script” jobs in the background in Fusion
 - Complex aggregation jobs
 - Unsupervised learning (LDA topic modeling)
 - Re-train supervised models as new training data flows in



Getting started with spark-solr

- Import package via maven. Available at spark-packages.org

```
./bin/spark-shell --packages "com.lucidworks.spark:spark-solr:2.0.1"
```

- Build from source

```
git clone https://github.com/Lucidworks/spark-solr
```

```
cd spark-solr
```

```
mvn clean package -DskipTests
```

```
./bin/spark-shell --jars 2.1.0-SNAPSHOT.jar
```



Example : Deep paging via shards

```
// Connect to Solr  
val opts = Map(  
  "zkhost" -> "localhost:9983",  
  "collection" -> "nyc_trips")  
  
val solrDF = sqlContext.read.format("solr").options(opts).load  
  
// Register DF as temp table  
solrDF.registerTempTable("trips")  
  
sqlContext.sql("SELECT * FROM trips LIMIT 2").show()
```


Example : Deep paging with intra shard splitting

```
// Connect to Solr
val opts = Map(
  "zkhost" -> "localhost:9983",
  "collection" -> "nyc_trips",
  "splits" -> "true")
val solrDF = sqlContext.read.format("solr").options(opts).load

// Register DF as temp table
solrDF.registerTempTable("trips")

sqlContext.sql("SELECT * FROM trips").count()
```


Example : Streaming API (/export handler)

```
// Connect to Solr
val opts = Map(
  "zkhost" -> "localhost:9983",
  "collection" -> "nyc_trips")
val solrDF = sqlContext.read.format("solr").options(opts).load

// Register DF as temp table
solrDF.registerTempTable("trips")

sqlContext.sql("SELECT avg(tip_amount), avg(fare_amount) FROM
trips").show()
```

Performance test

- NYC taxi data (30 months - **91.7M** rows)
- Dataset loaded in to AWS RDS instance (Postgres)
- 3 EC2 nodes of r3.2x large instances
- Solr and Spark instances co-located together
- Collection 'nyc-taxi' created with 6 shards, 1 replication
- Deployed using solr-scale-tk (<https://github.com/LucidWorks/solr-scale-tk>)
- Dataset link: <https://github.com/toddwschneider/nyc-taxi-data>
- More details: <https://gist.github.com/kiranchitturi/0be62fc13e4ec7f9ae5def53180ed181>



Index performance

- **91.4M** rows indexed to Solr in 49 minutes
- Docs per second: **31K**
- JDBC batch size: 5000
- Indexing batch size: 50000
- Partitions: 200



Query performance

- Query - simple aggregation query to calculate averages
- Streaming expressions took **2.3** mins across 6 tasks
- Deep paging took **20** minutes across 120 tasks



Roadmap

- Implement Datasource with Catalyst scan.
- Fallback to different replica if the preferred replica is down
- HashQParserPlugin for intra-shard splits
- Output PreAnalyzedField compatible JSON for Solr

THANK YOU.

kiran.chitturi@lucidworks.com / @chitturikiran

<https://github.com/Lucidworks/spark-solr/>

Download Fusion: <http://lucidworks.com/fusion/download/>



SPARK SUMMIT 2016
DATA SCIENCE AND ENGINEERING AT SCALE